

TRASH TALK



EXPLORING THE MEMORY MANAGEMENT IN THE JVM

ABOUT ME.



Gerrit Grunwald | Developer Advocate | Azul

MEMORY MANAGEMENT

IN THE JVM...

IS
AUTOMATIC
RIGHT...?

**SO...WHY
CARE...?**

MEMORY MANAGEMENT

Why you should care...



Impact on application performance

MEMORY MANAGEMENT

Why you should care...

Impact on application performance

Impact on application responsiveness

MEMORY MANAGEMENT

Why you should care...



Impact on application performance



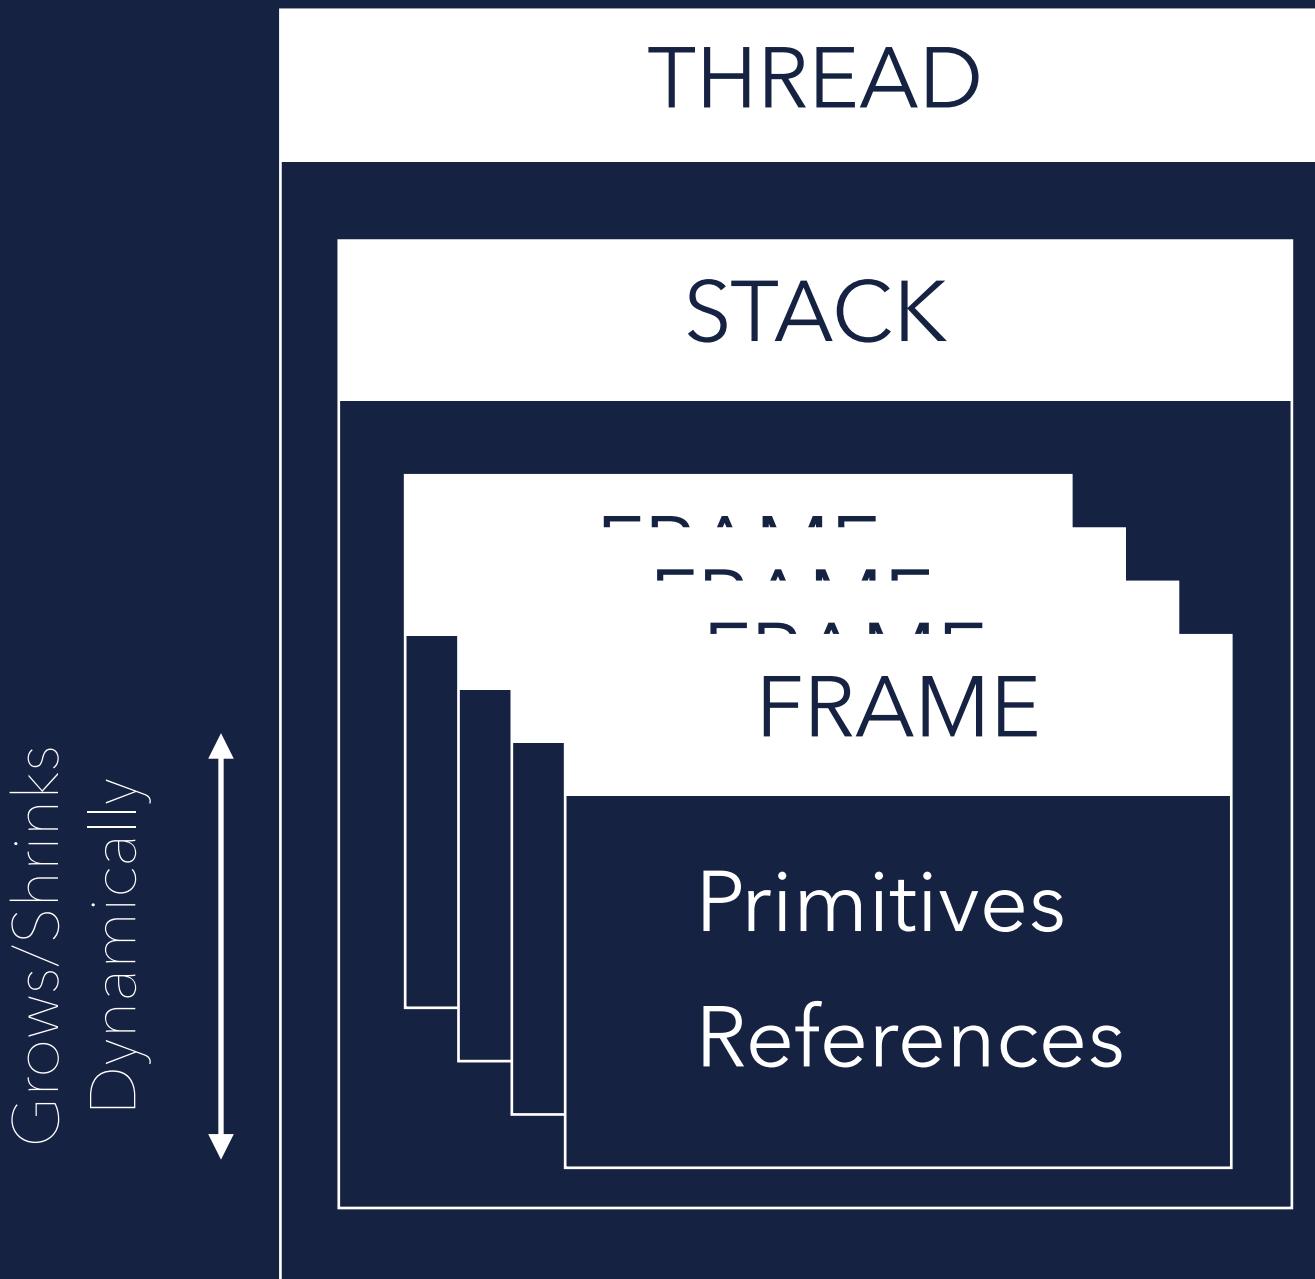
Impact on application responsiveness



Impact on system requirements

MEMORY MANAGEMENT

Stack, Heap and Metaspace



Local access -> thread safe

MEMORY MANAGEMENT

Stack, Heap and Metaspace

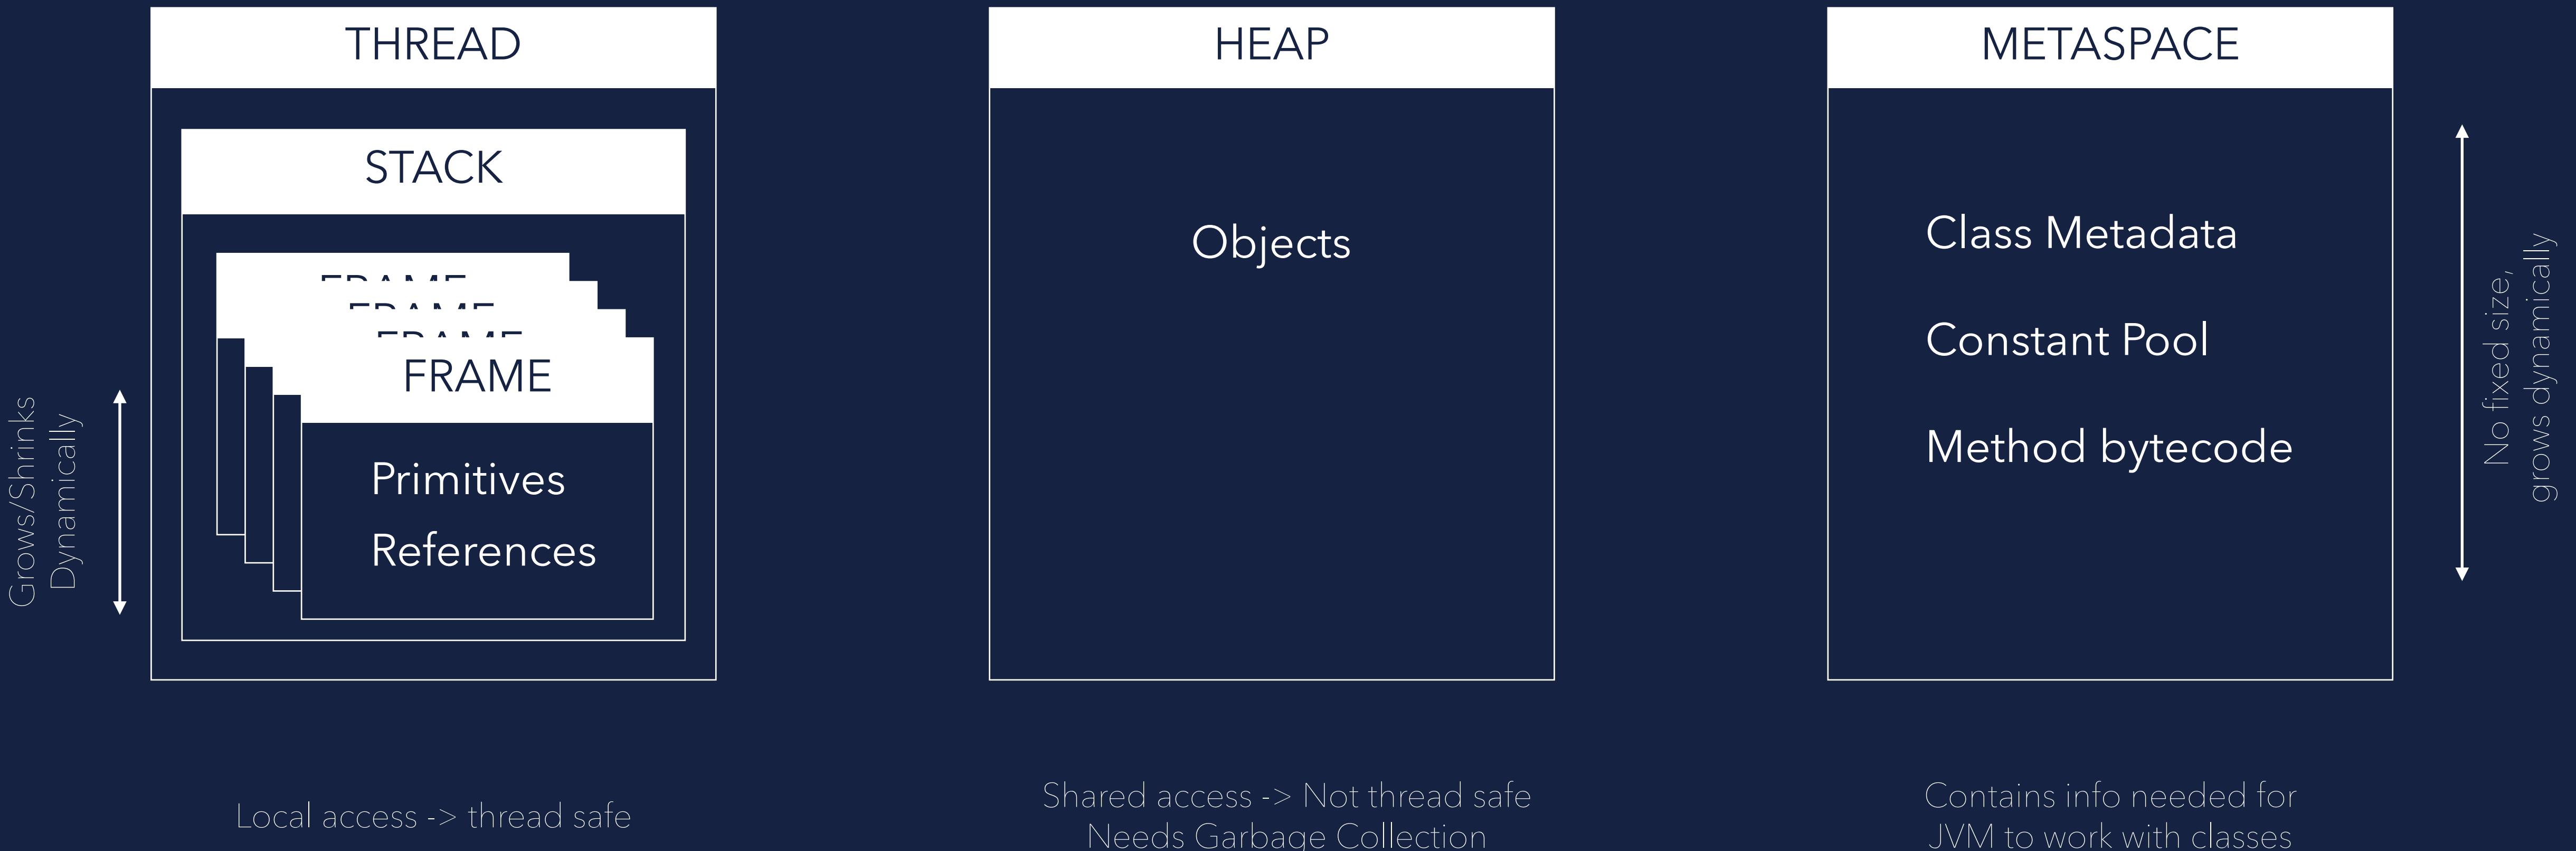


Local access -> thread safe

Shared access -> Not thread safe
Needs Garbage Collection

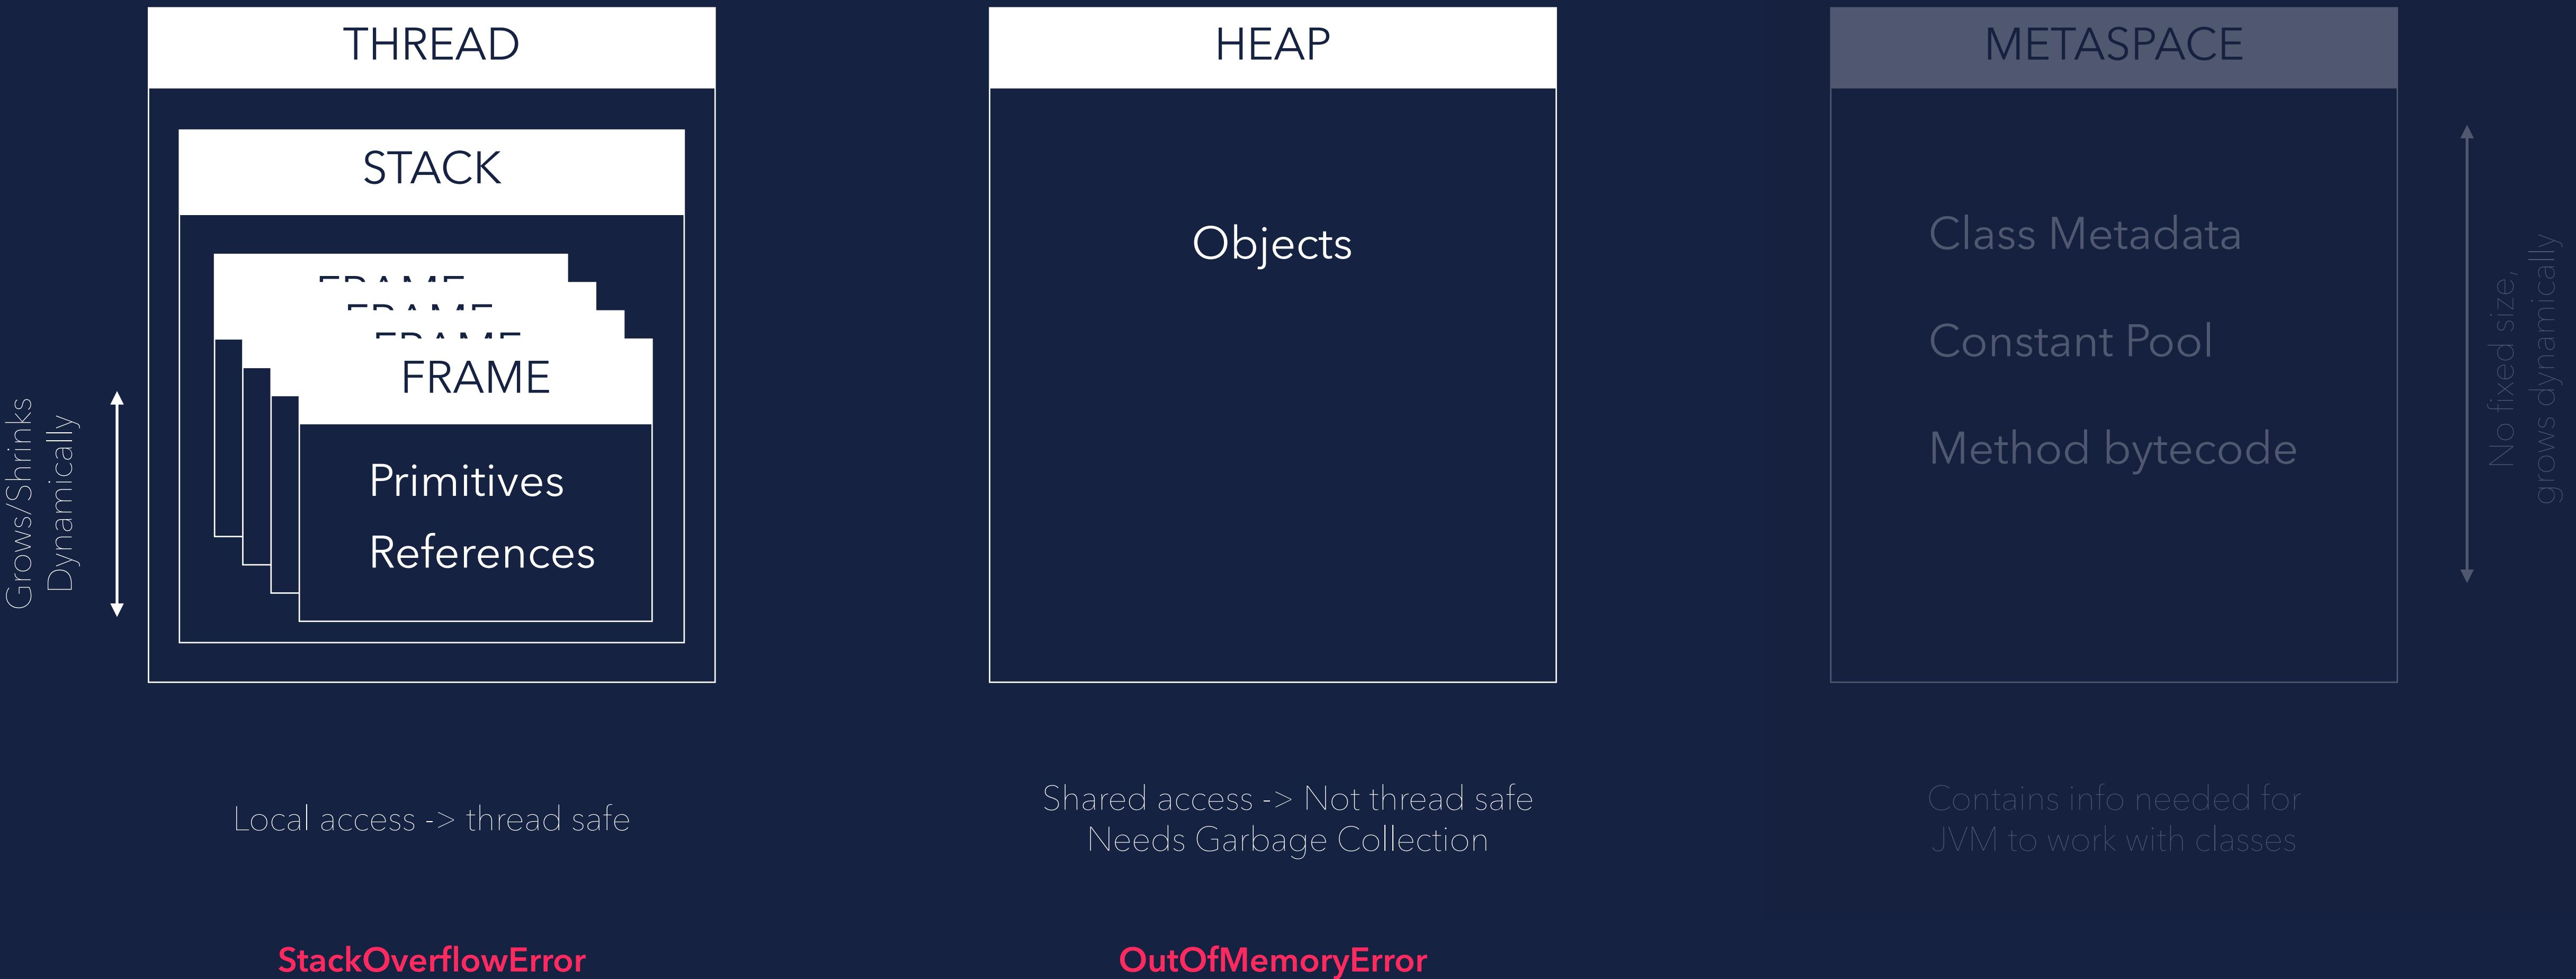
MEMORY MANAGEMENT

Stack, Heap and Metaspace



MEMORY MANAGEMENT

Stack, Heap and Metaspace



MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
}
```

MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
}
```

Stack for thread 1

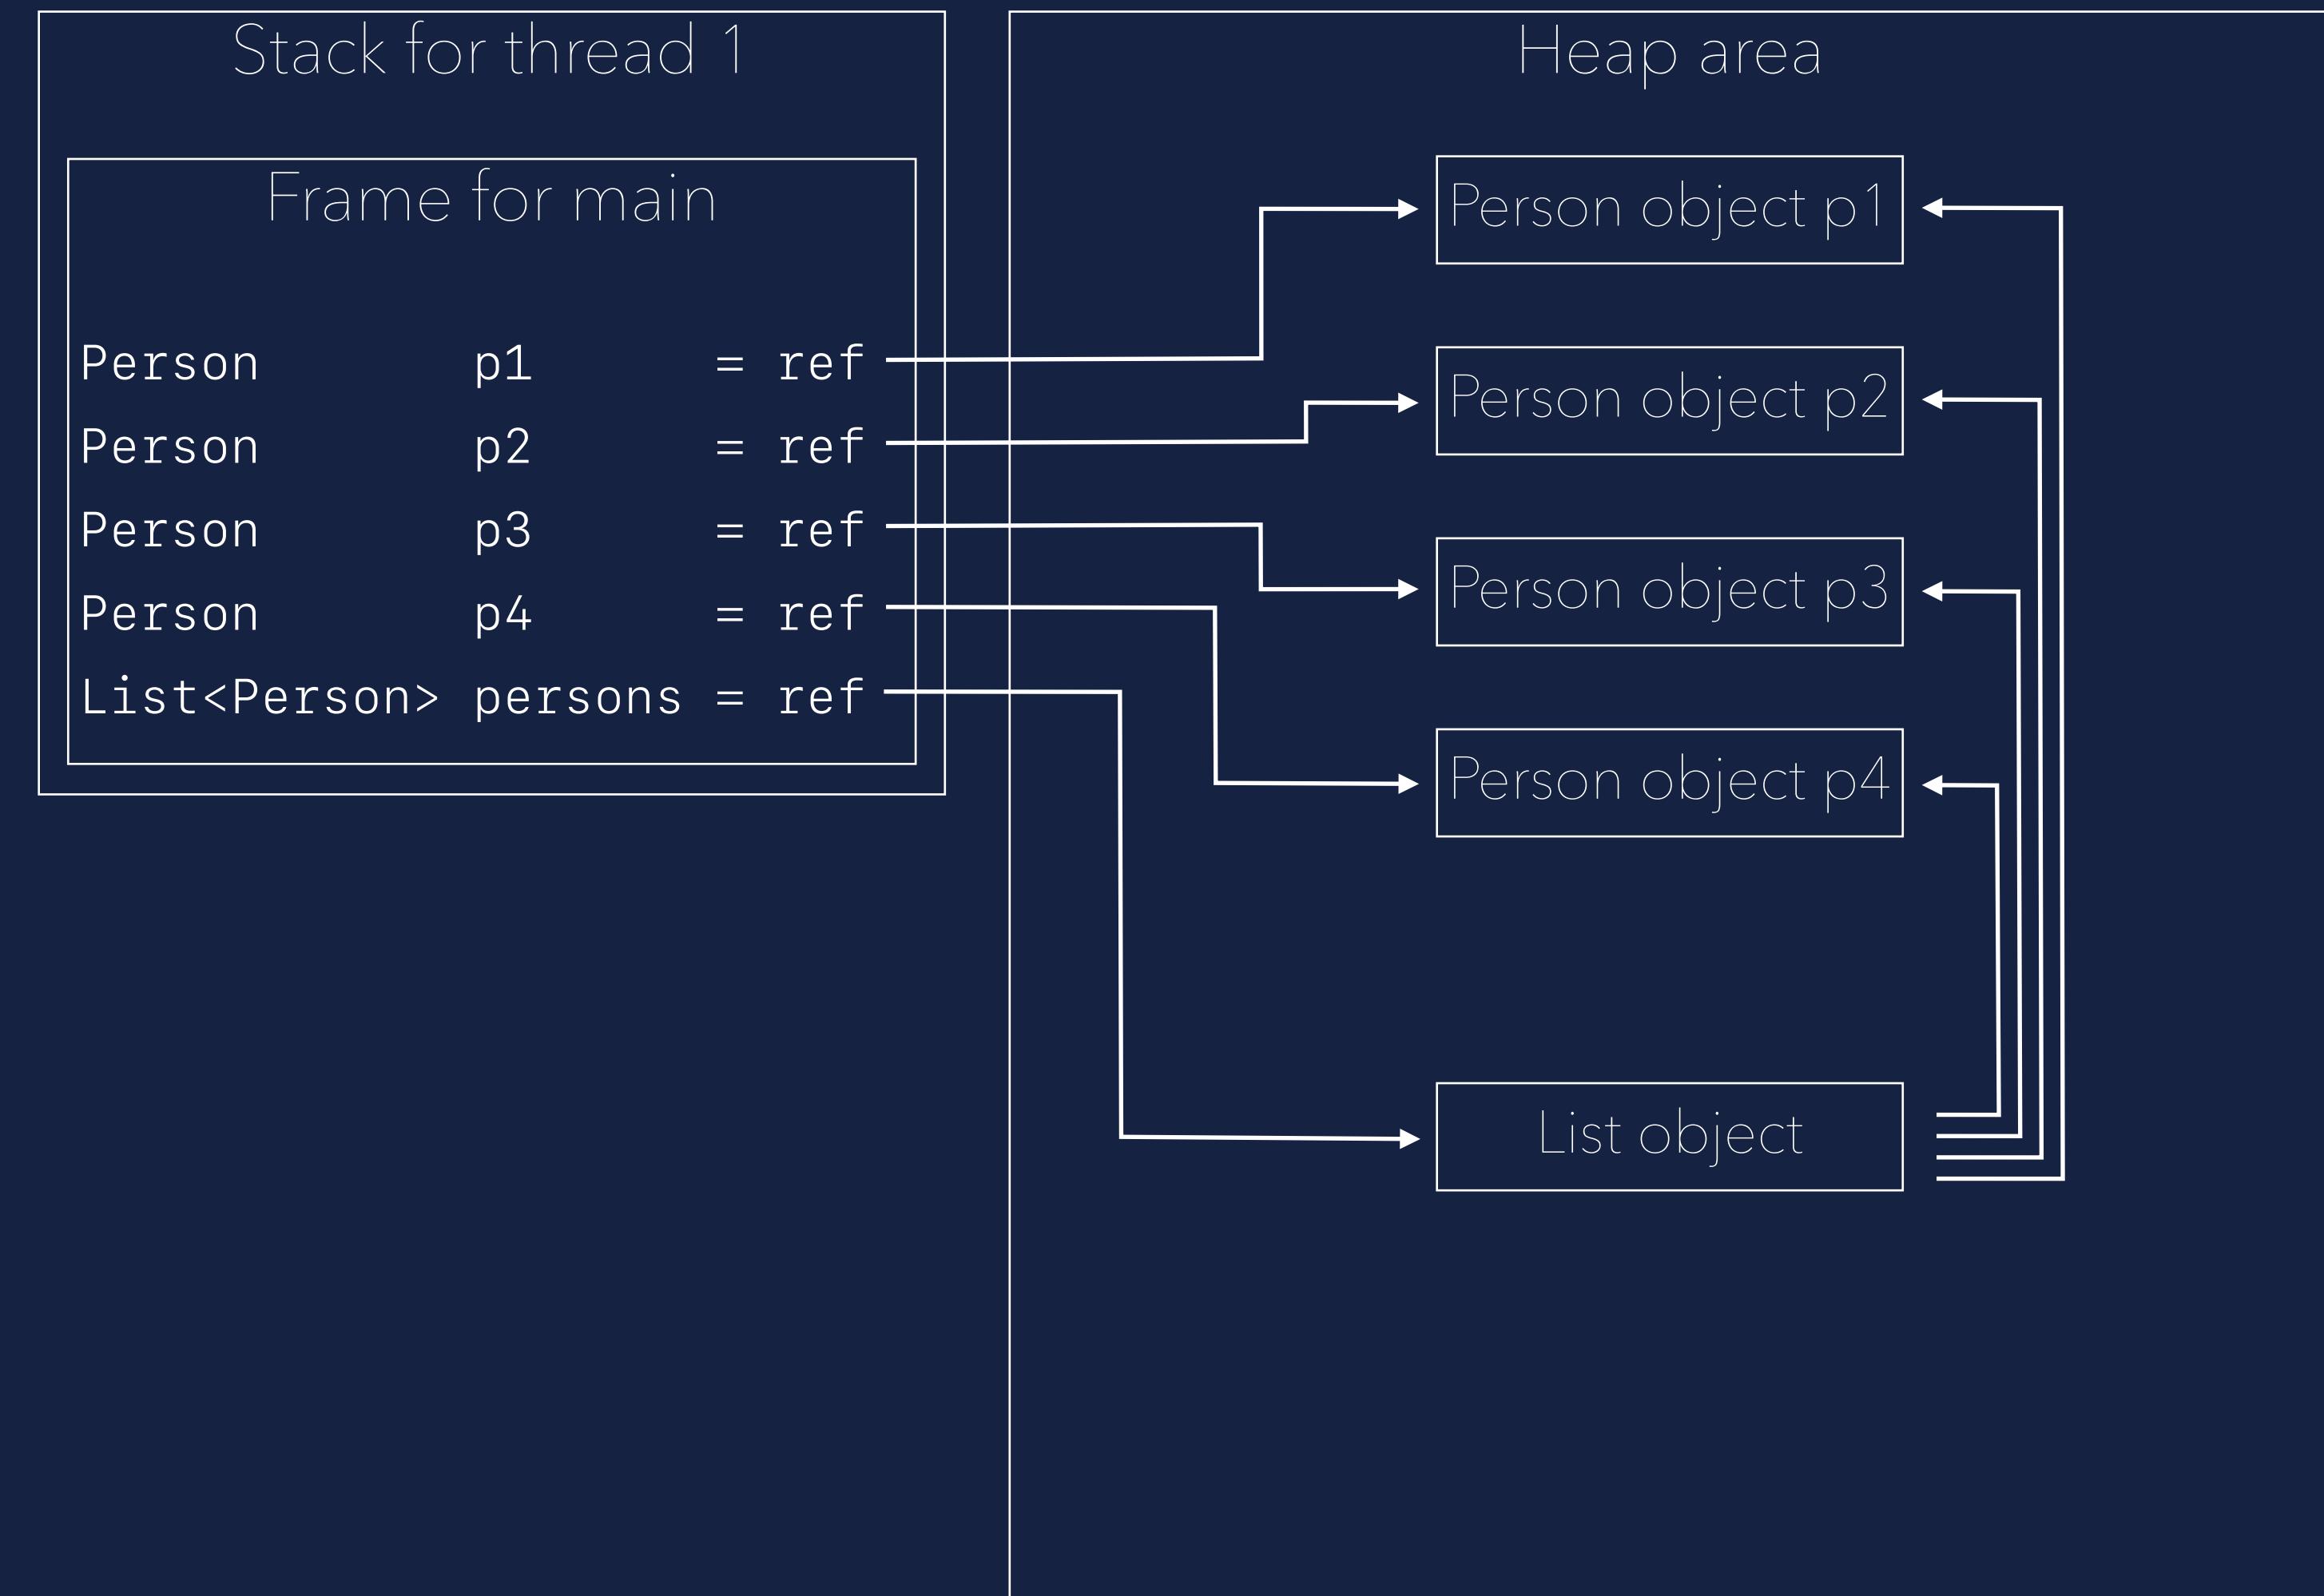
Frame for main

Person	p1	= ref
Person	p2	= ref
Person	p3	= ref
Person	p4	= ref
List<Person>	persons	= ref

MEMORY MANAGEMENT

In the JVM...

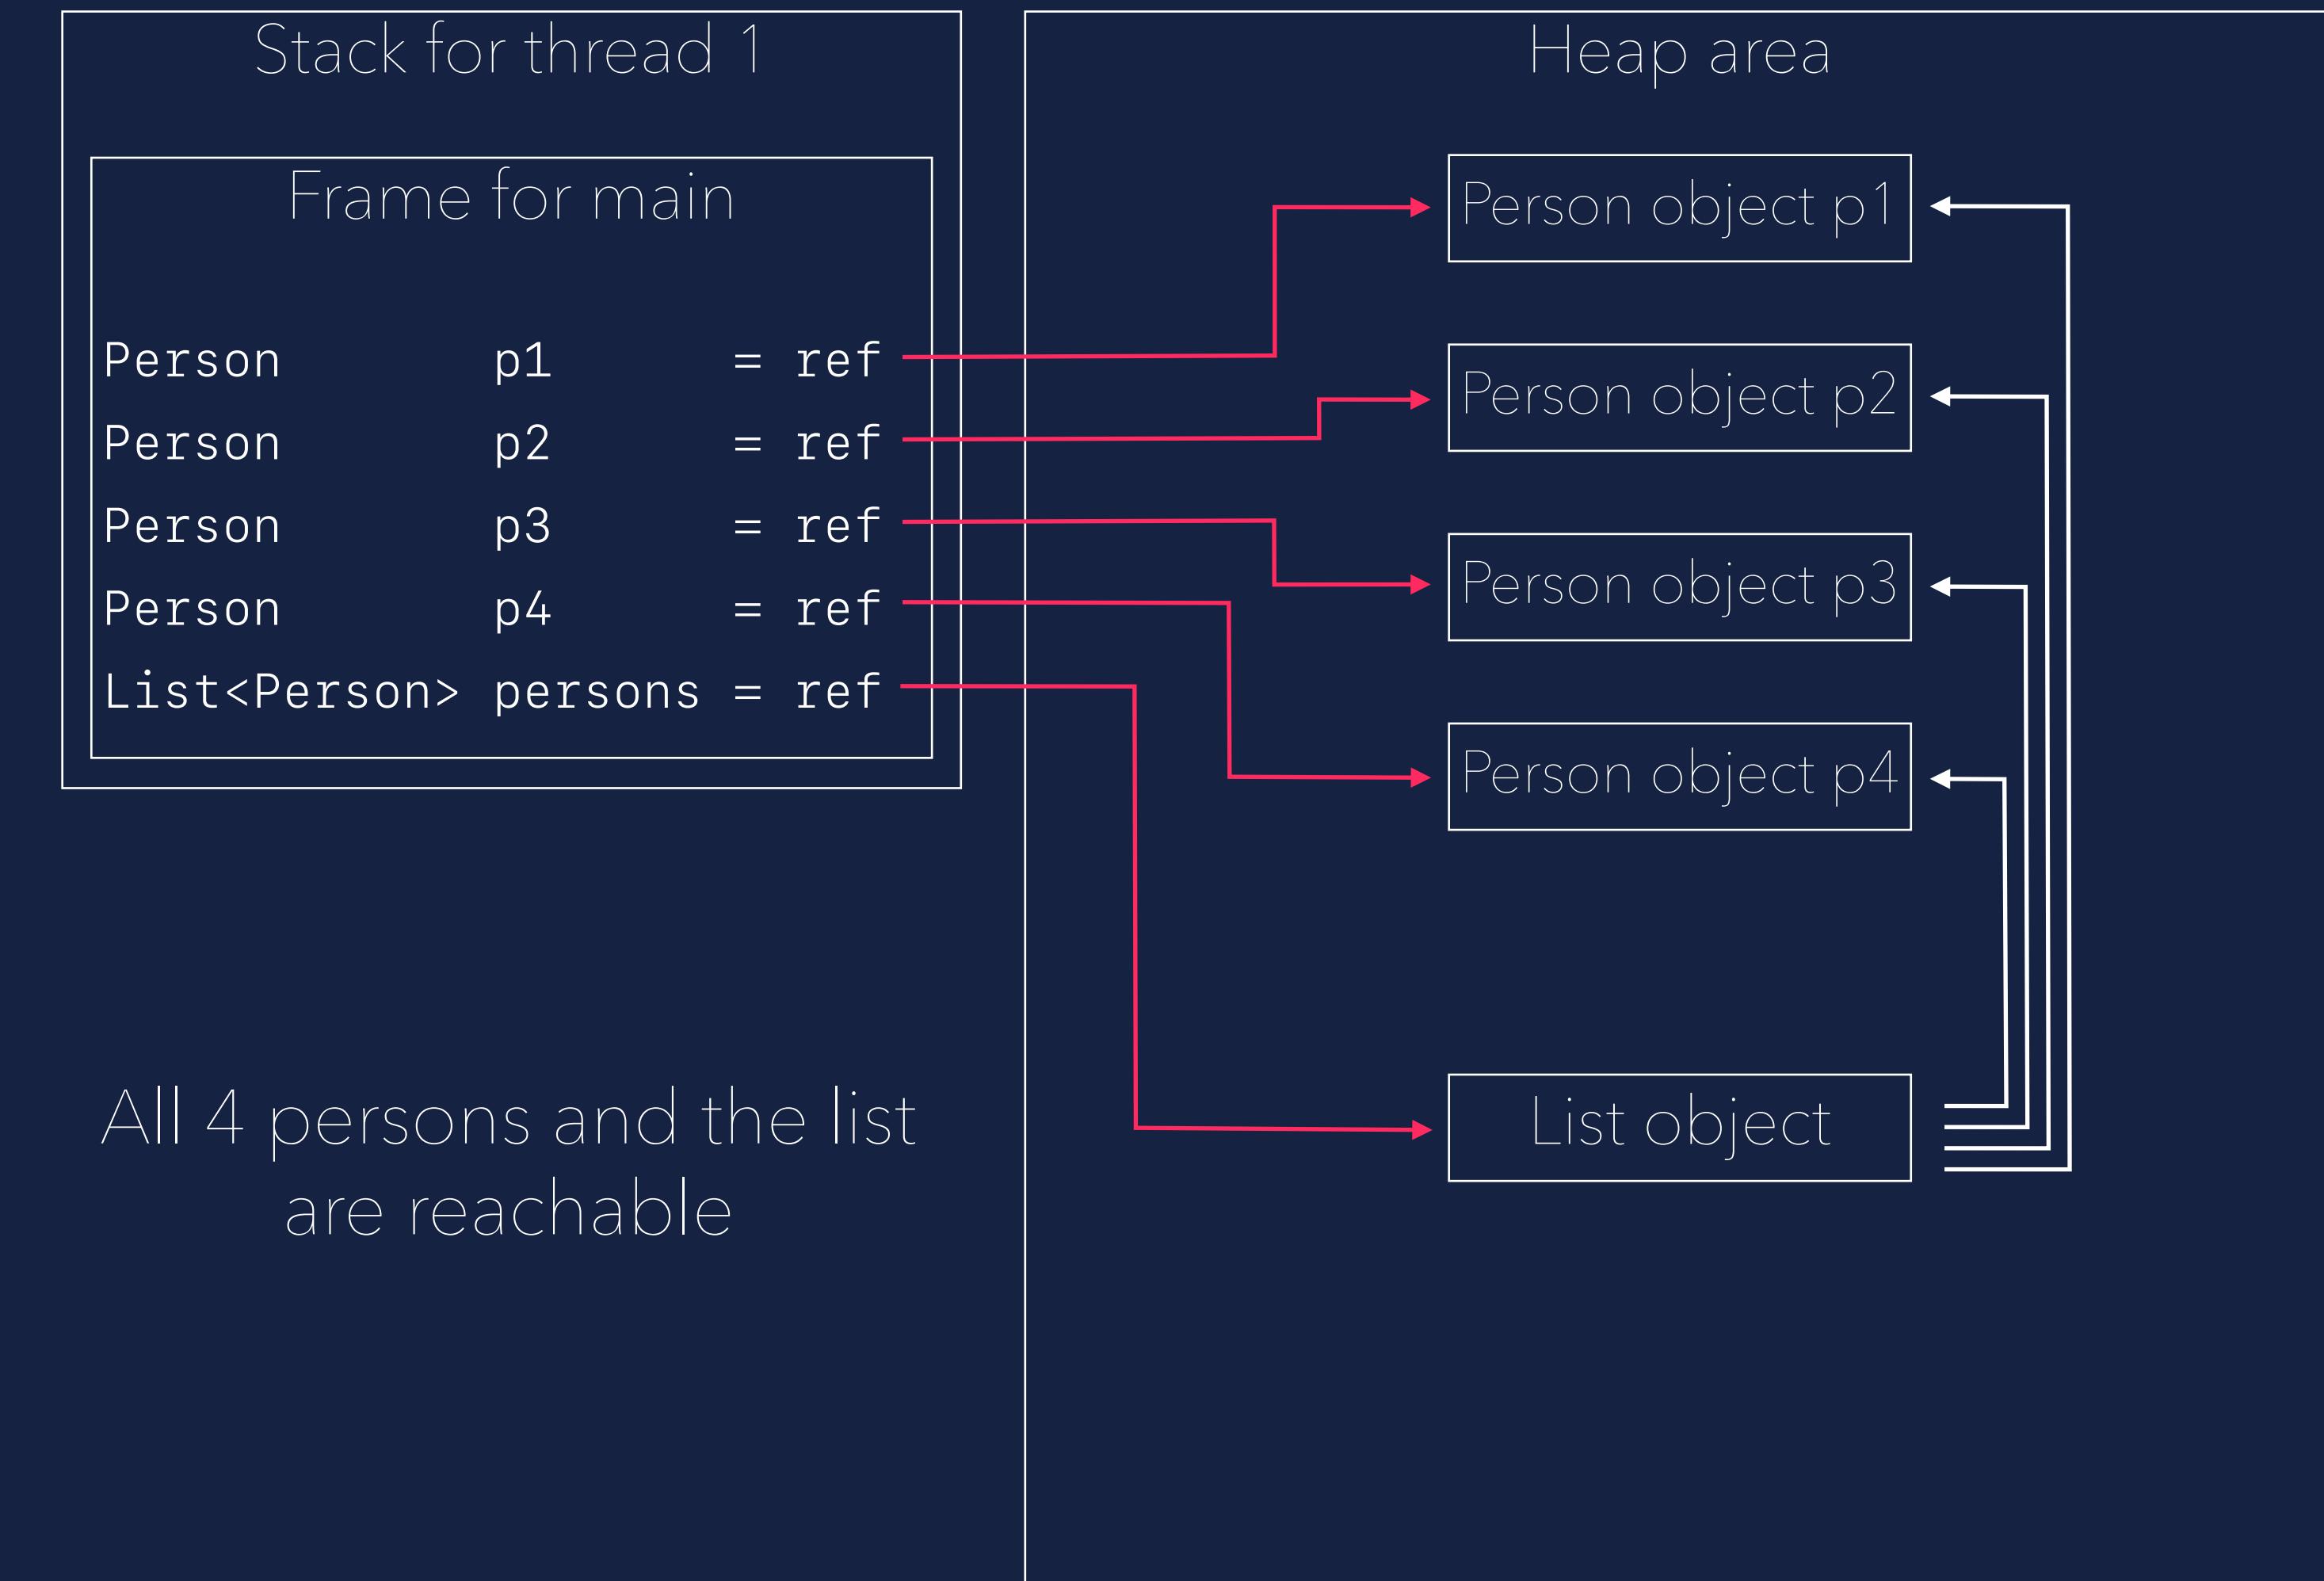
```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
}
```



MEMORY MANAGEMENT

In the JVM...

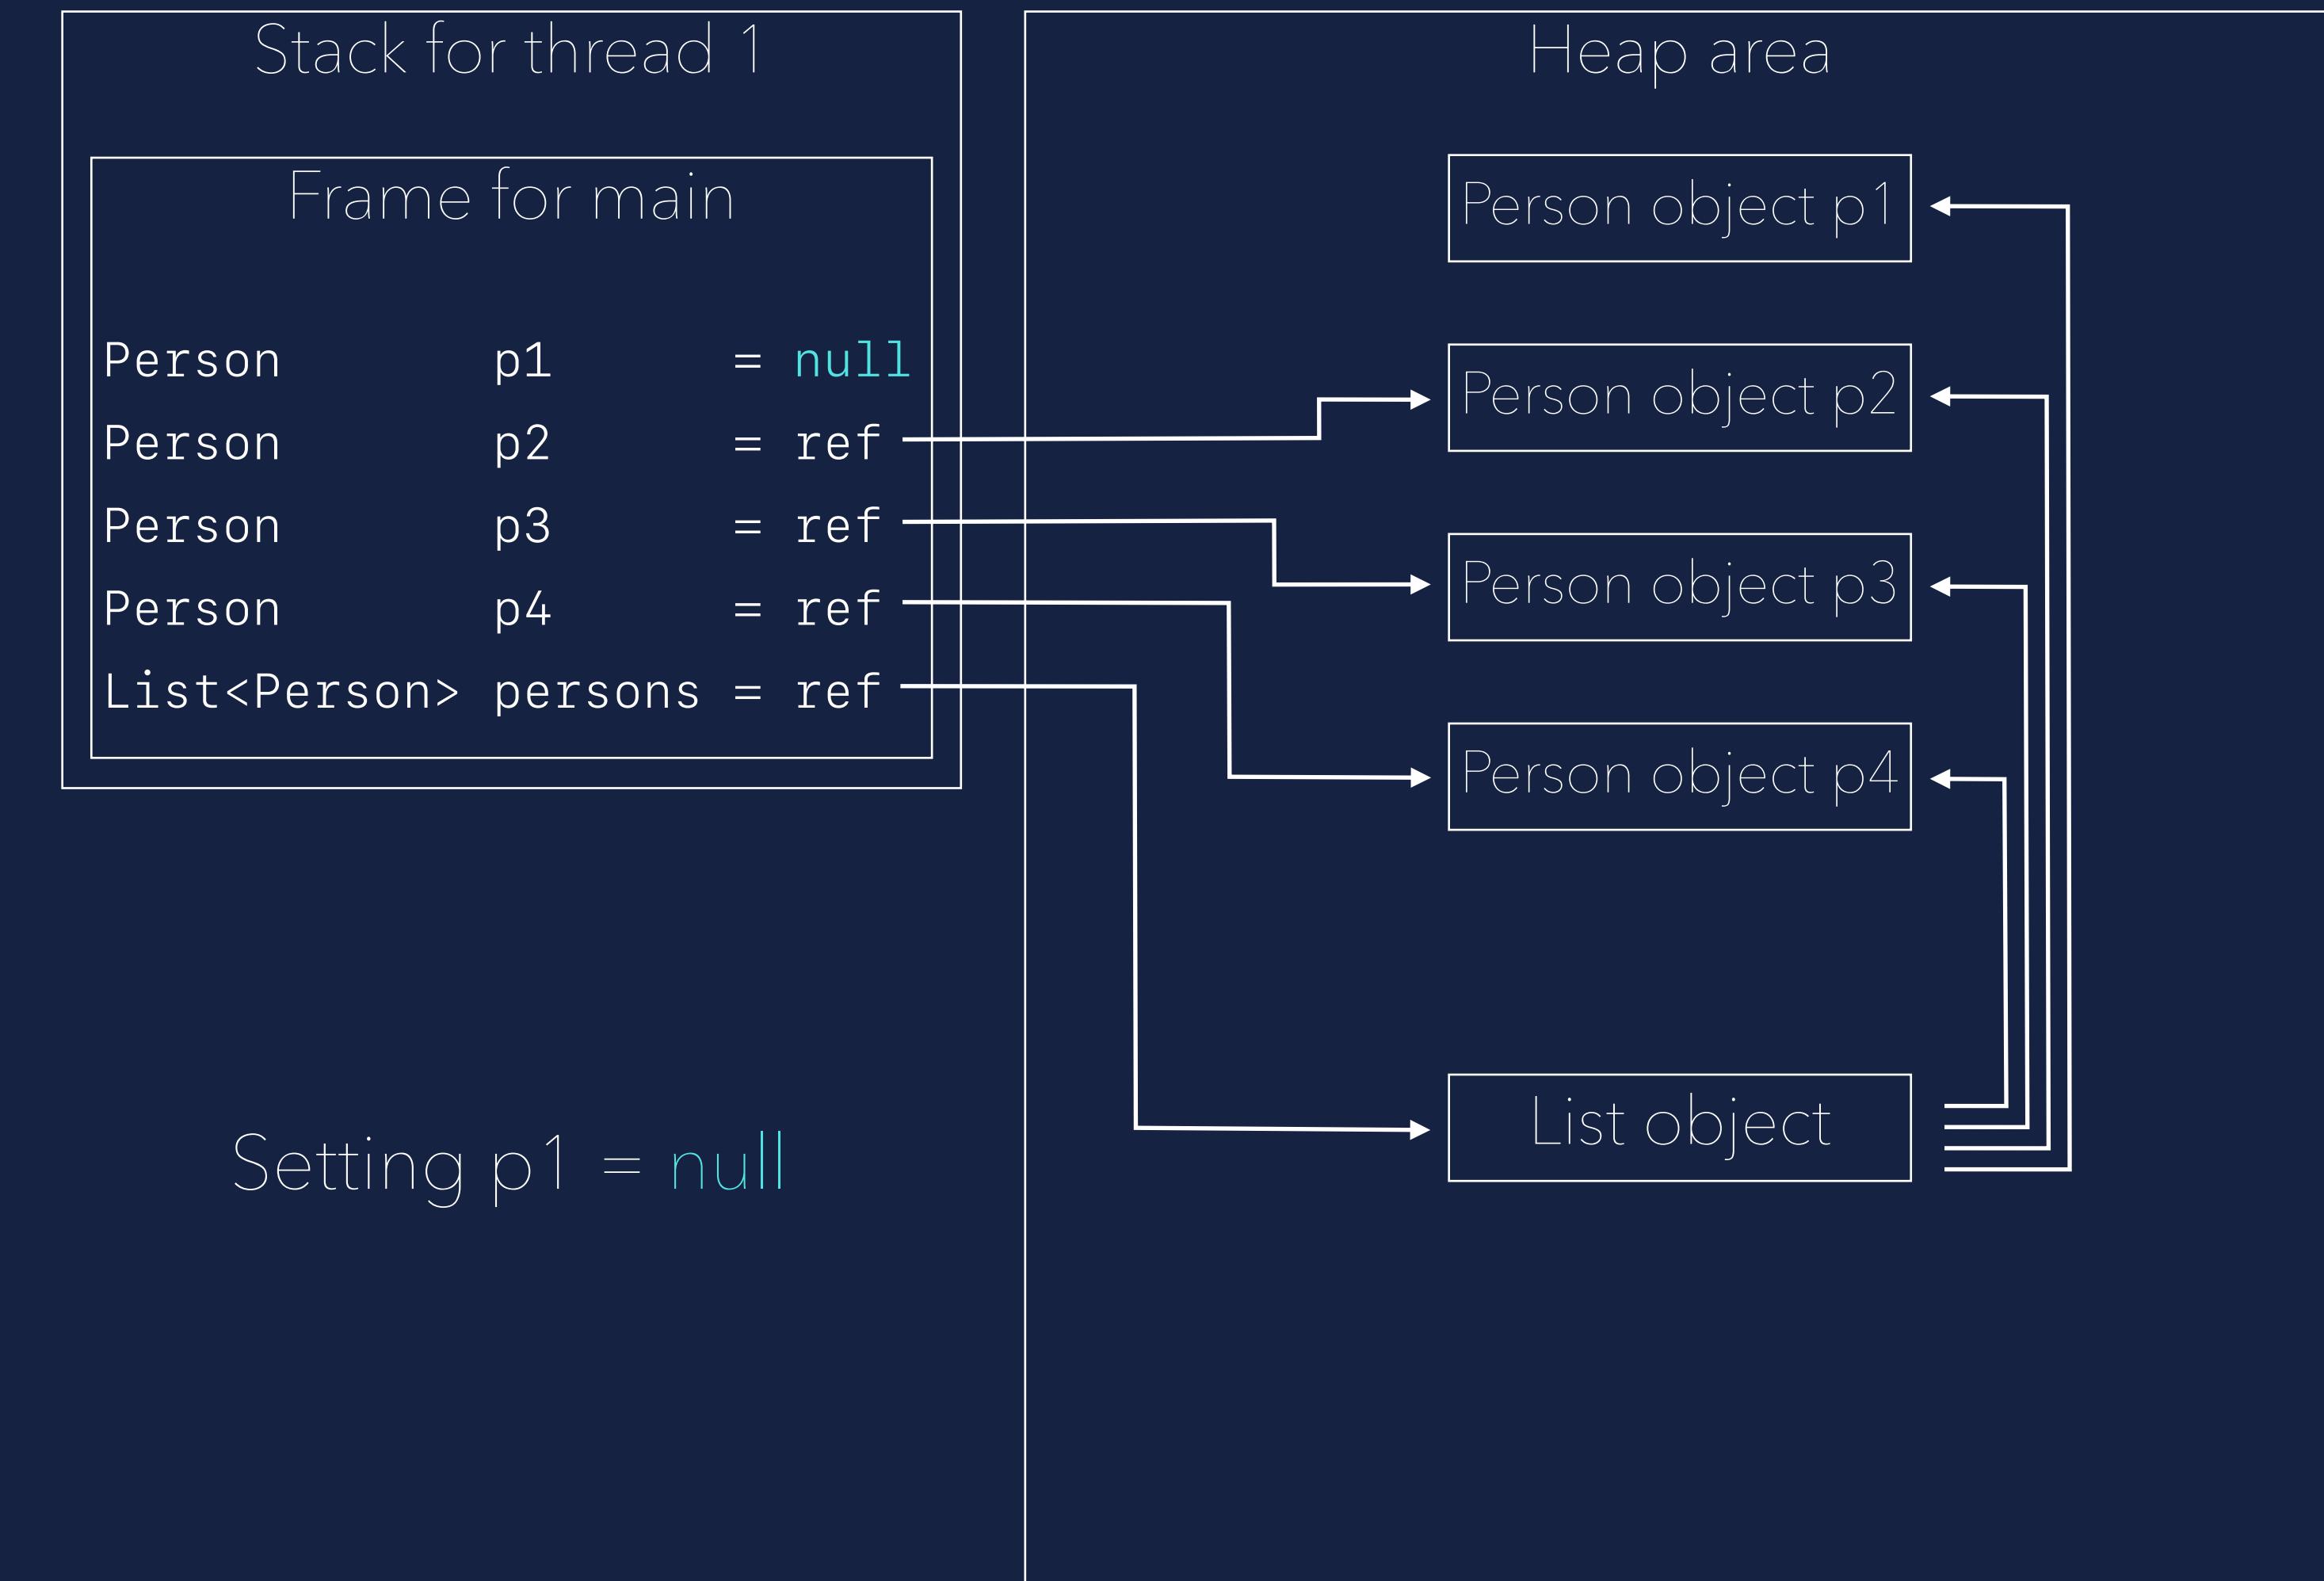
```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
}
```



MEMORY MANAGEMENT

In the JVM...

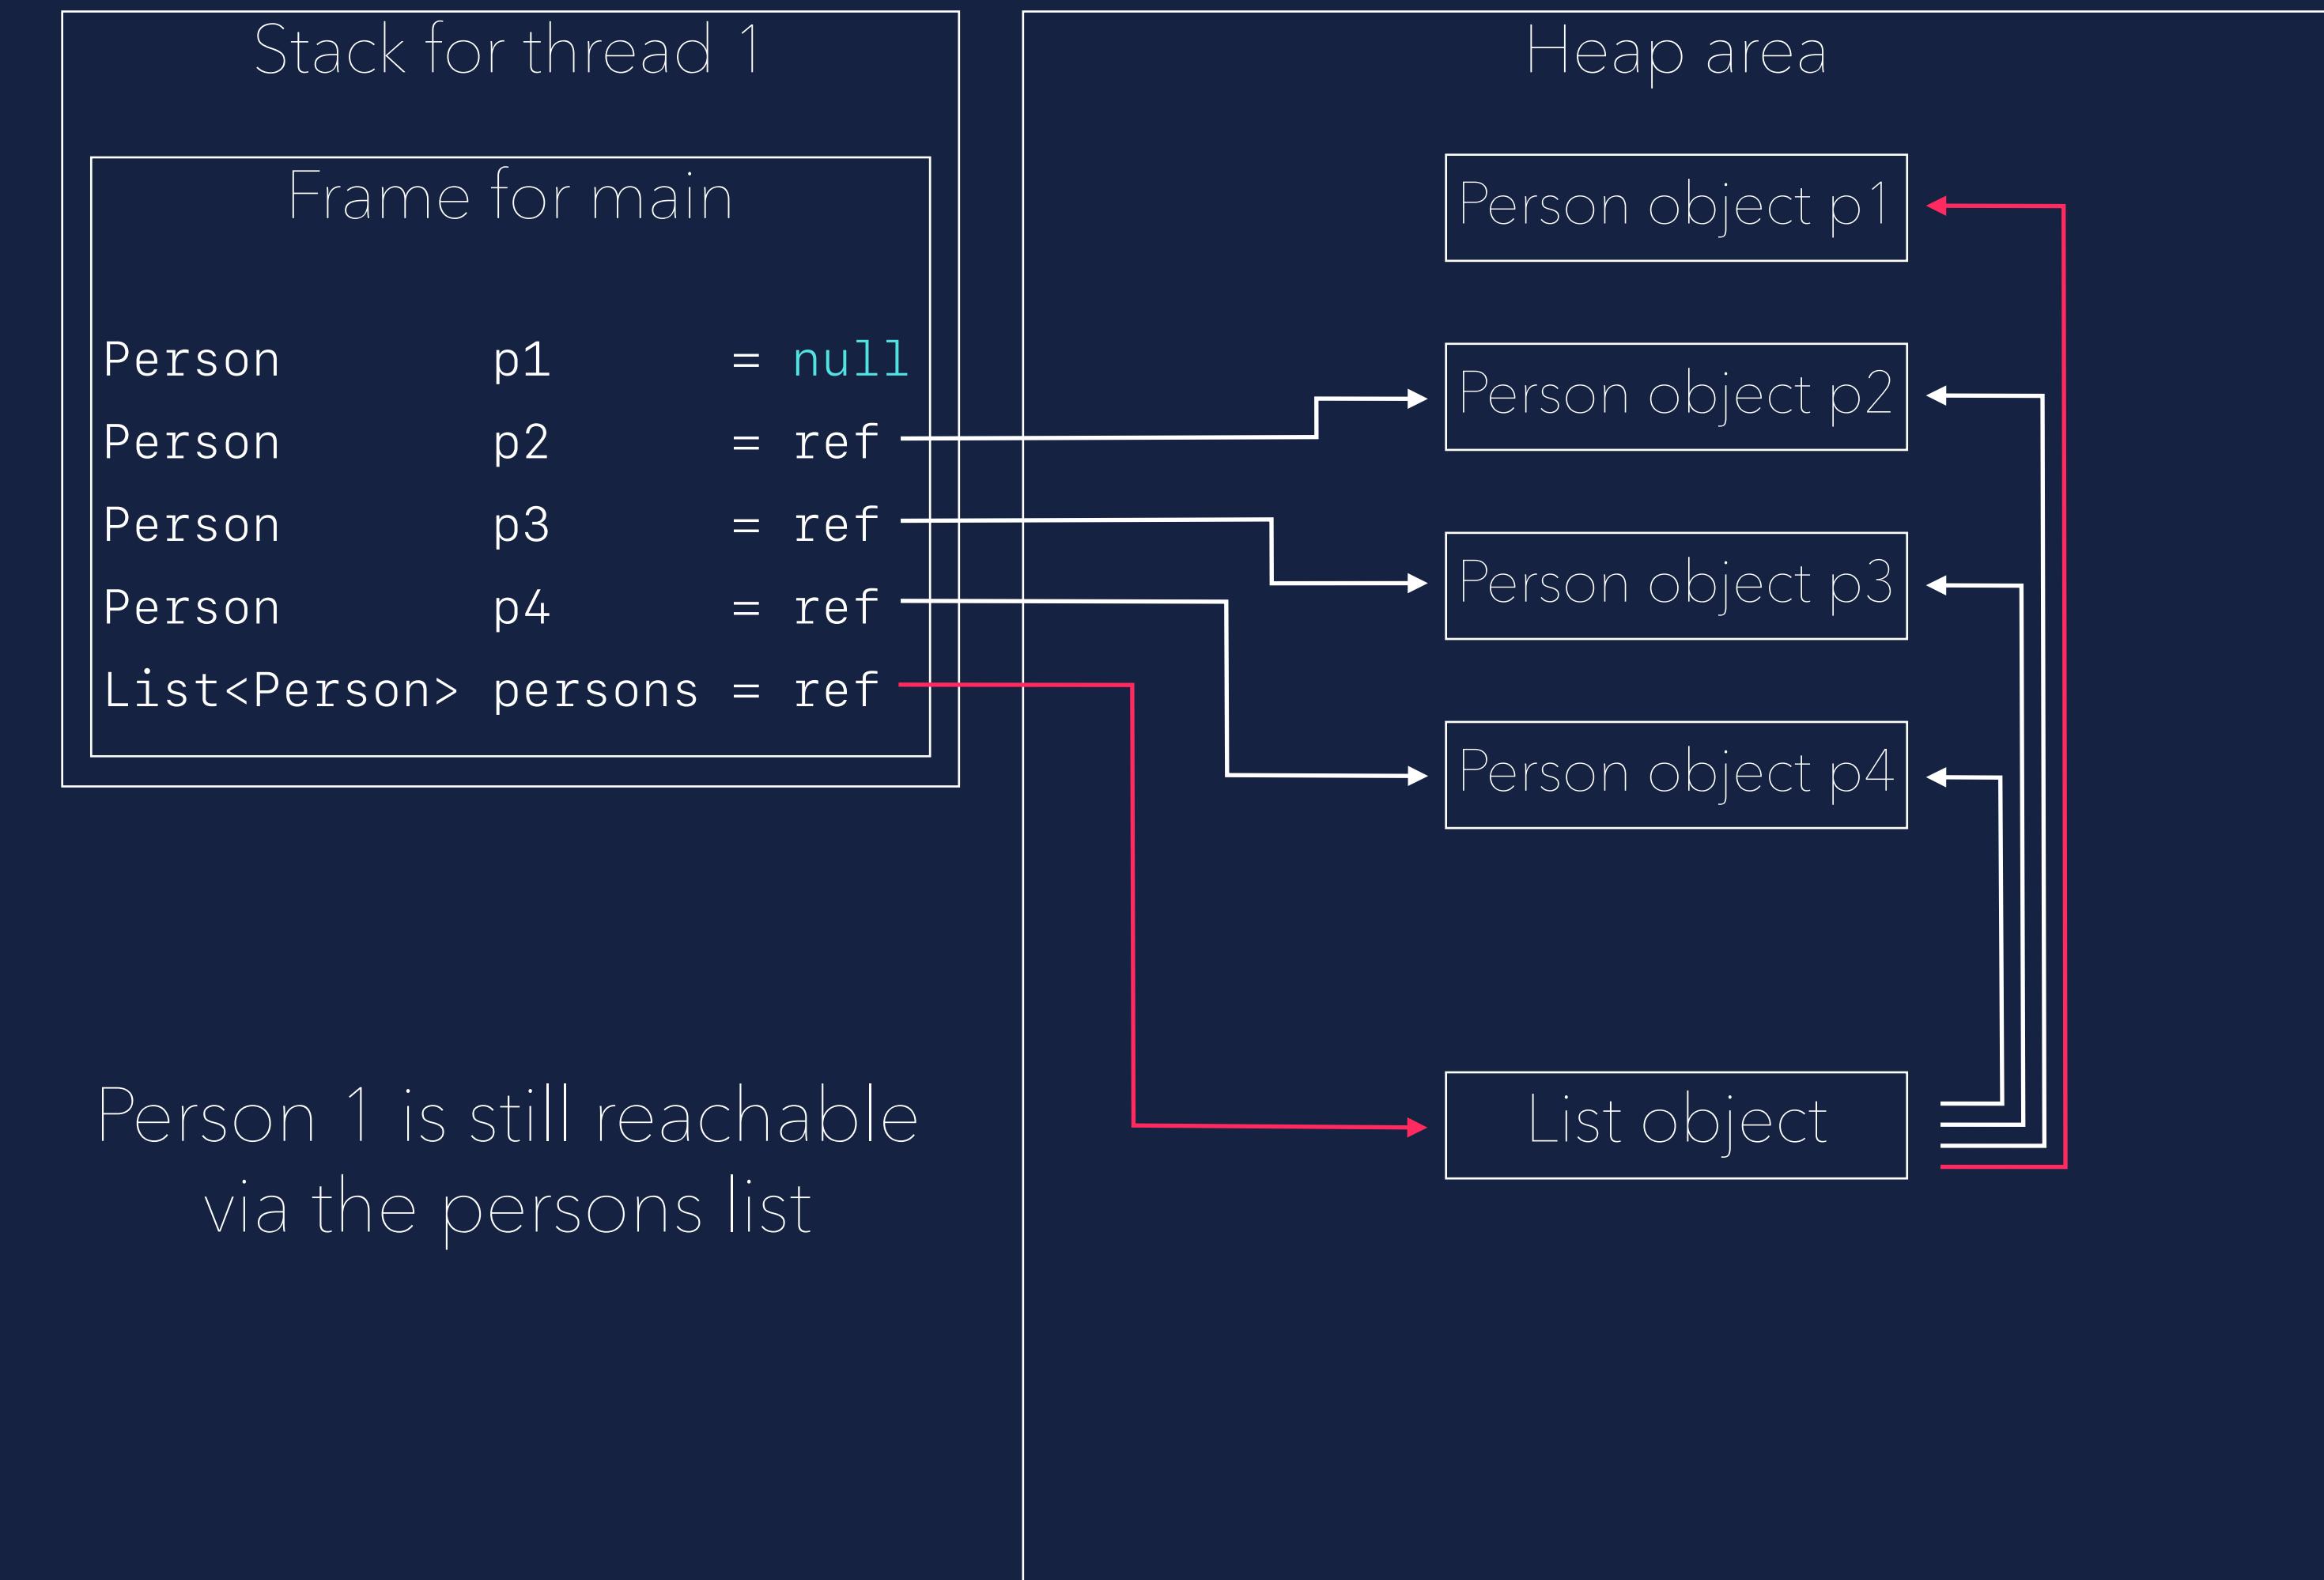
```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
}
```



MEMORY MANAGEMENT

In the JVM...

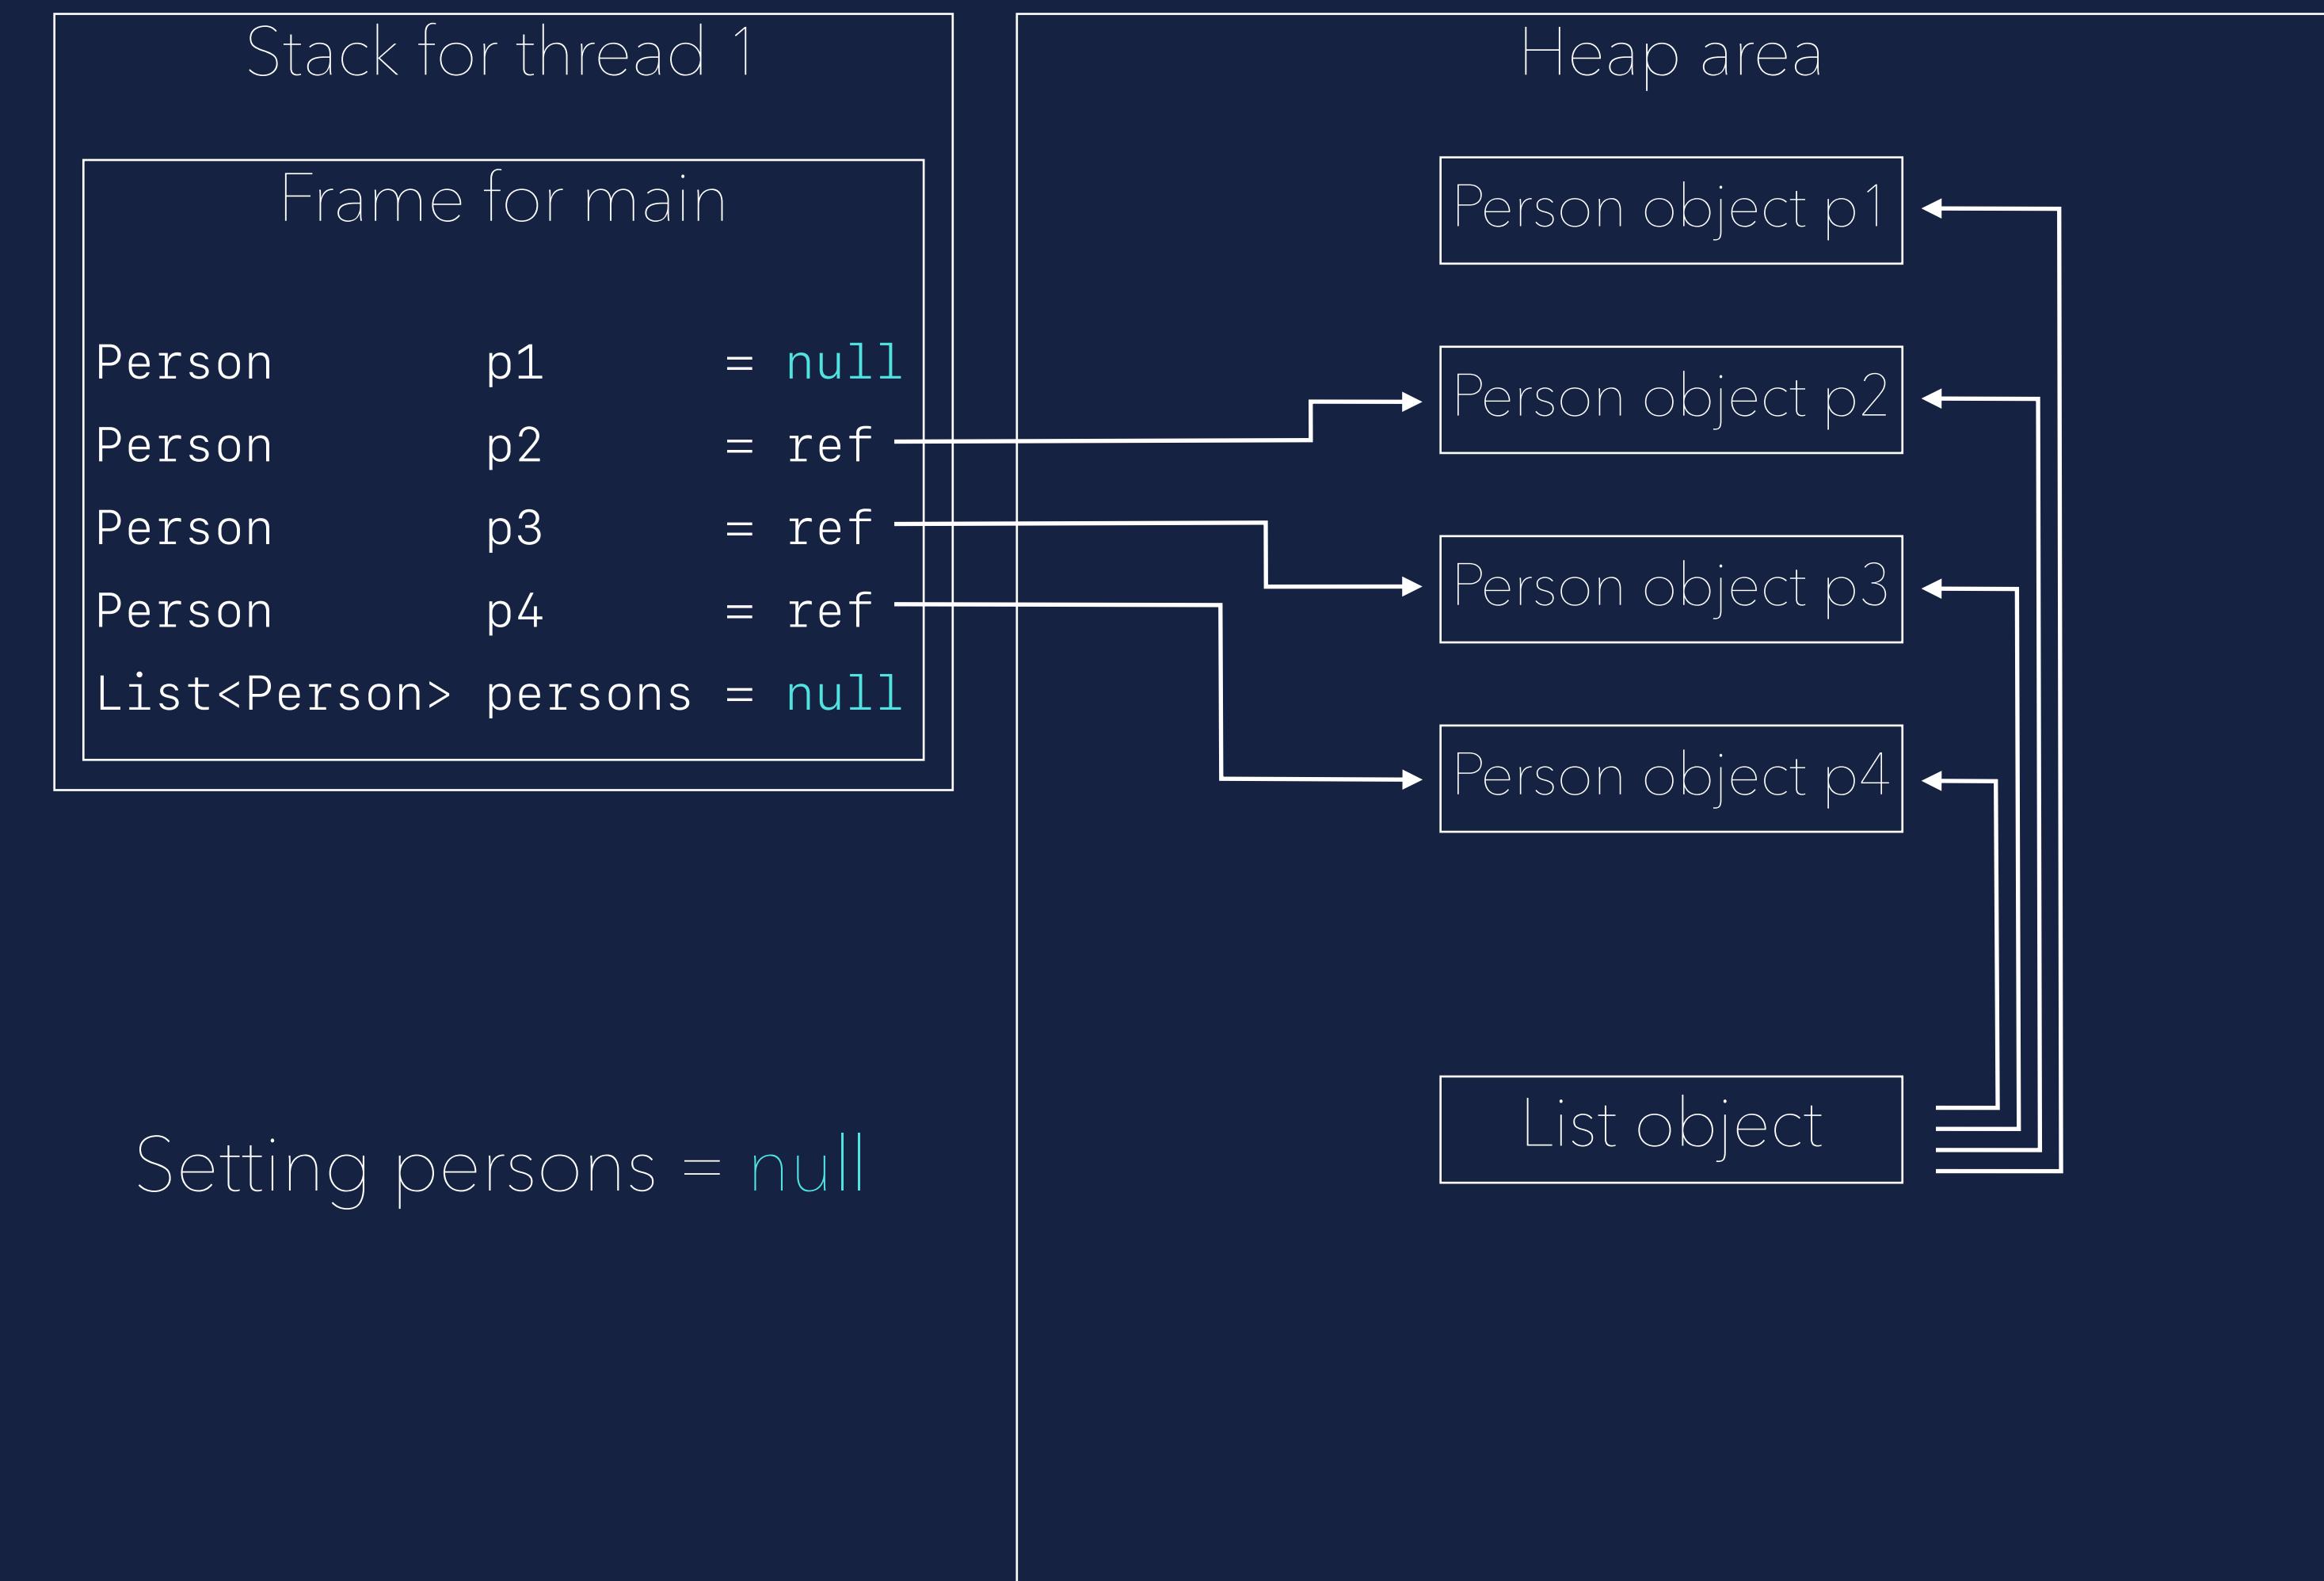
```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
  
    System.out.println(persons.get(0)); // -> Gerrit  
}
```



MEMORY MANAGEMENT

In the JVM...

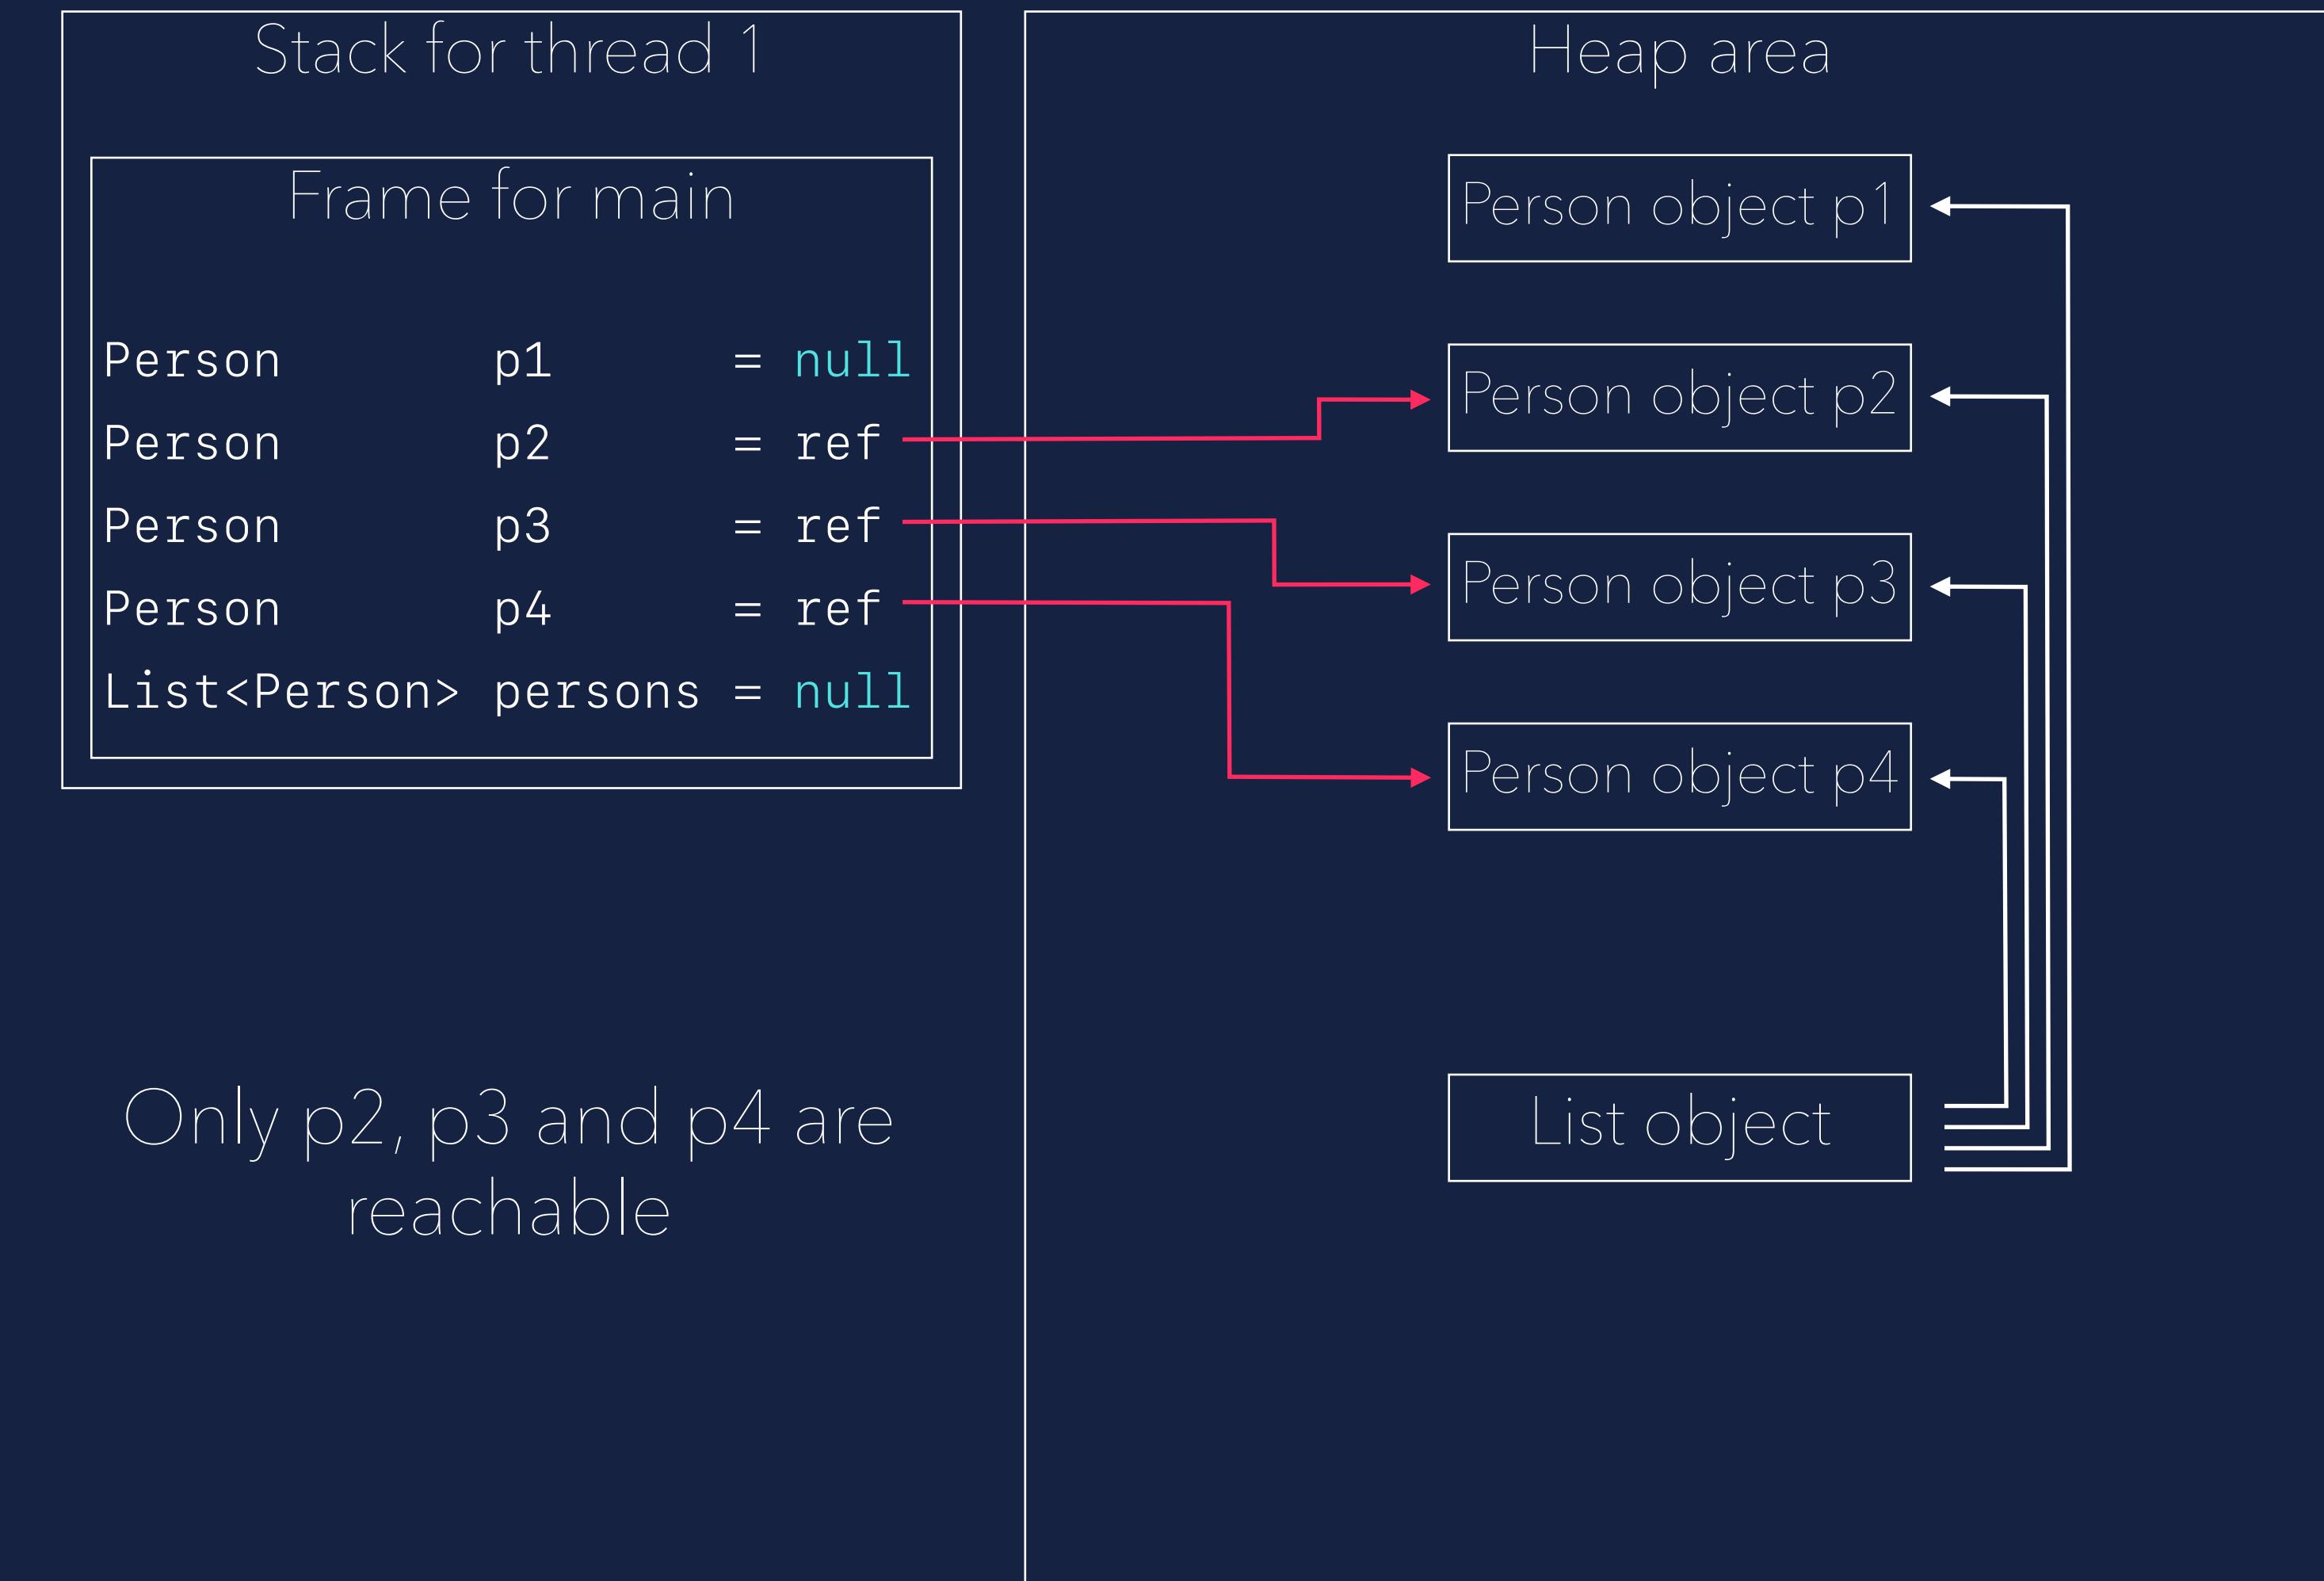
```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
  
    System.out.println(persons.get(0)); // -> Gerrit  
  
    persons = null;  
}
```



MEMORY MANAGEMENT

In the JVM...

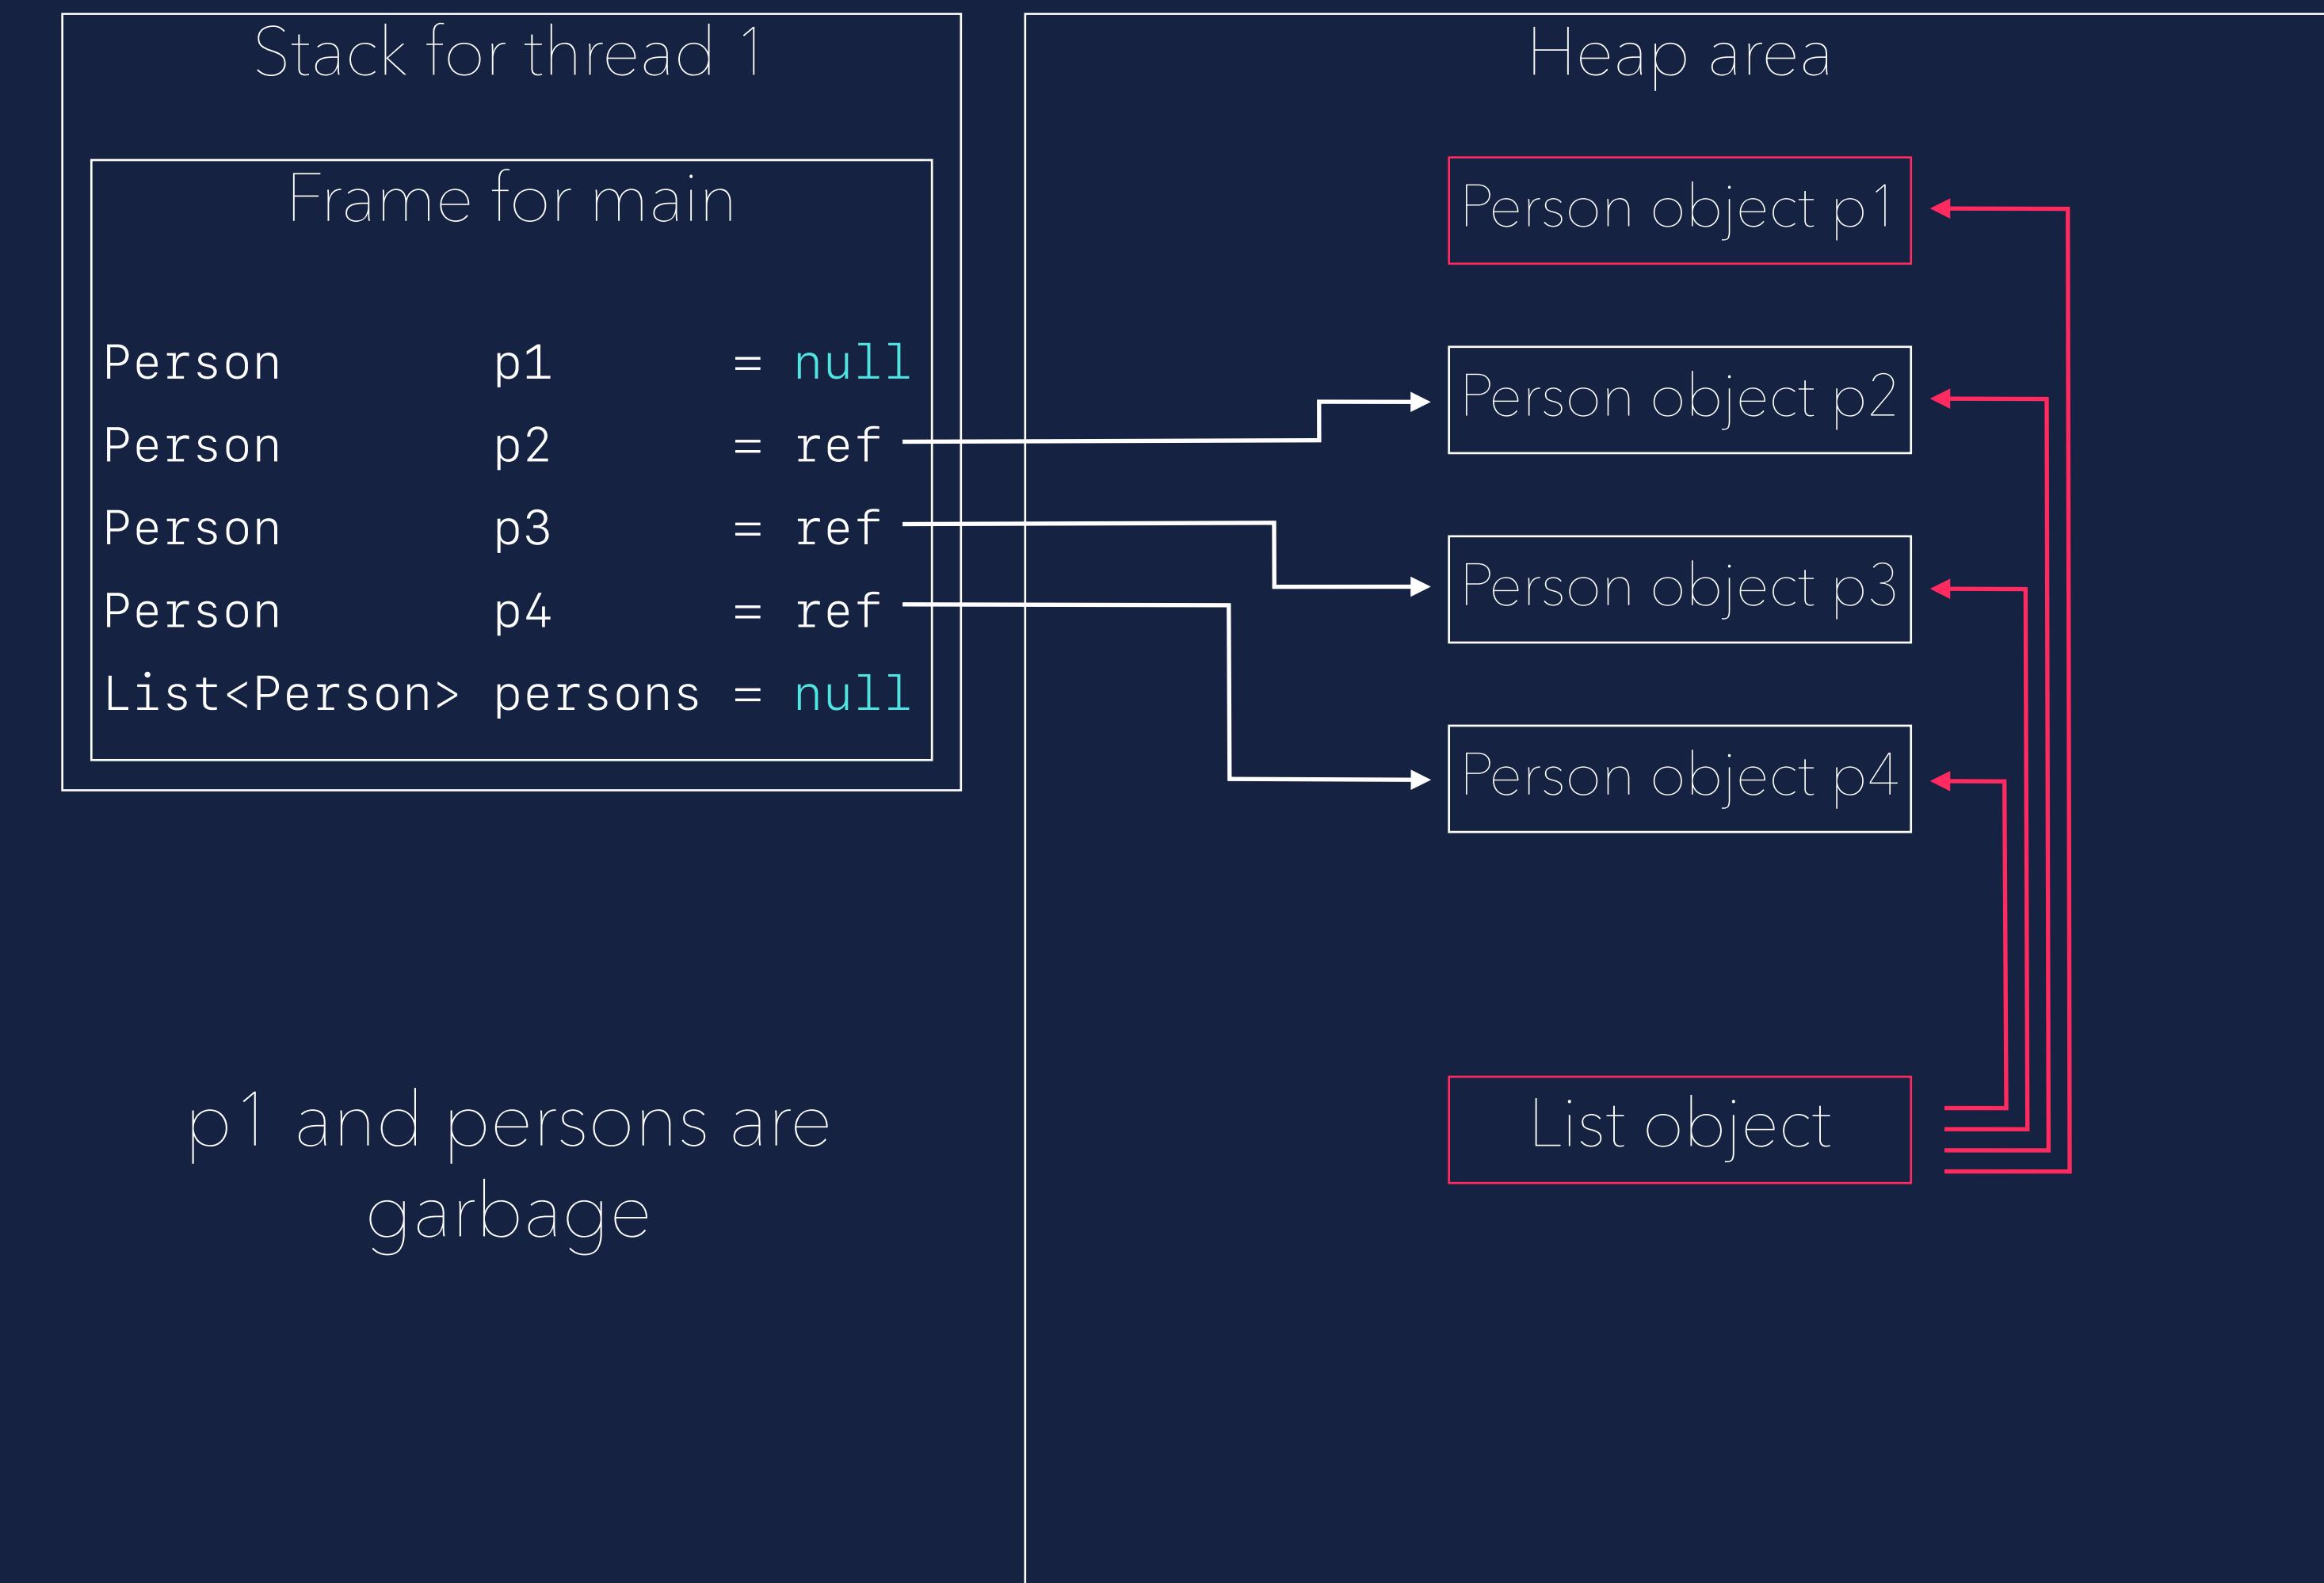
```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
  
    System.out.println(persons.get(0)); // -> Gerrit  
  
    persons = null;  
}
```



MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
  
    System.out.println(persons.get(0)); // -> Gerrit  
  
    persons = null;  
}
```



**HOW TO GET
RID OF IT...?**

**GARBAGE
COLLECTION**

GARBAGE COLLECTION

What is it...



Form of automatic memory management

GARBAGE COLLECTION

What is it...

- Form of automatic memory management
- Identifies and reclaims no longer used memory

GARBAGE COLLECTION

What is it...

- Form of automatic memory management
- Identifies and reclaims no longer used memory
- Ensures efficient memory utilisation

GARBAGE COLLECTION

What is it...

- Form of automatic memory management
- Identifies and reclaims no longer used memory
- Ensures efficient memory utilisation
- Frees user from managing the memory manually

**CONSERVATIVE
AND
PRECISE**

GARBAGE COLLECTION

Conservative and Precise



Conservative does not fully identify all object references
(assumes any bit pattern in memory could be a reference, lead to more false positives)

GARBAGE COLLECTION

Conservative and Precise

-  Conservative does not fully identify all object references
(assumes any bit pattern in memory could be a reference, lead to more false positives)
-  Precise correctly identifies all references in an object
(needed in order to move objects)

PHASES

(precise collectors)

GARBAGE COLLECTION

Phases (precise collectors)



Tracing

Identify live objects on the heap

GARBAGE COLLECTION

Phases (precise collectors)



Tracing

Identify live objects on the heap



Freeing

Reclaim resources held by dead objects

GARBAGE COLLECTION

Phases (precise collectors)



Tracing

Identify live objects on the heap



Freeing

Reclaim resources held by dead objects



Compaction

Periodically relocate live objects

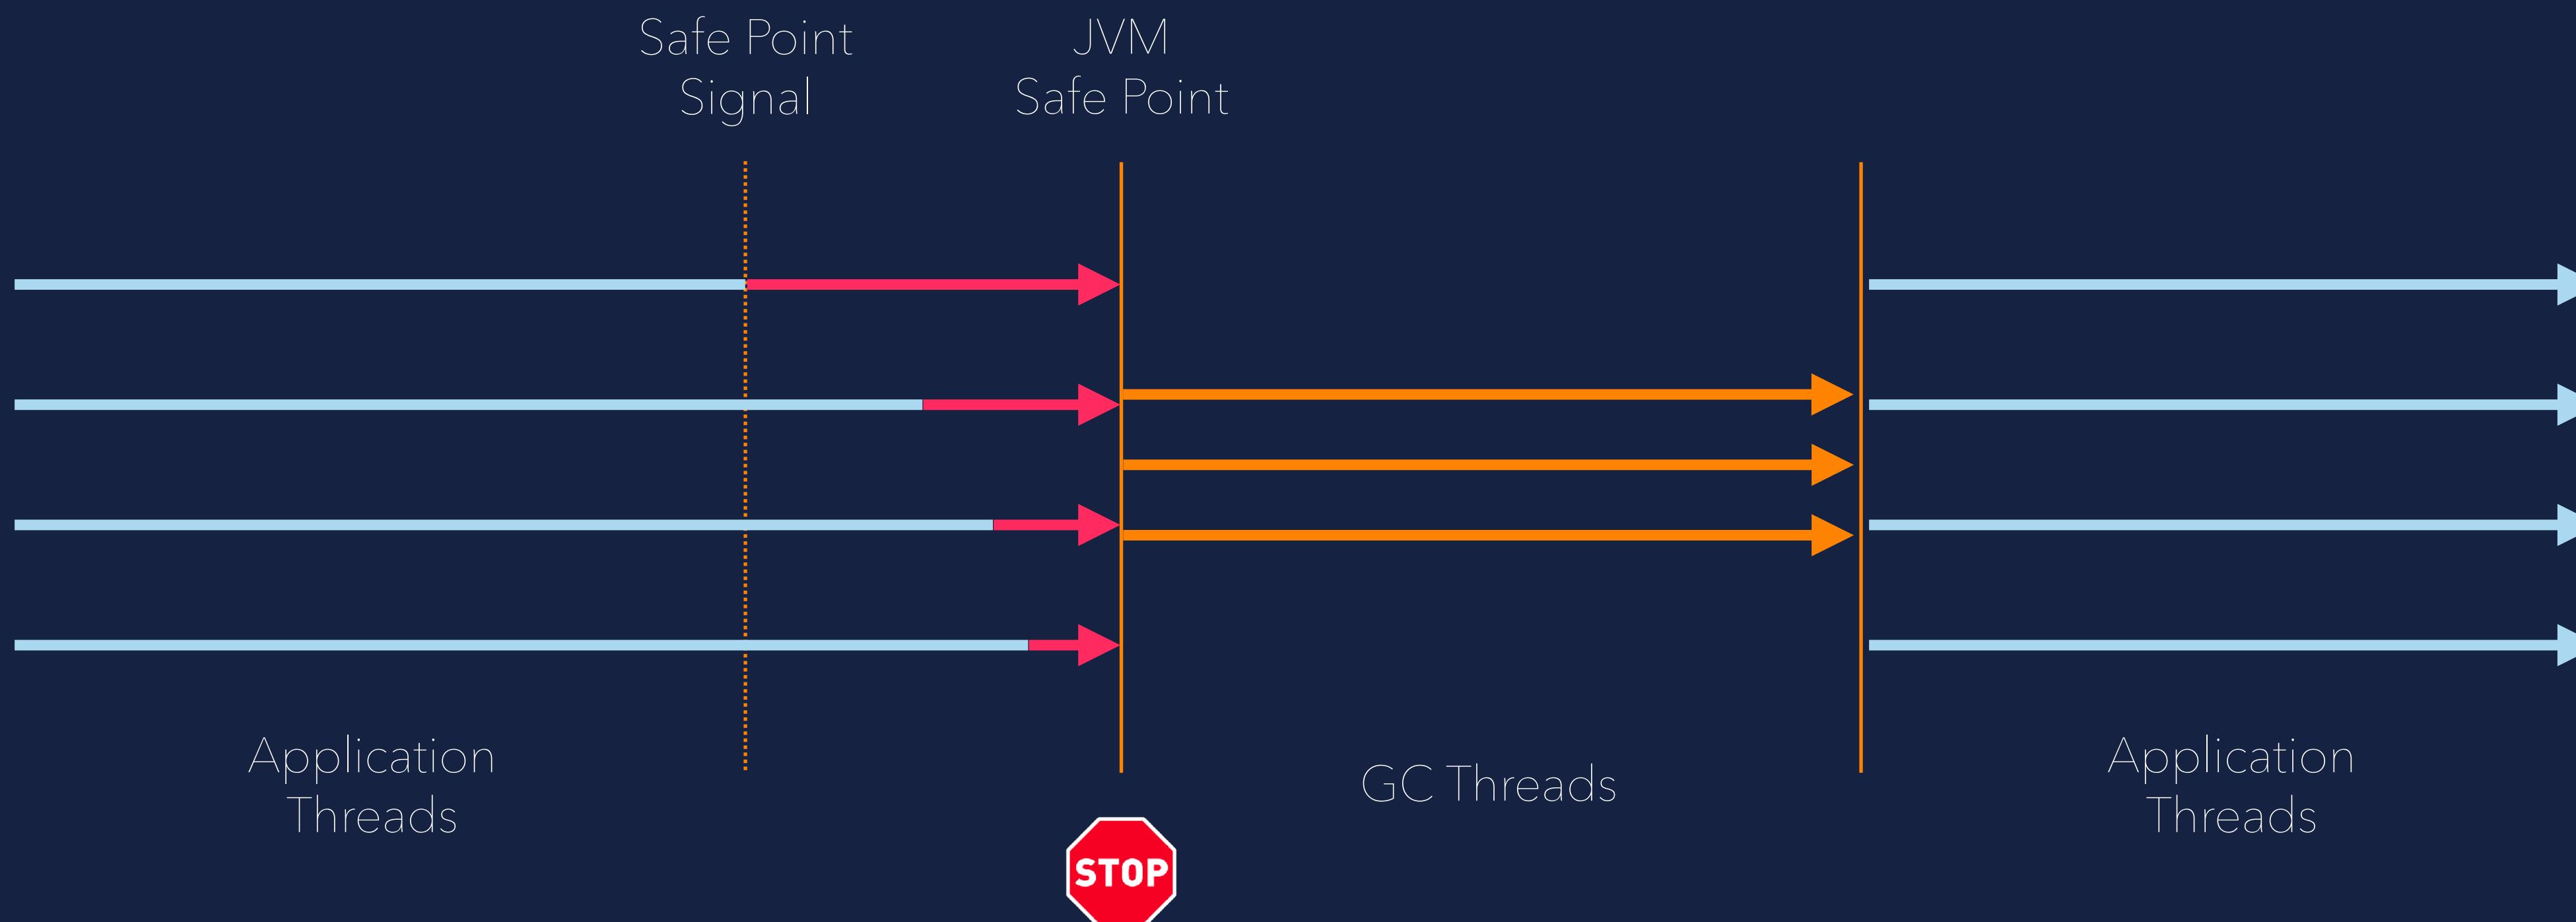
STOPPING THE
WORLD



STOPPING THE WORLD



Halt of all application threads



COLLECTORS

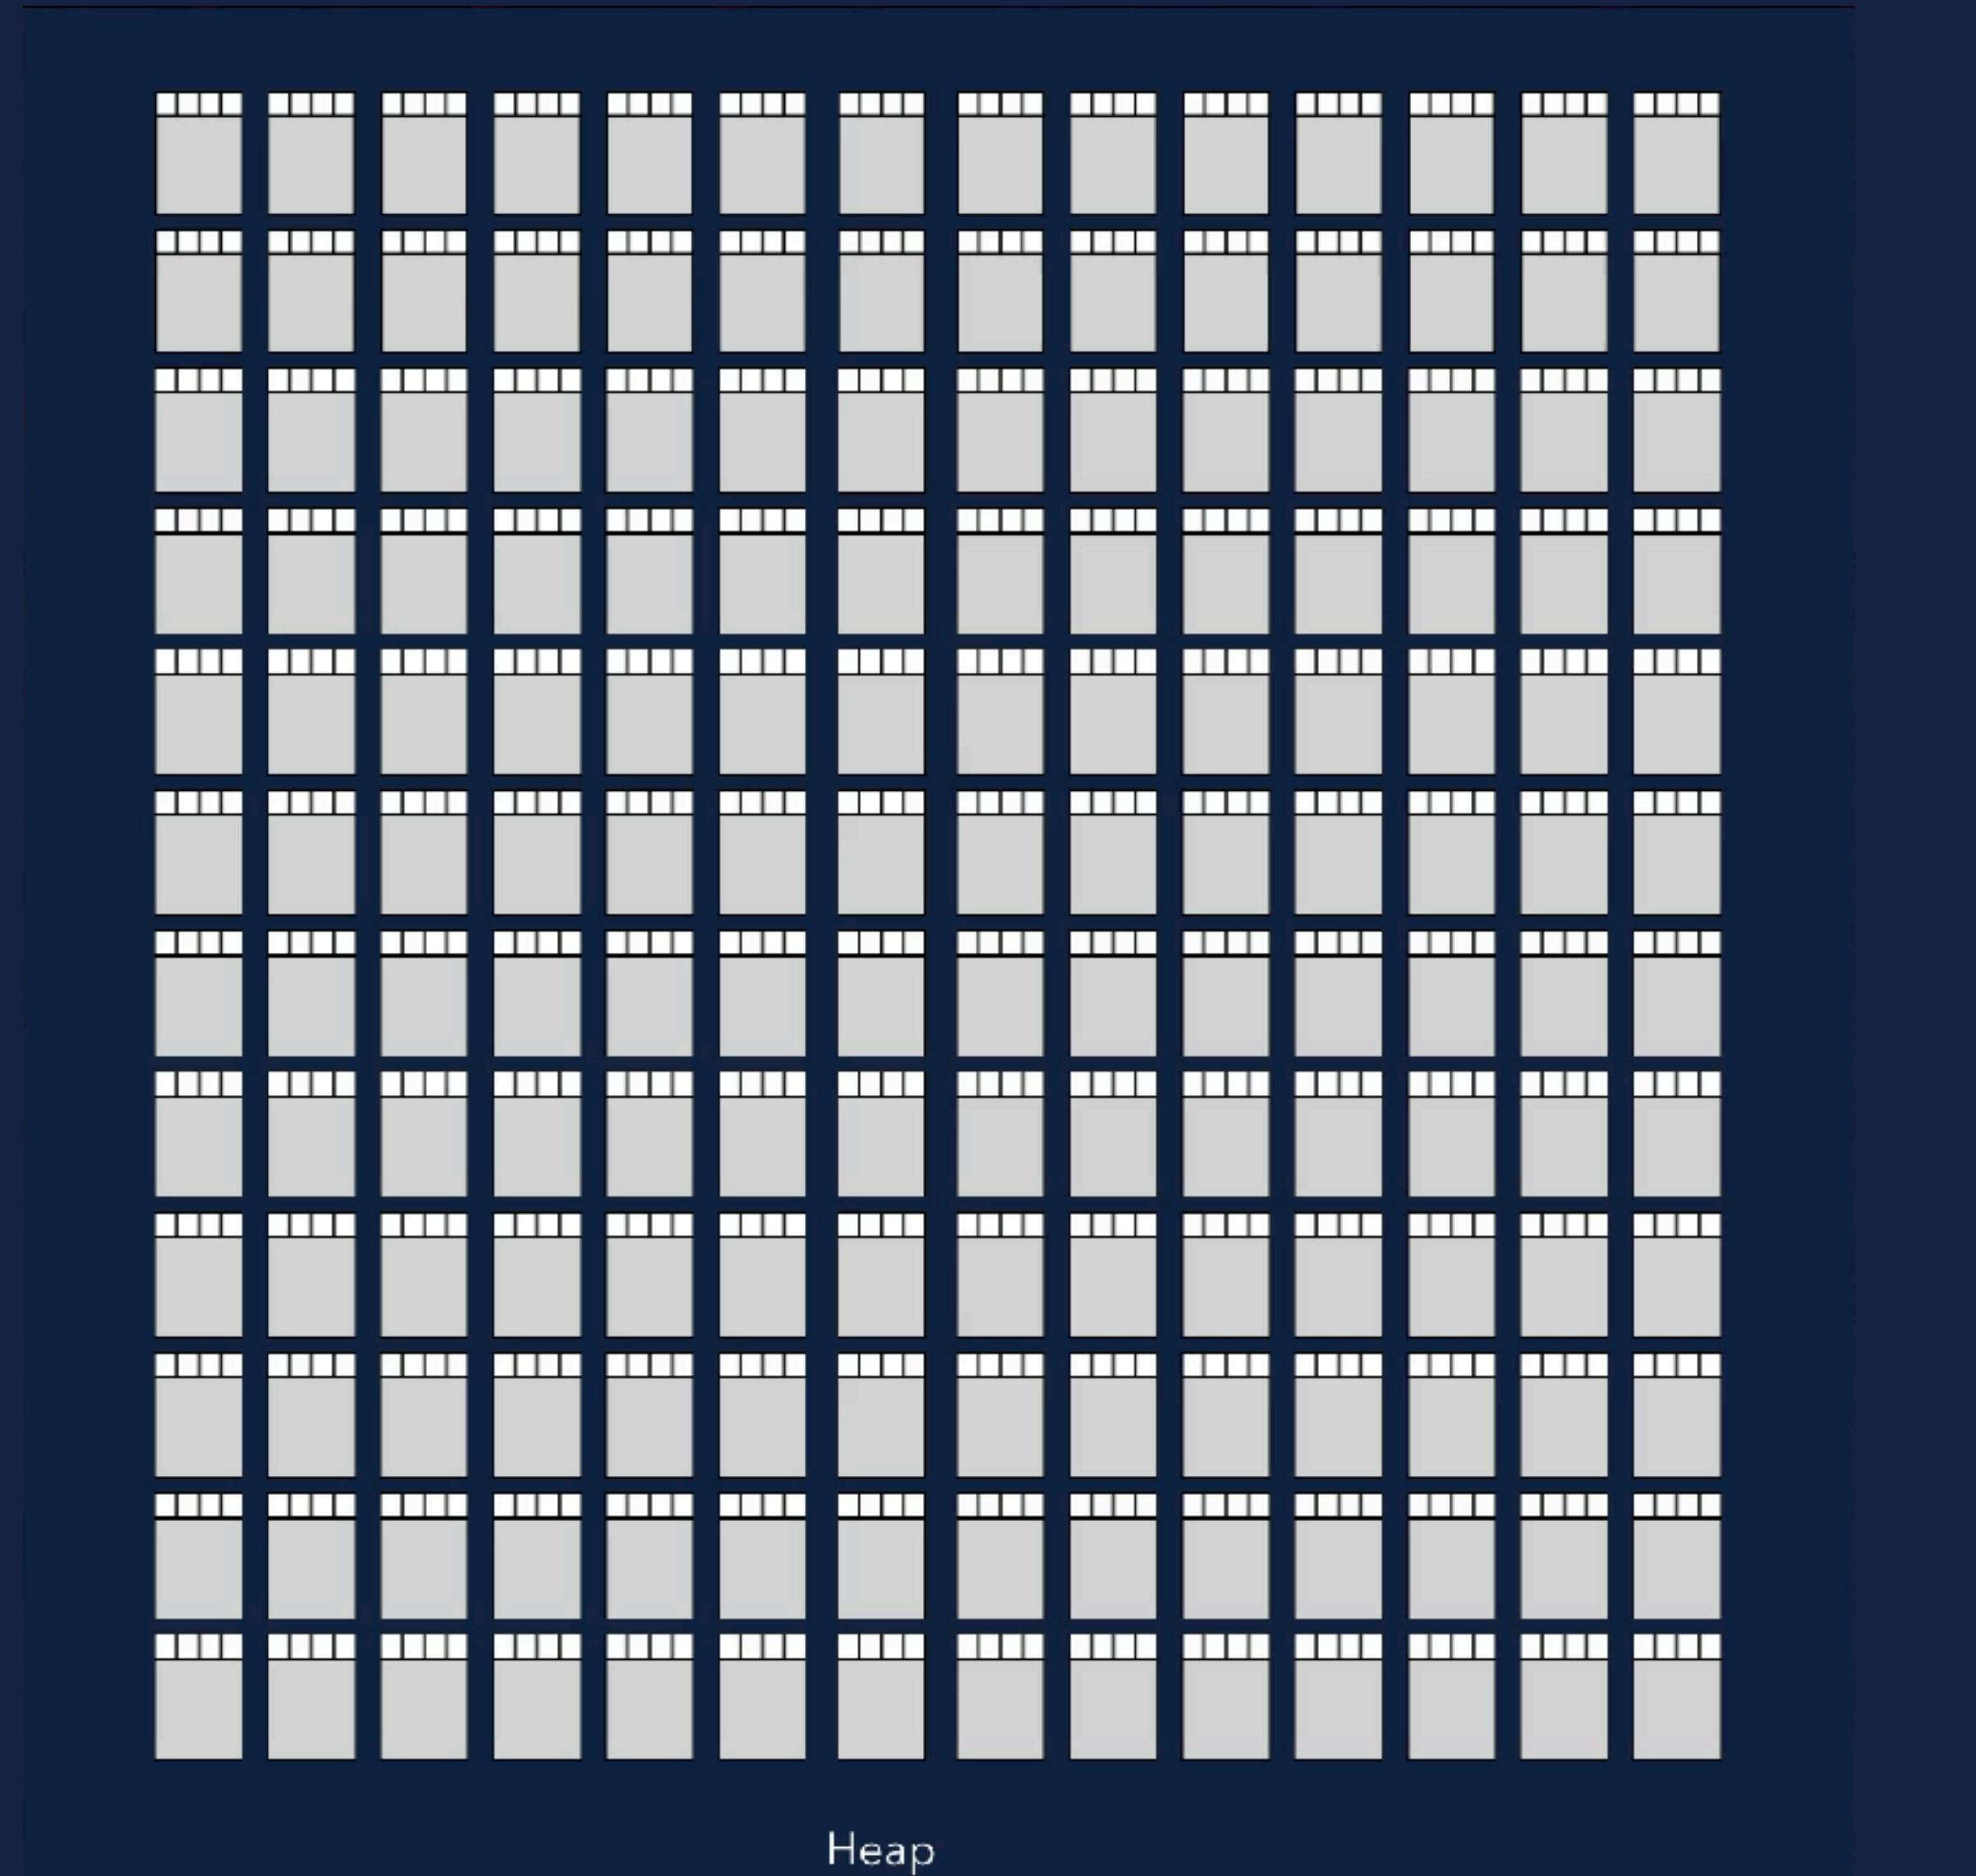
NON MOVING COLLECTOR

Mark & Sweep

NON MOVING COLLECTOR

Demo

1. Mutator allocates cells in Heap
2. Heap is out of memory -> GC
3. Mark all live cells
4. Free all dead cells
5. Unmark all live cells
6. Resume Mutator



Fragmentation

MOVING COLLECTORS

Compacting Collector & Copy Collector

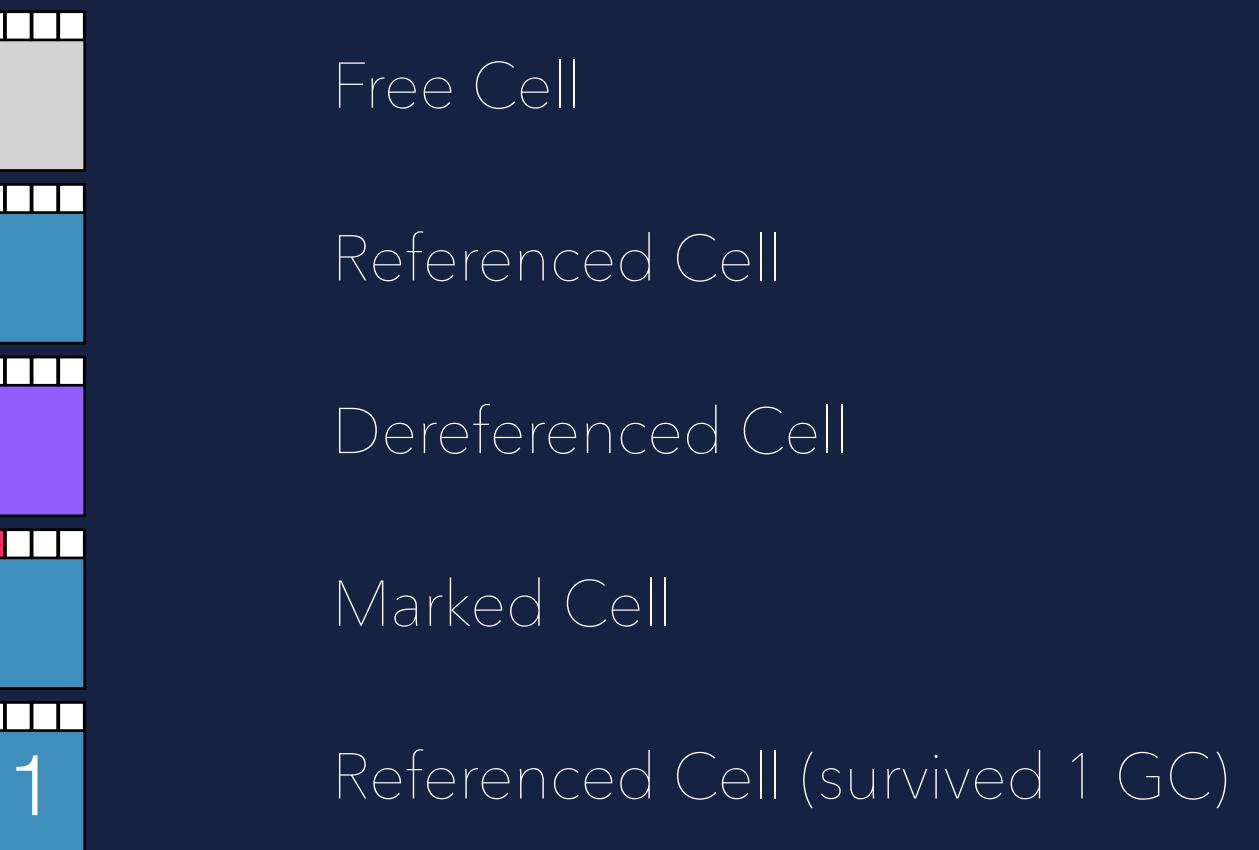
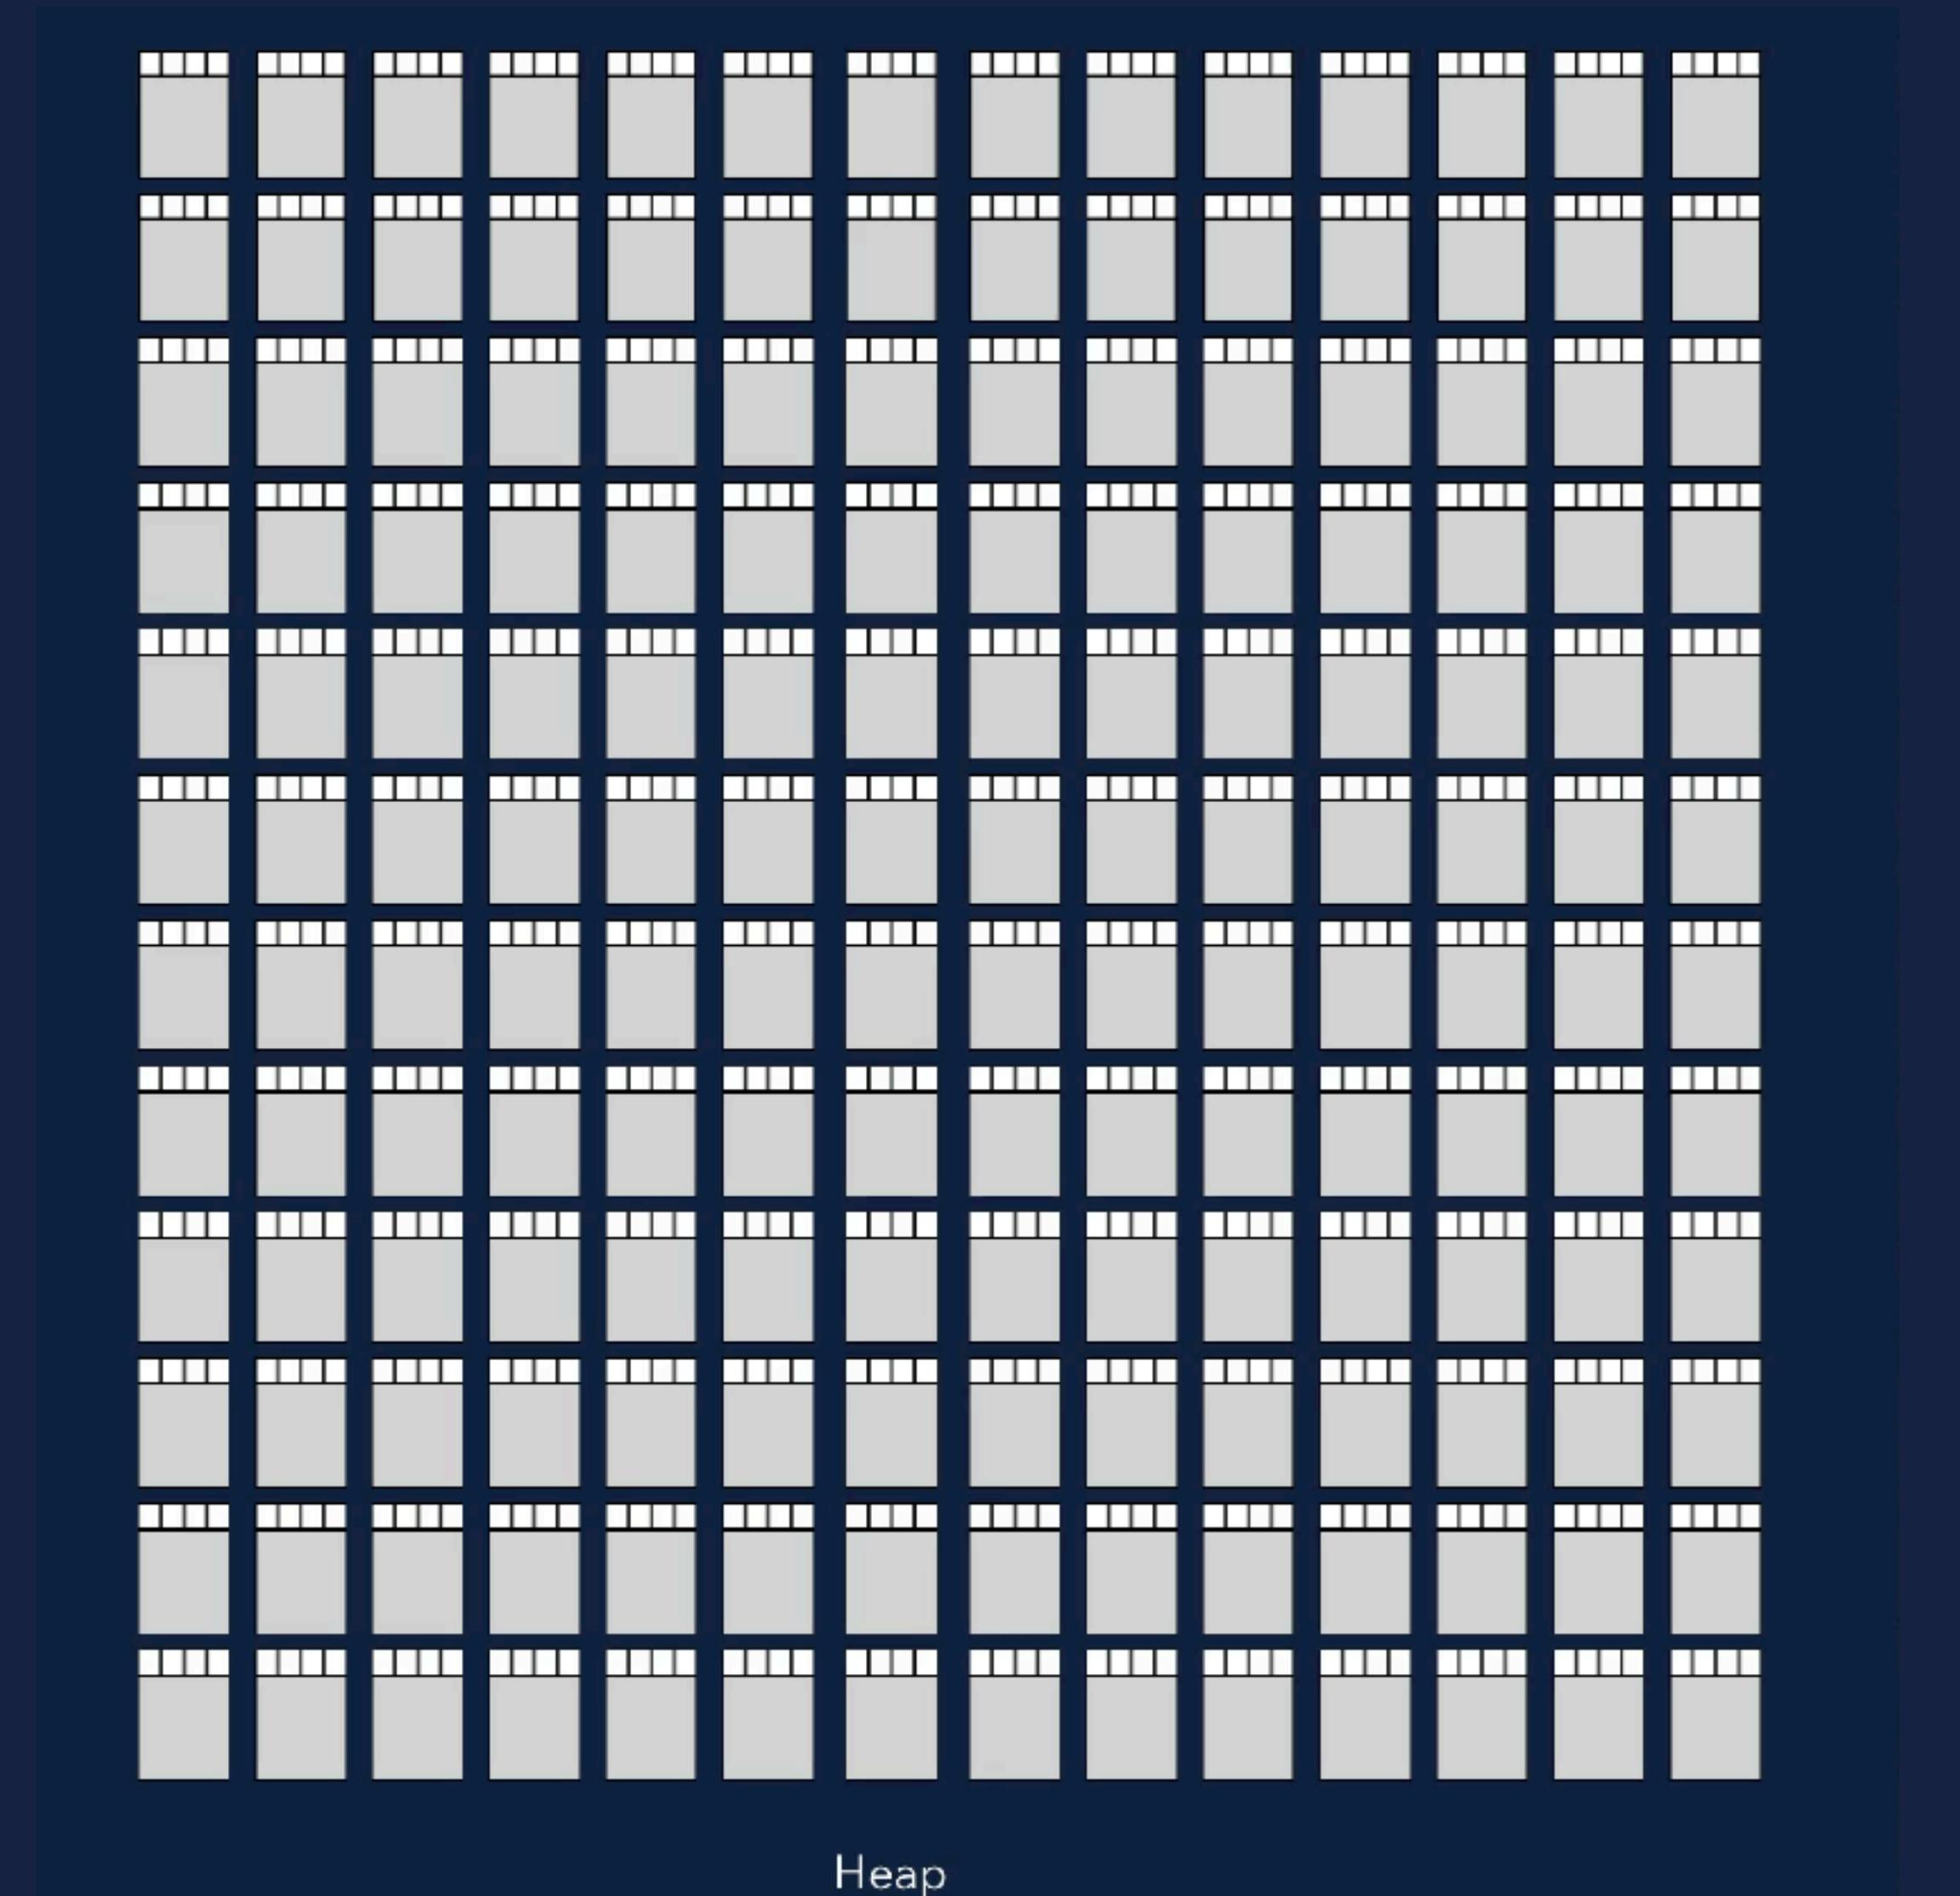
COMPACTING COLLECTOR

Mark & Compact

COMPACTING COLLECTOR

Demo

1. Mutator allocates cells in Heap
2. Heap is out of memory -> GC
3. Mark all live cells
4. Free all dead cells
5. Unmark all live cells
6. Compact all live cells
7. Resume Mutator



Easy allocation through
"Bump the pointer" technique



Headroom
20-50%

COPY COLLECTOR

Mark & Copy

COPY COLLECTOR

Demo

1. Allocating in ToSpace
2. ToSpace is out of memory -> GC
3. Toggle To- and FromSpace
4. Mark live cells in FromSpace
5. Copy live cells to ToSpace
6. Free all cells in FromSpace
7. Resume Mutator



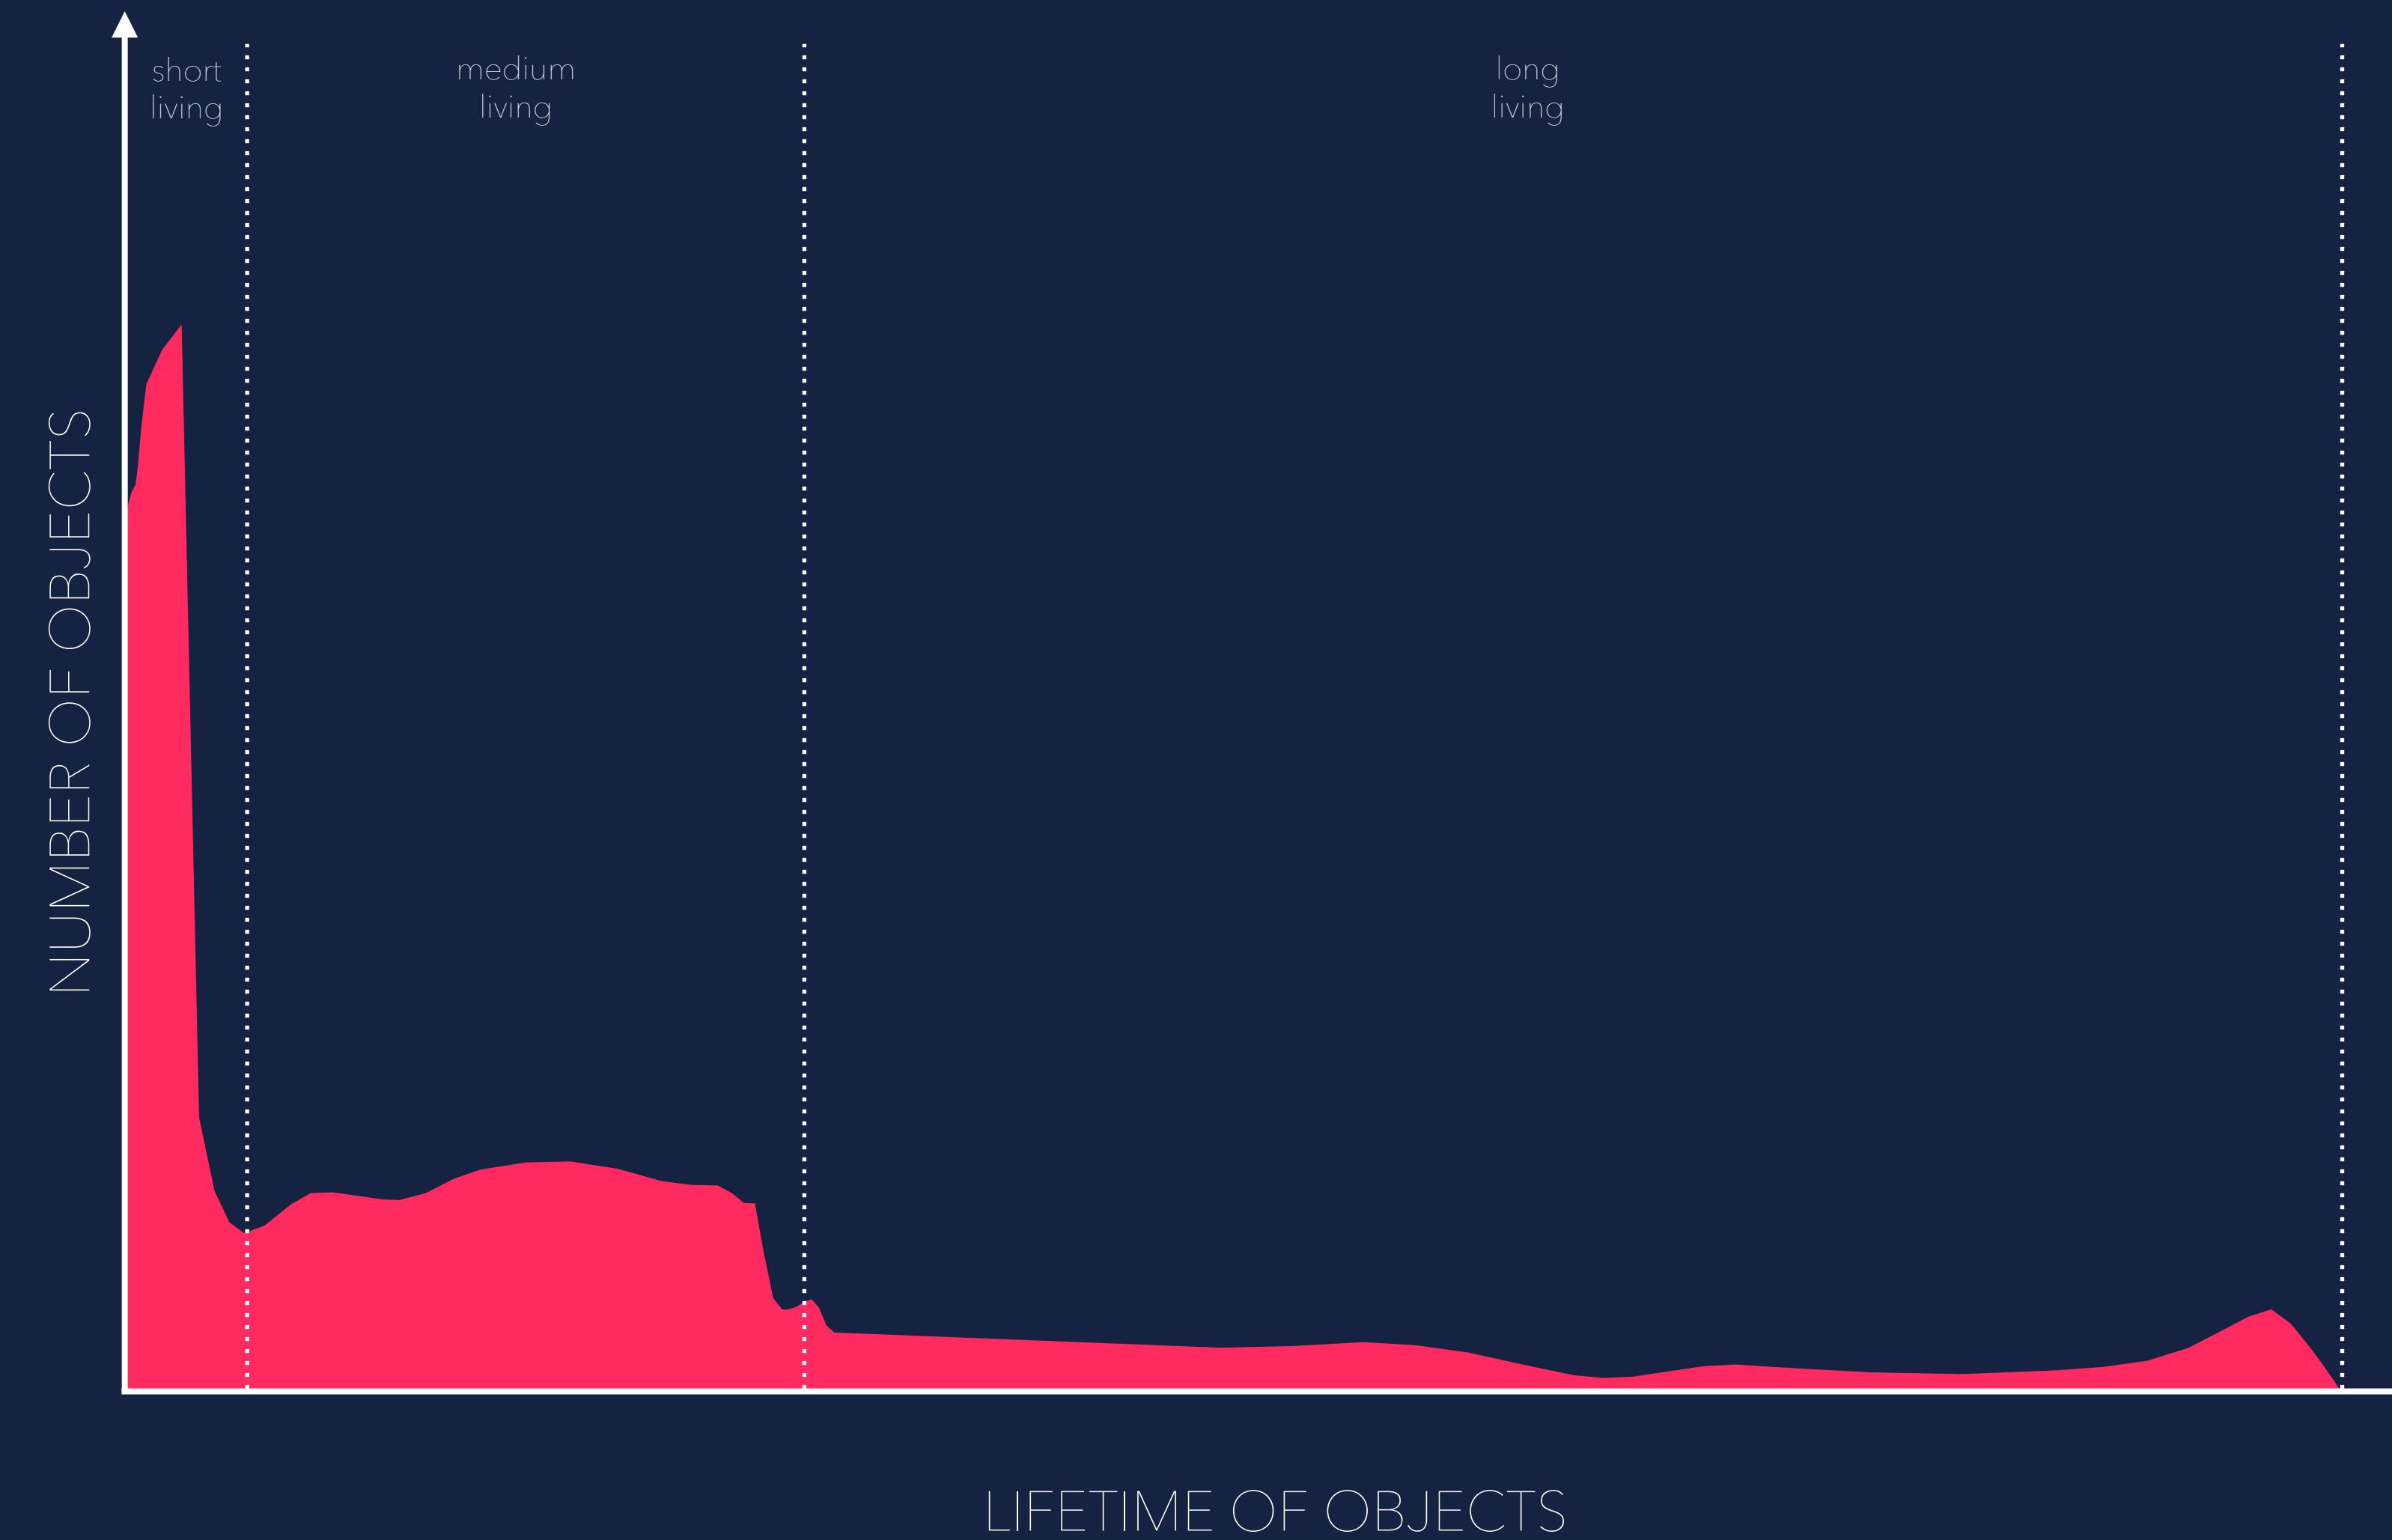
⚠️ Long living
objects and twice
as much memory

GENERATIONAL COLLECTOR

Generational Mark & Compact

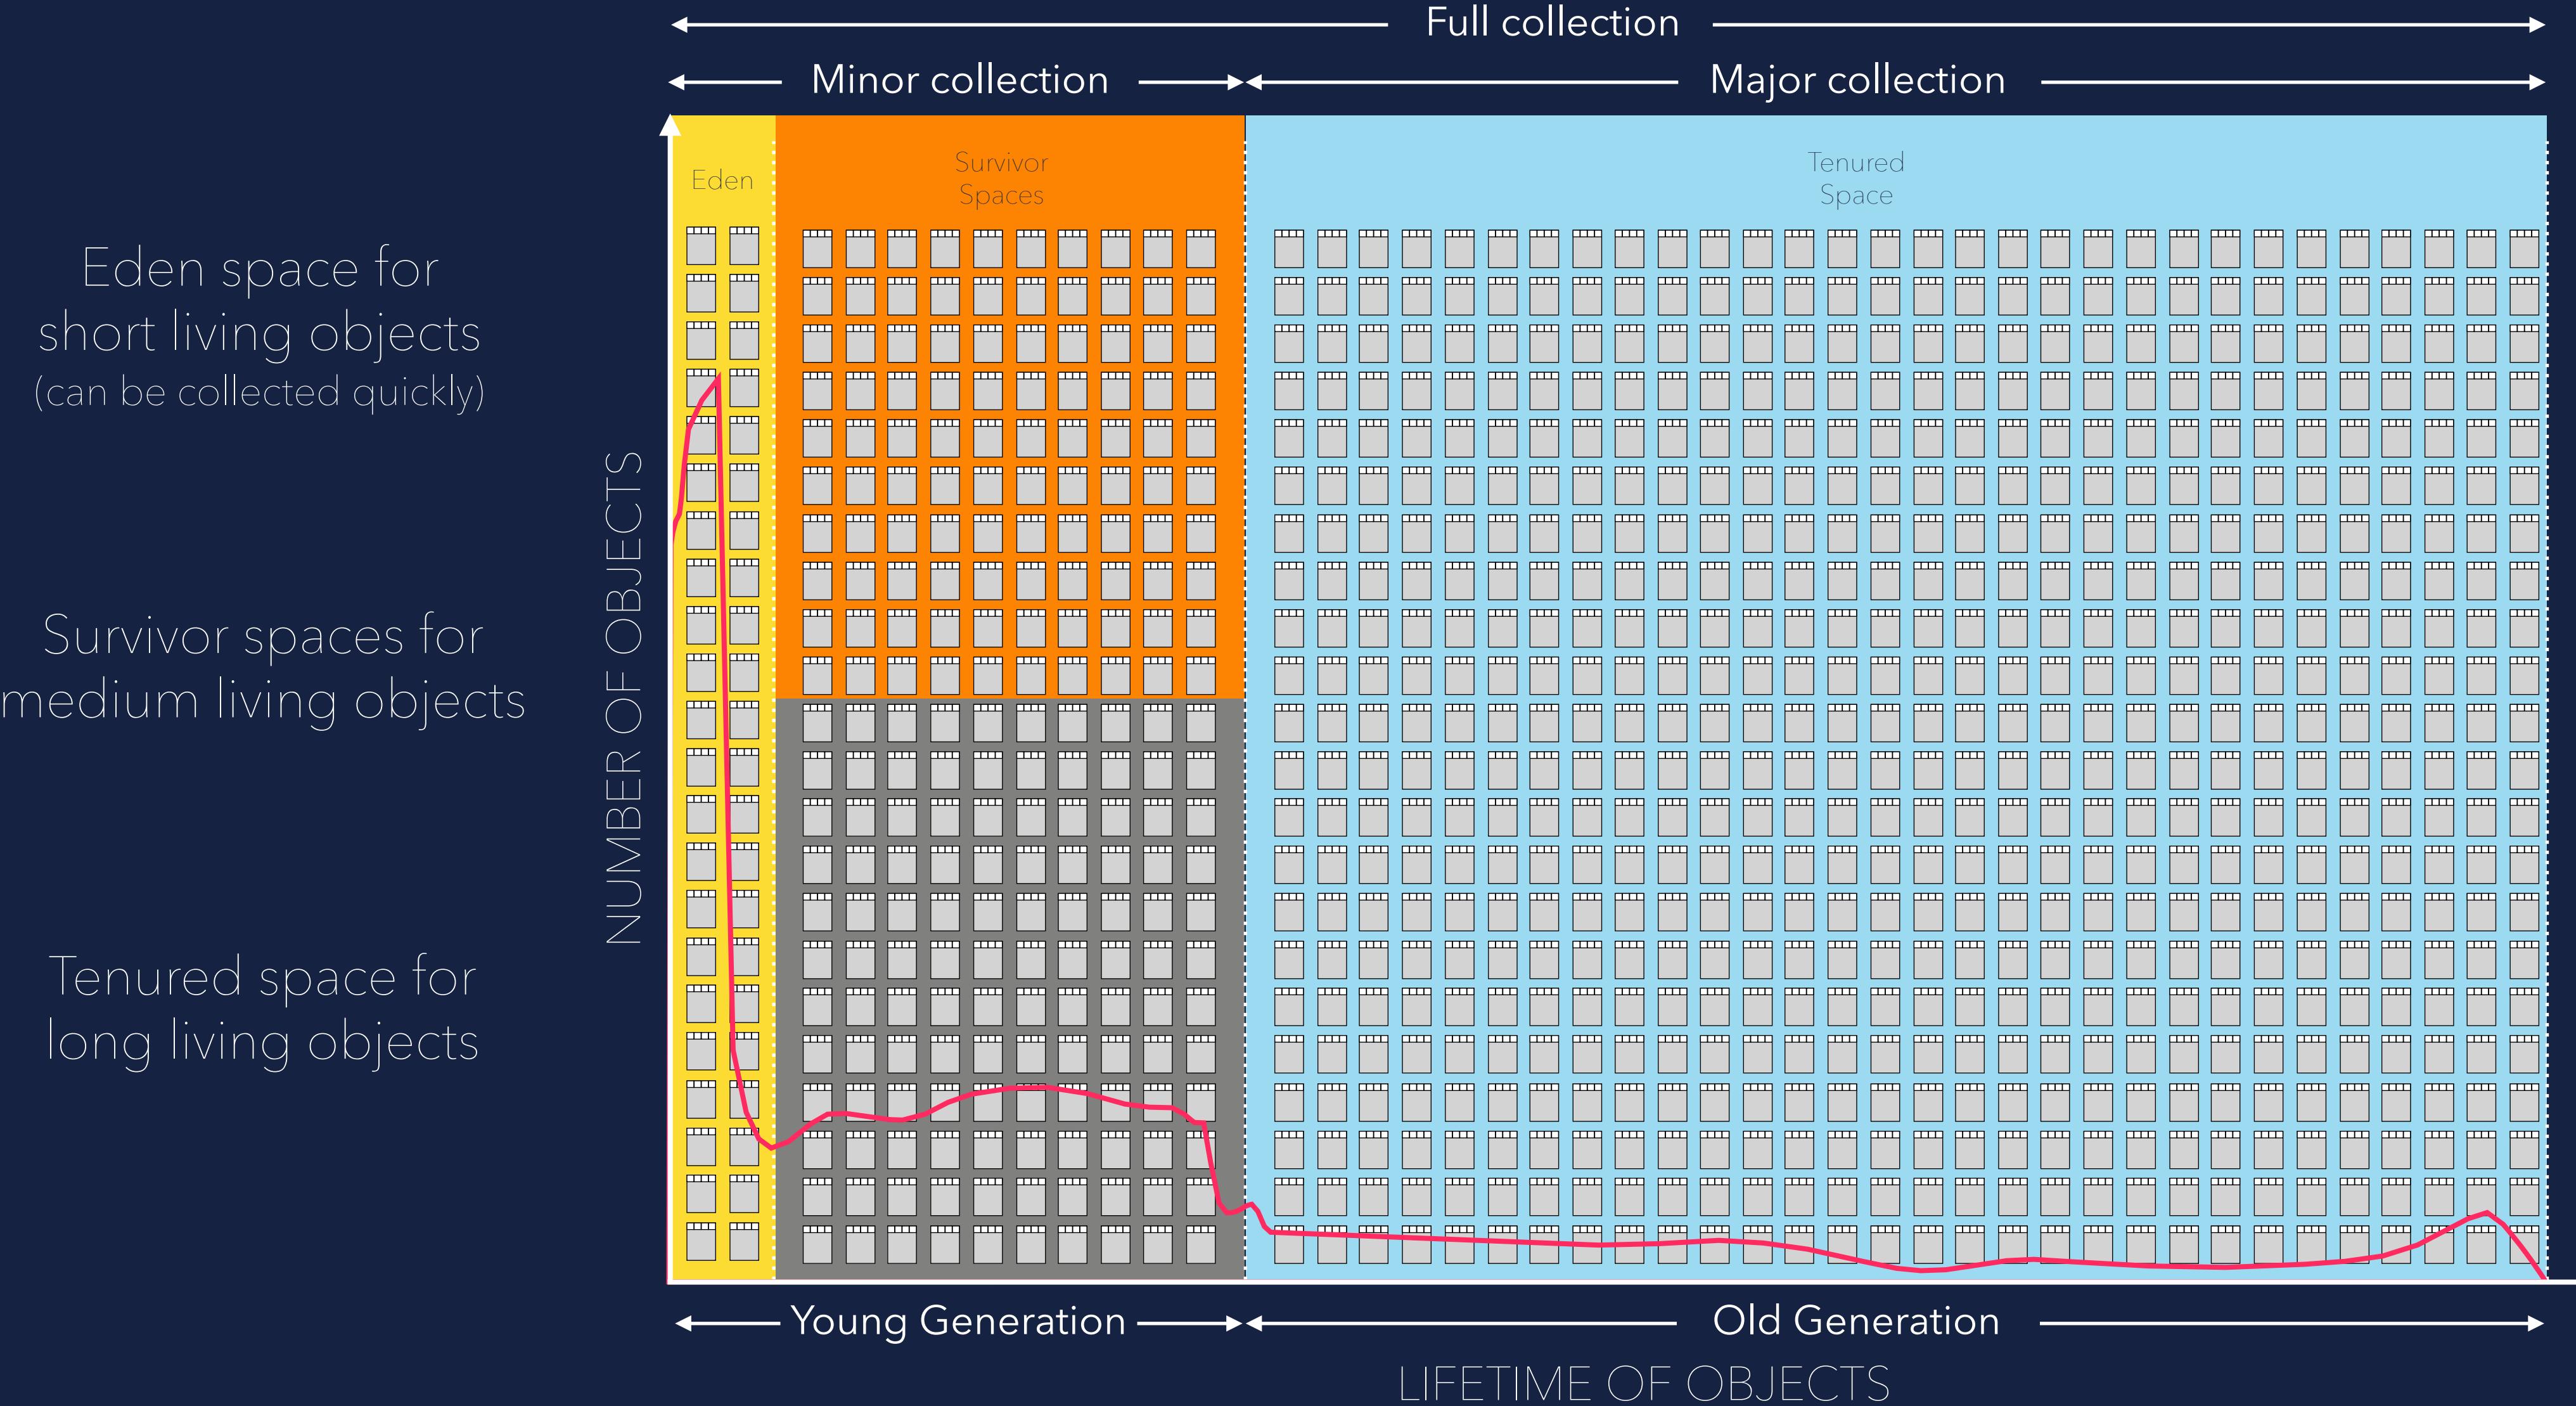
GENERATIONAL COLLECTOR

Weak Generational Hypothesis (Most objects die young)



GENERATIONAL COLLECTOR

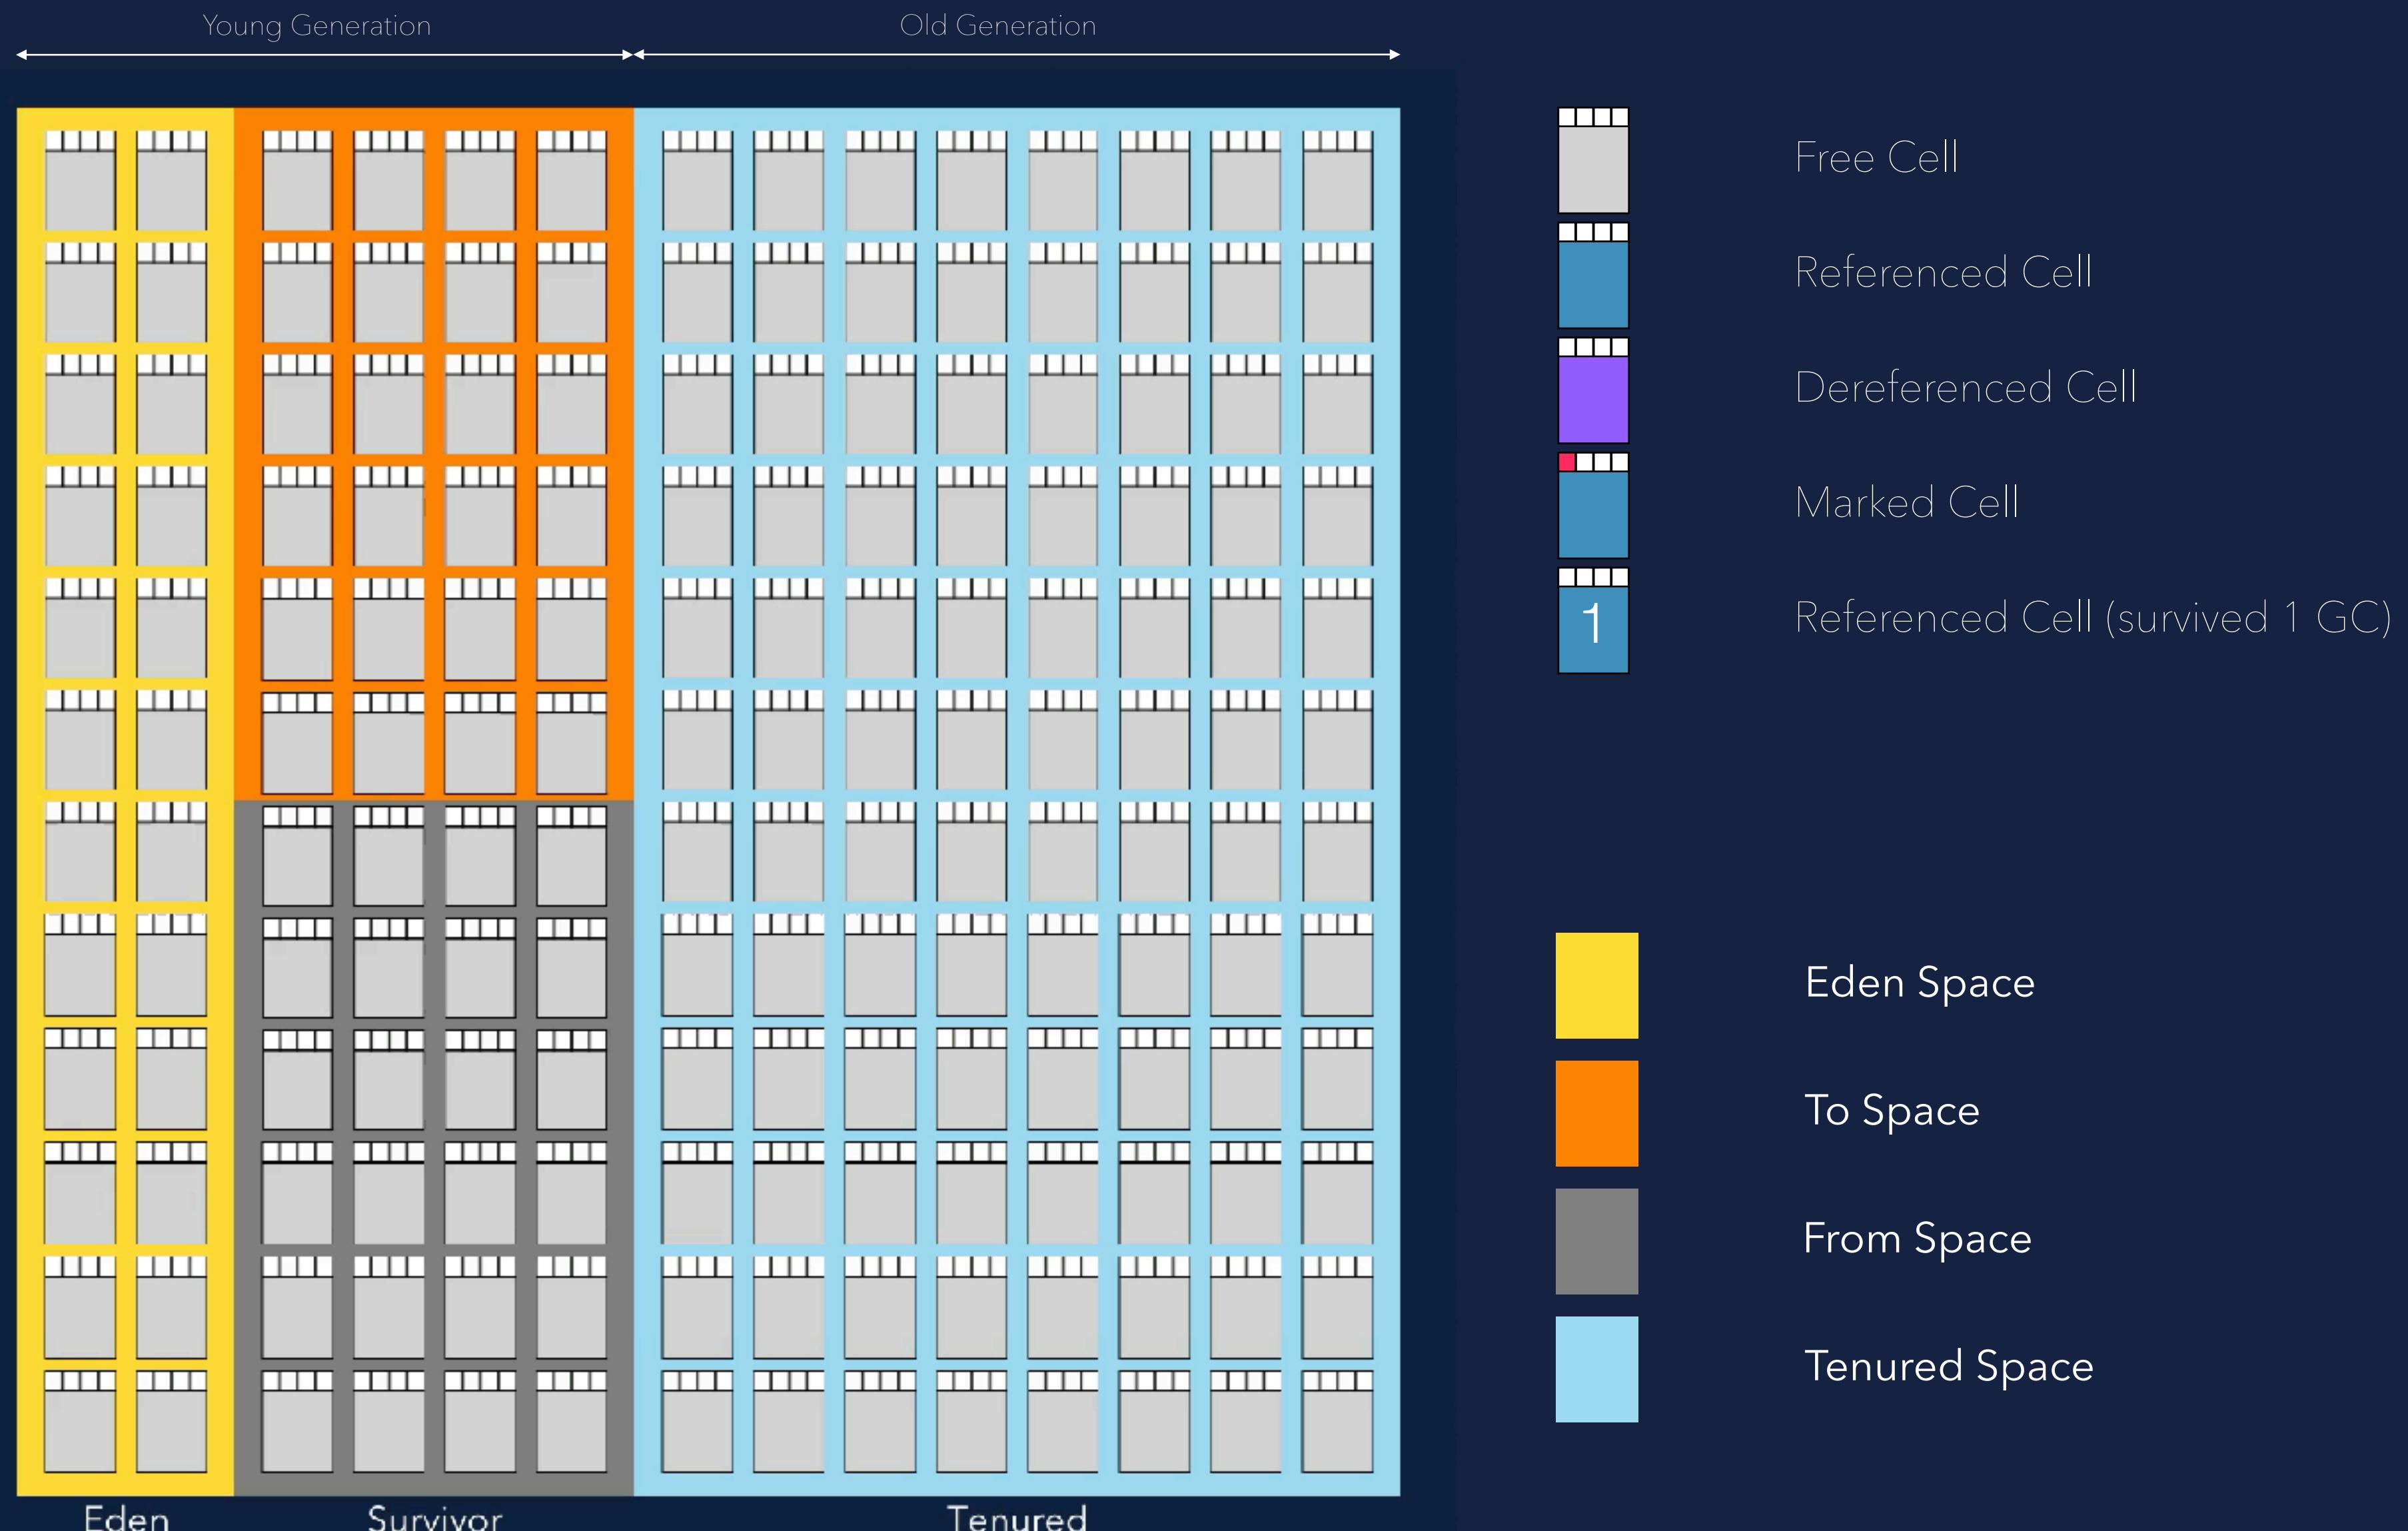
Weak Generational Hypothesis (Most objects die young)



GENERATIONAL COLLECTOR

Demo

1. Mutator allocates cells in Eden
2. Eden is out of memory -> GC
3. Toggle To- and FromSpace
4. Copy all live cells from FromSpace to ToSpace
5. Copy all live cells from Eden to ToSpace
6. Promote live cells from FromSpace to TenuredSpace
7. Free all dead cells
8. Resume Mutator



REMEMBERED SET

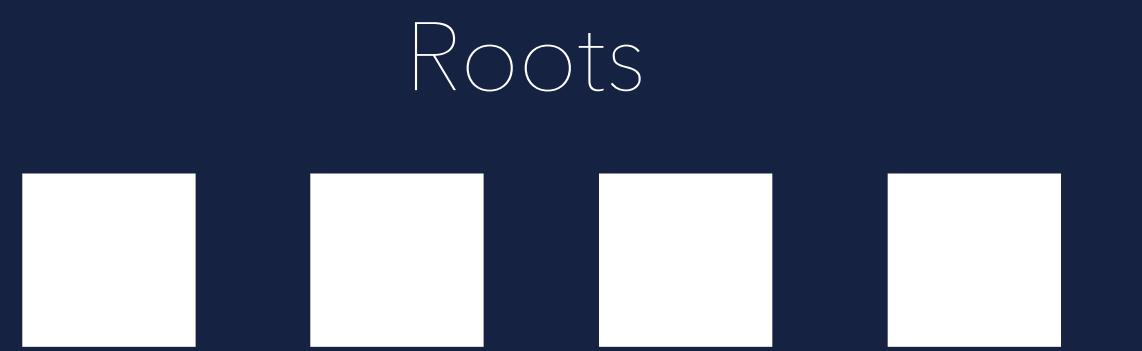
Intergenerational References

REMEMBERED SET

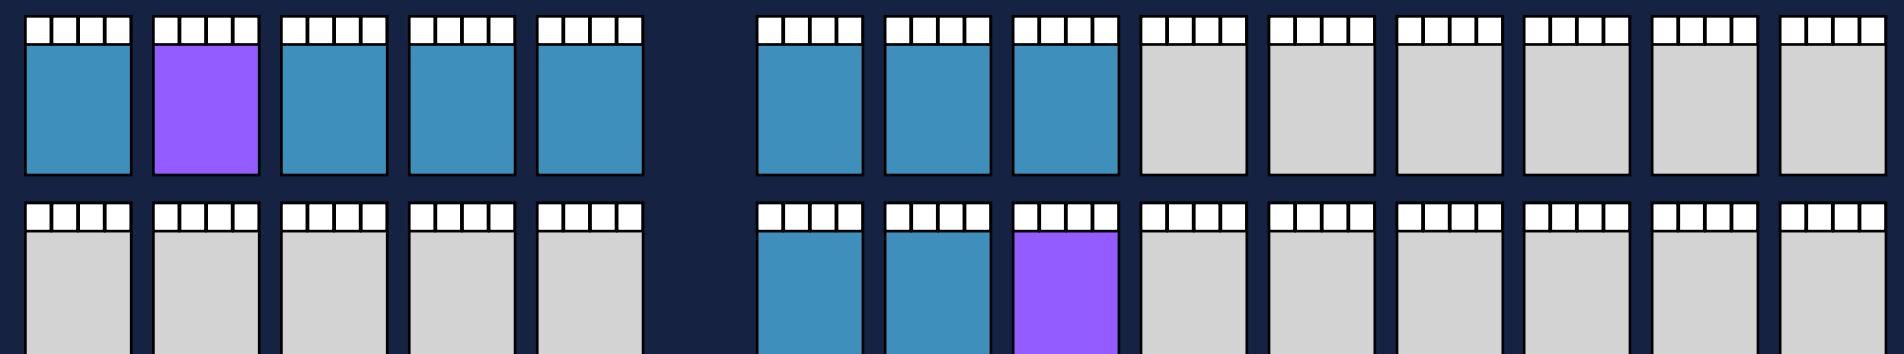
How to do a minor collection
with references from old to
young generation...?

REMEMBERED SET

Also known as Card Table



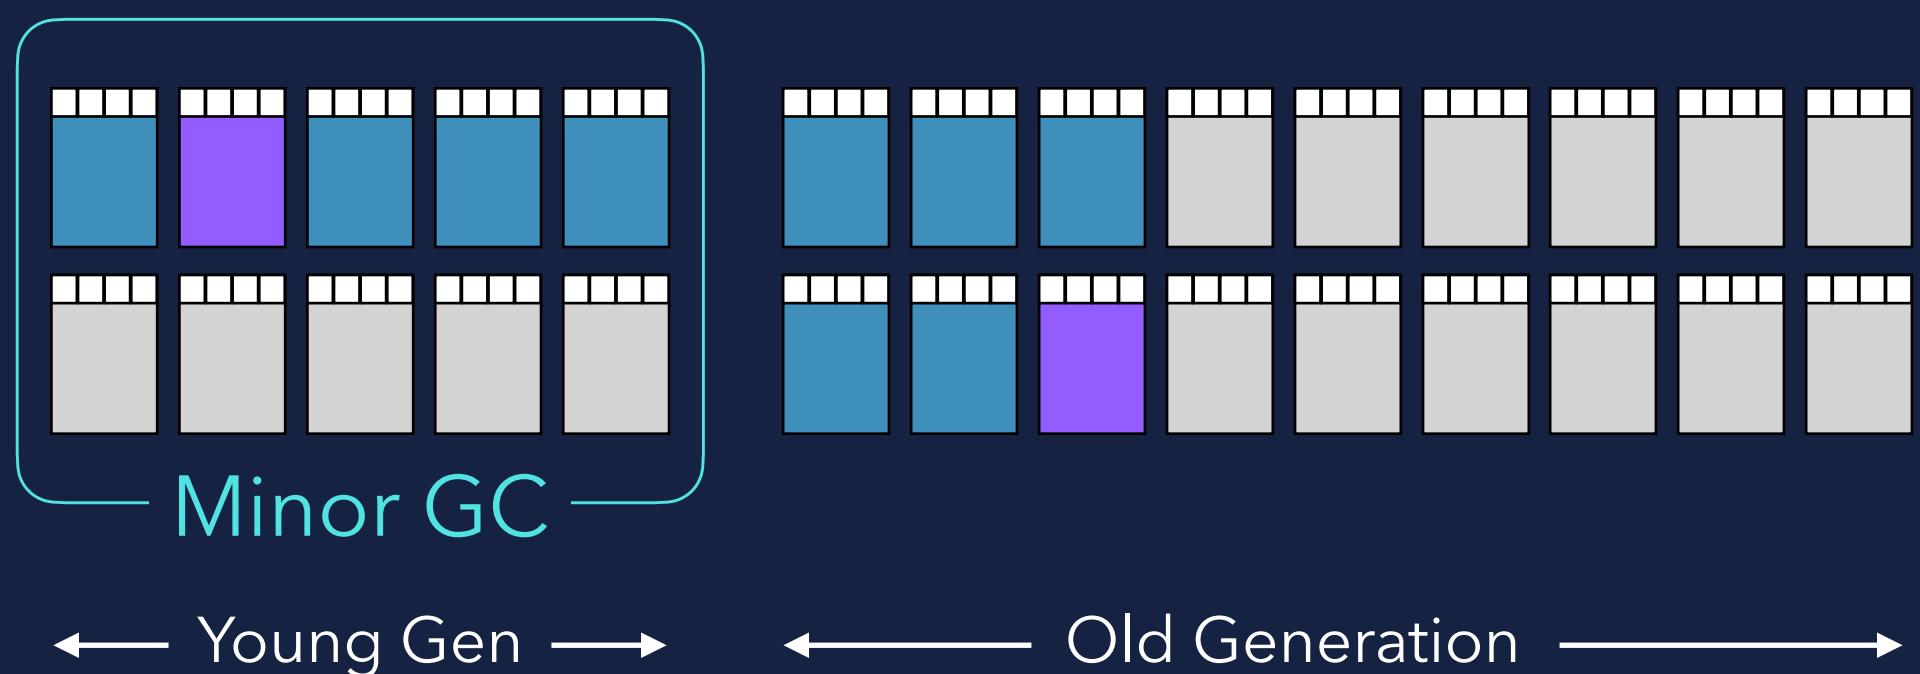
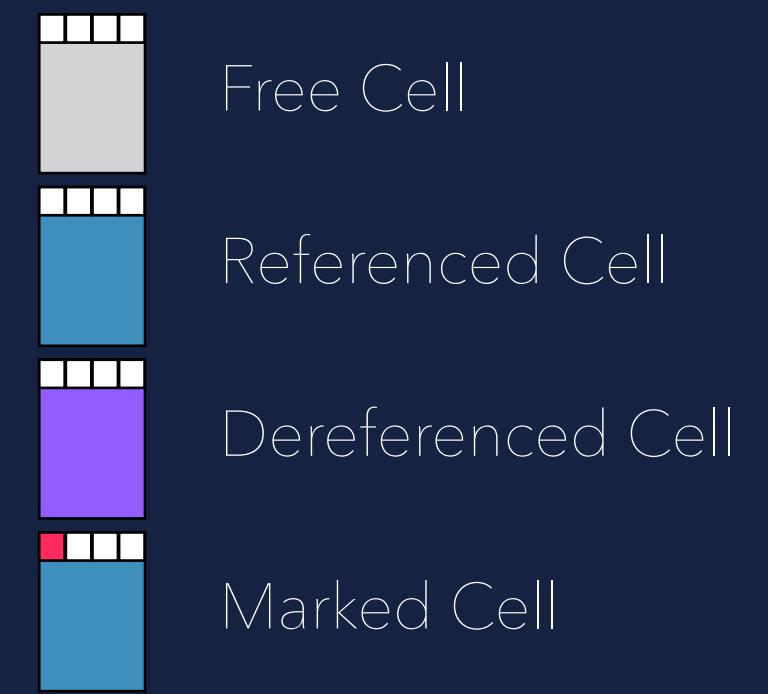
	Free Cell
	Referenced Cell
	Dereferenced Cell
	Marked Cell



← Young Gen → ← Old Generation →

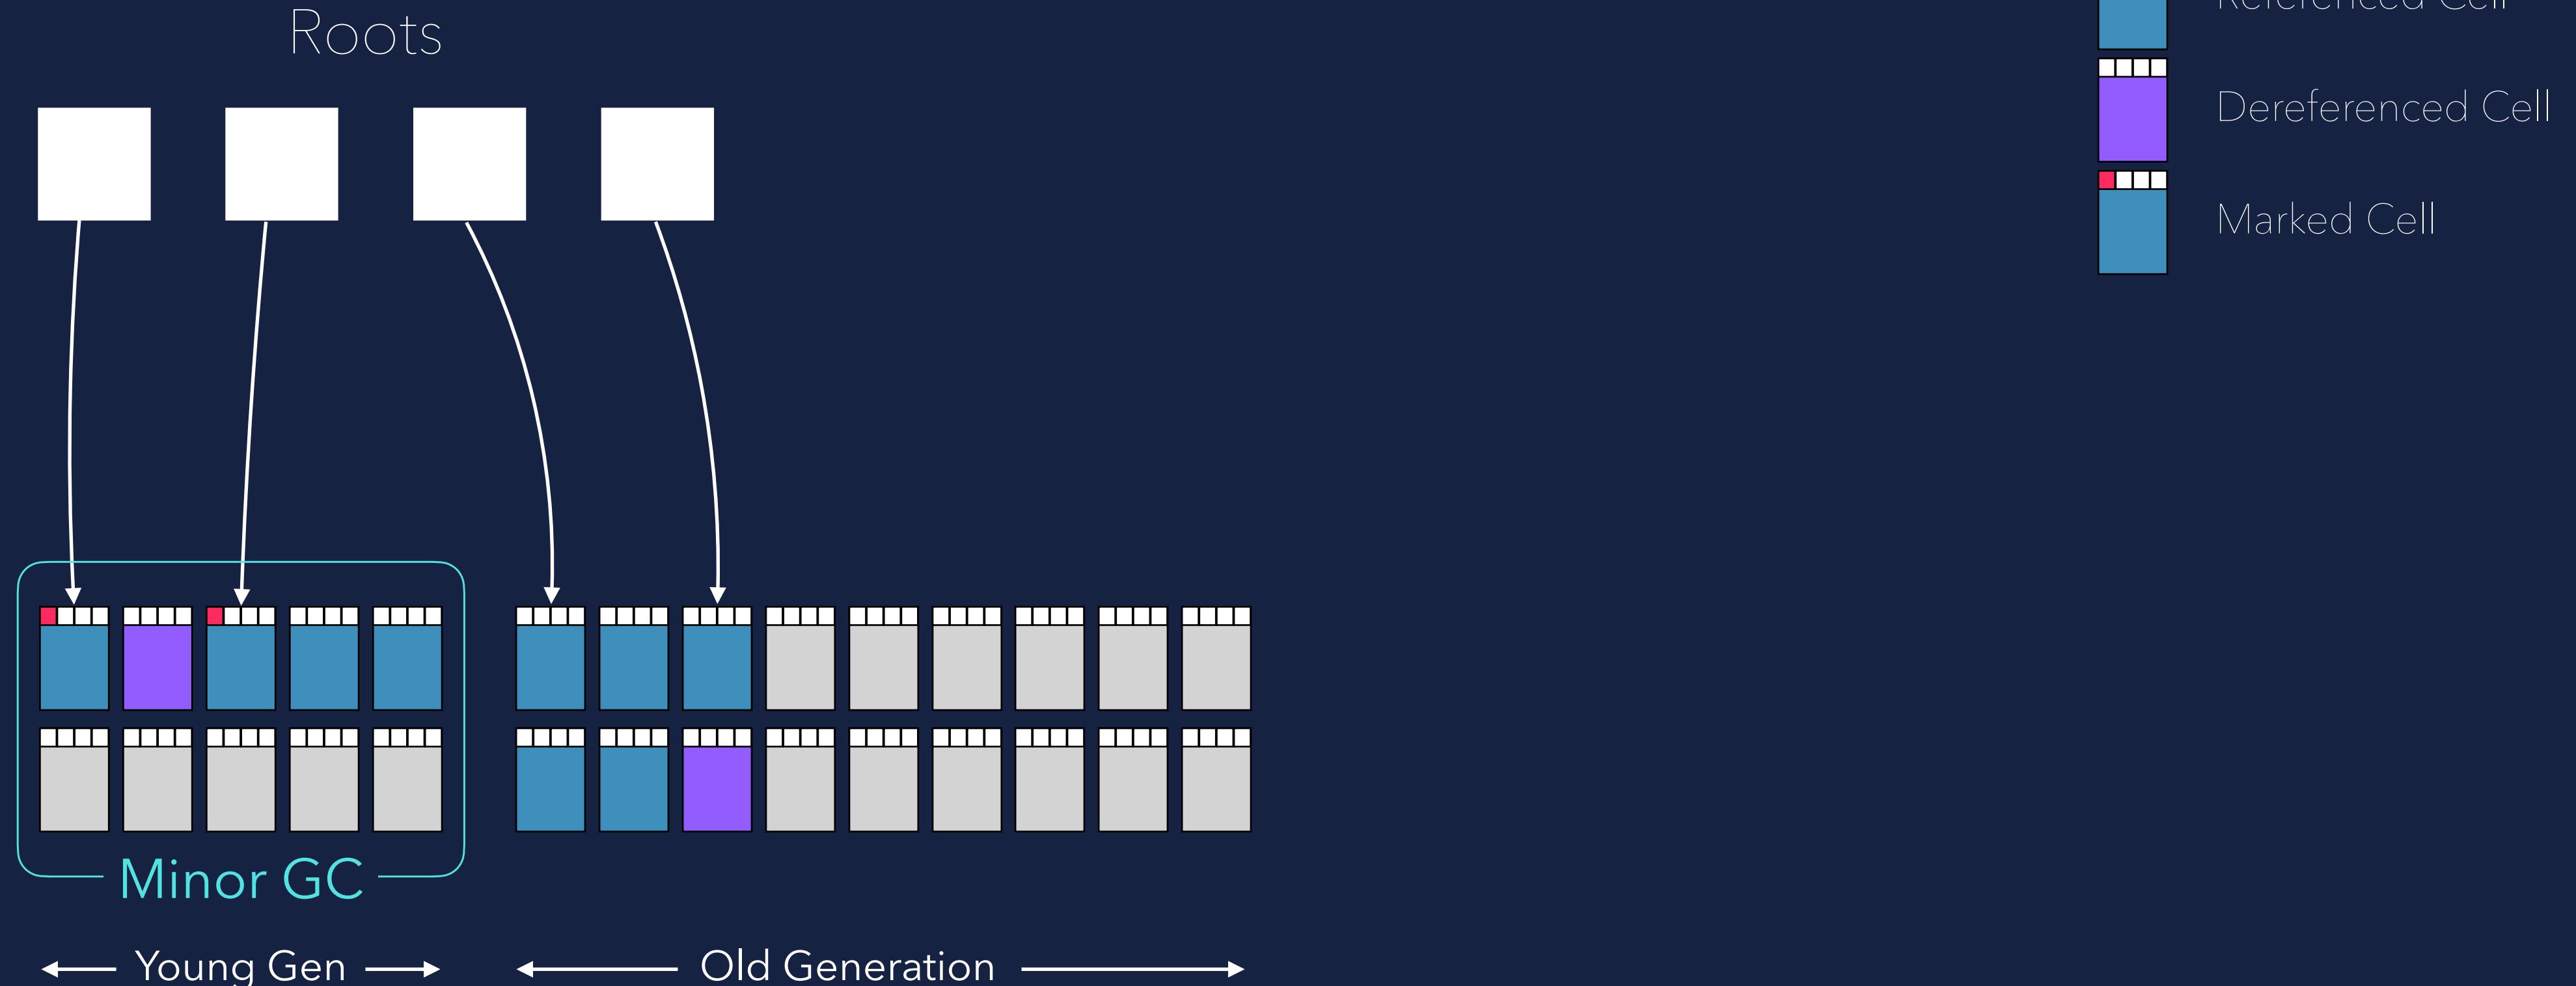
REMEMBERED SET

Also known as Card Table



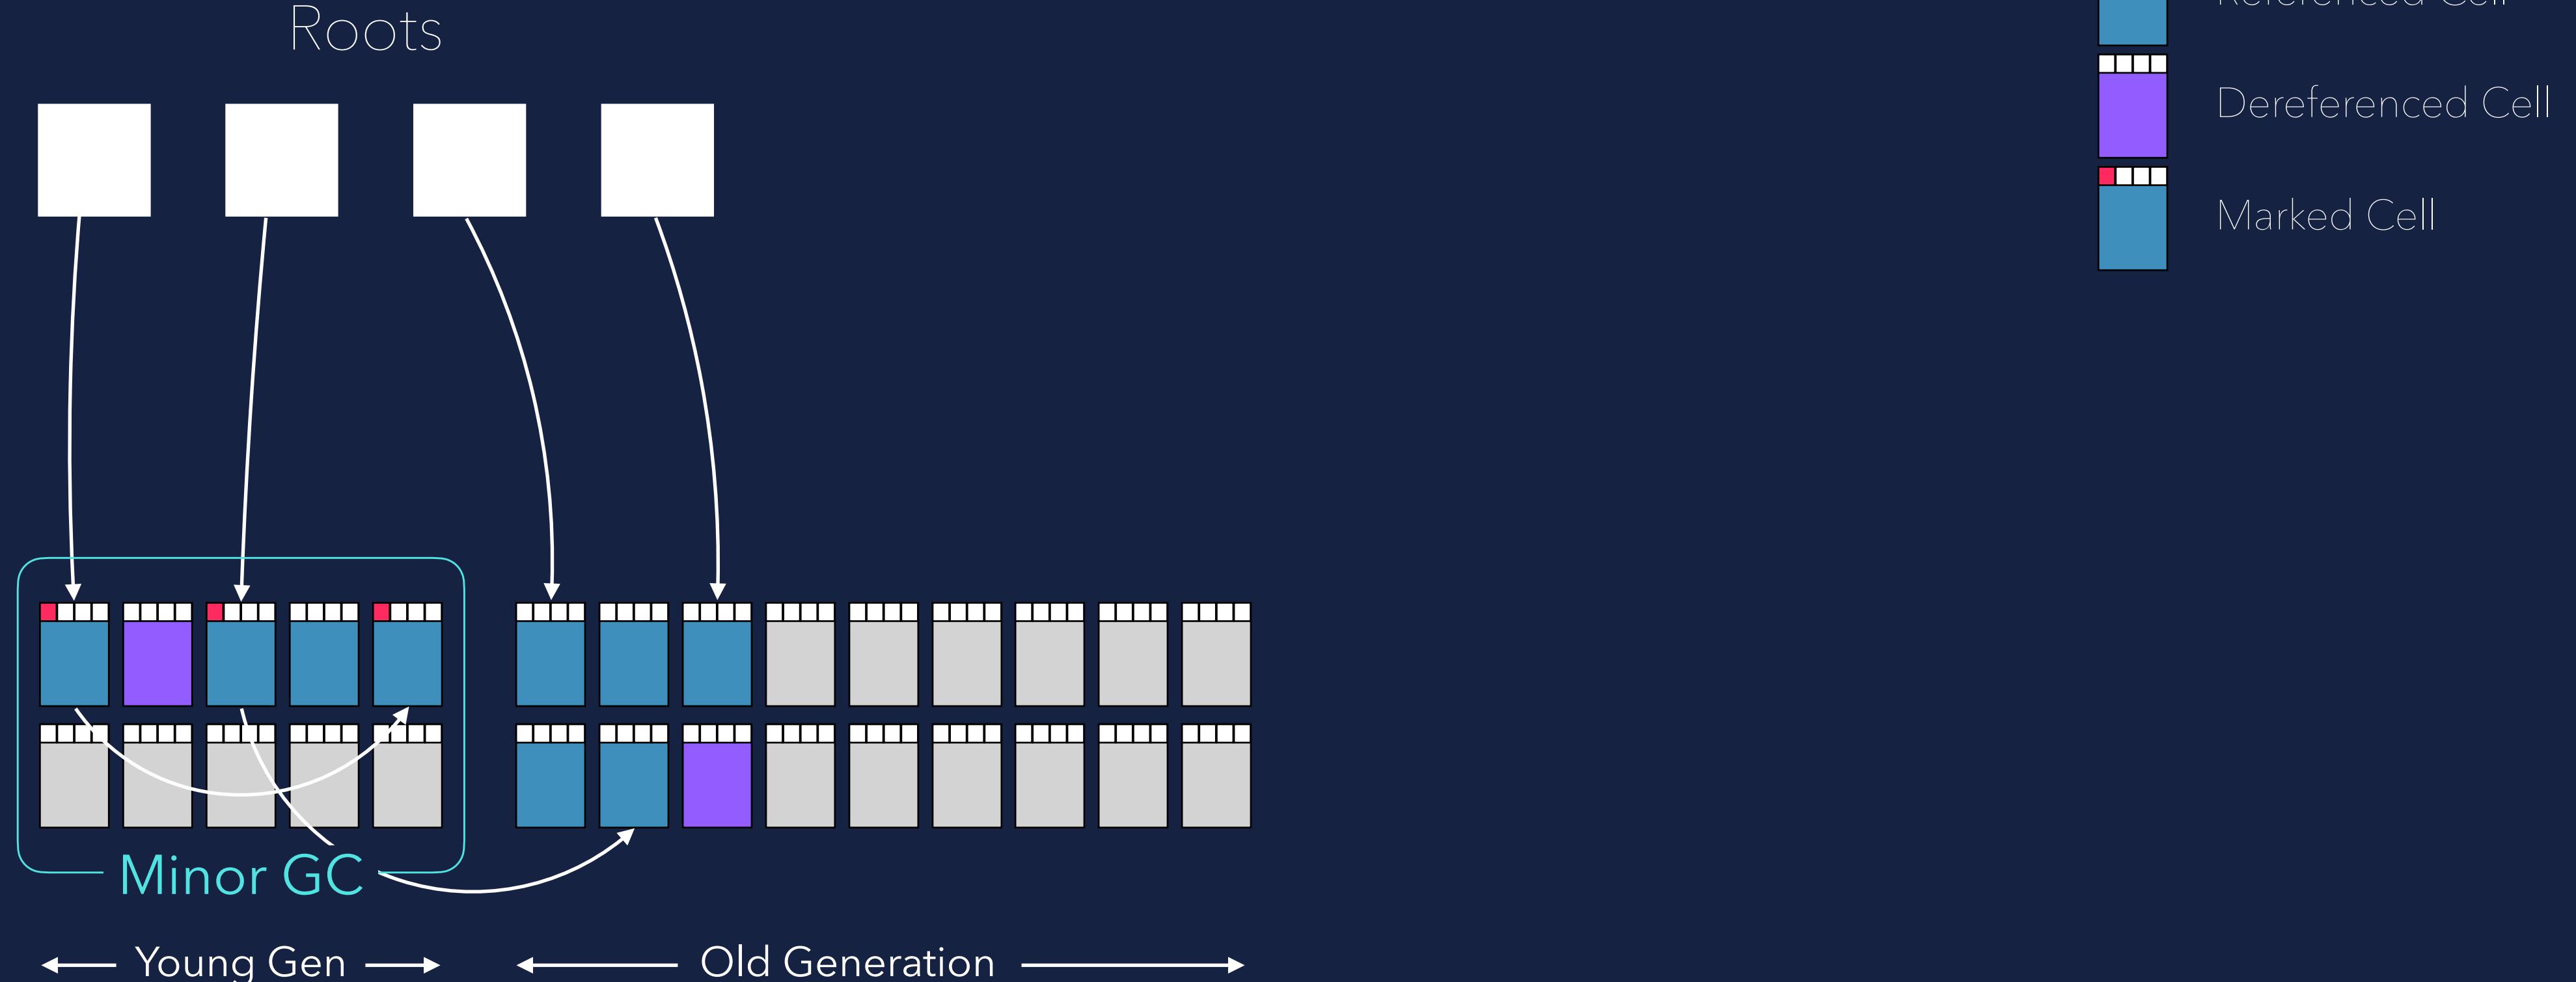
REMEMBERED SET

Also known as Card Table



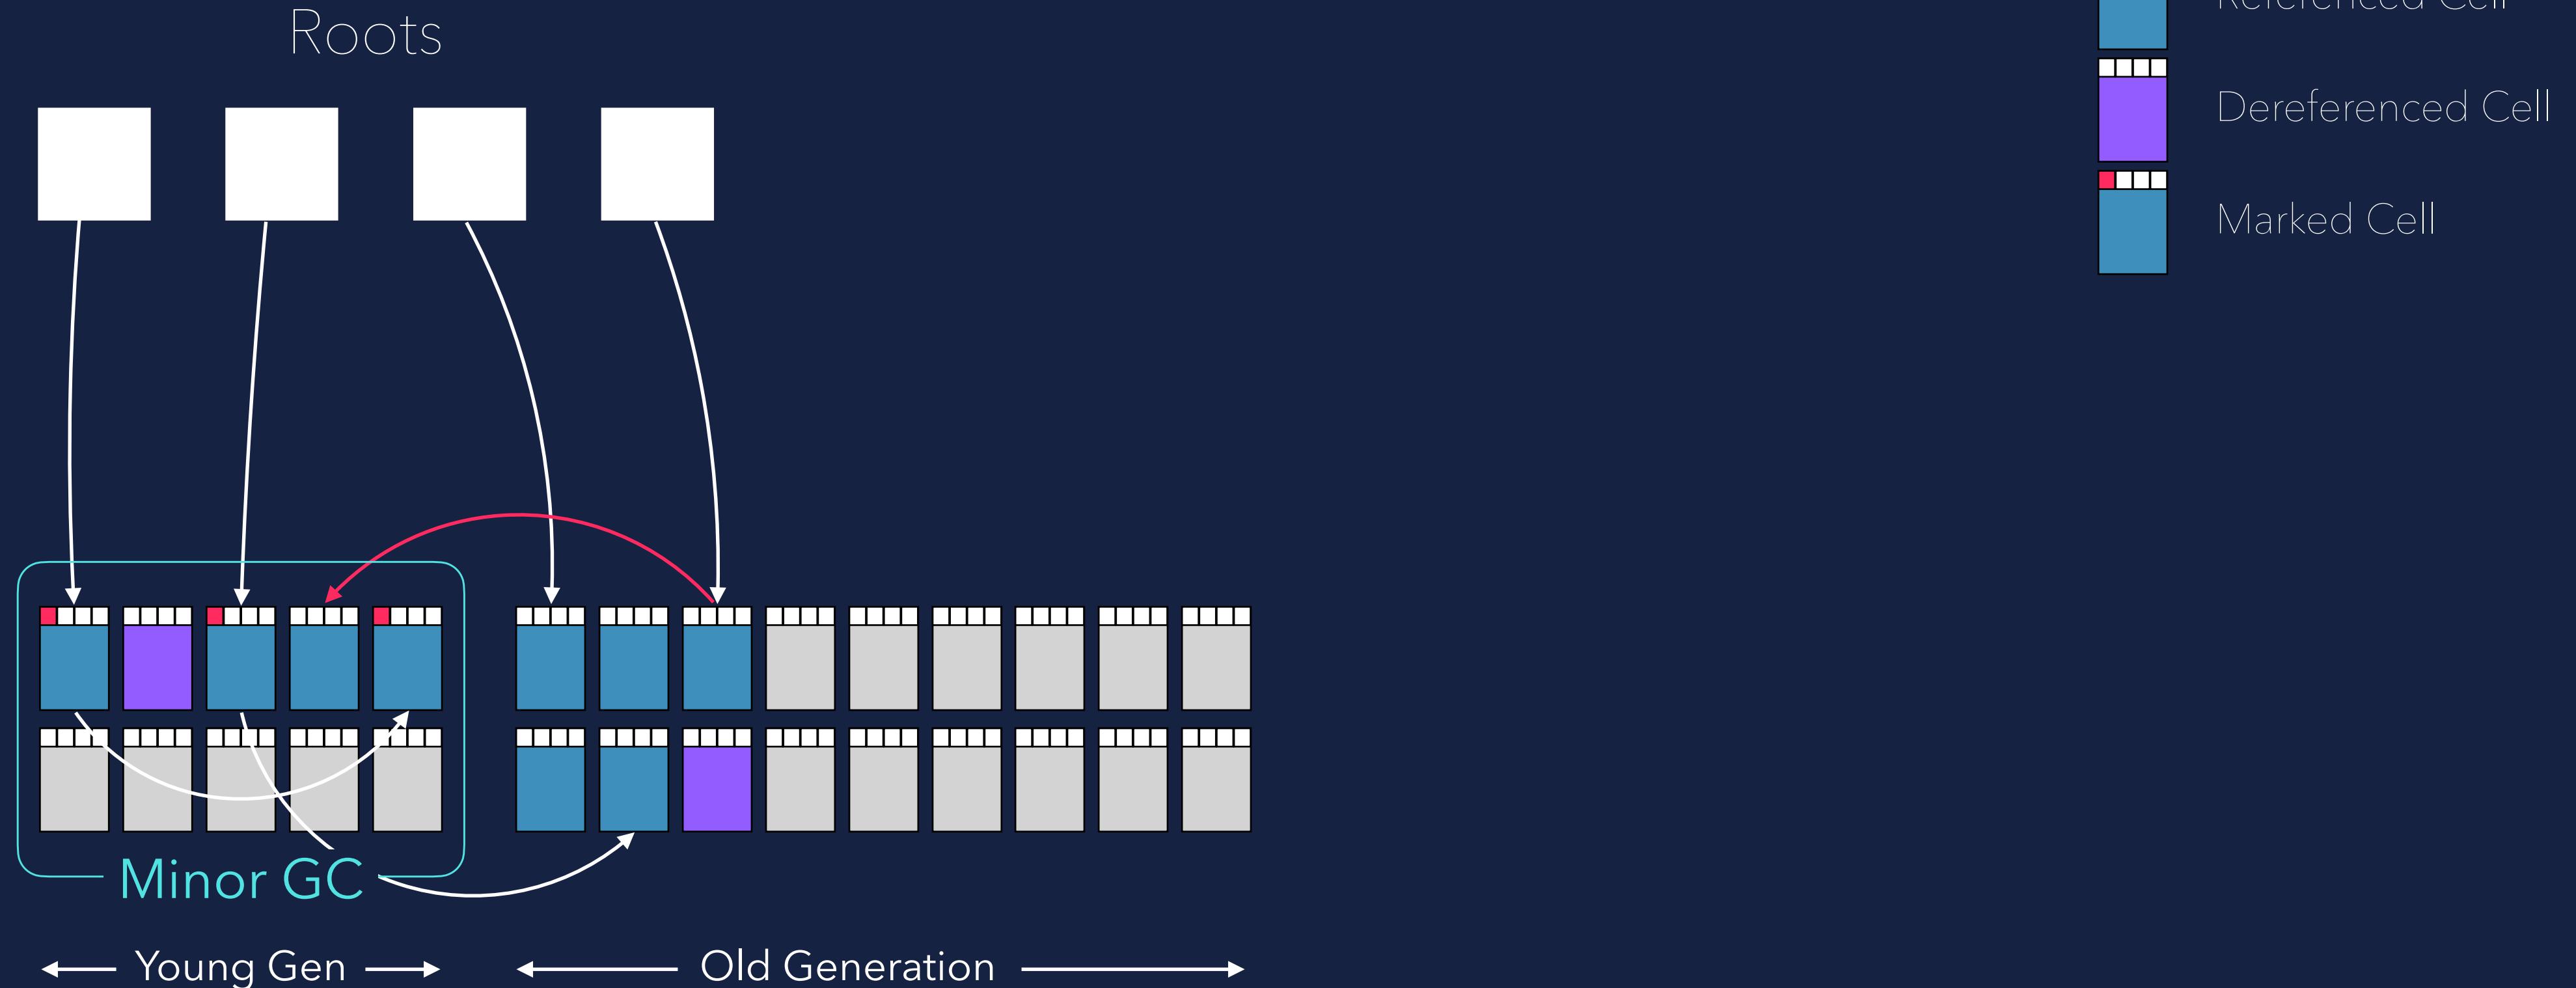
REMEMBERED SET

Also known as Card Table



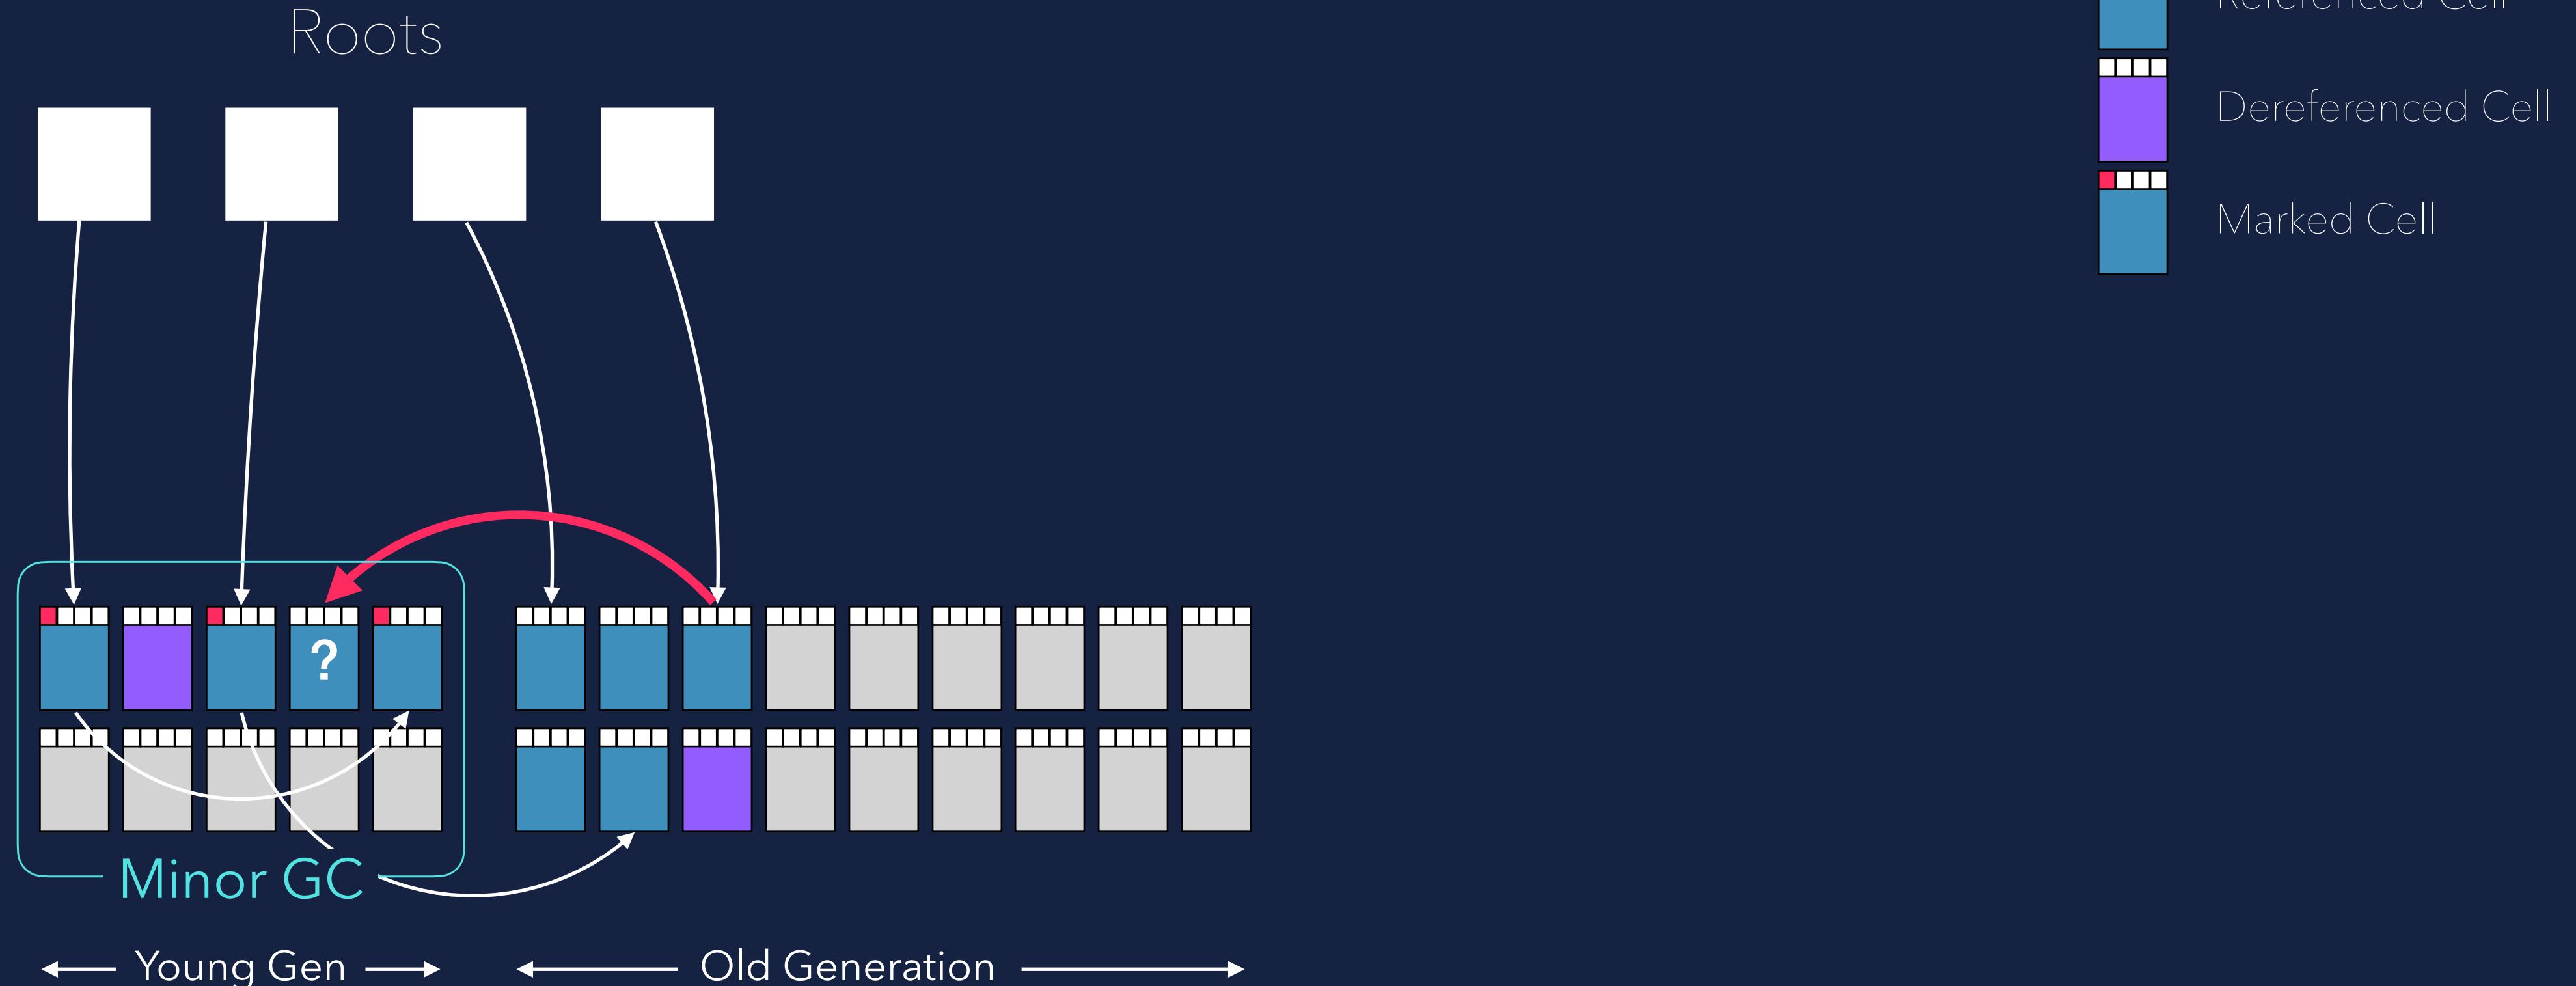
REMEMBERED SET

Also known as Card Table



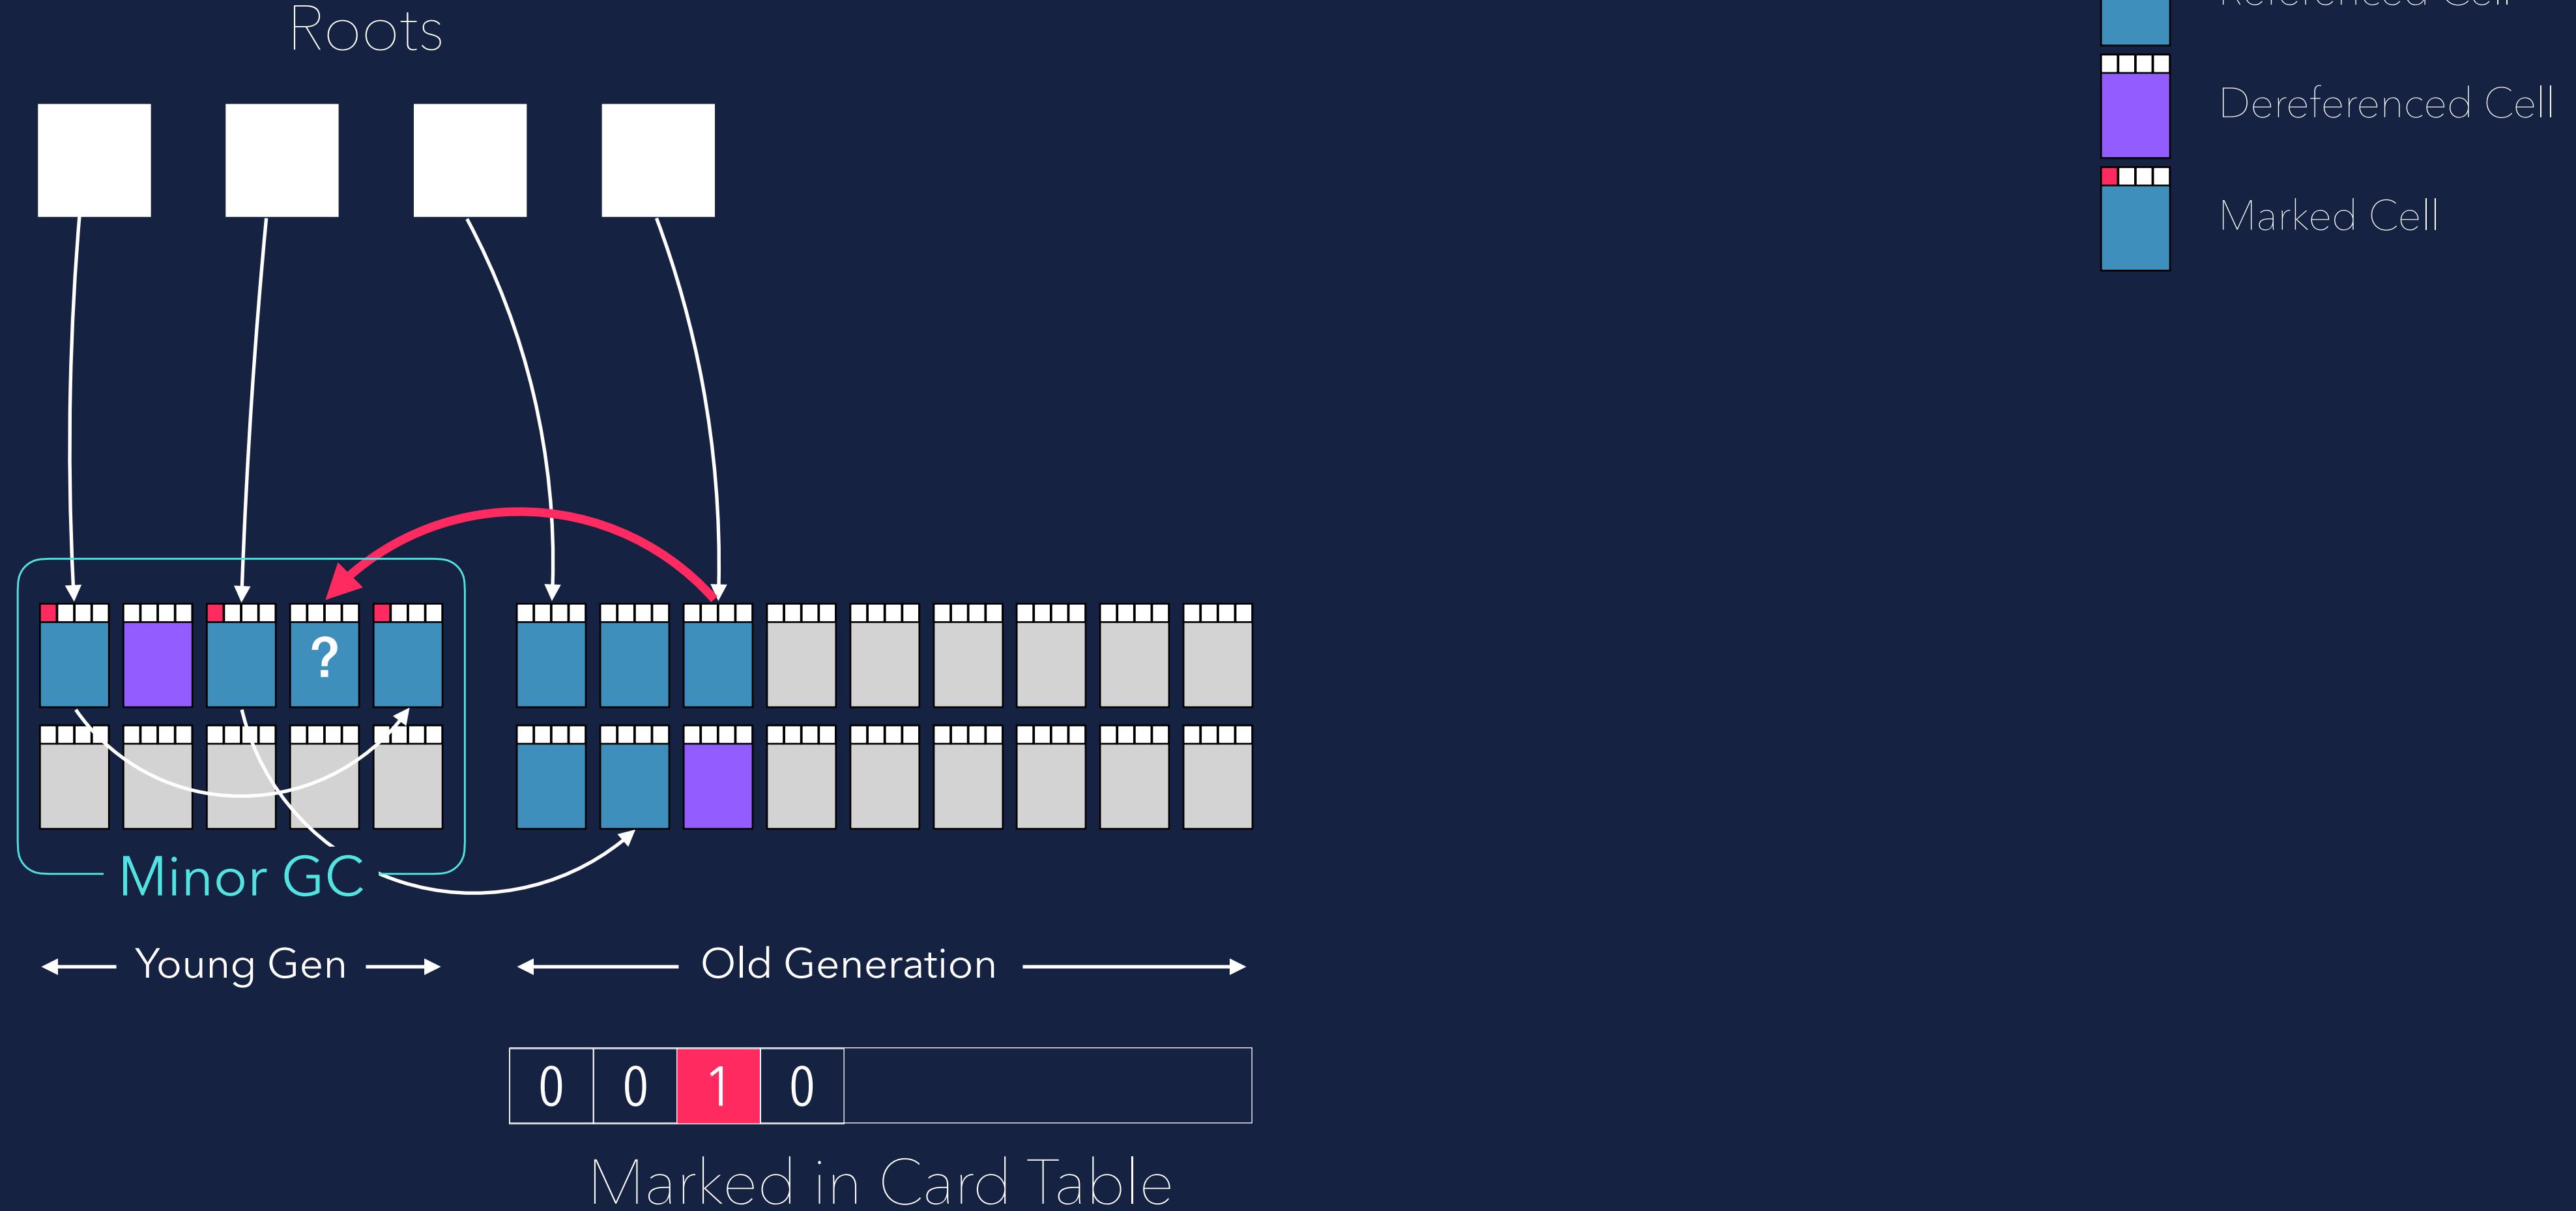
REMEMBERED SET

Also known as Card Table



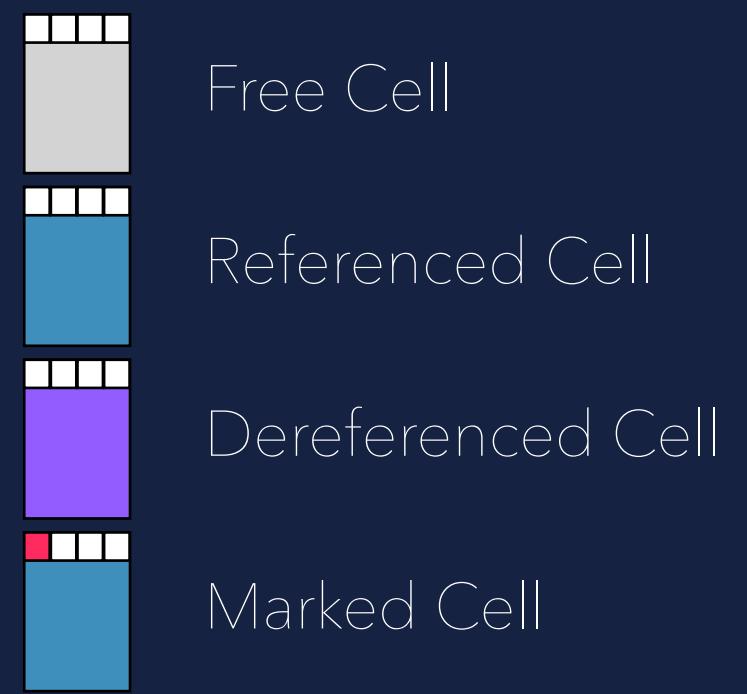
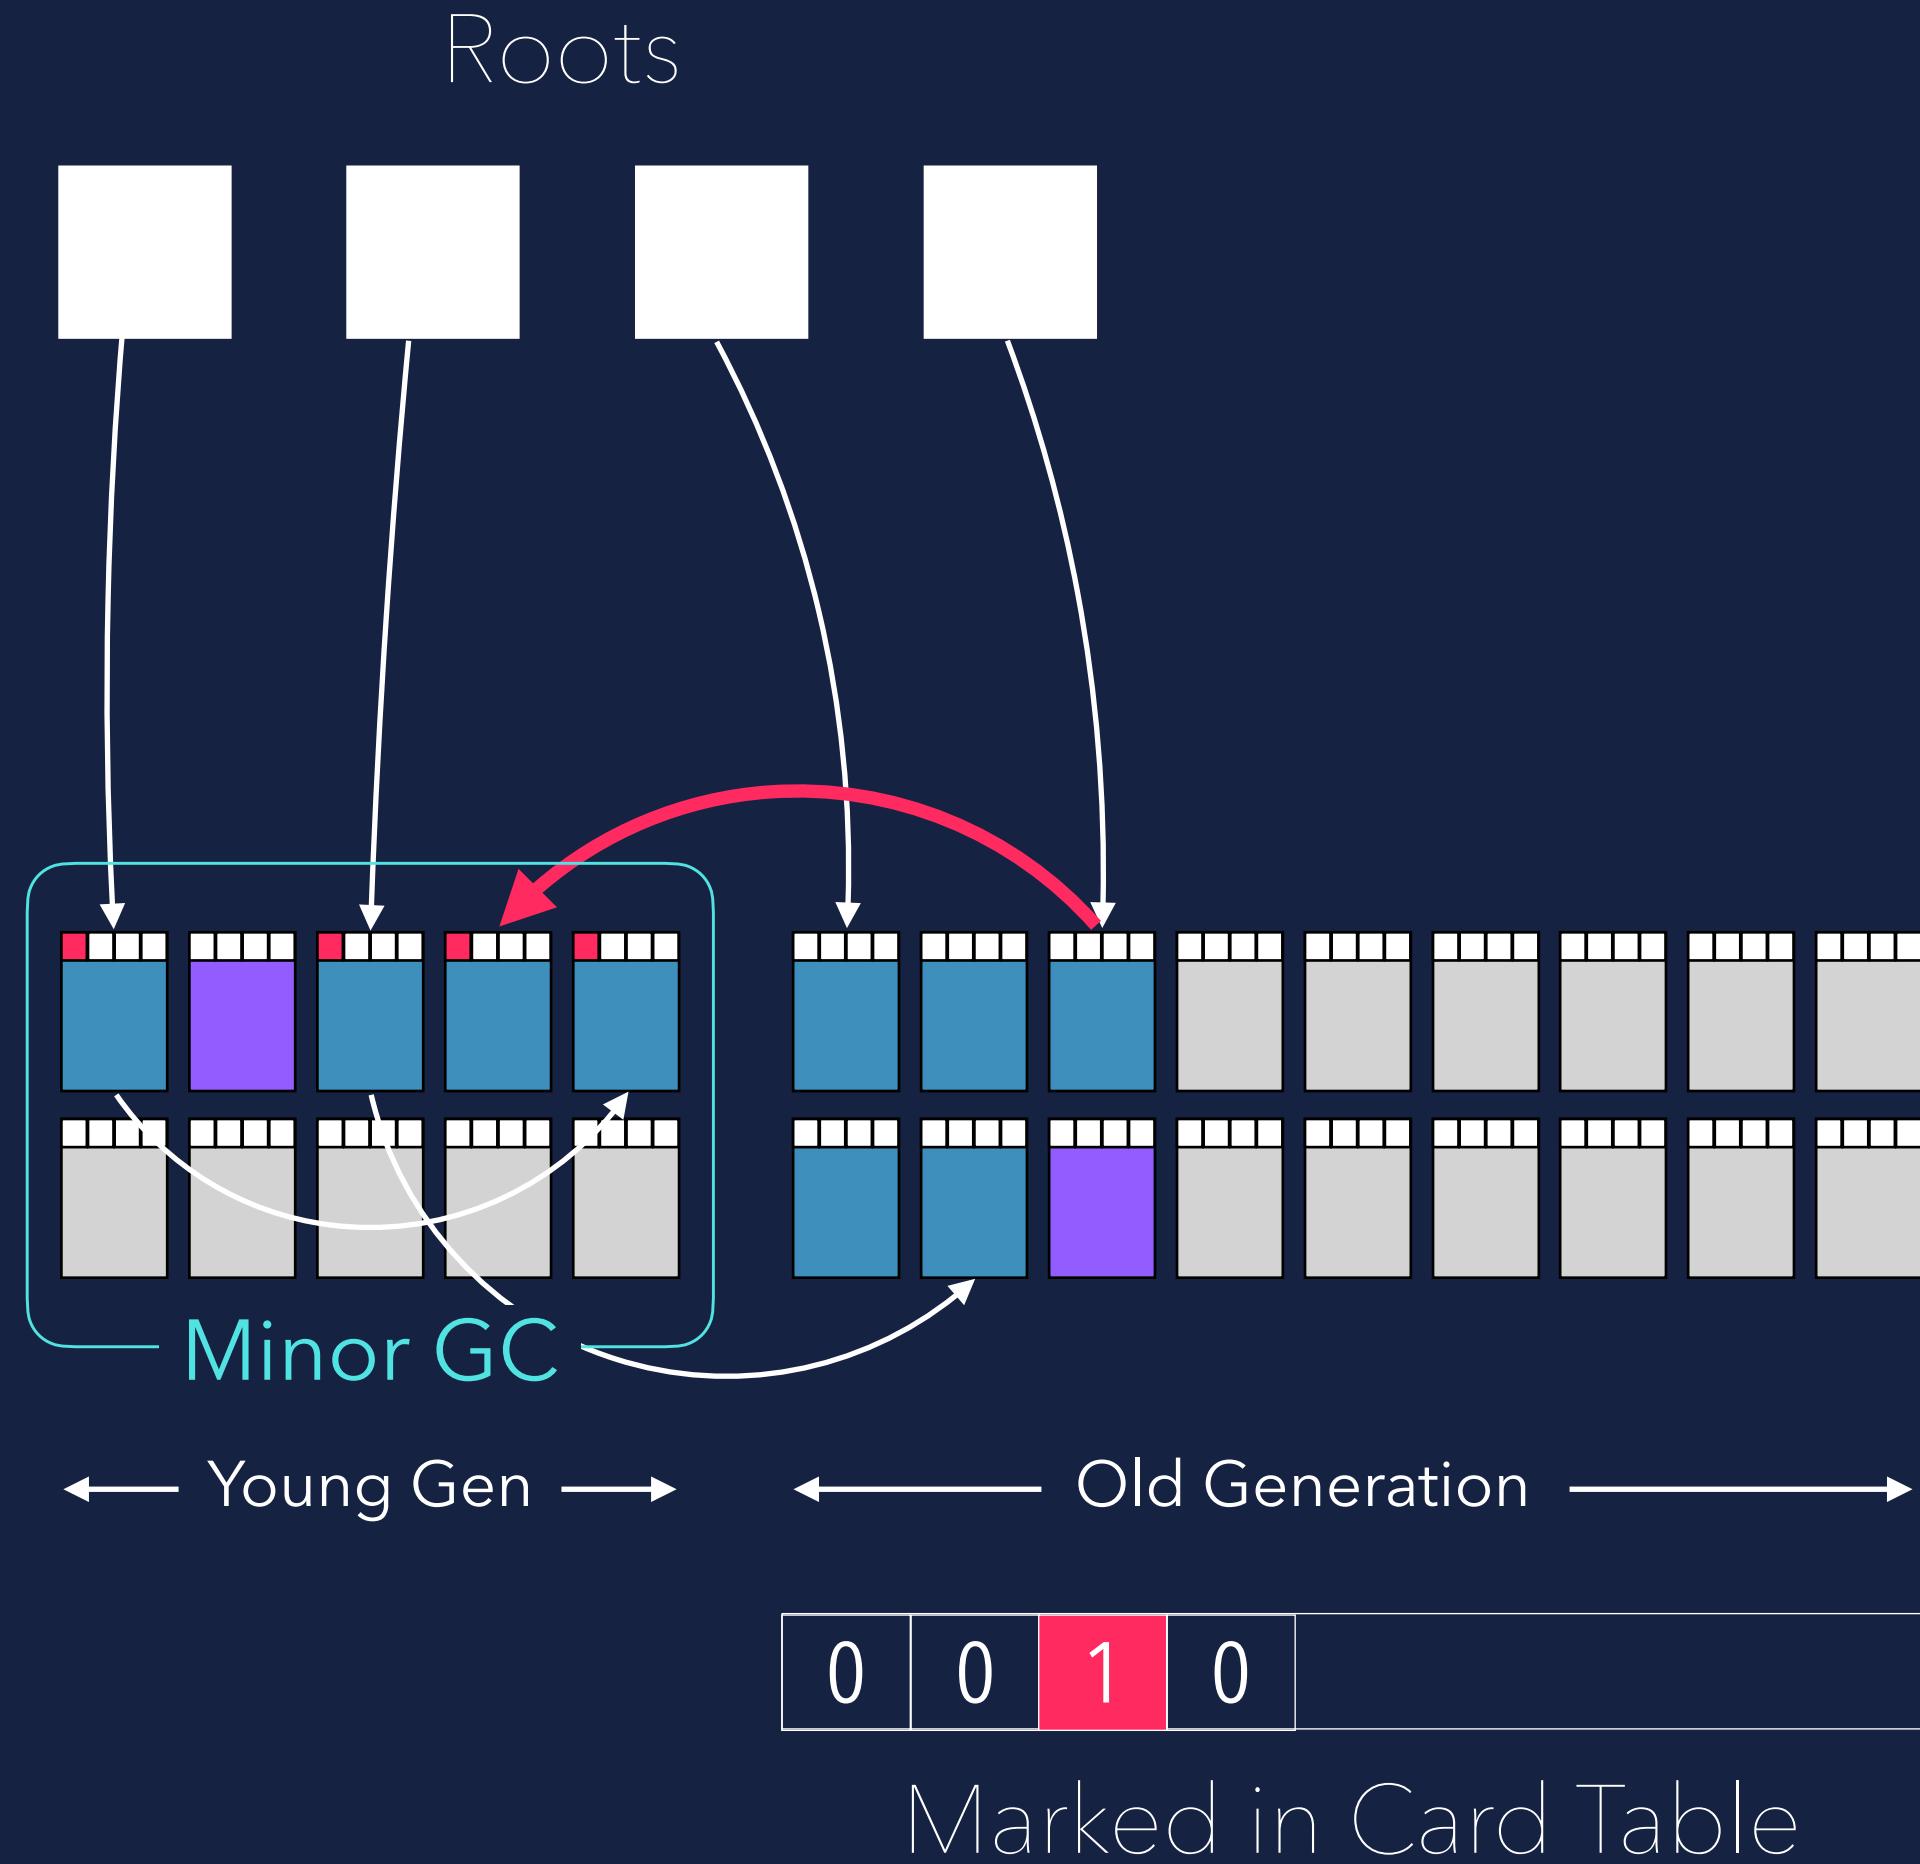
REMEMBERED SET

Also known as Card Table



REMEMBERED SET

Also known as Card Table



GC looks up Card Table,
finds the reference and
marks it as live

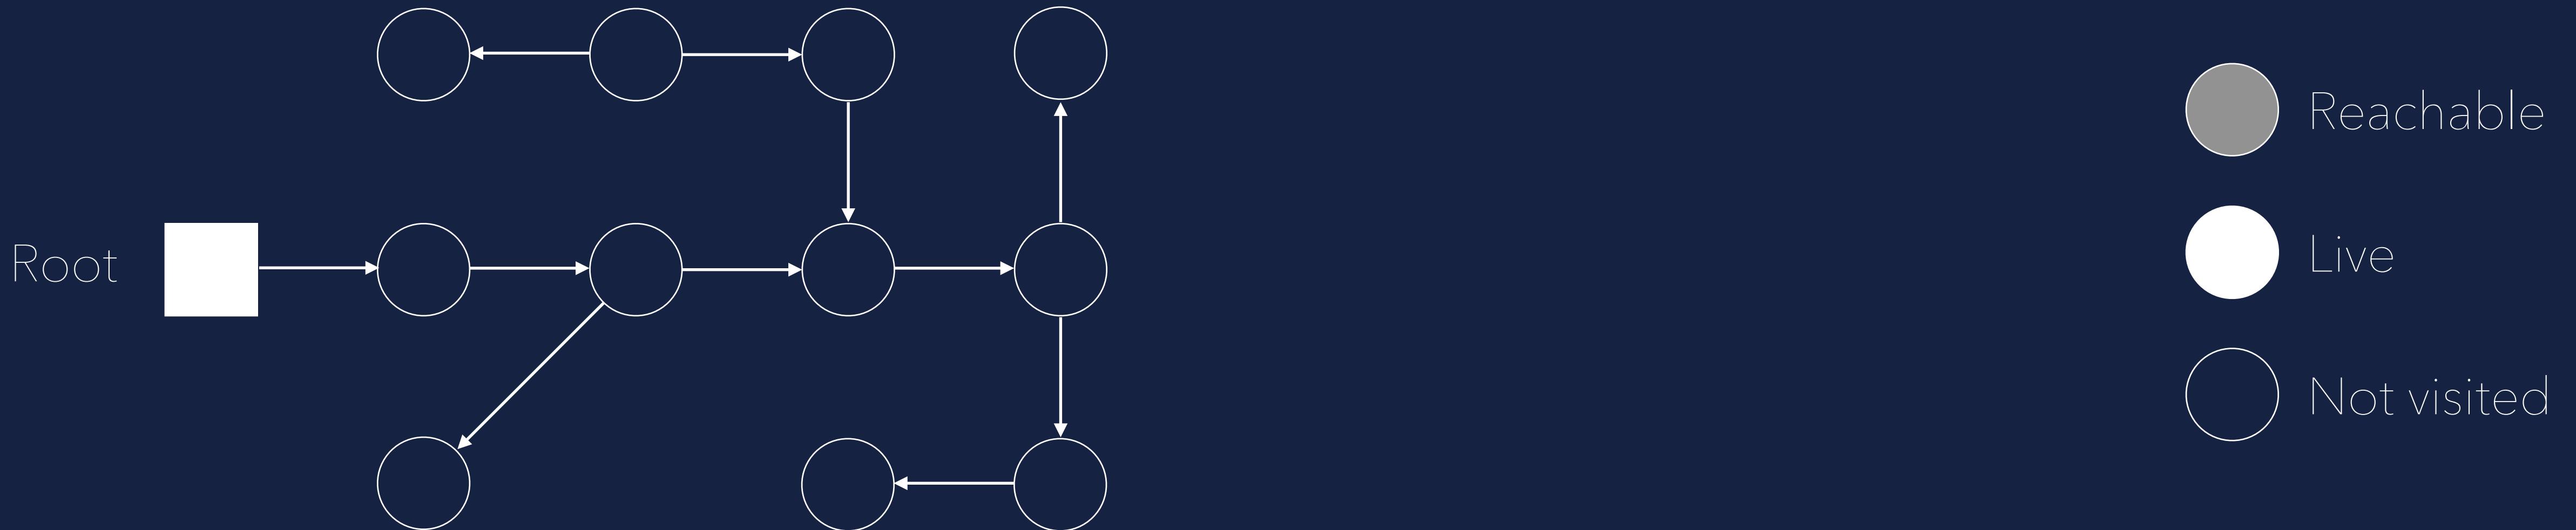
**CONCURRENT
COLLECTION?**

**CONCURRENCY
IS
HARD...**

CONCURRENT MARKING

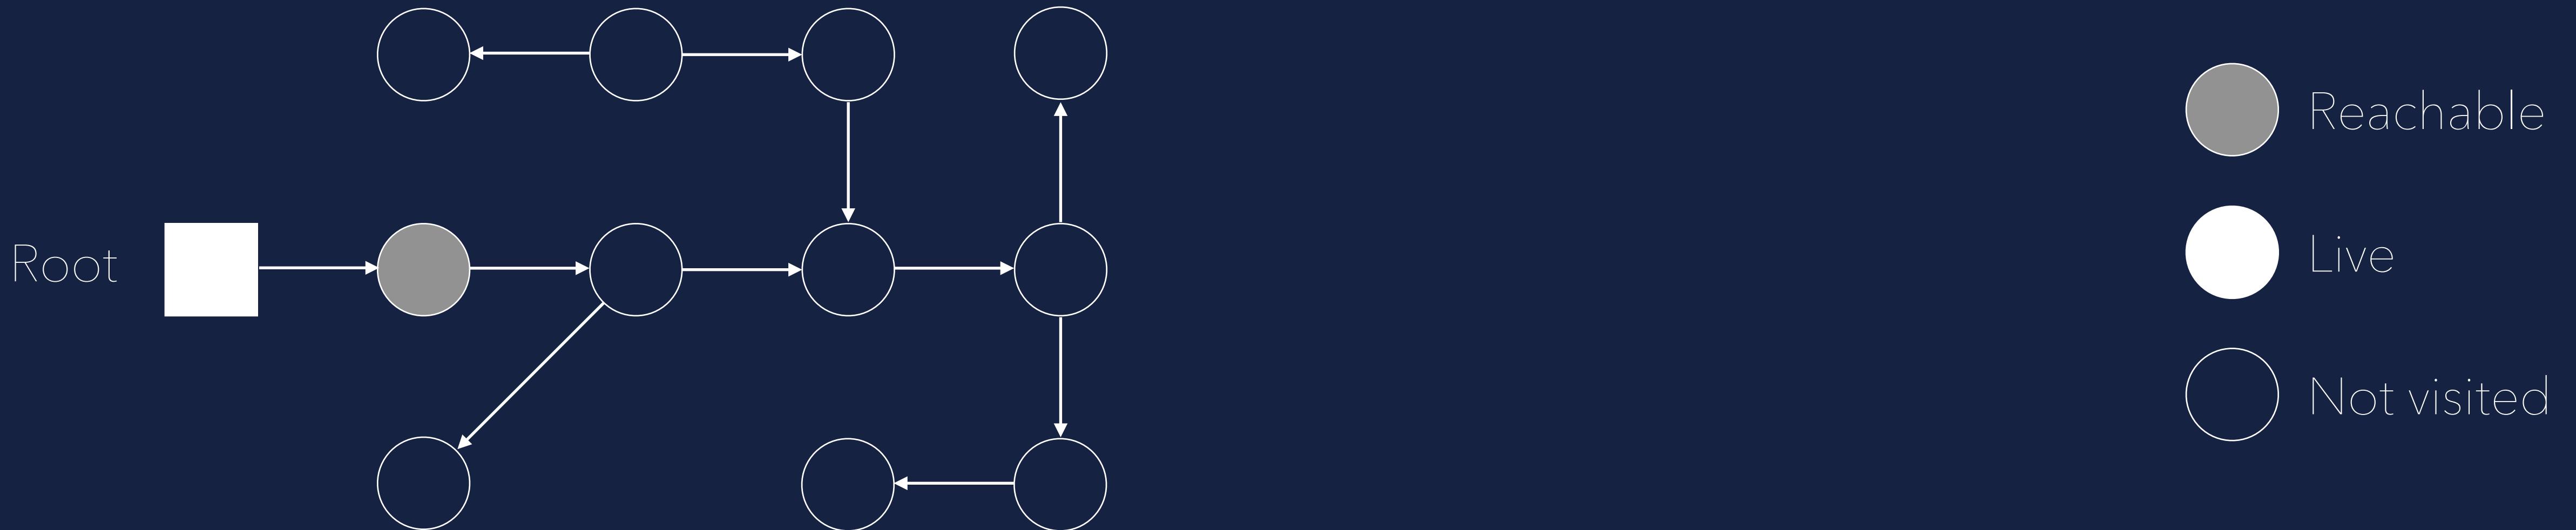
CONCURRENCY IS HARD...

Concurrent Marking



CONCURRENCY IS HARD...

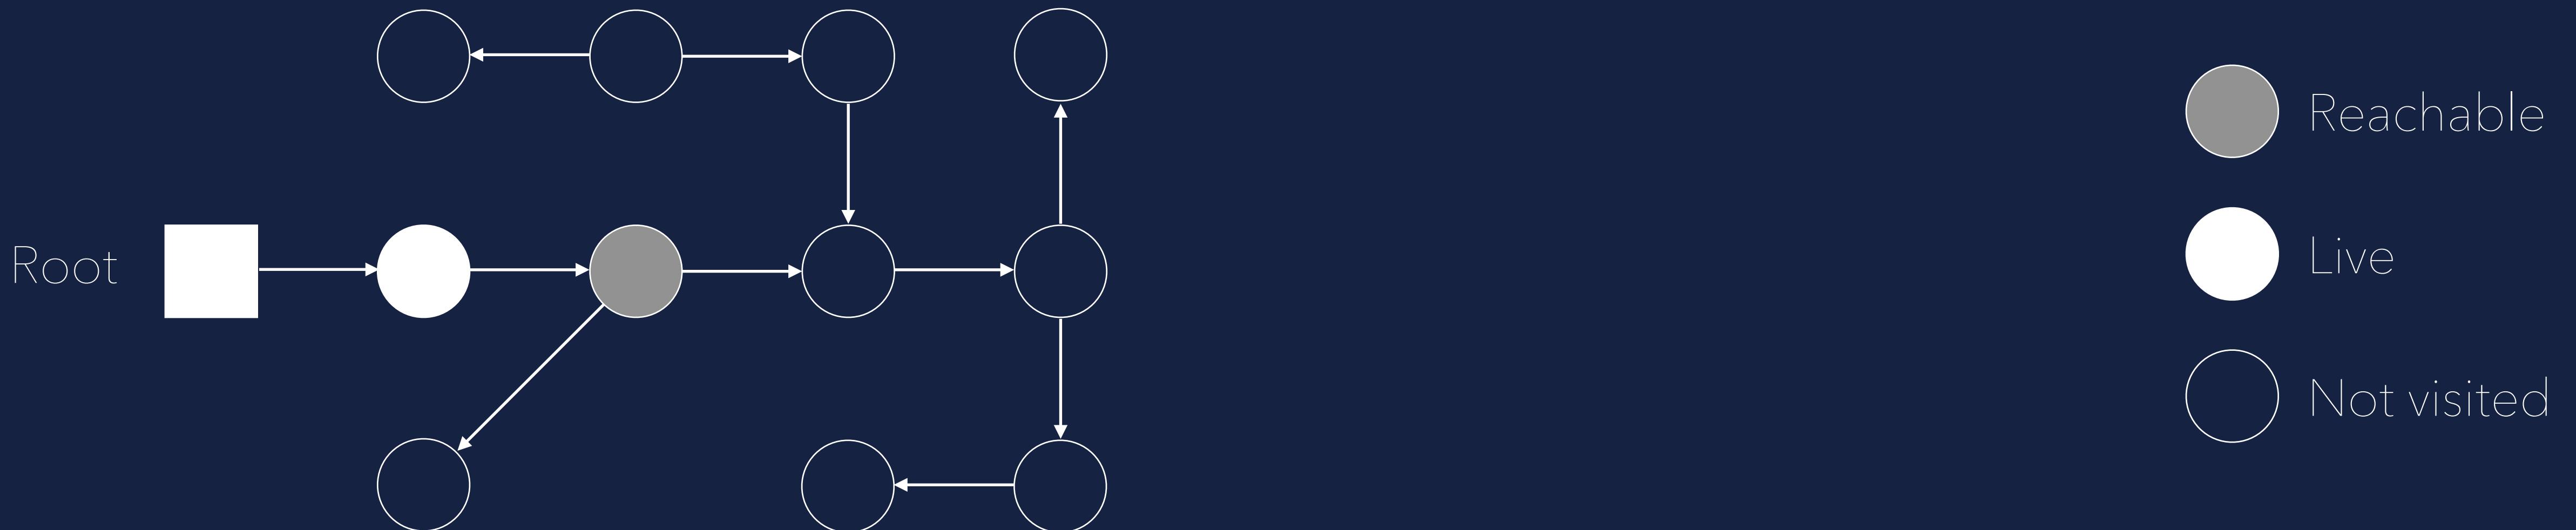
Concurrent Marking



Collector starts marking objects

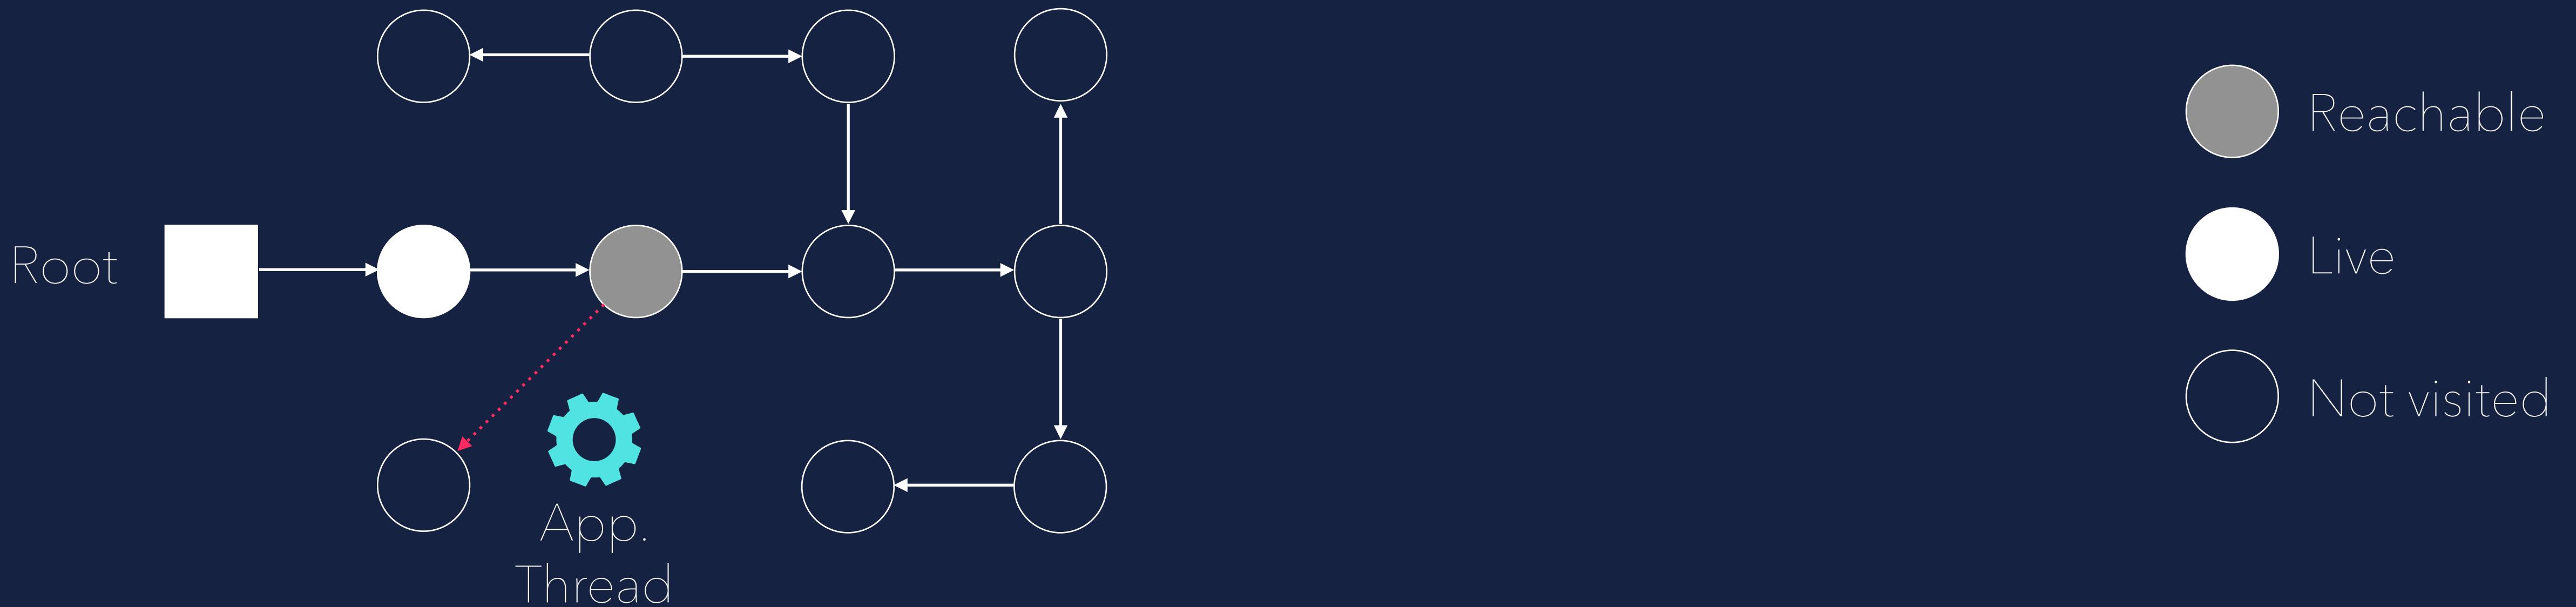
CONCURRENCY IS HARD...

Concurrent Marking



CONCURRENCY IS HARD...

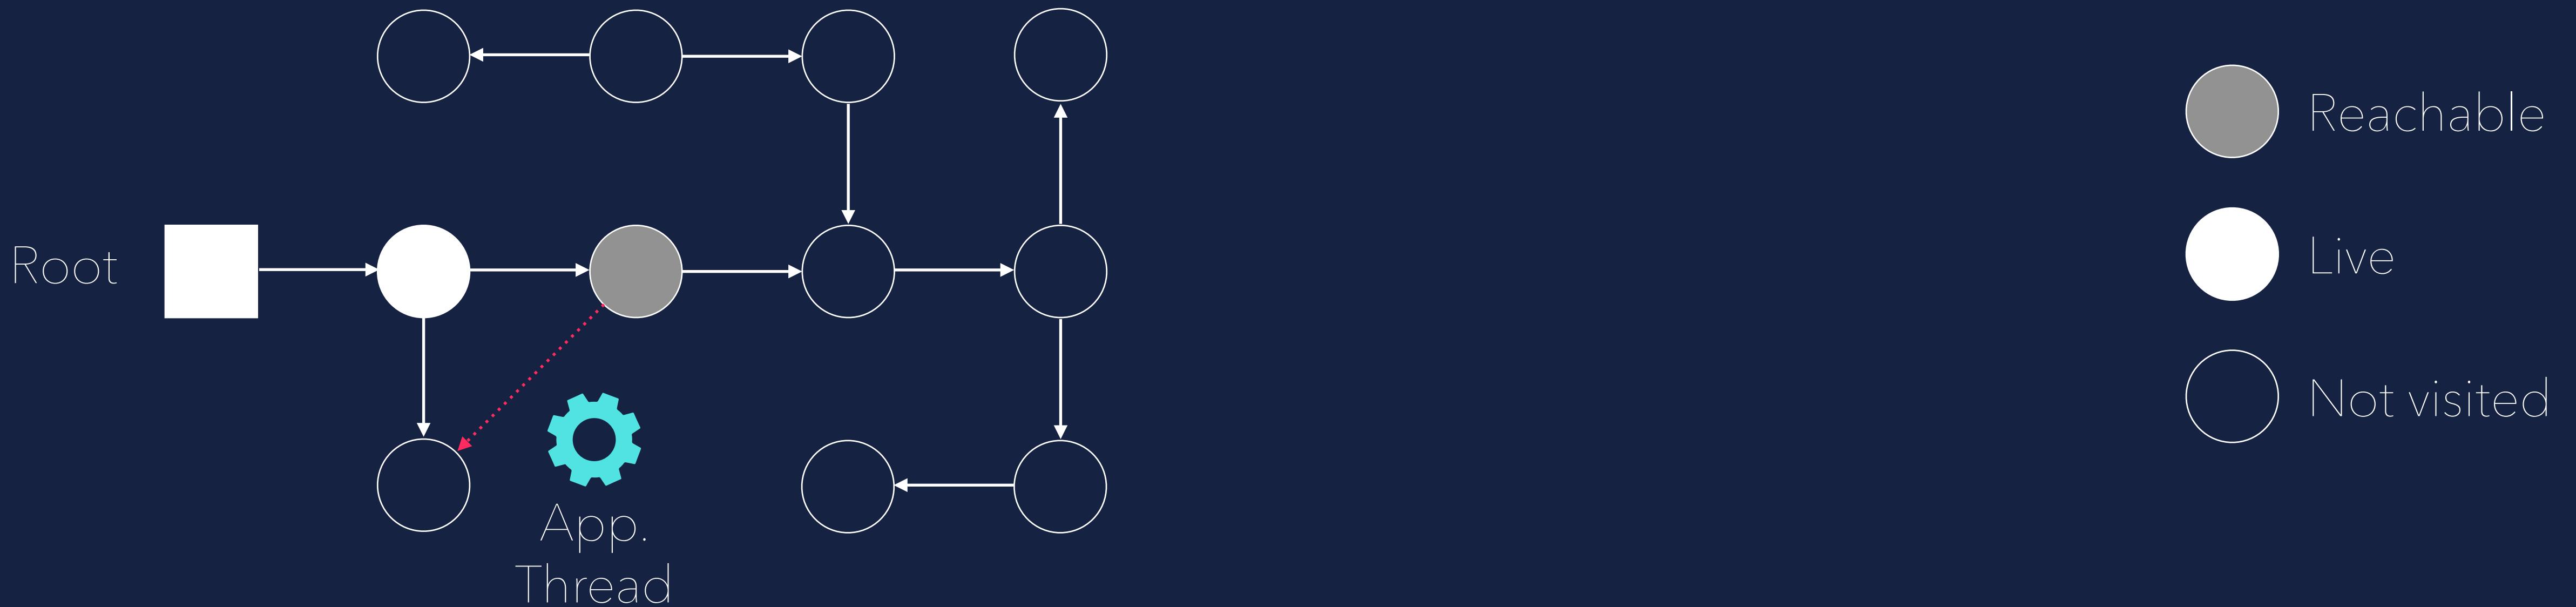
Concurrent Marking



Mutator removes reference and creates a new one from an already visited cell !

CONCURRENCY IS HARD...

Concurrent Marking



Mutator removes reference and creates a new one from an already visited cell !

CONCURRENCY IS HARD...

Concurrent Marking



CONCURRENCY IS HARD...

Concurrent Marking



Won't be detected by the Garbage Collector !

PPP

azul

BARRIERS TO THE RESCUE

BARRIERS

Read / Write Barriers



Mechanisms to execute memory management code when a read/write on some object takes place

BARRIERS

Read / Write Barriers

- Mechanisms to execute memory management code when a read/write on some object takes place
- Used to keep track of inter-generational references.
(references from old generation to young generation, the so called Rembered Set)

BARRIERS

Read / Write Barriers

- Mechanisms to execute memory management code when a read/write on some object takes place
- Used to keep track of inter-generational references.
(references from old generation to young generation, the so called Rembered Set)
- Used to synchronize action between mutator and collector
(allocation concurrent to collection)

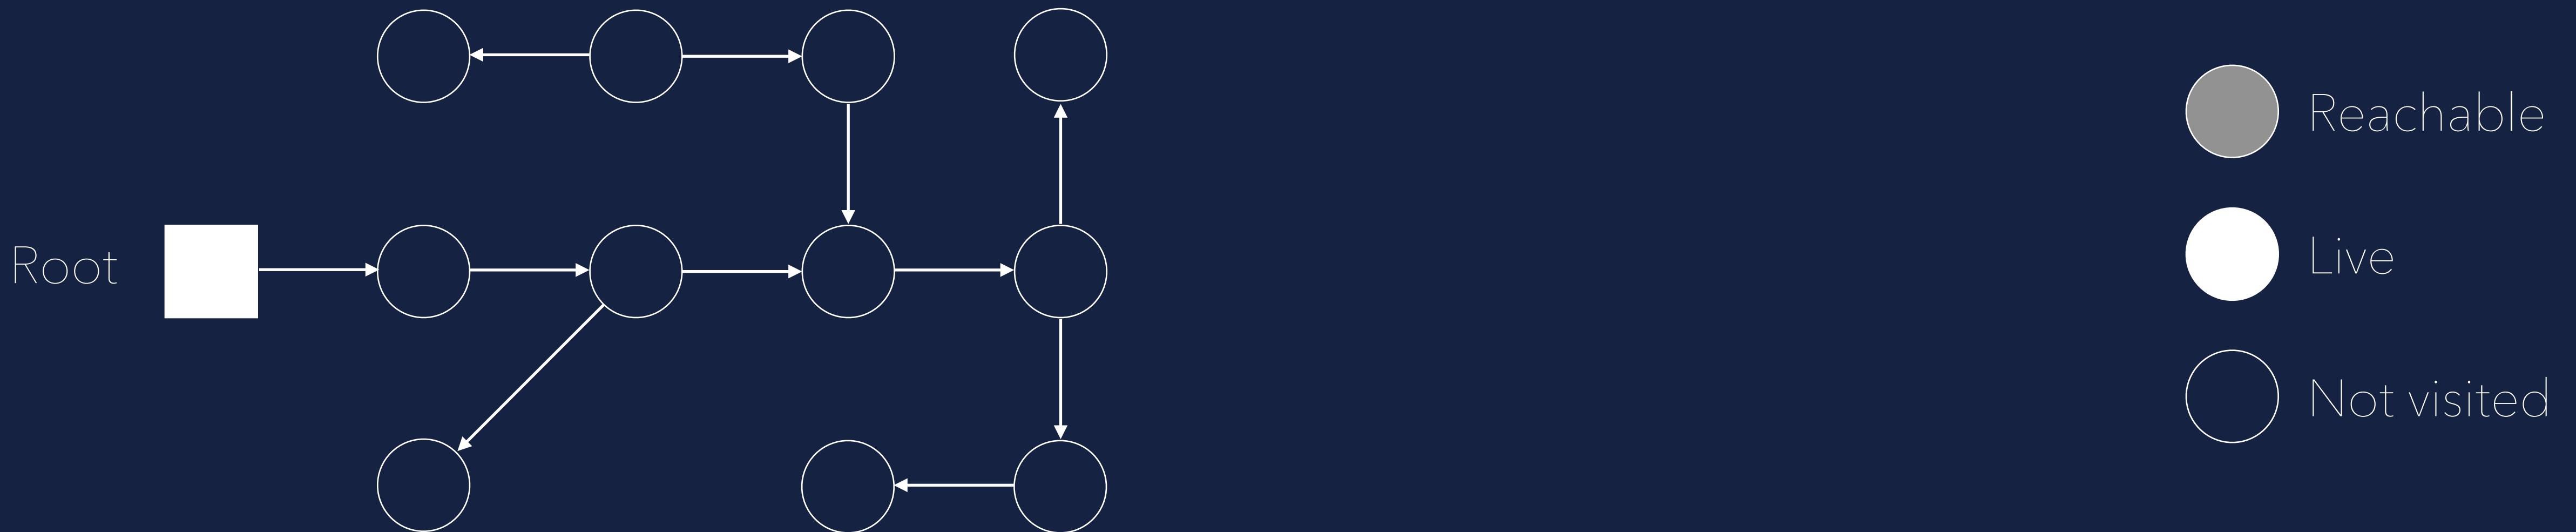
BARRIERS

Read / Write Barriers

- Mechanisms to execute memory management code when a read/write on some object takes place
- Used to keep track of inter-generational references.
(references from old generation to young generation, the so called Rembered Set)
- Used to synchronize action between mutator and collector
(allocation concurrent to collection)
- Read Barriers are usually more expensive
(reads 75% to writes 25% -> Read Barriers must be very efficient)

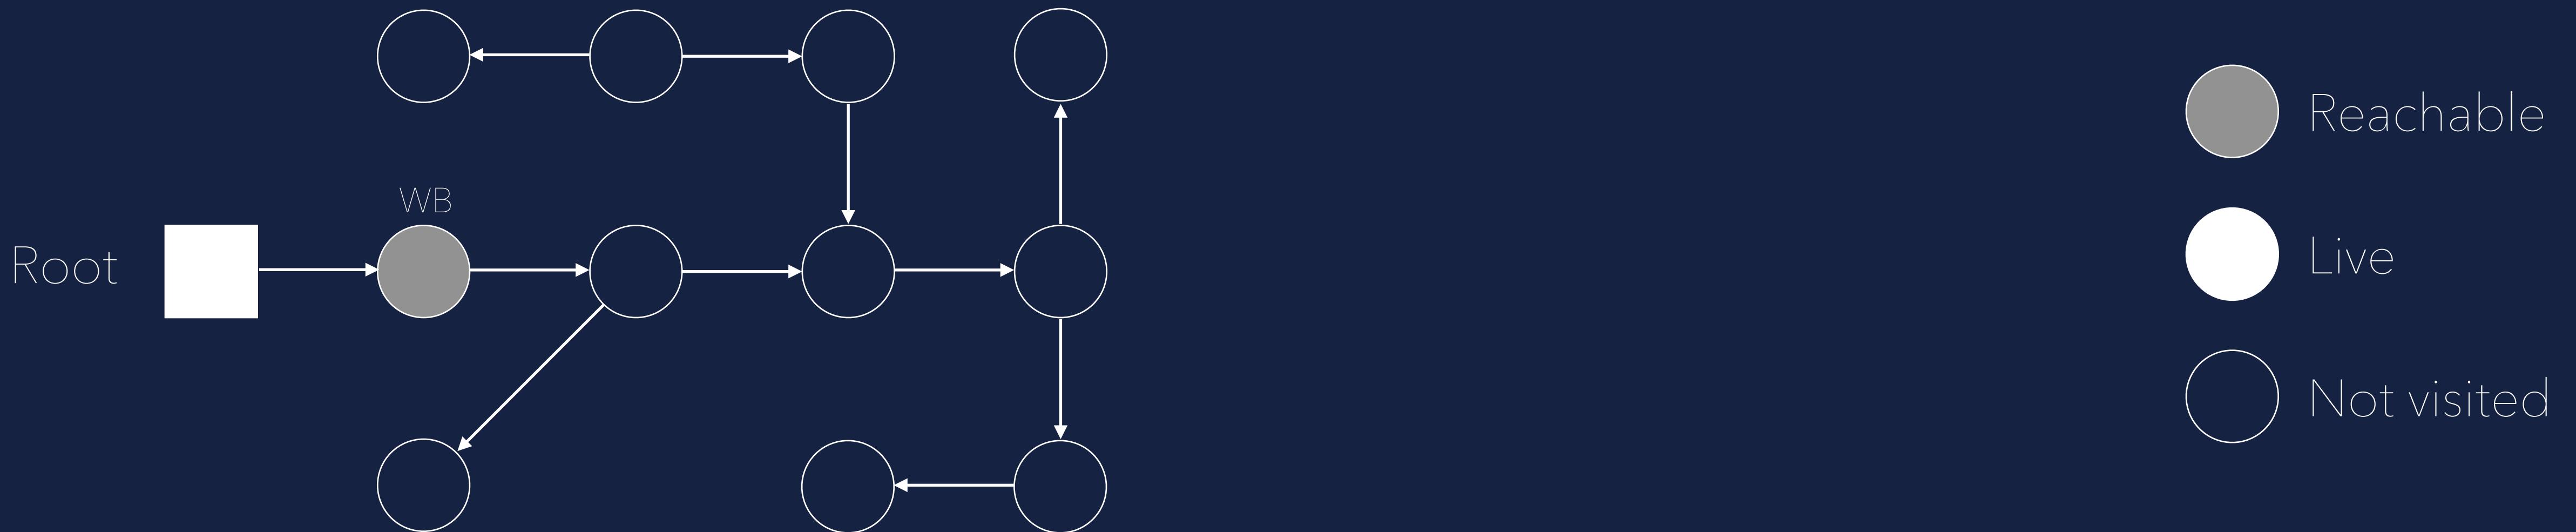
CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers



CONCURRENCY IS HARD...

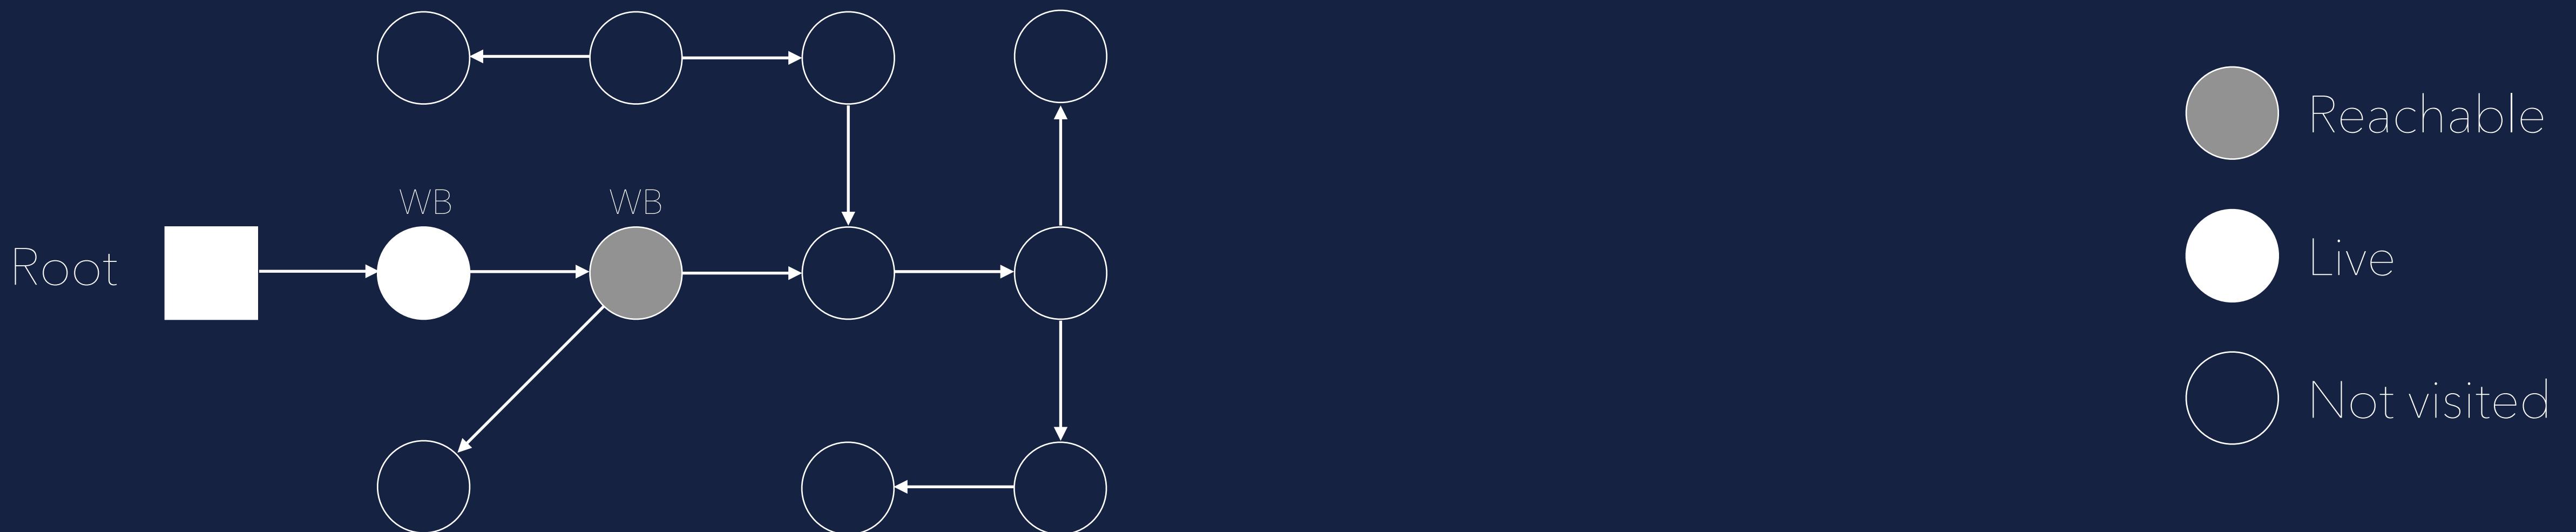
Concurrent Marking using Write Barriers



Collector starts marking objects

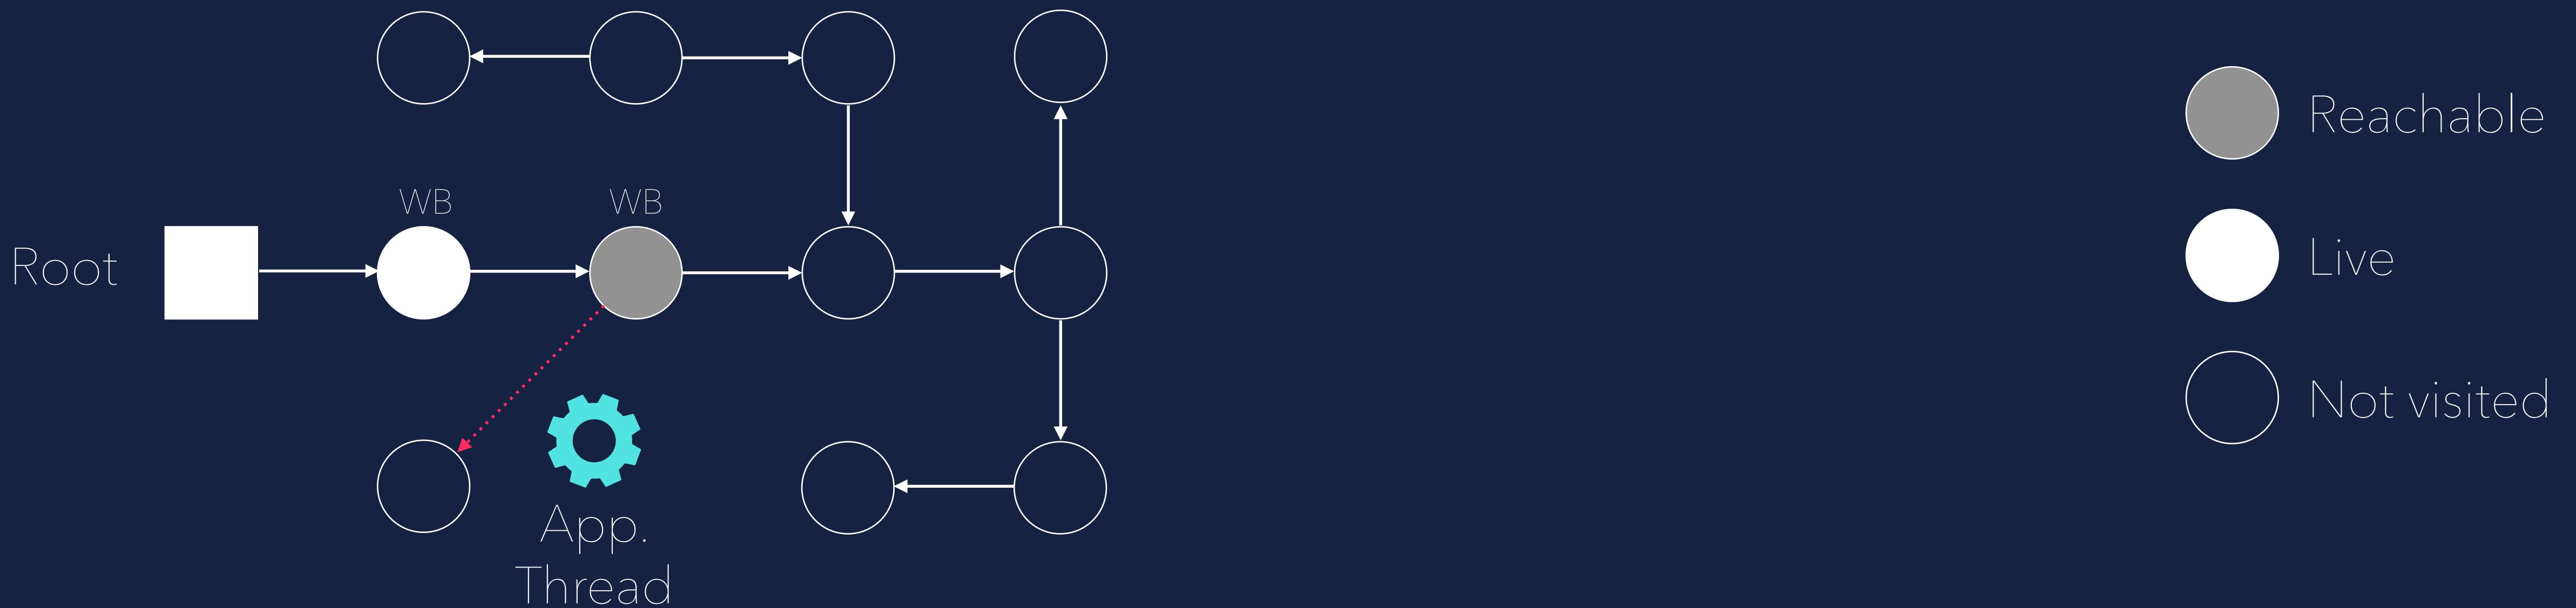
CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers



CONCURRENCY IS HARD...

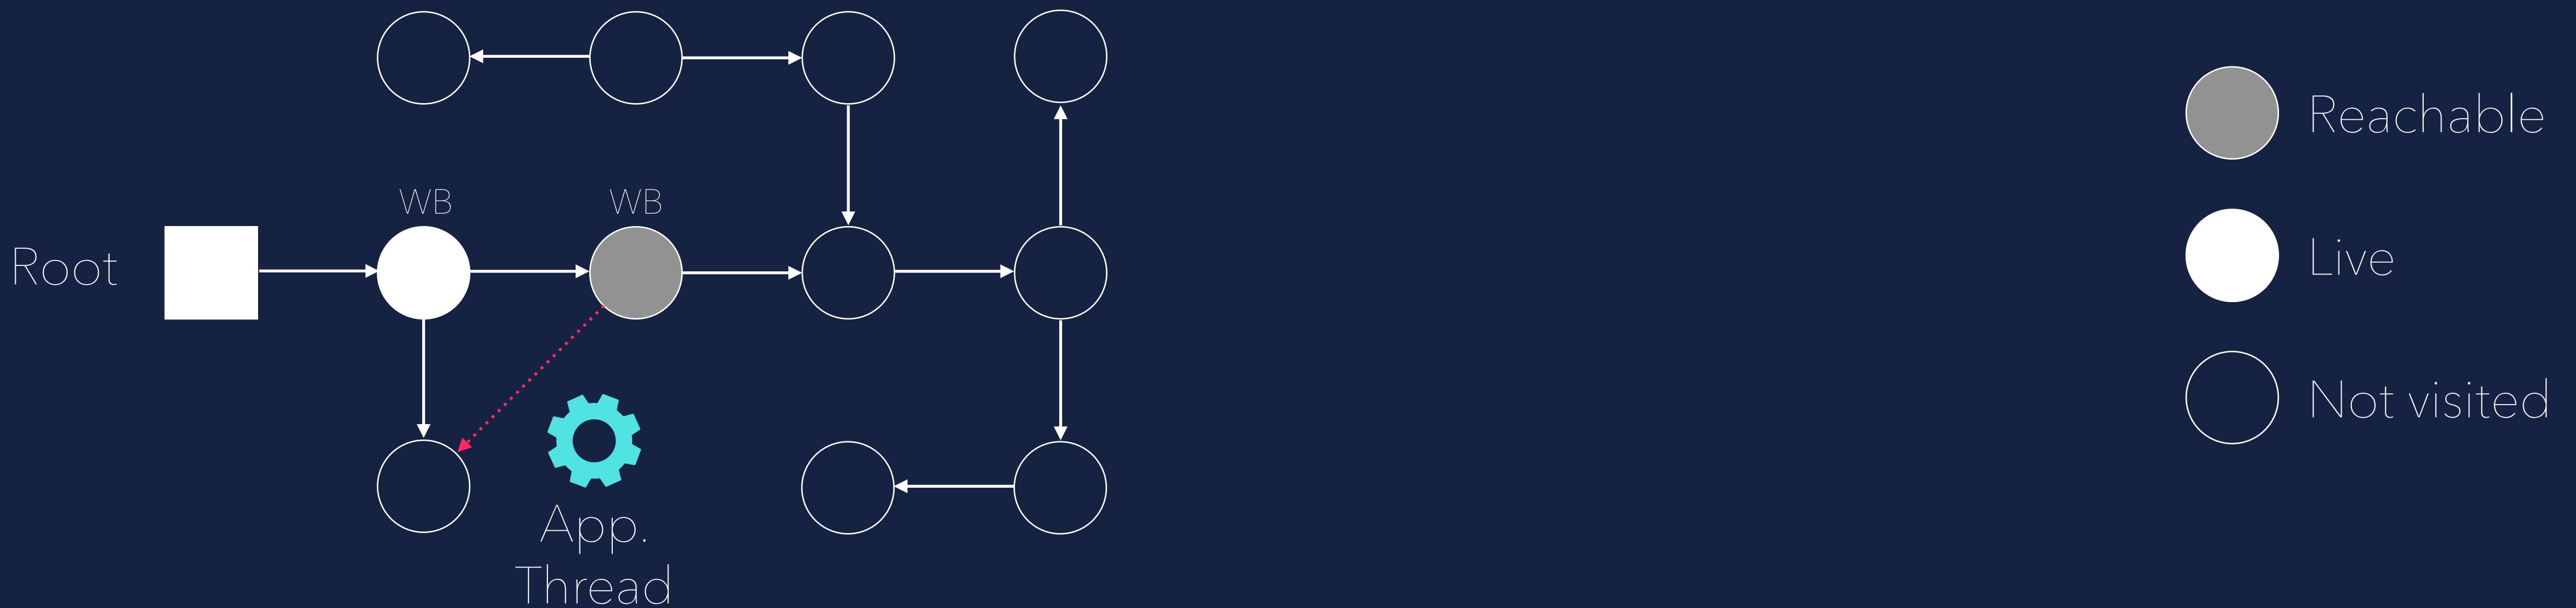
Concurrent Marking using Write Barriers



Mutator hits write barrier and removes reference and adds a new one

CONCURRENCY IS HARD...

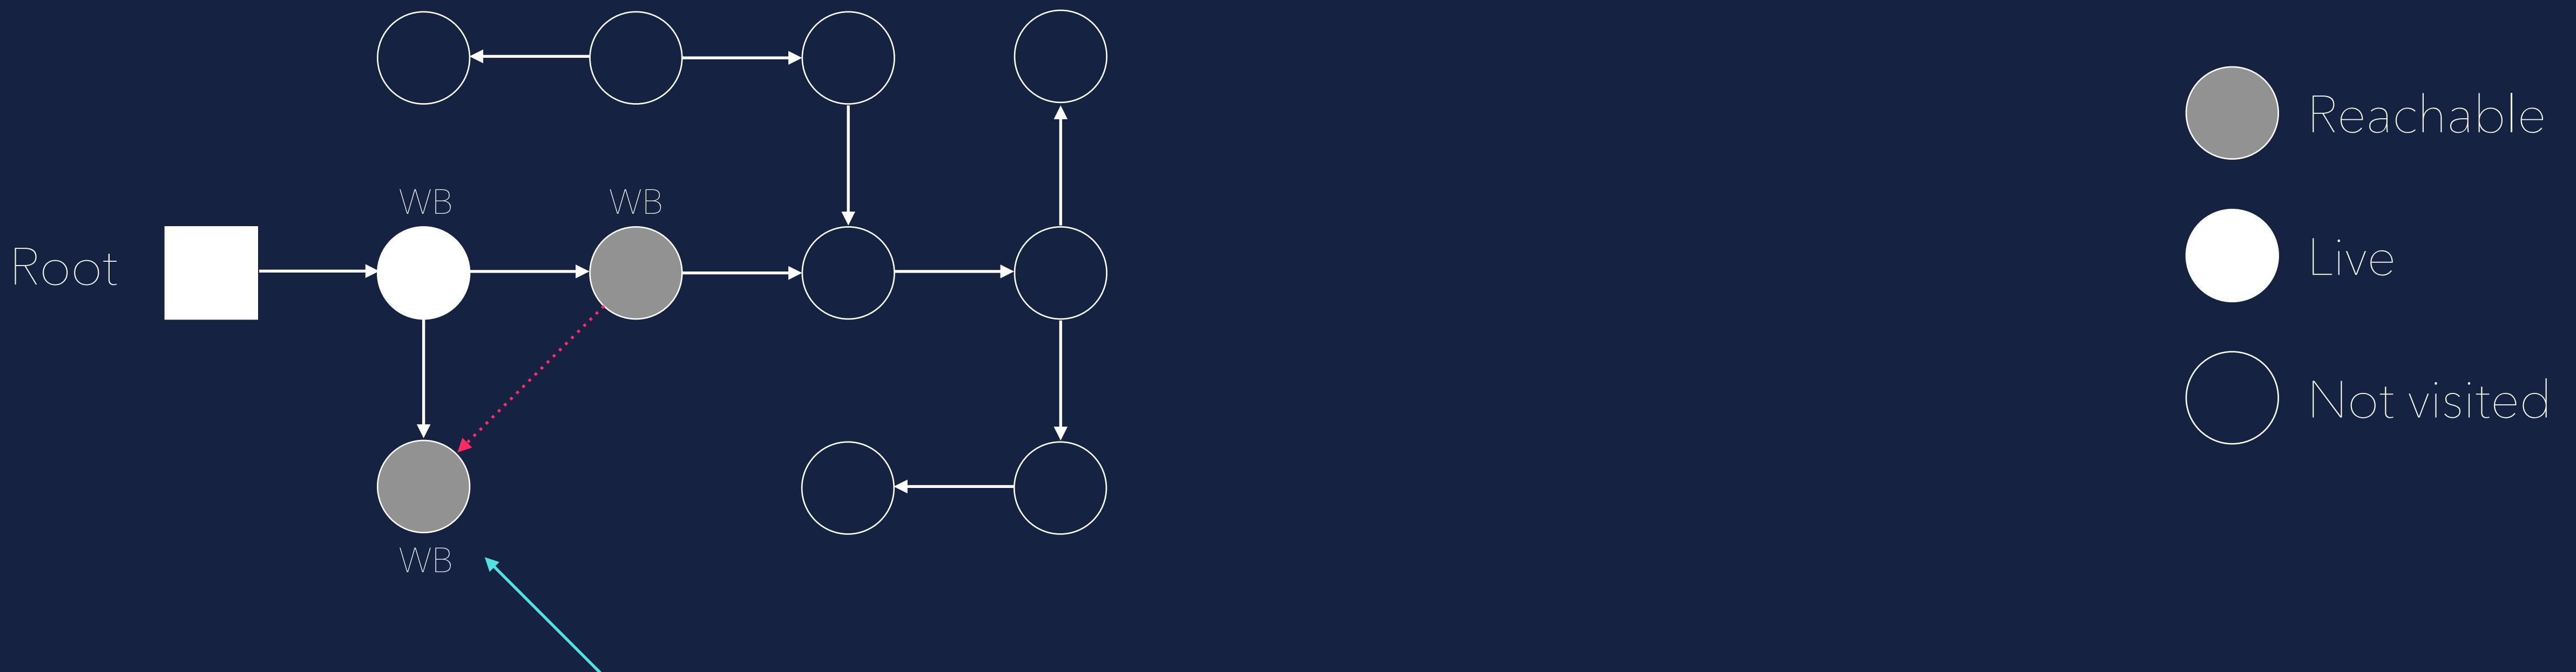
Concurrent Marking using Write Barriers



Mutator hits write barrier and removes reference and adds a new one

CONCURRENCY IS HARD...

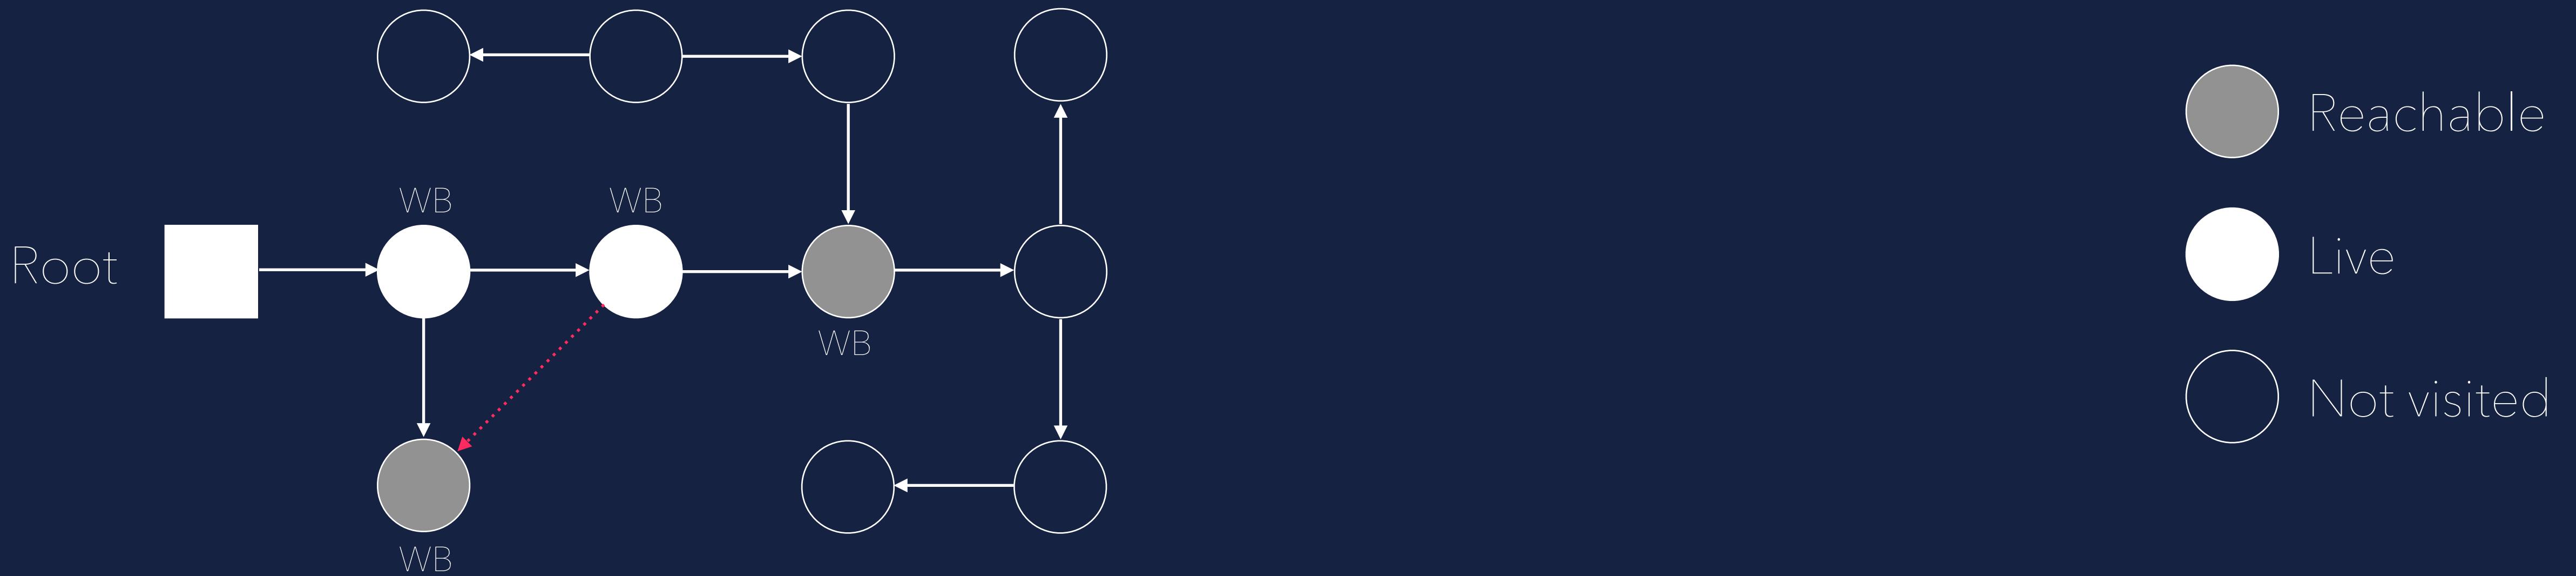
Concurrent Marking using Write Barriers



Removed references will be marked as reachable by Write Barrier

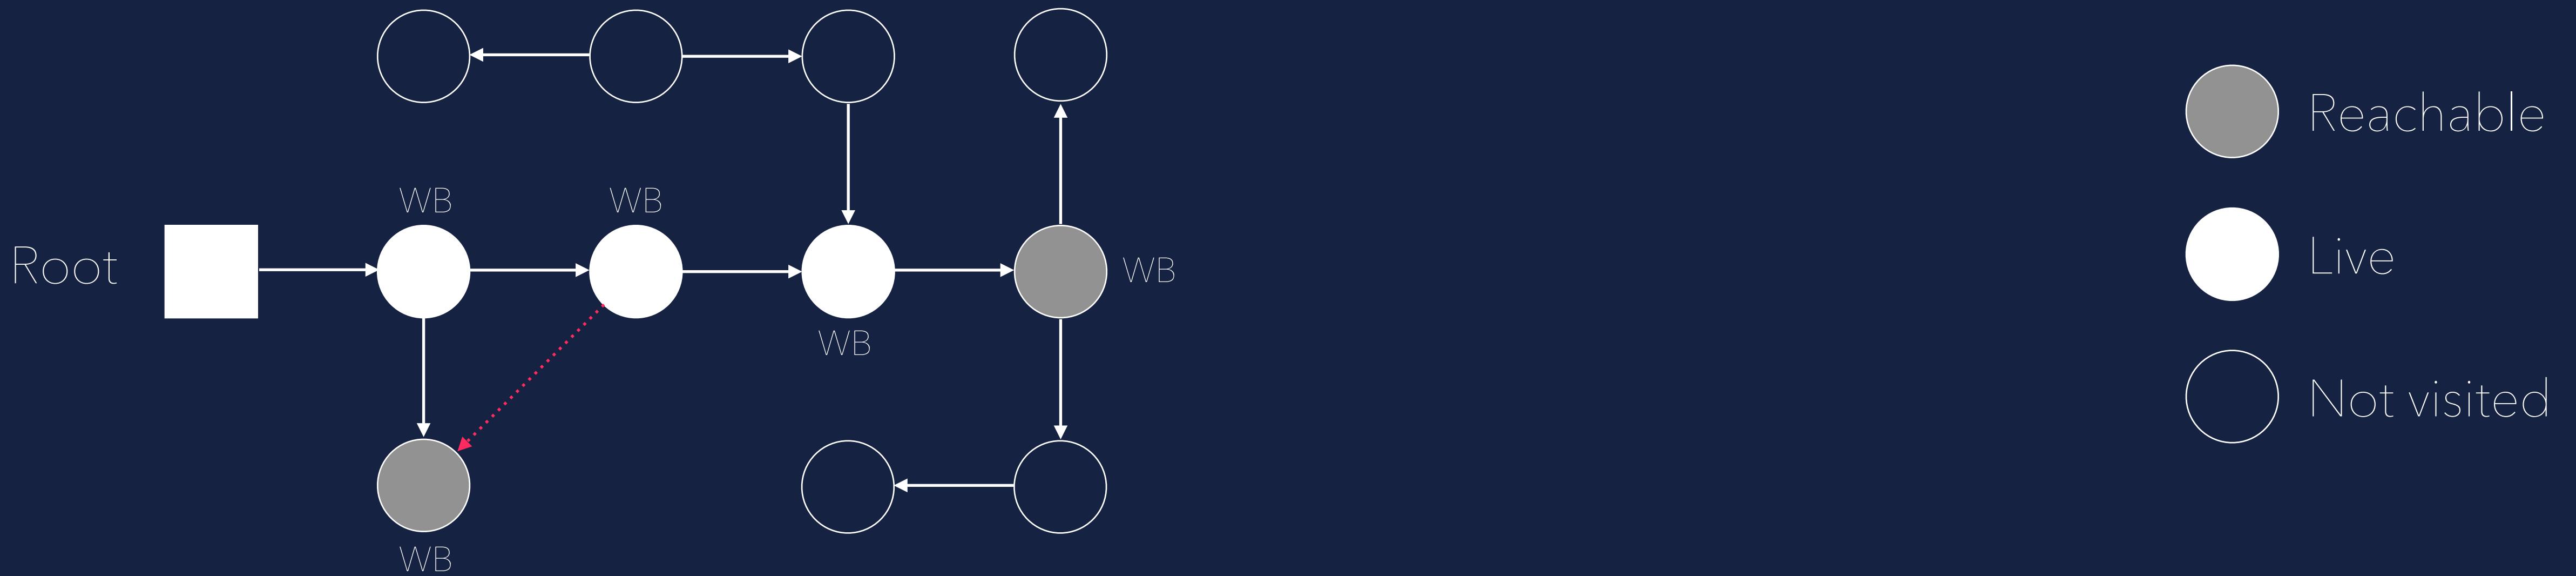
CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers



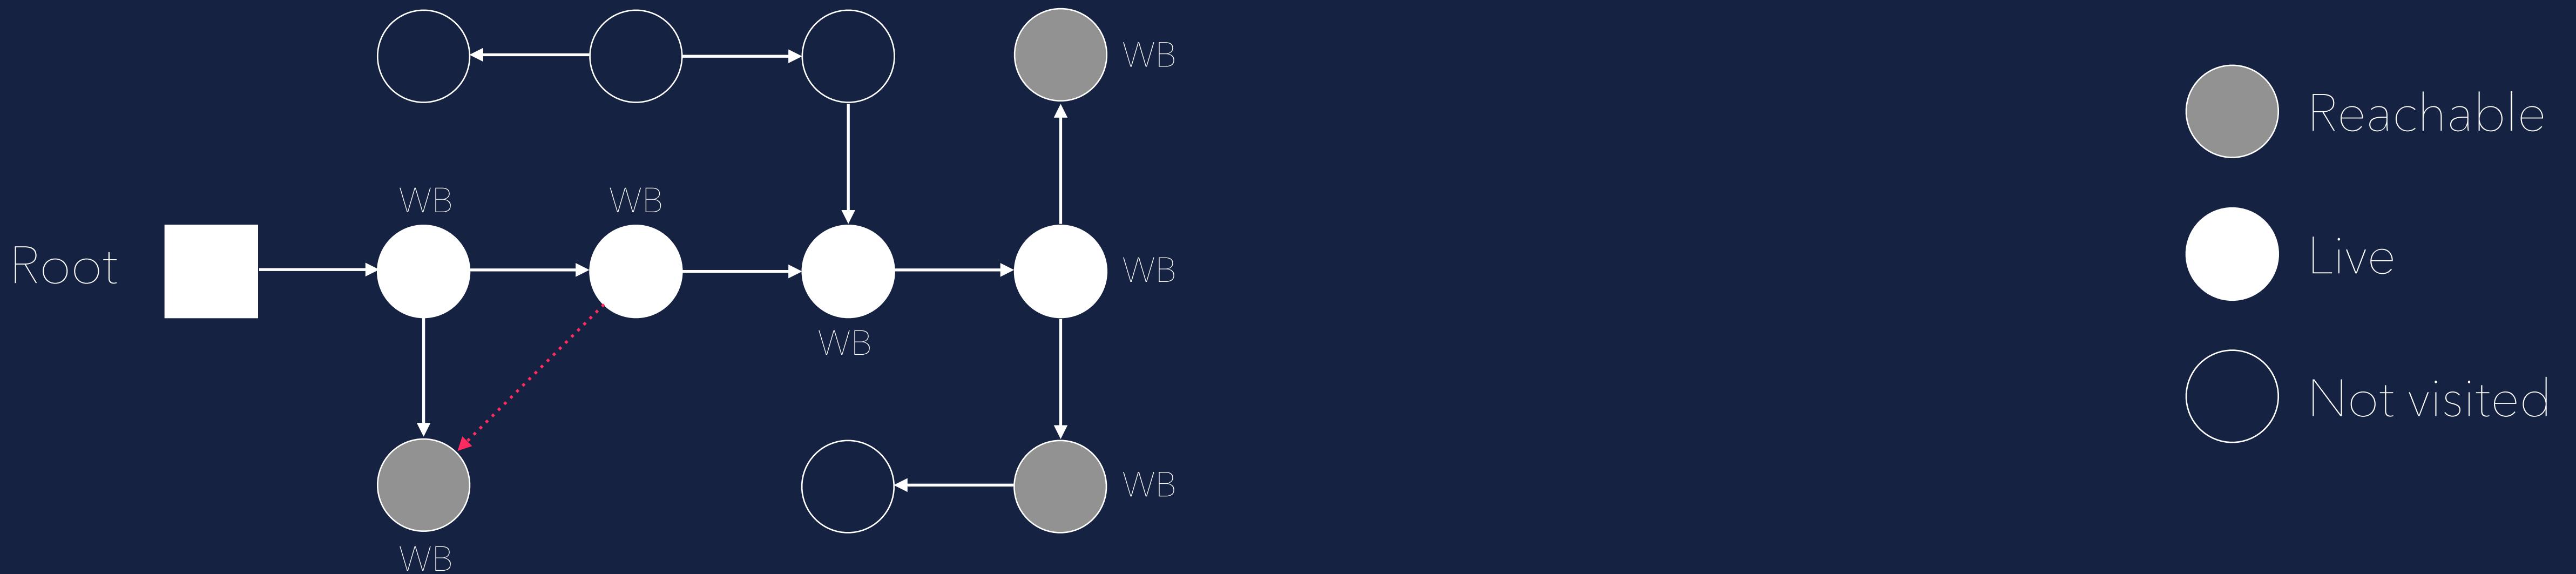
CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers



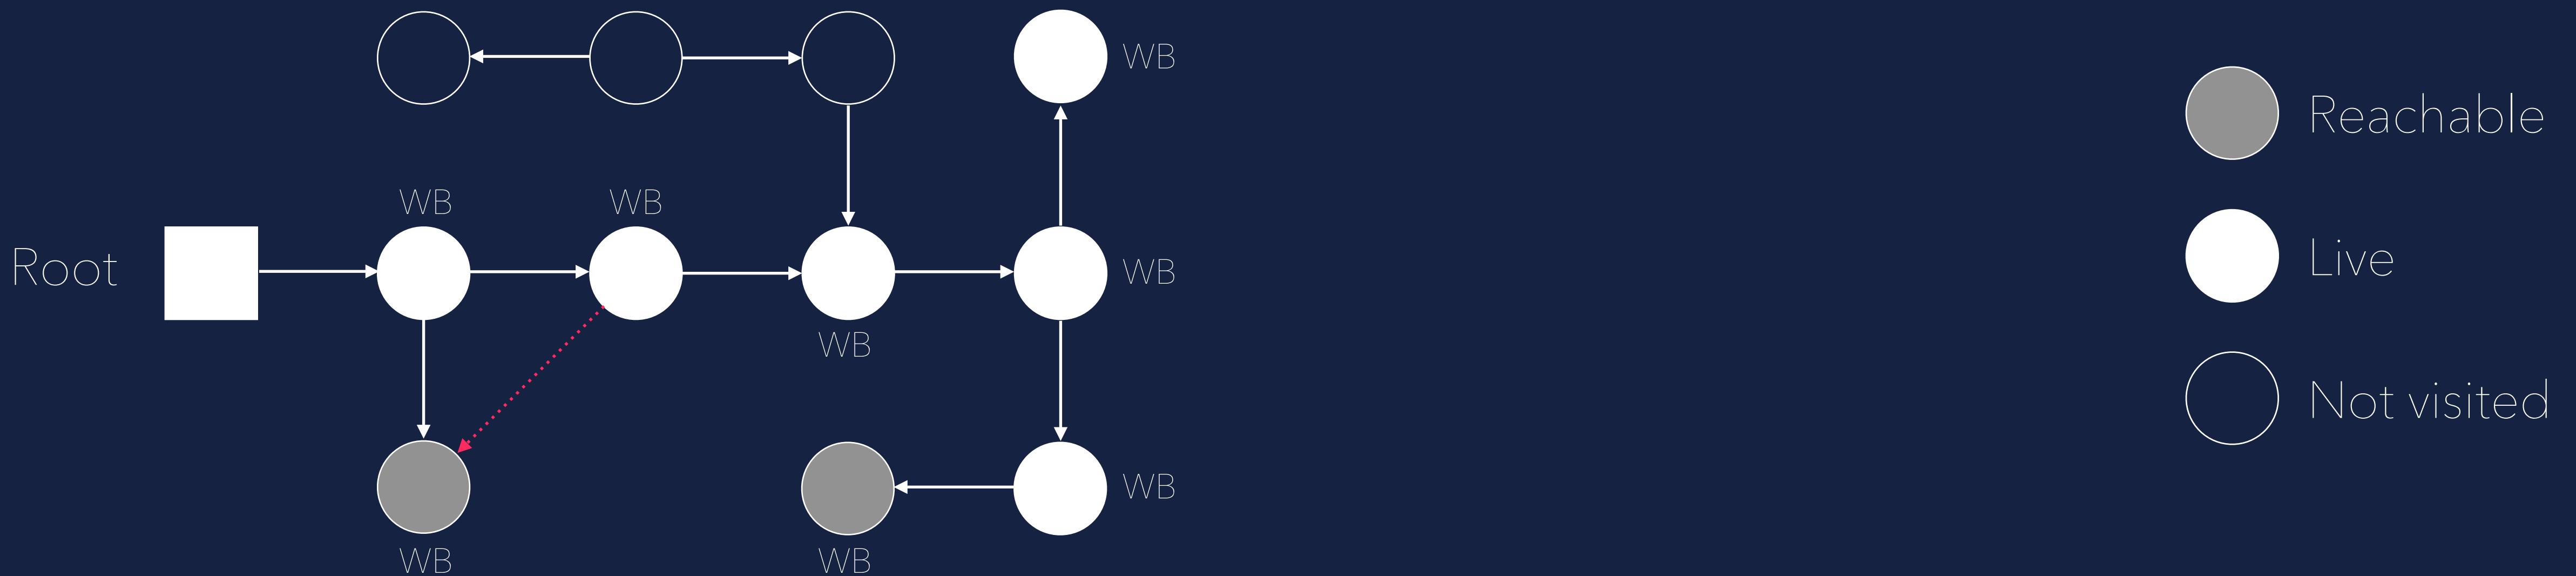
CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers



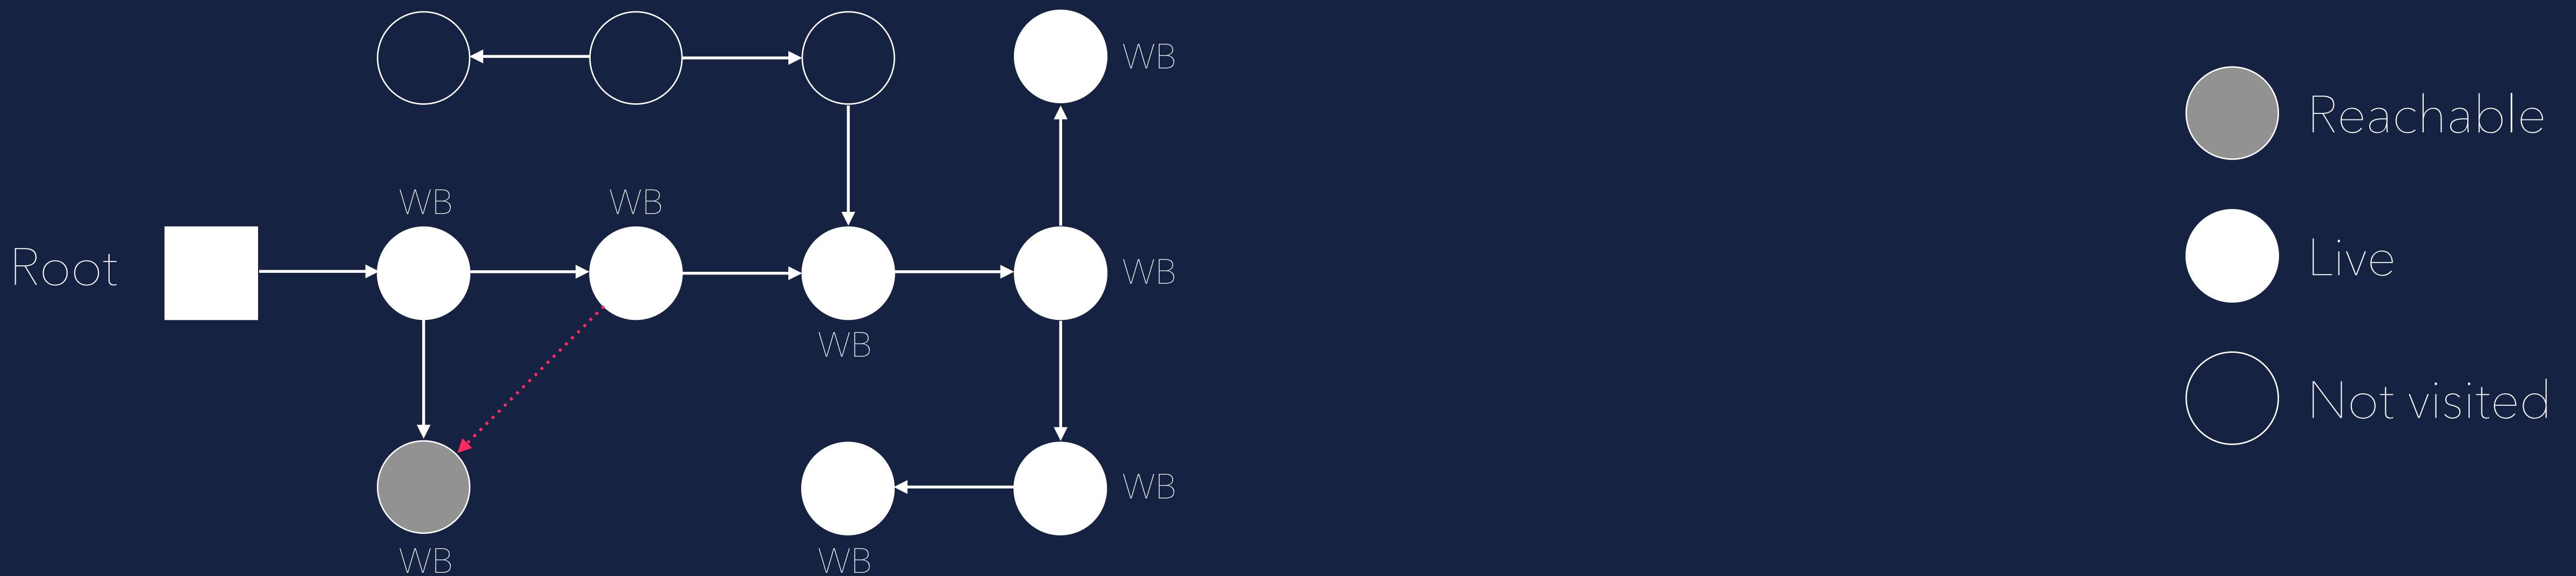
CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers



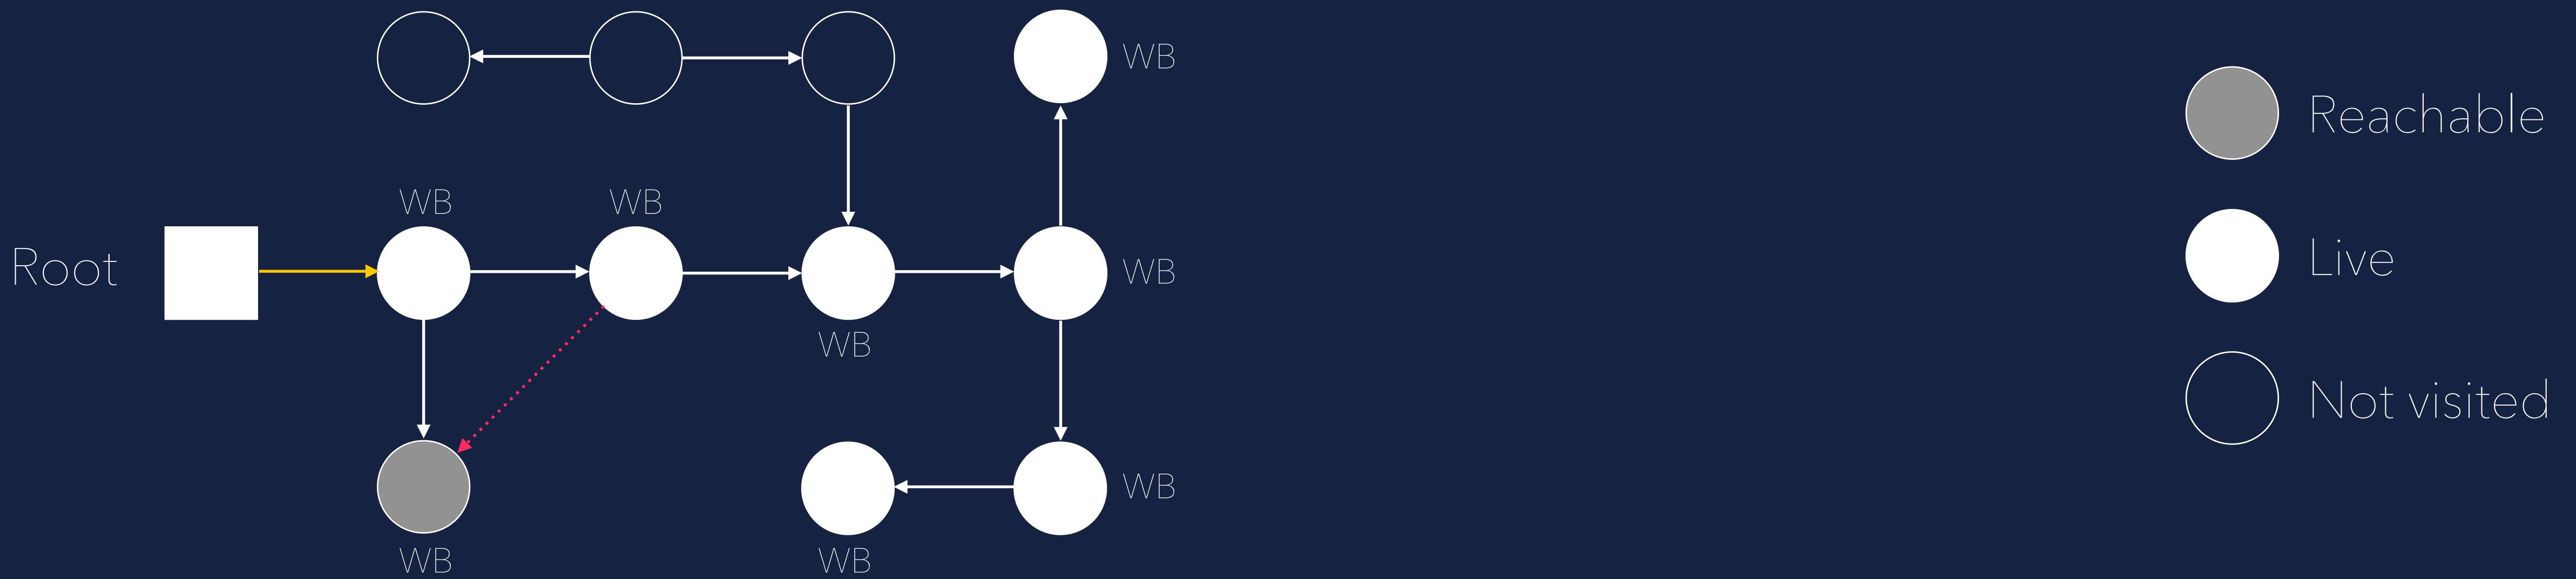
CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers



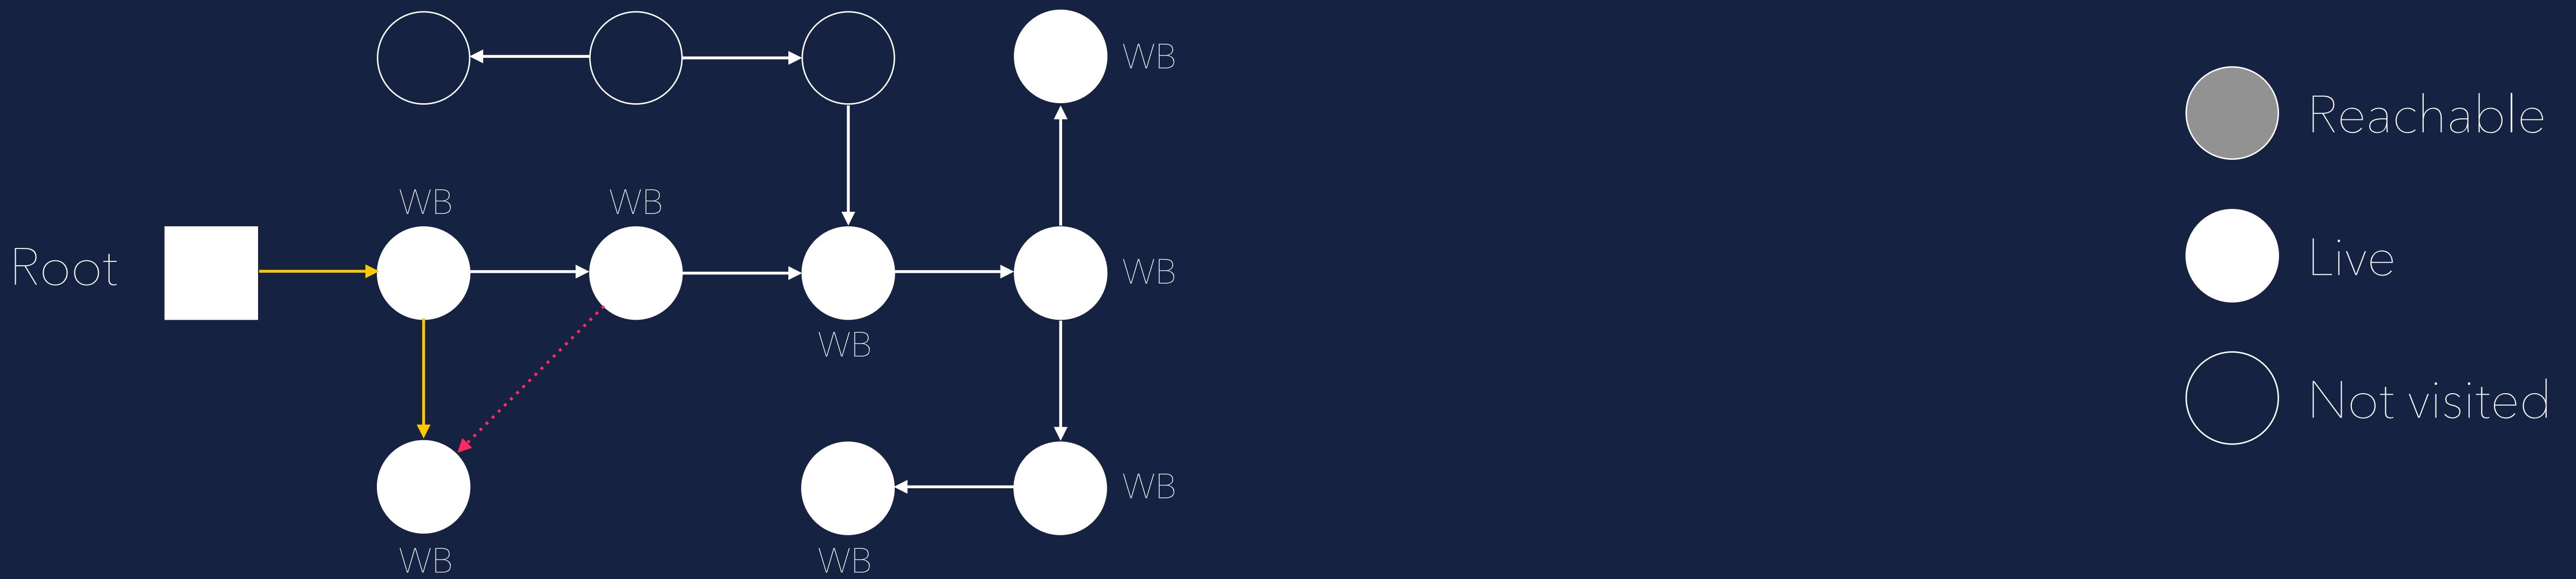
CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers



CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers

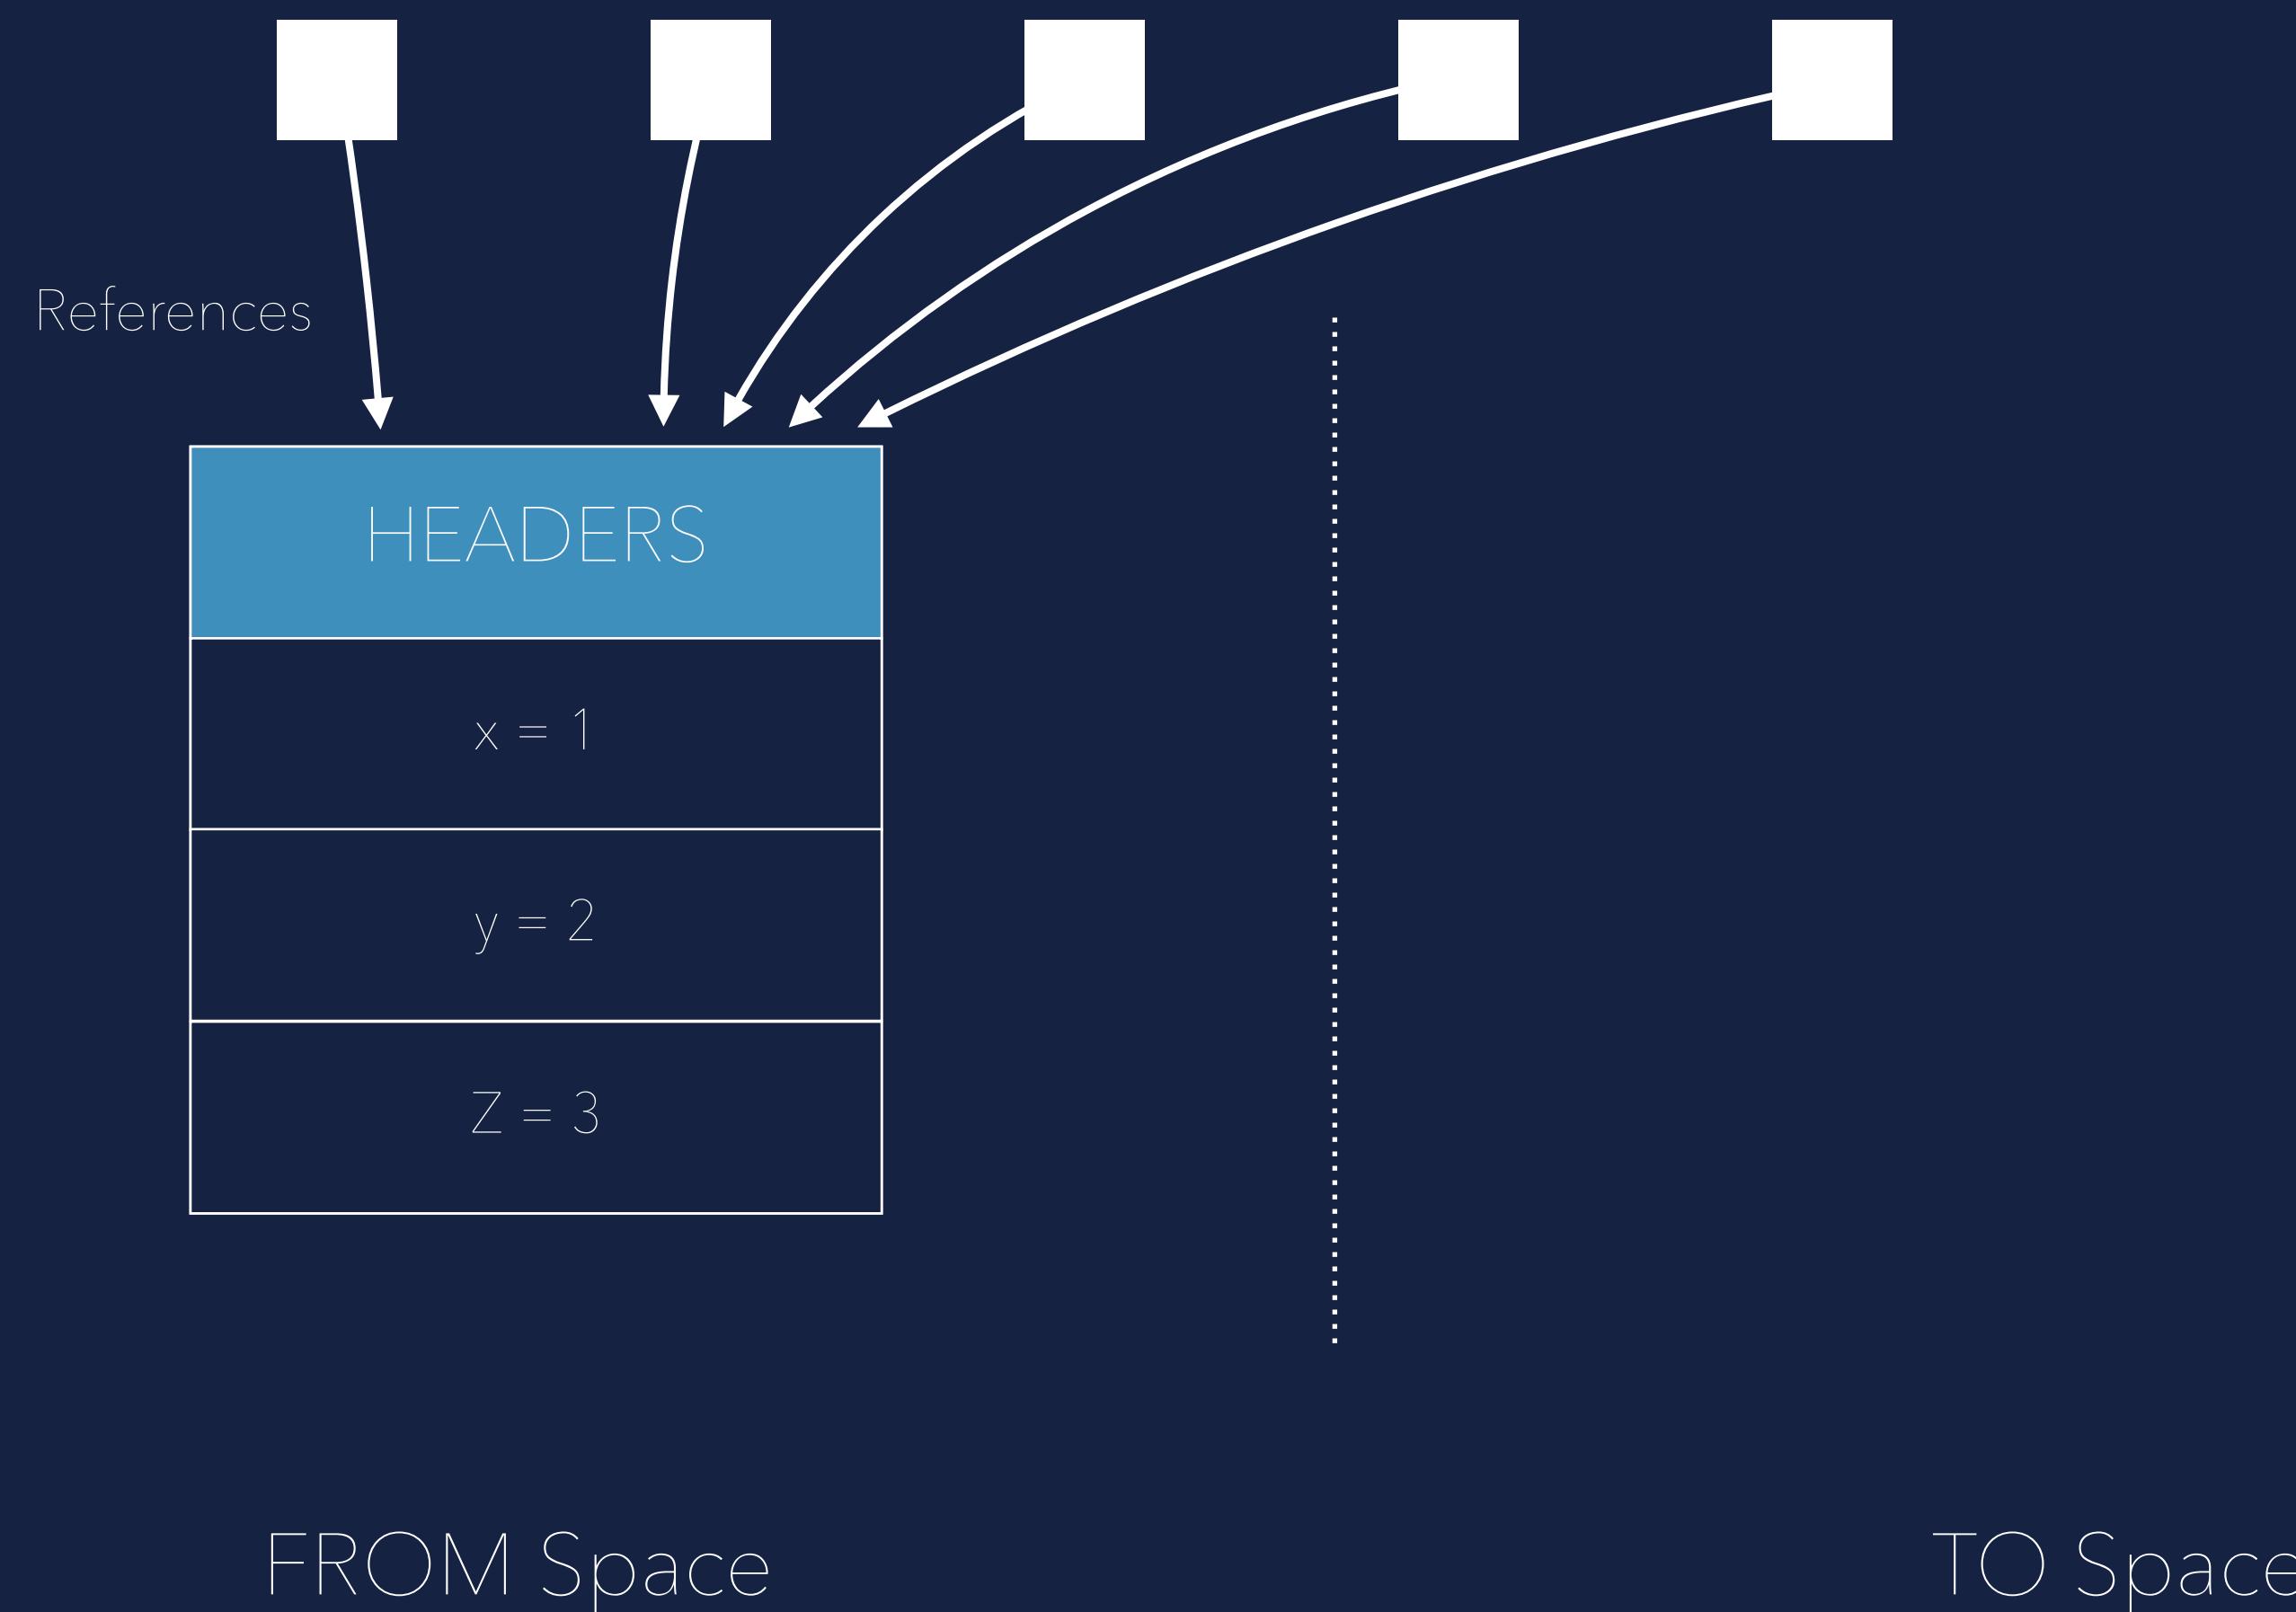


In the Re-Marking phase, in between marked references will be marked as live

**CONCURRENT
COPYING**

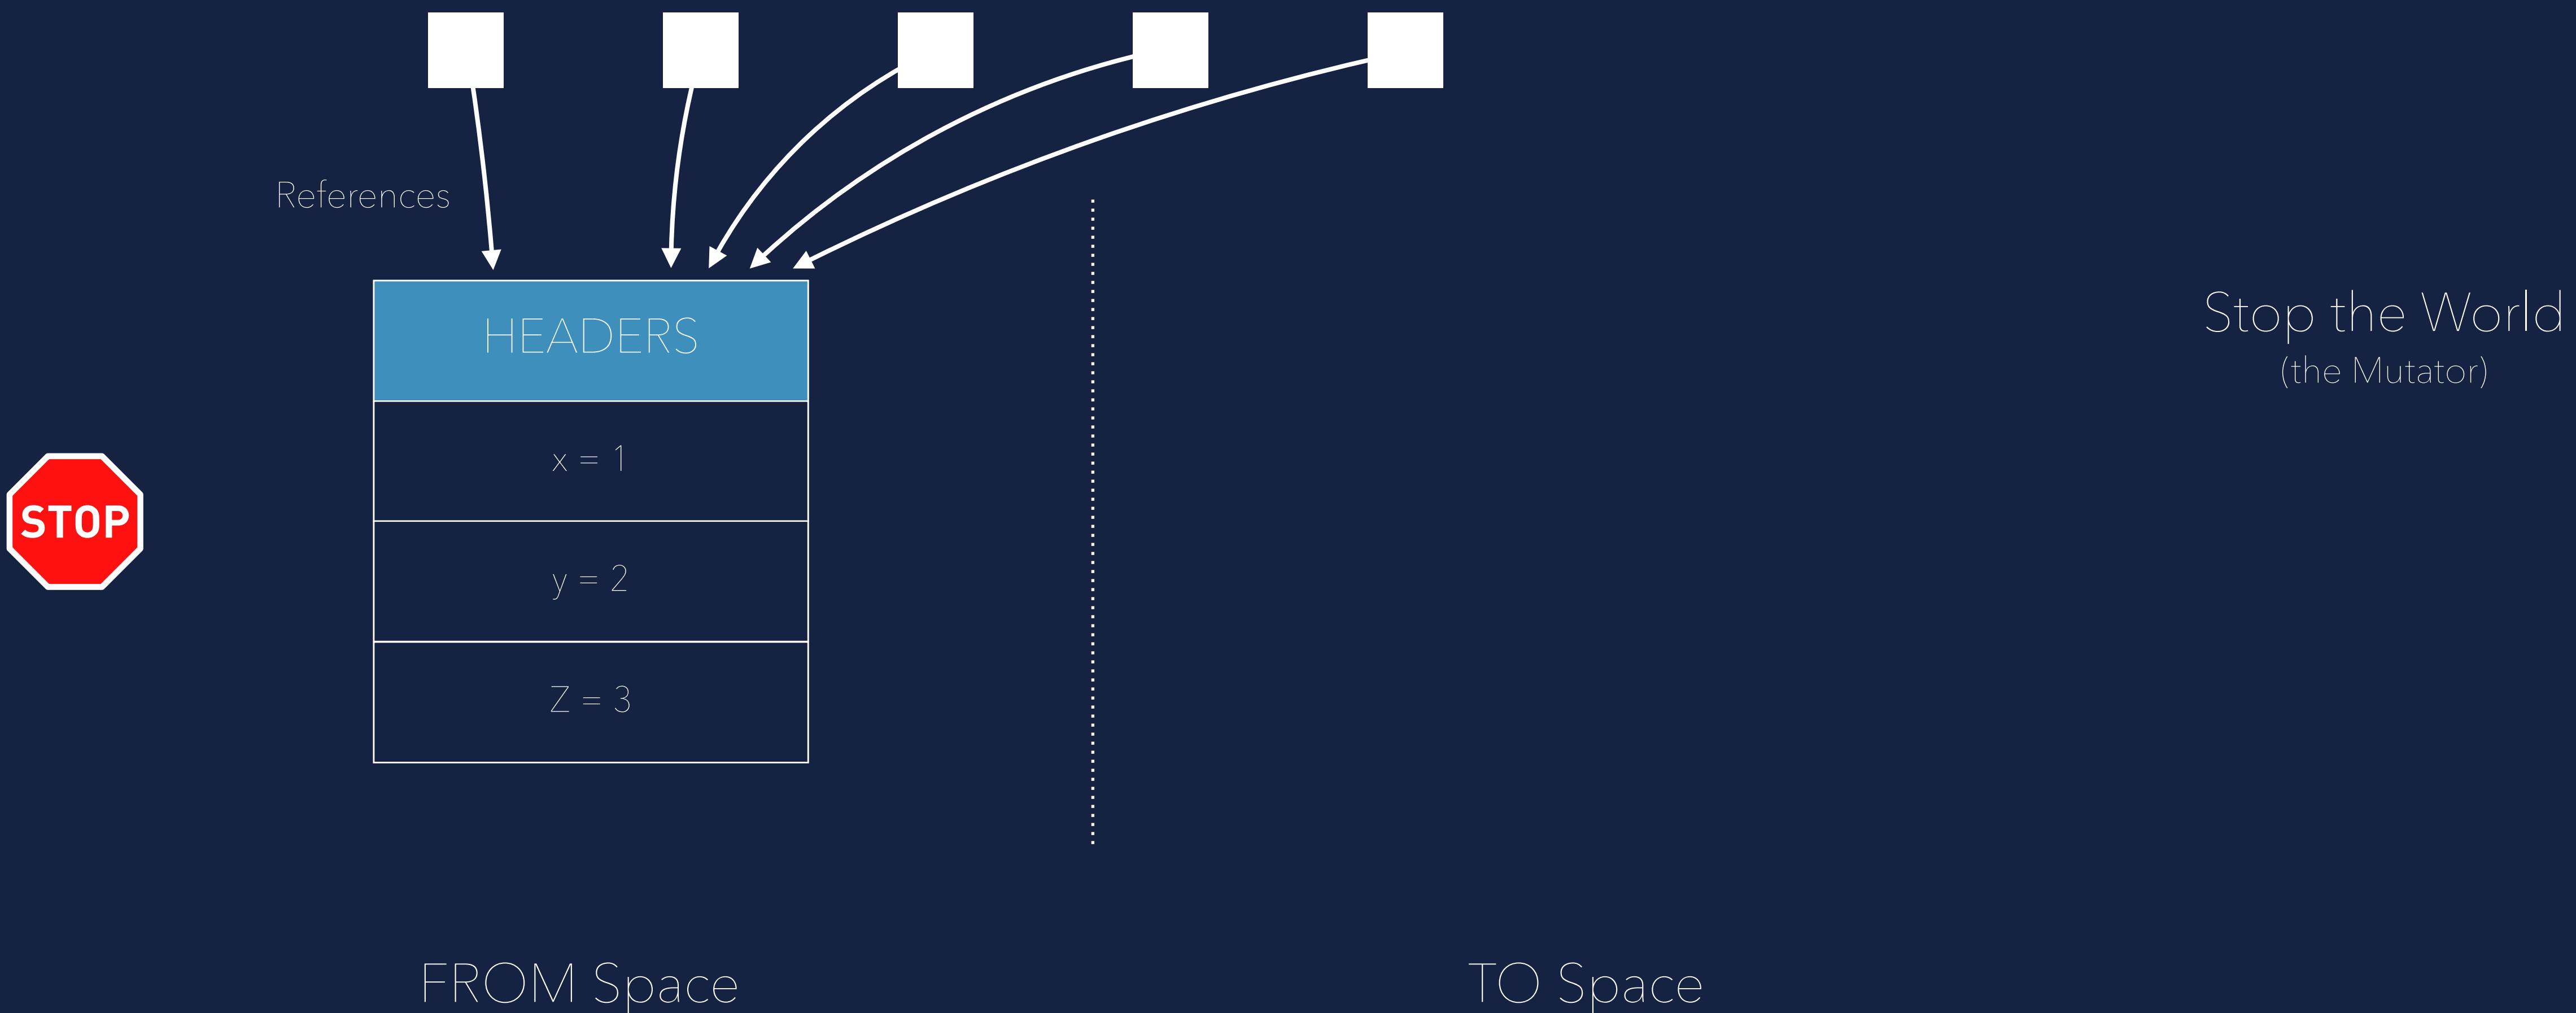
CONCURRENCY IS HARD...

Stop the world copying



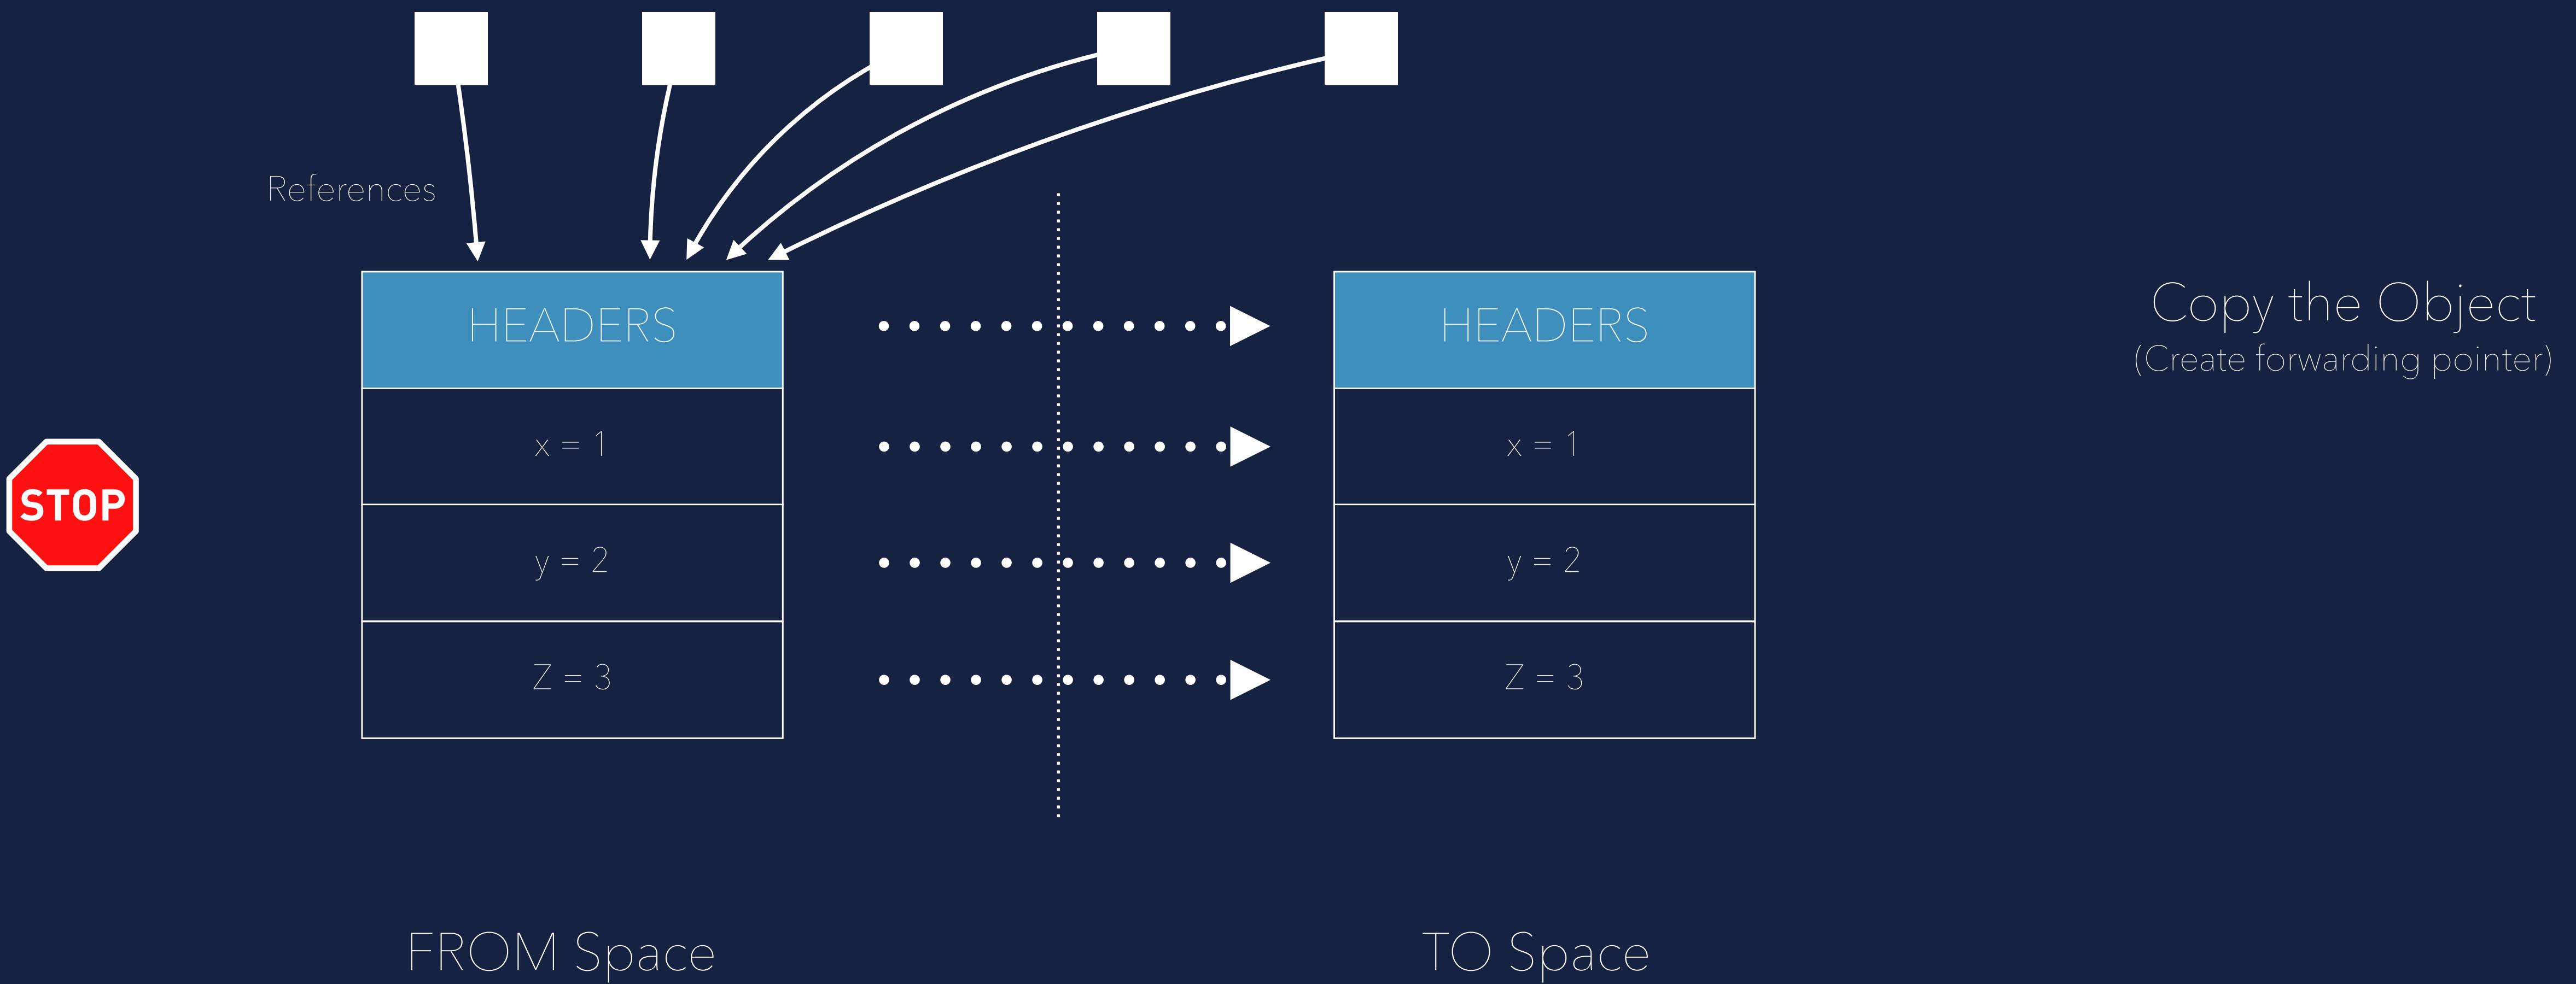
CONCURRENCY IS HARD...

Stop the world copying



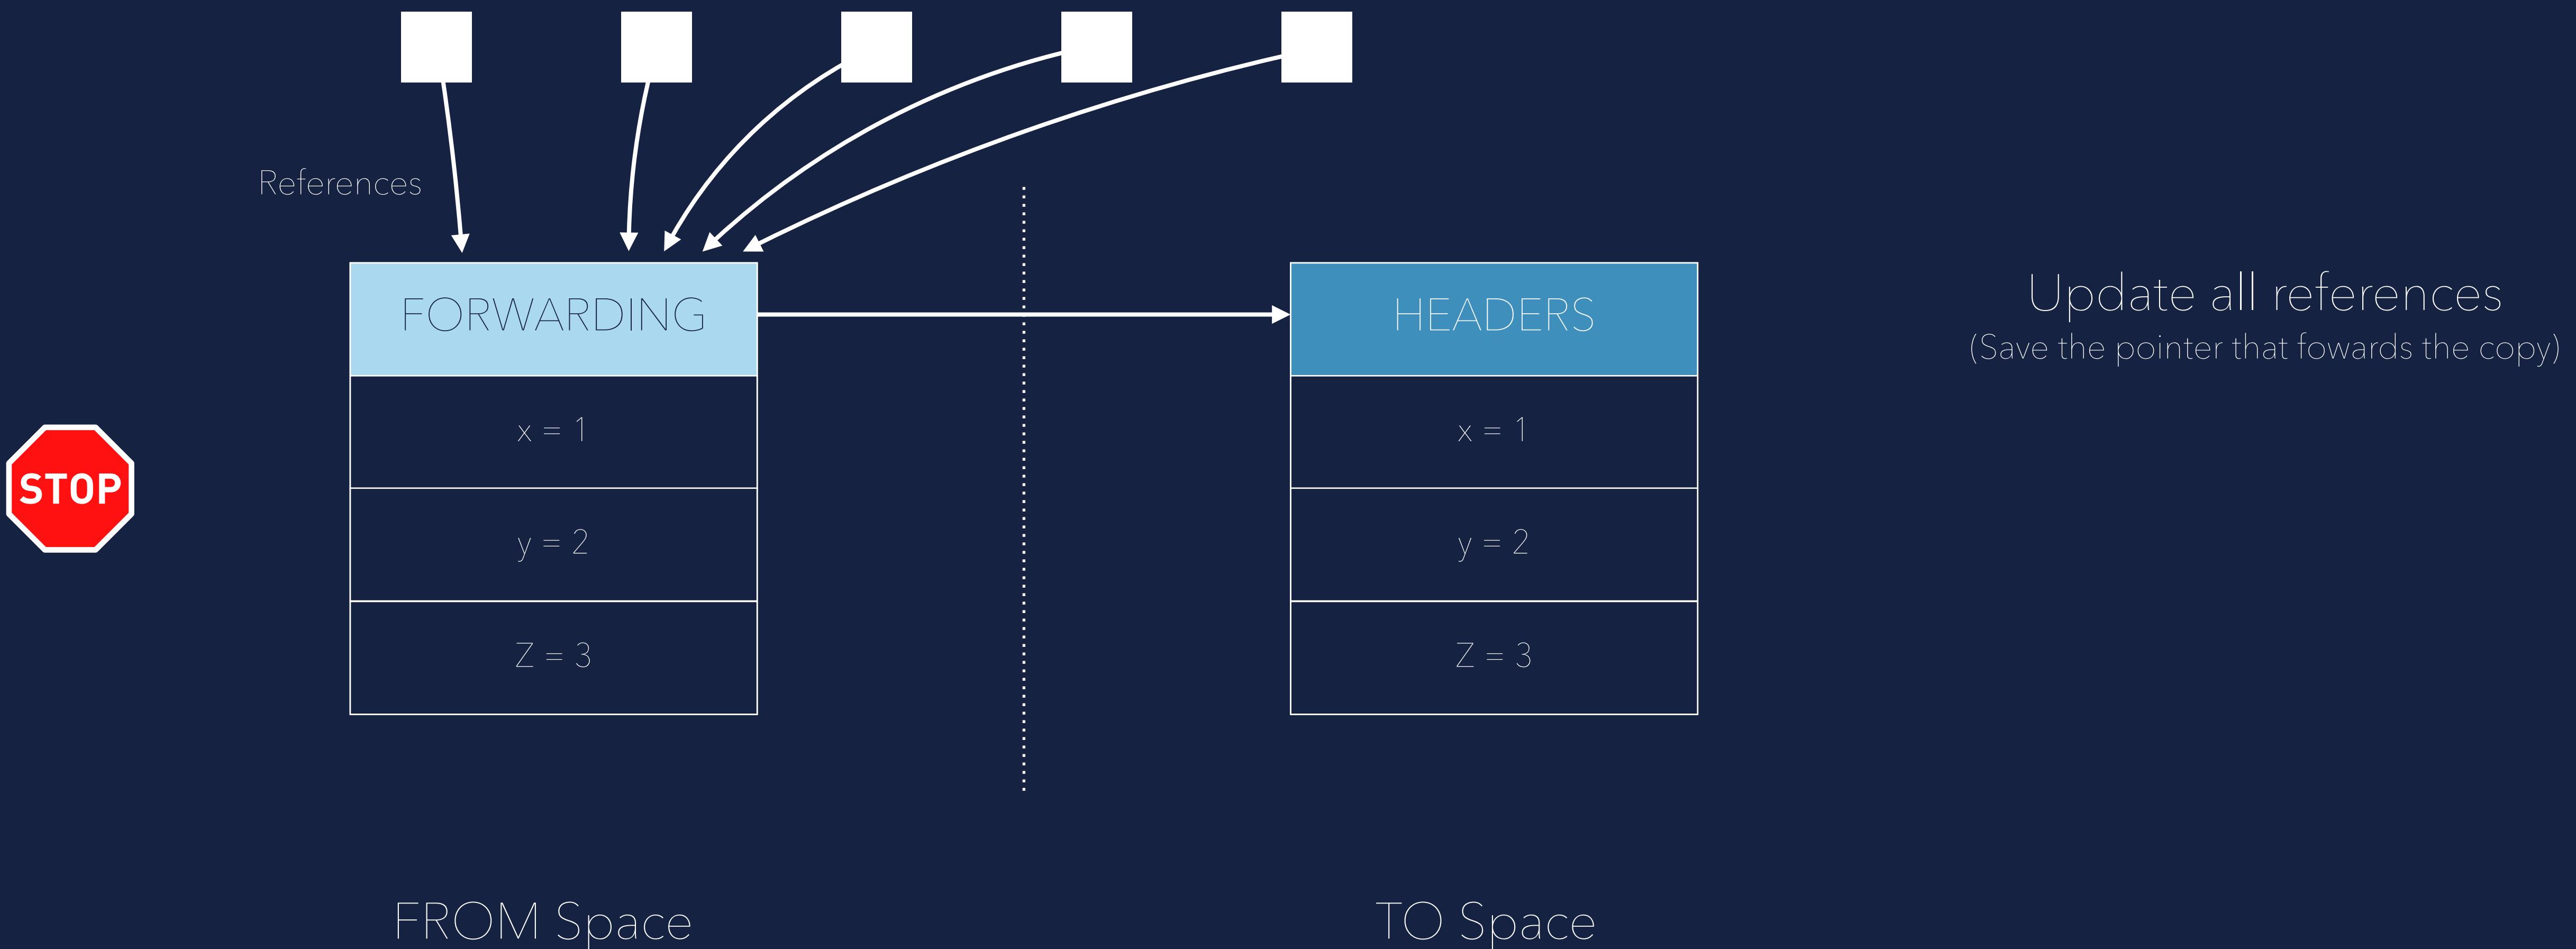
CONCURRENCY IS HARD...

Stop the world copying



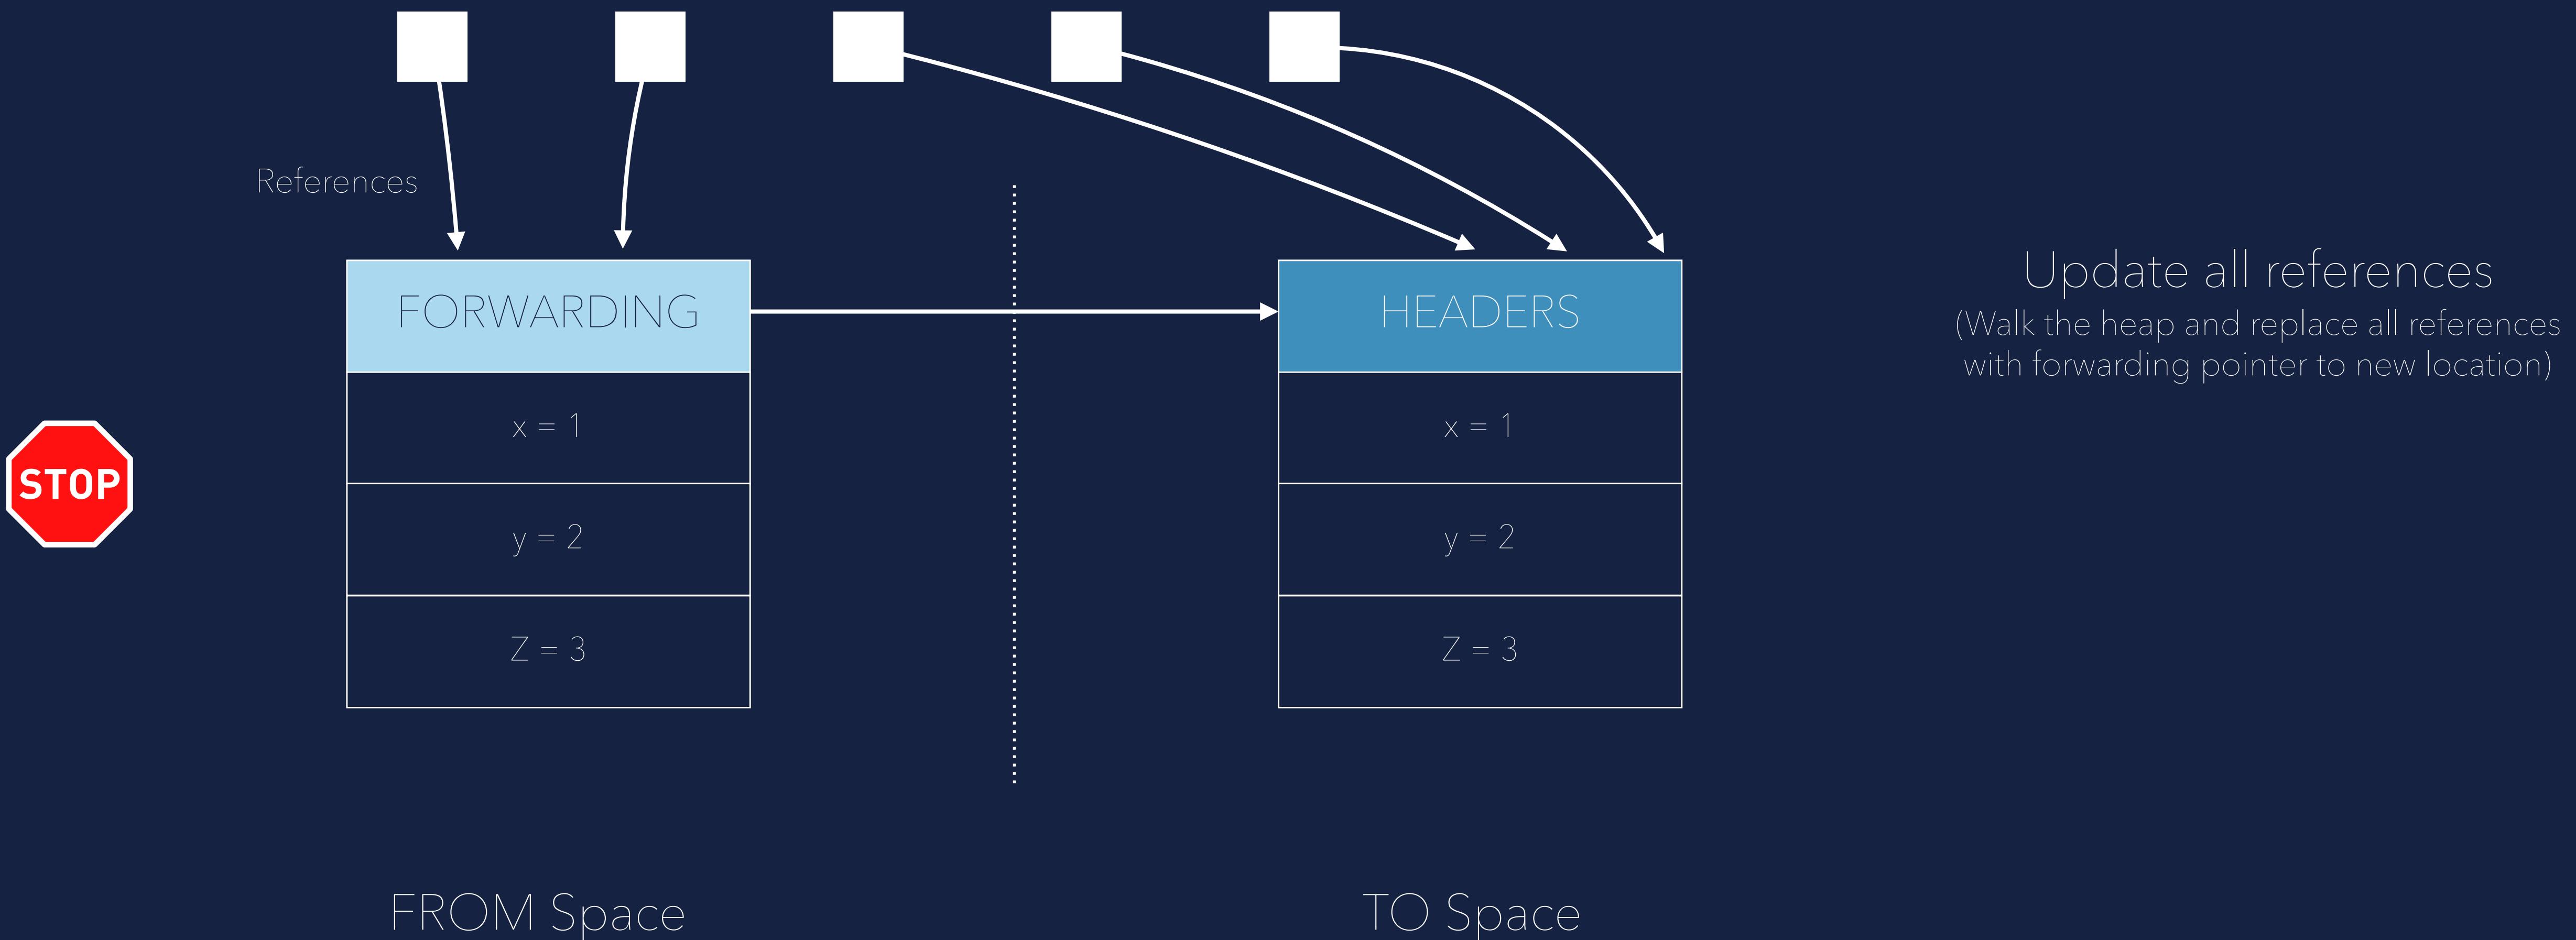
CONCURRENCY IS HARD...

Stop the world copying



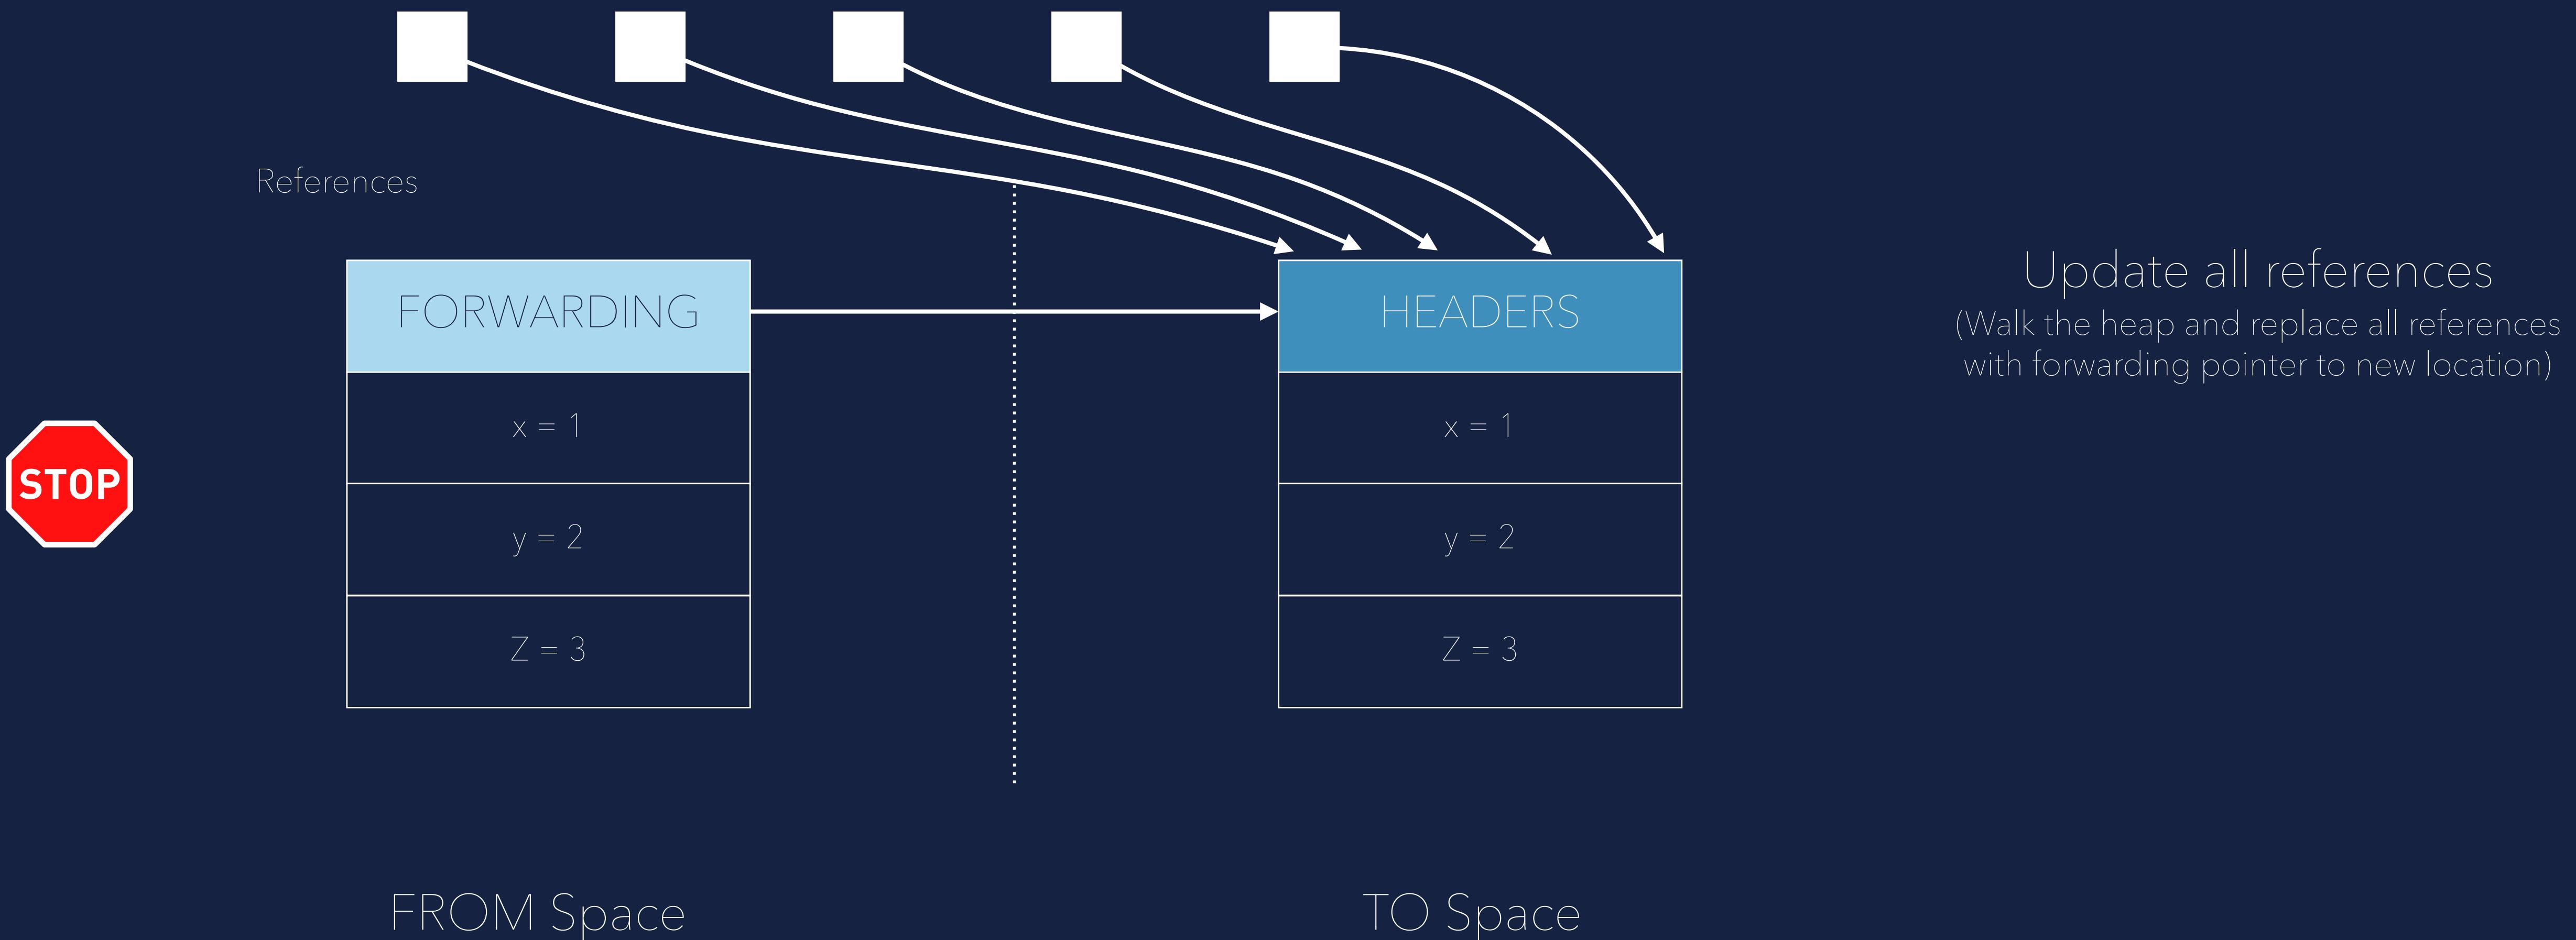
CONCURRENCY IS HARD...

Stop the world copying



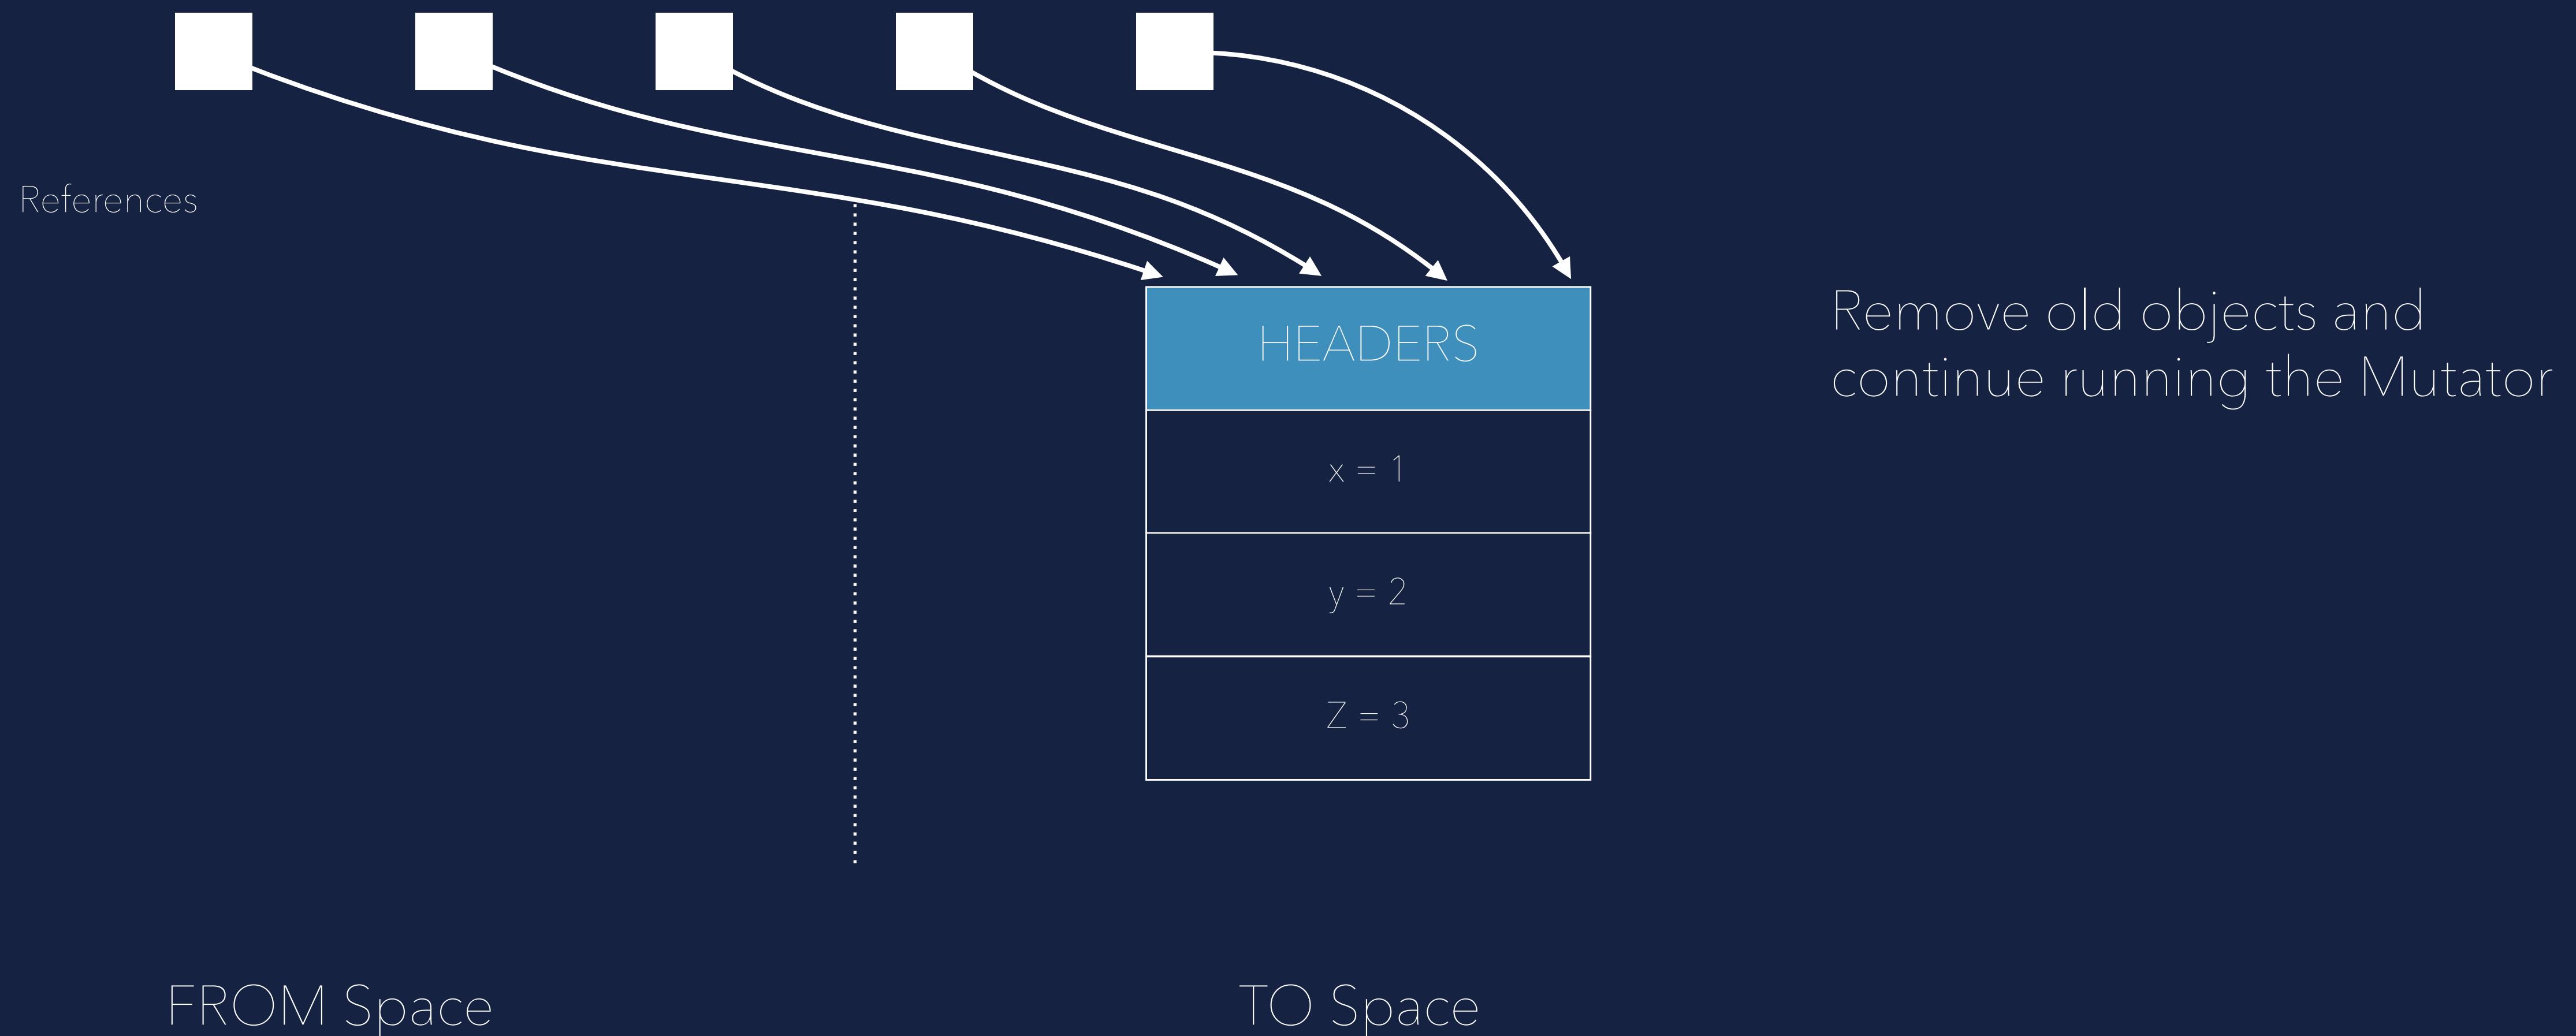
CONCURRENCY IS HARD...

Stop the world copying



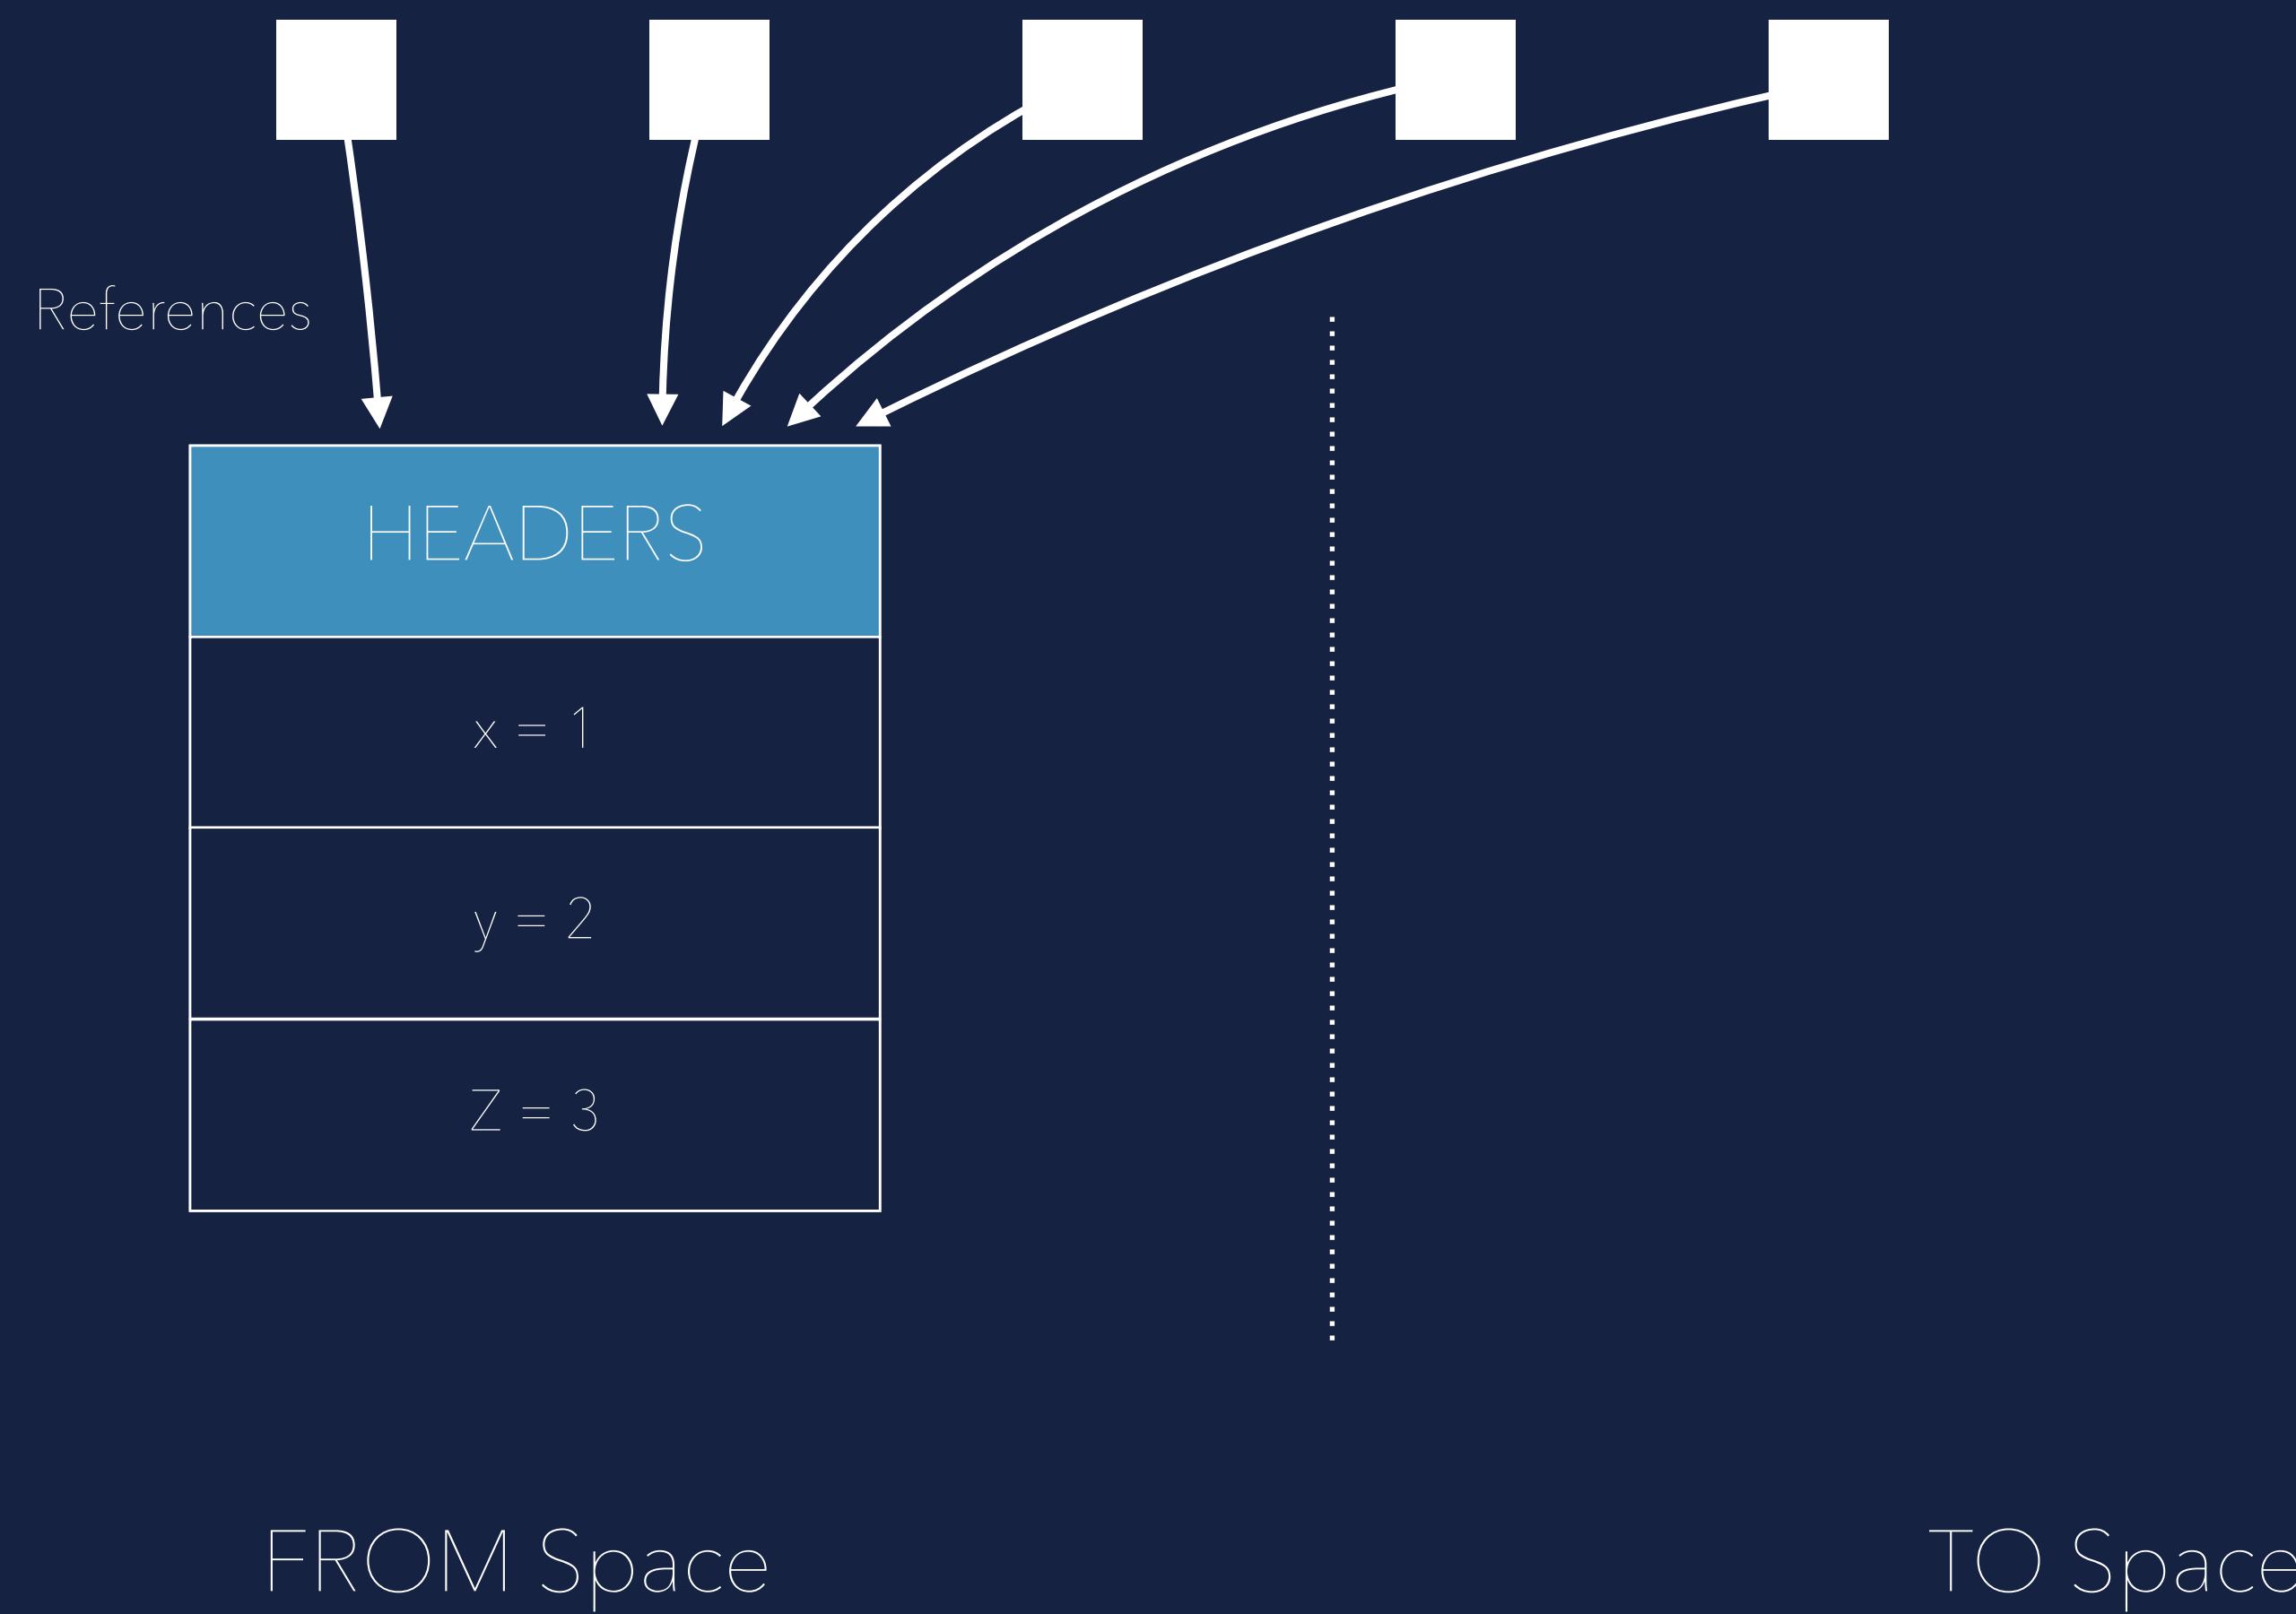
CONCURRENCY IS HARD...

Stop the world copying



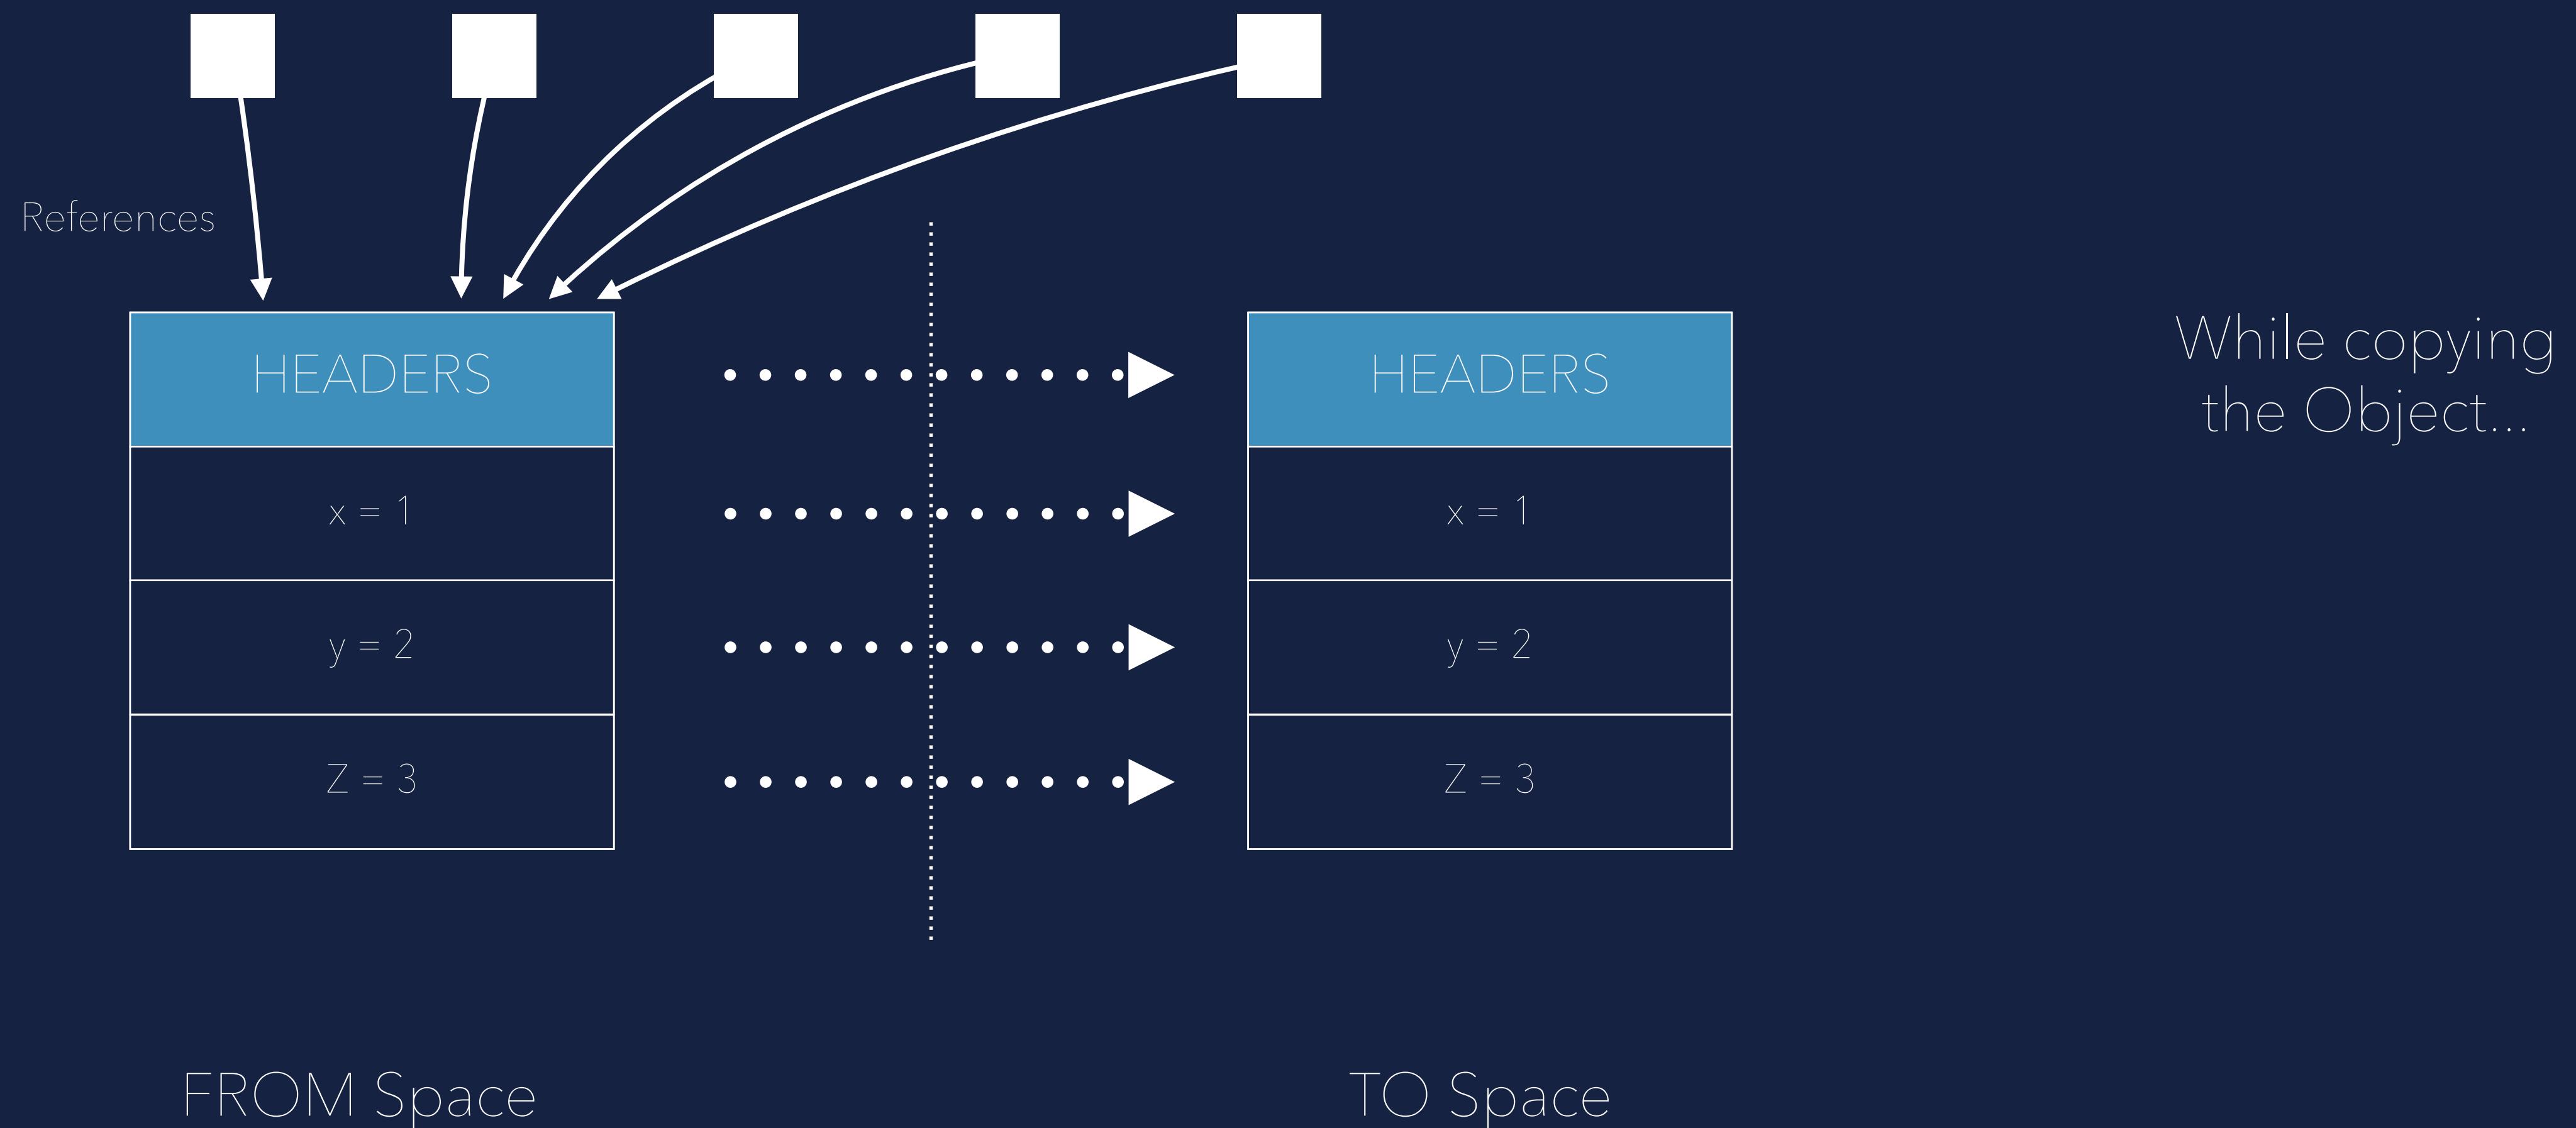
CONCURRENCY IS HARD...

Concurrent copying



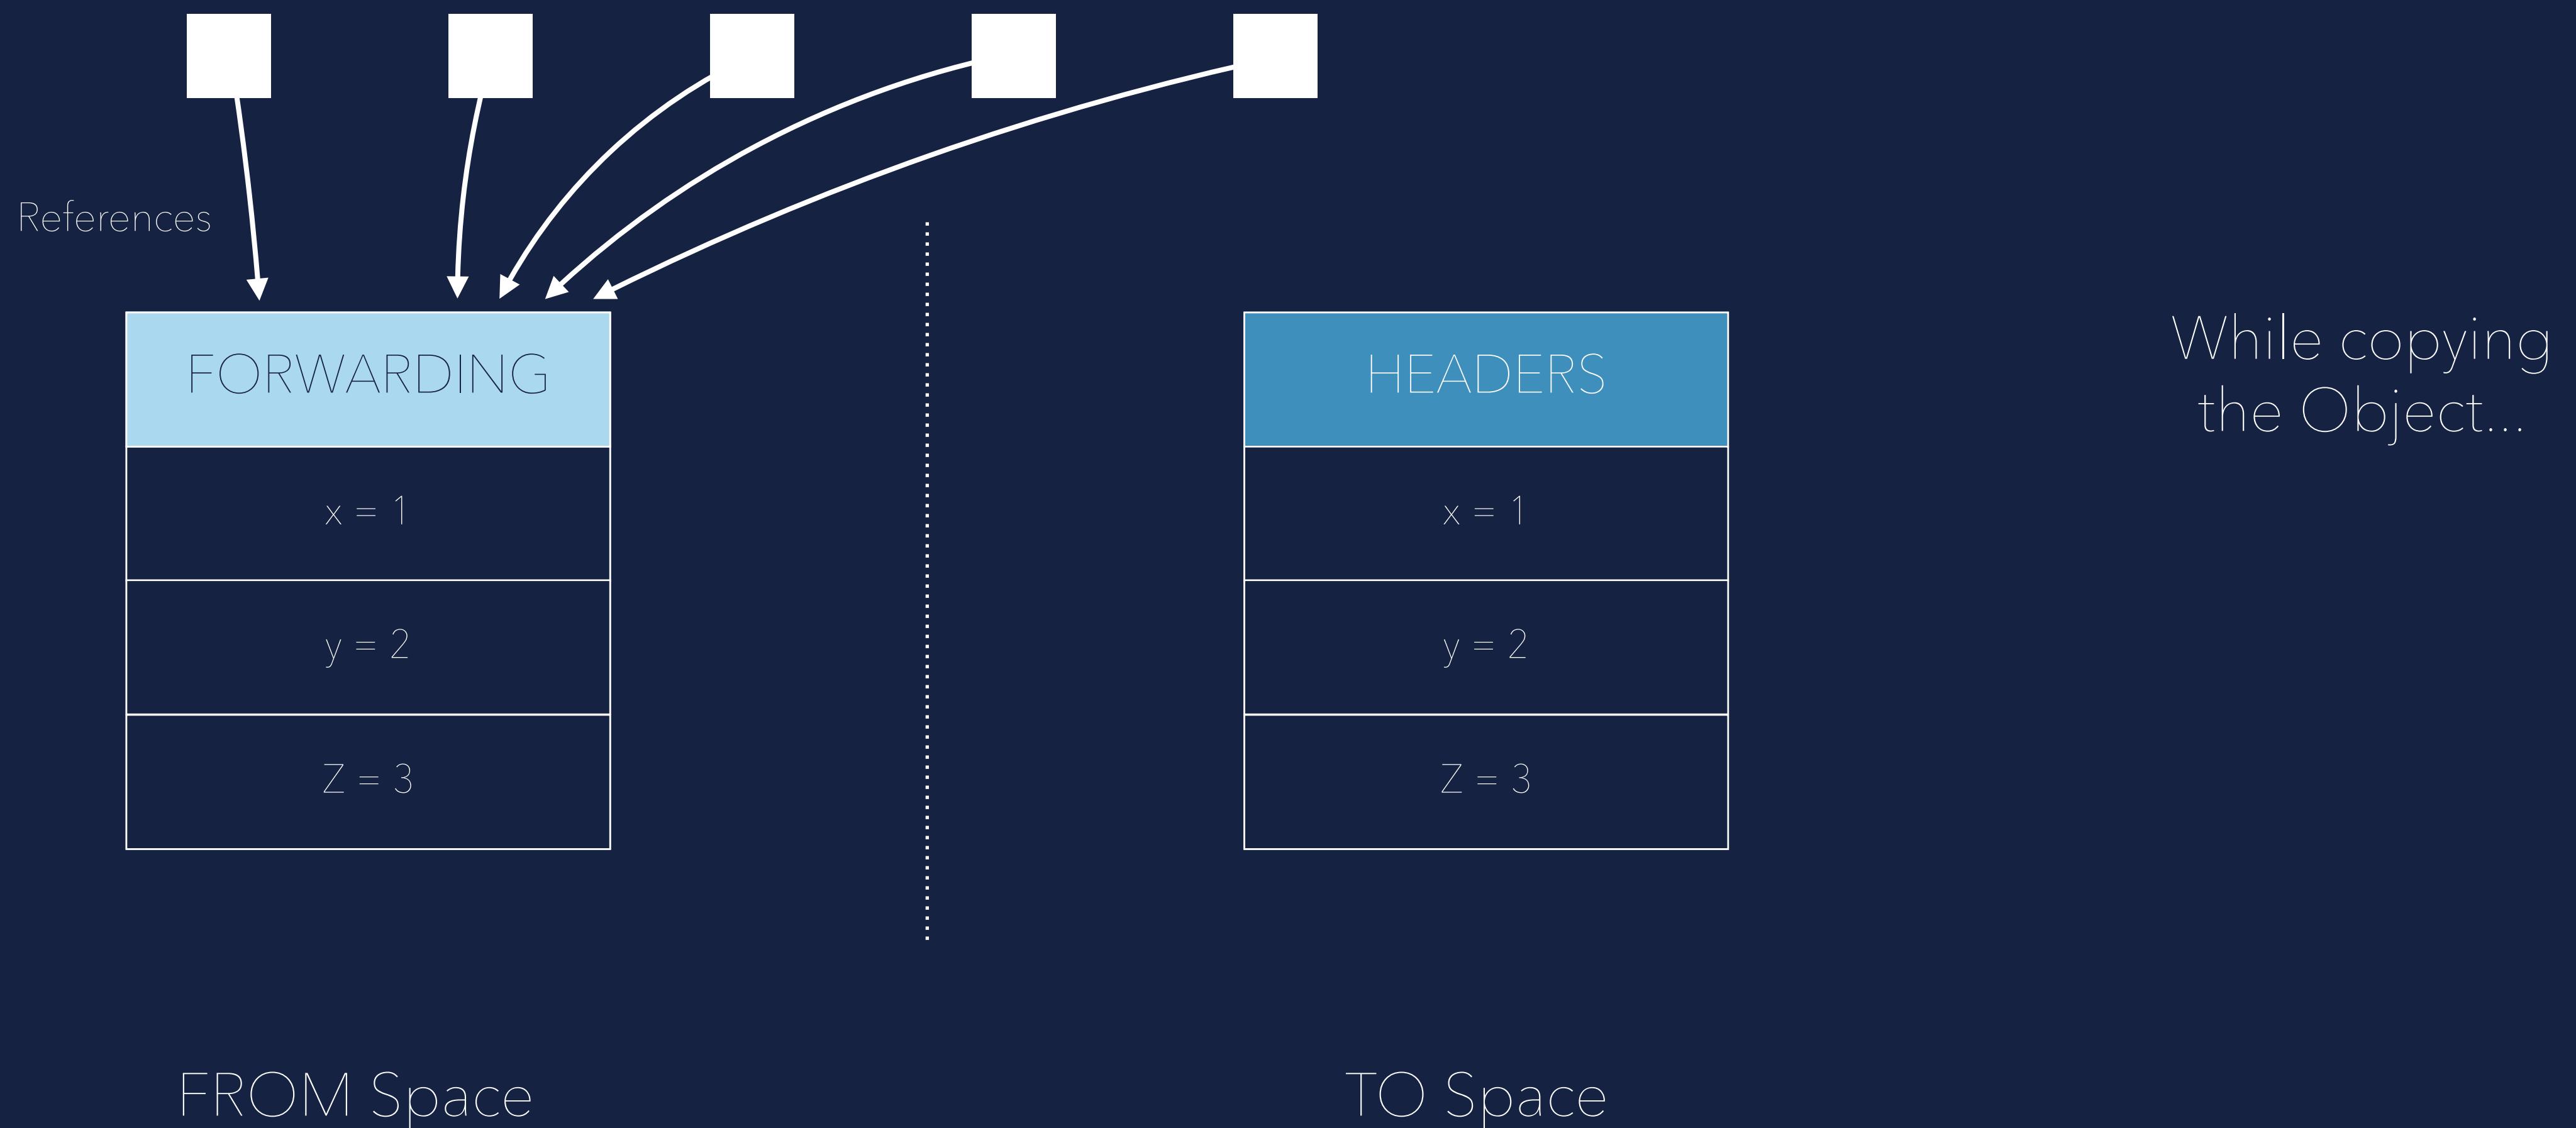
CONCURRENCY IS HARD...

Concurrent copying



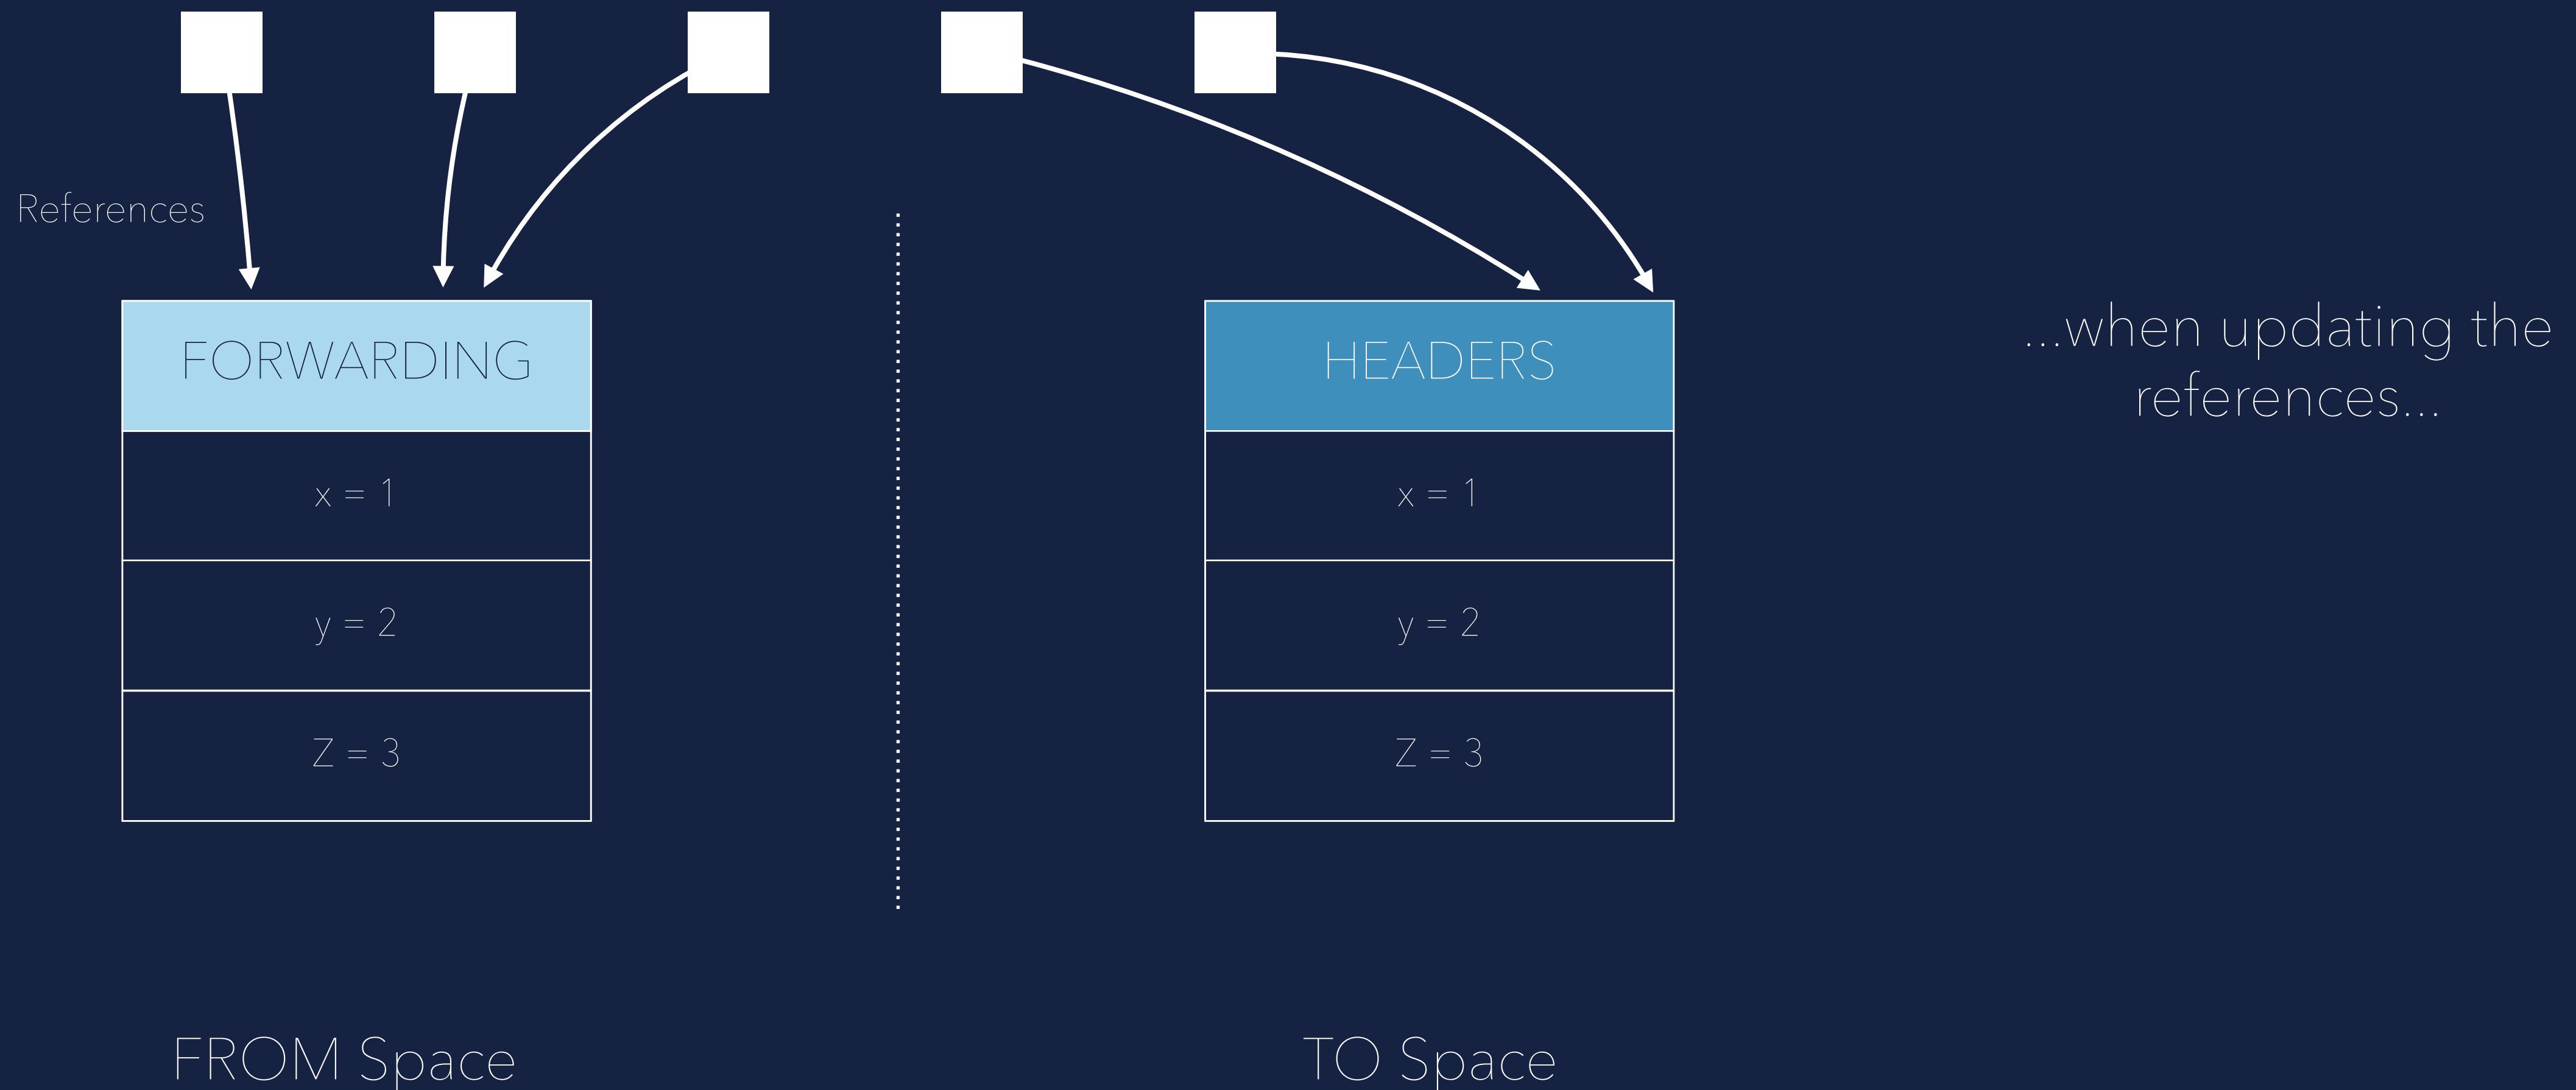
CONCURRENCY IS HARD...

Concurrent copying



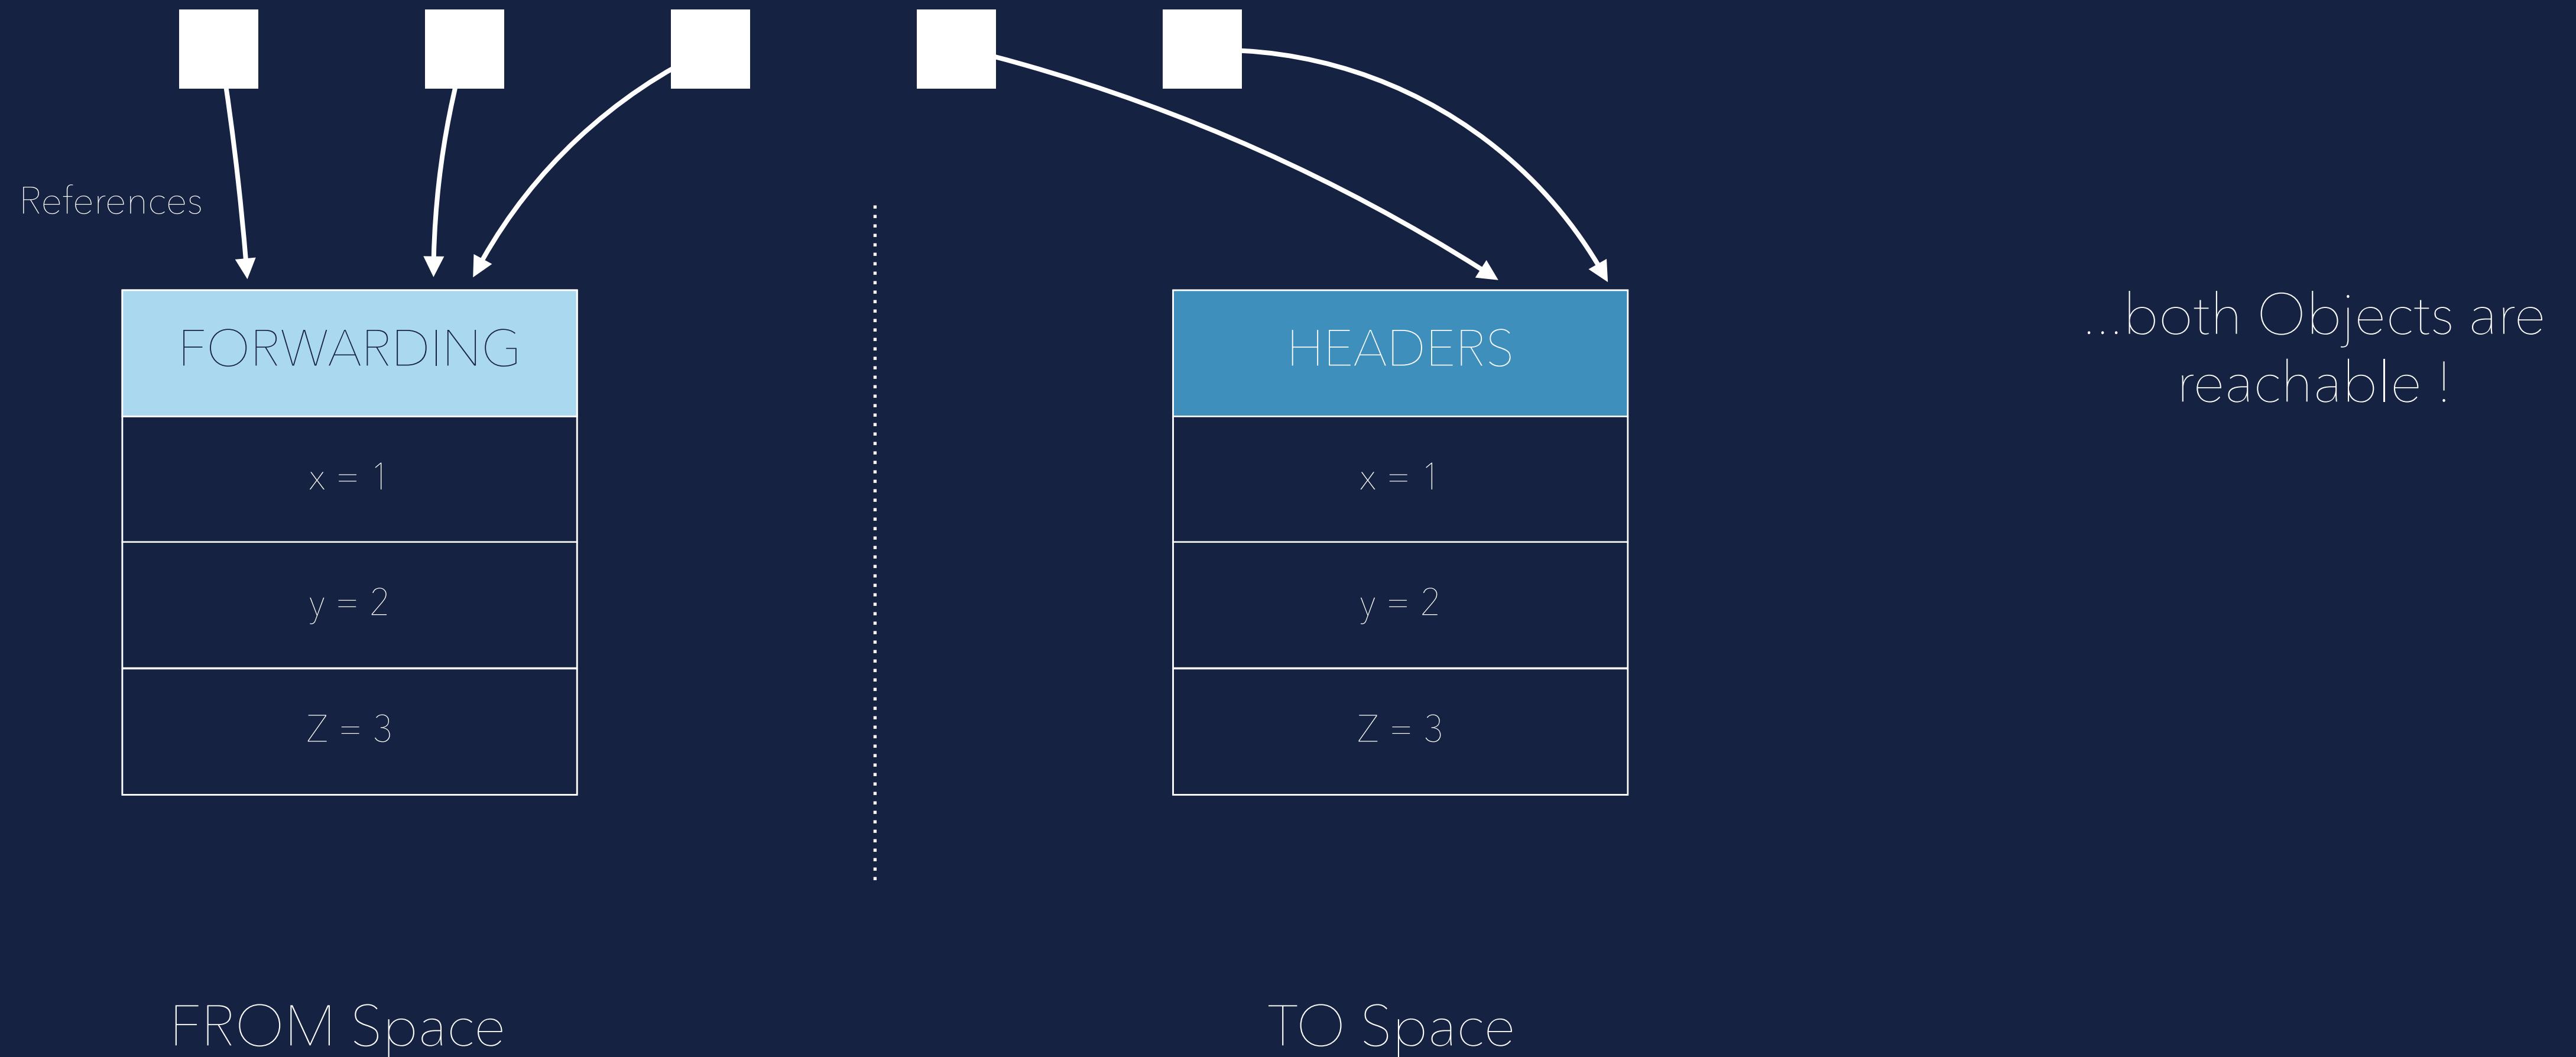
CONCURRENCY IS HARD...

Concurrent copying



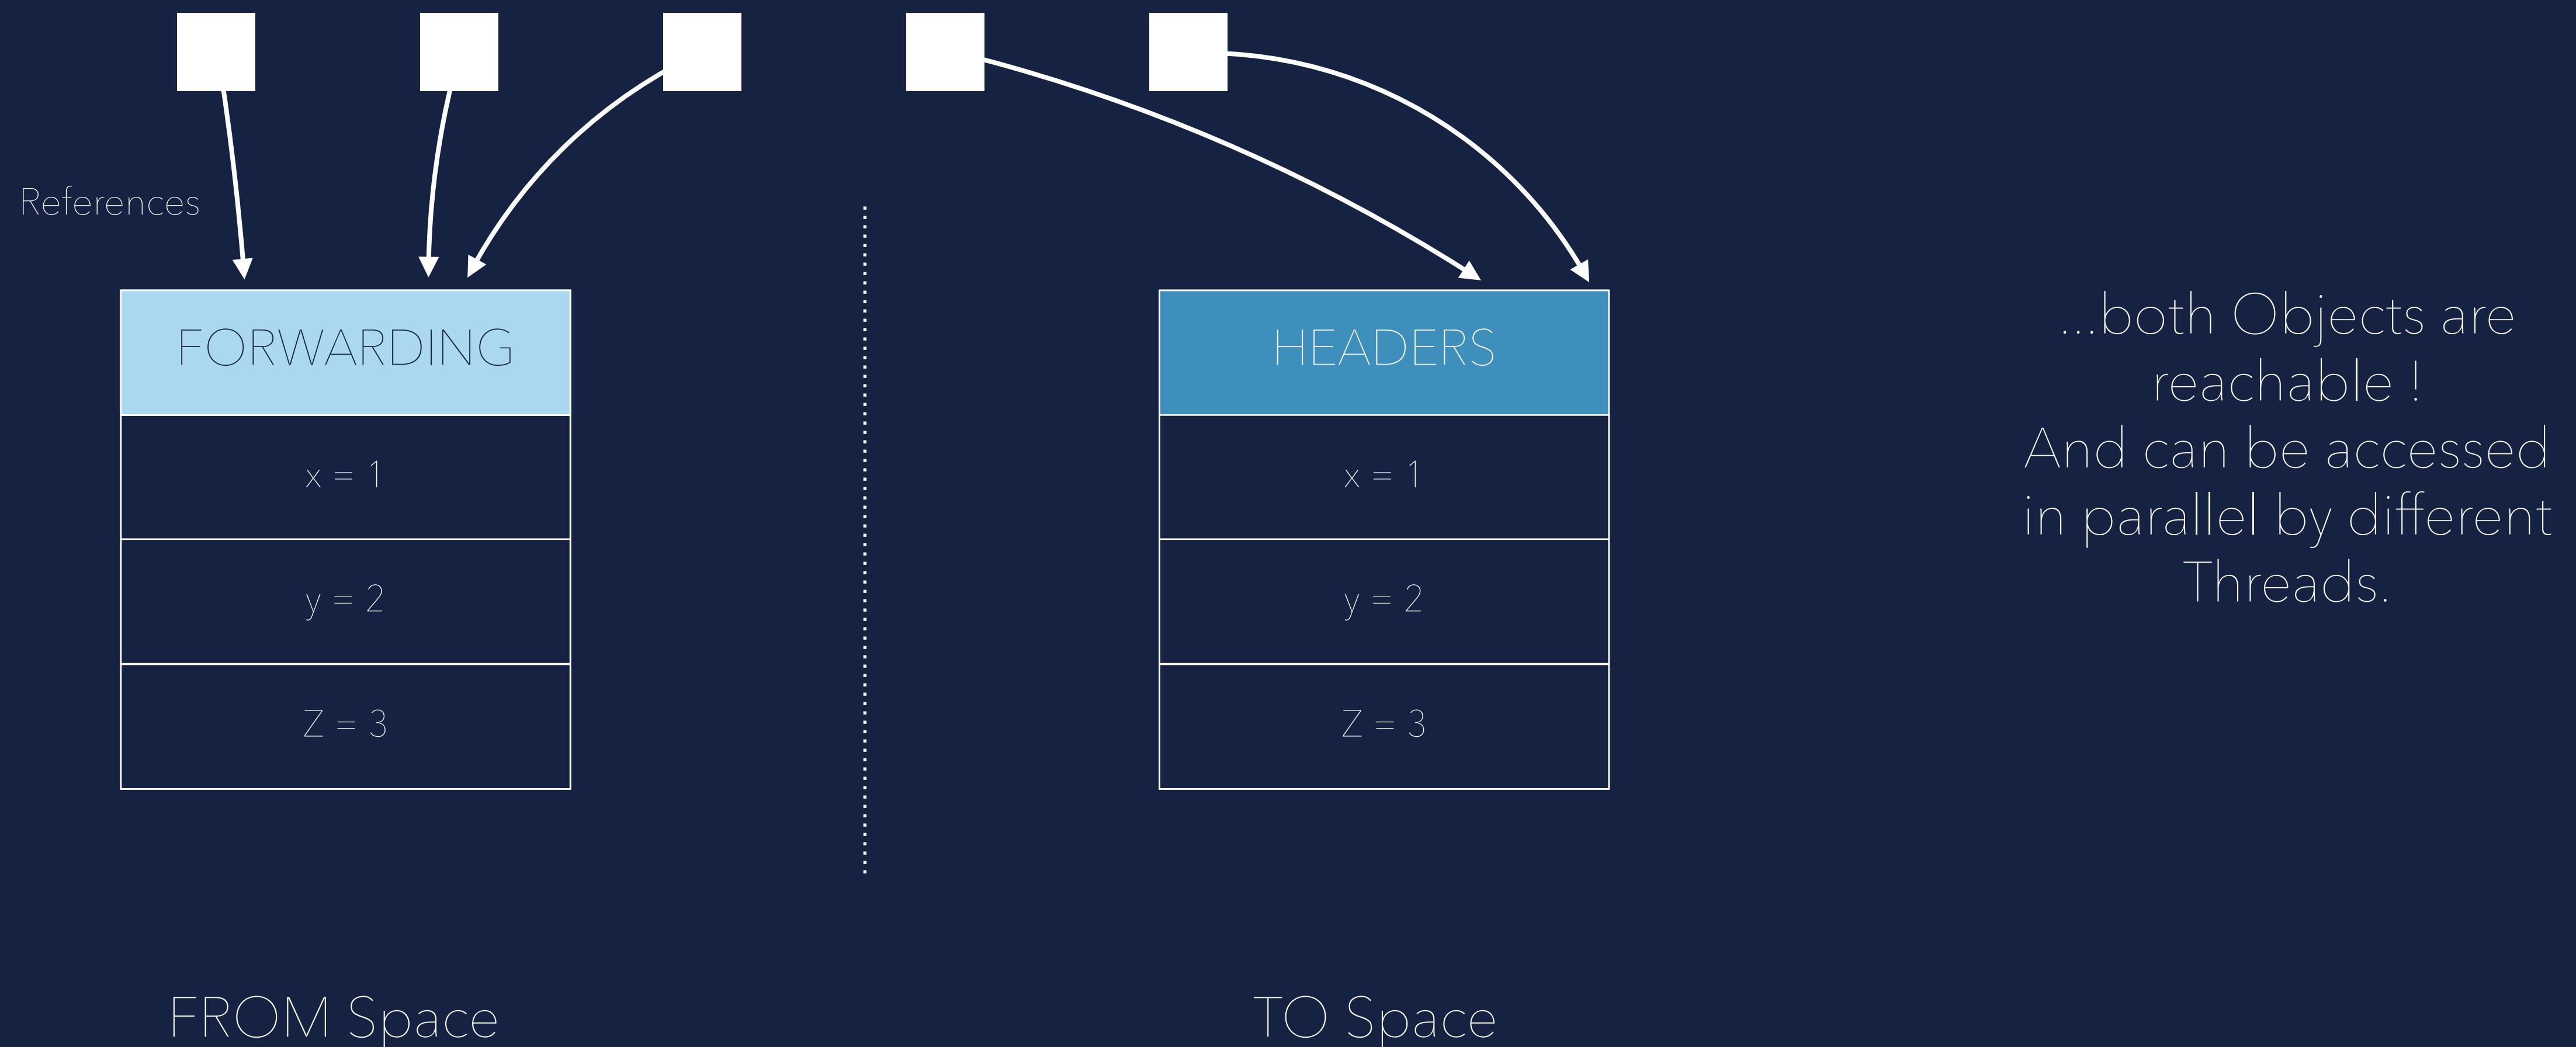
CONCURRENCY IS HARD...

Concurrent copying



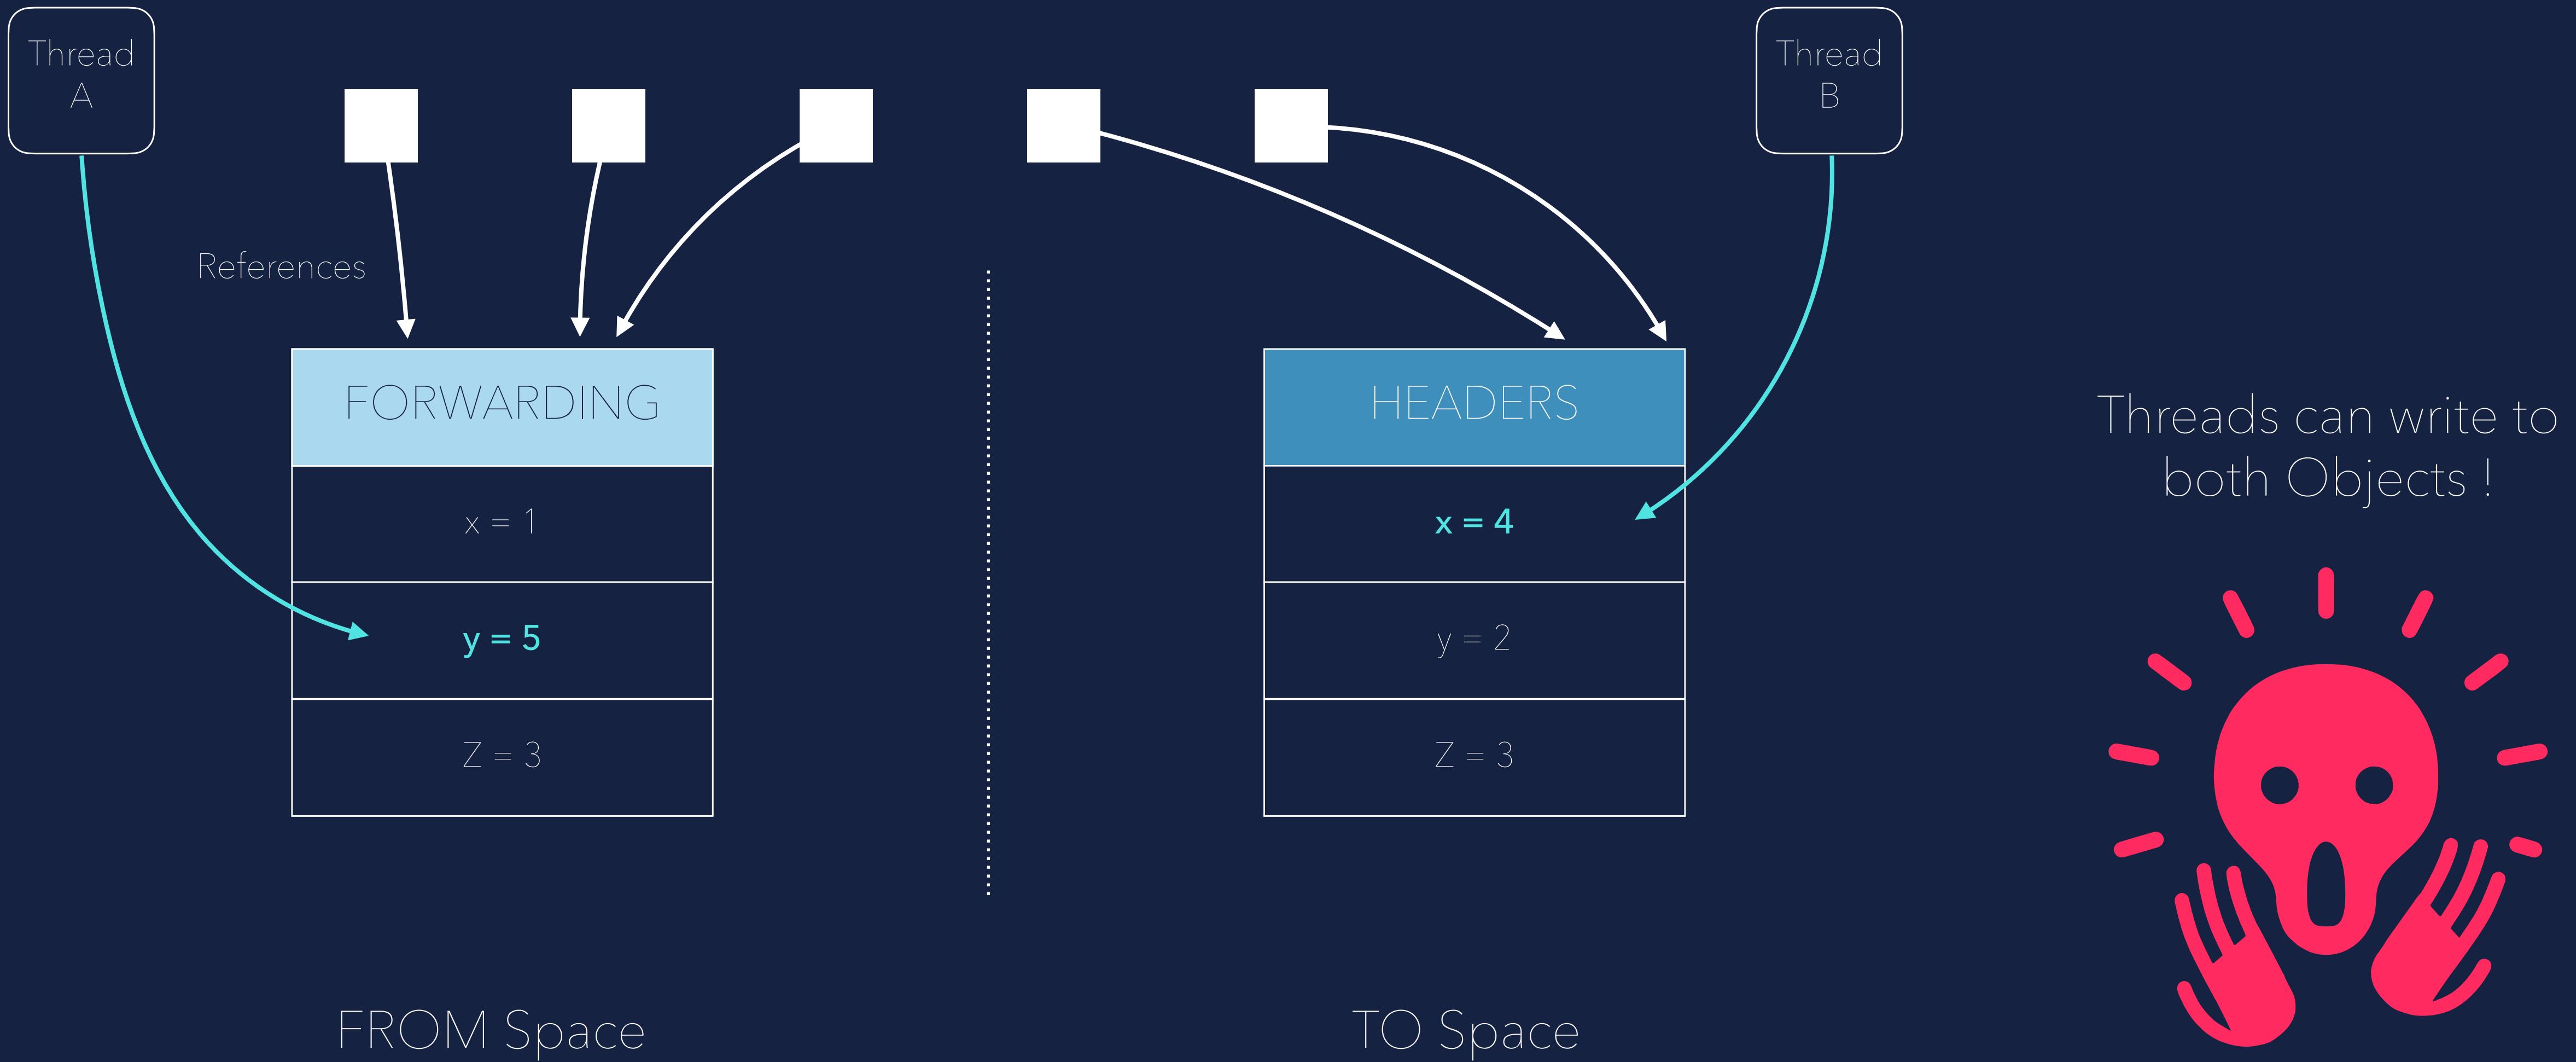
CONCURRENCY IS HARD...

Concurrent copying



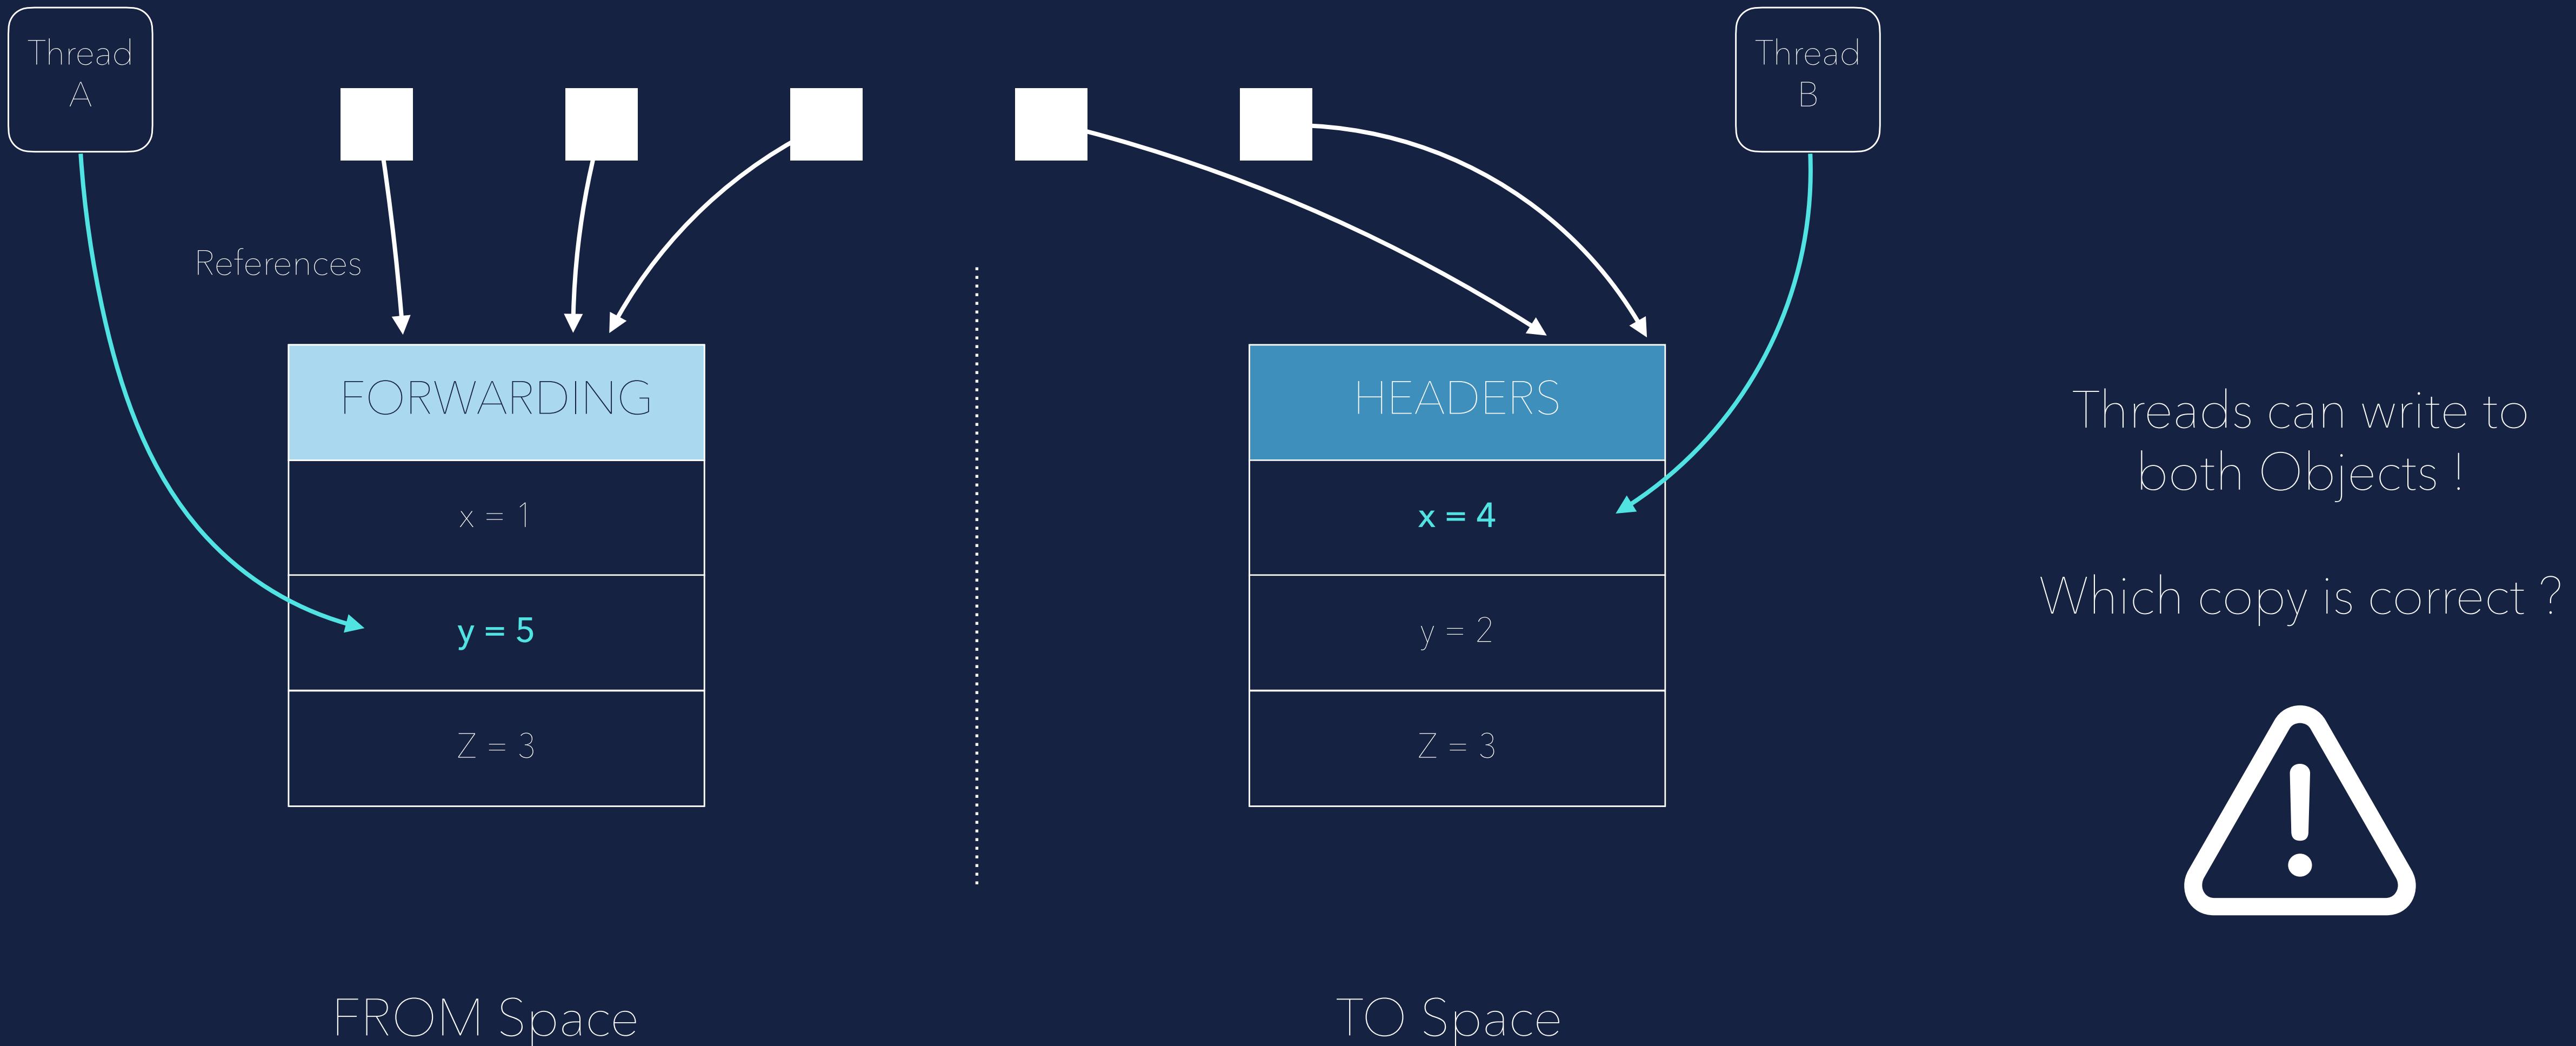
CONCURRENCY IS HARD...

Concurrent copying



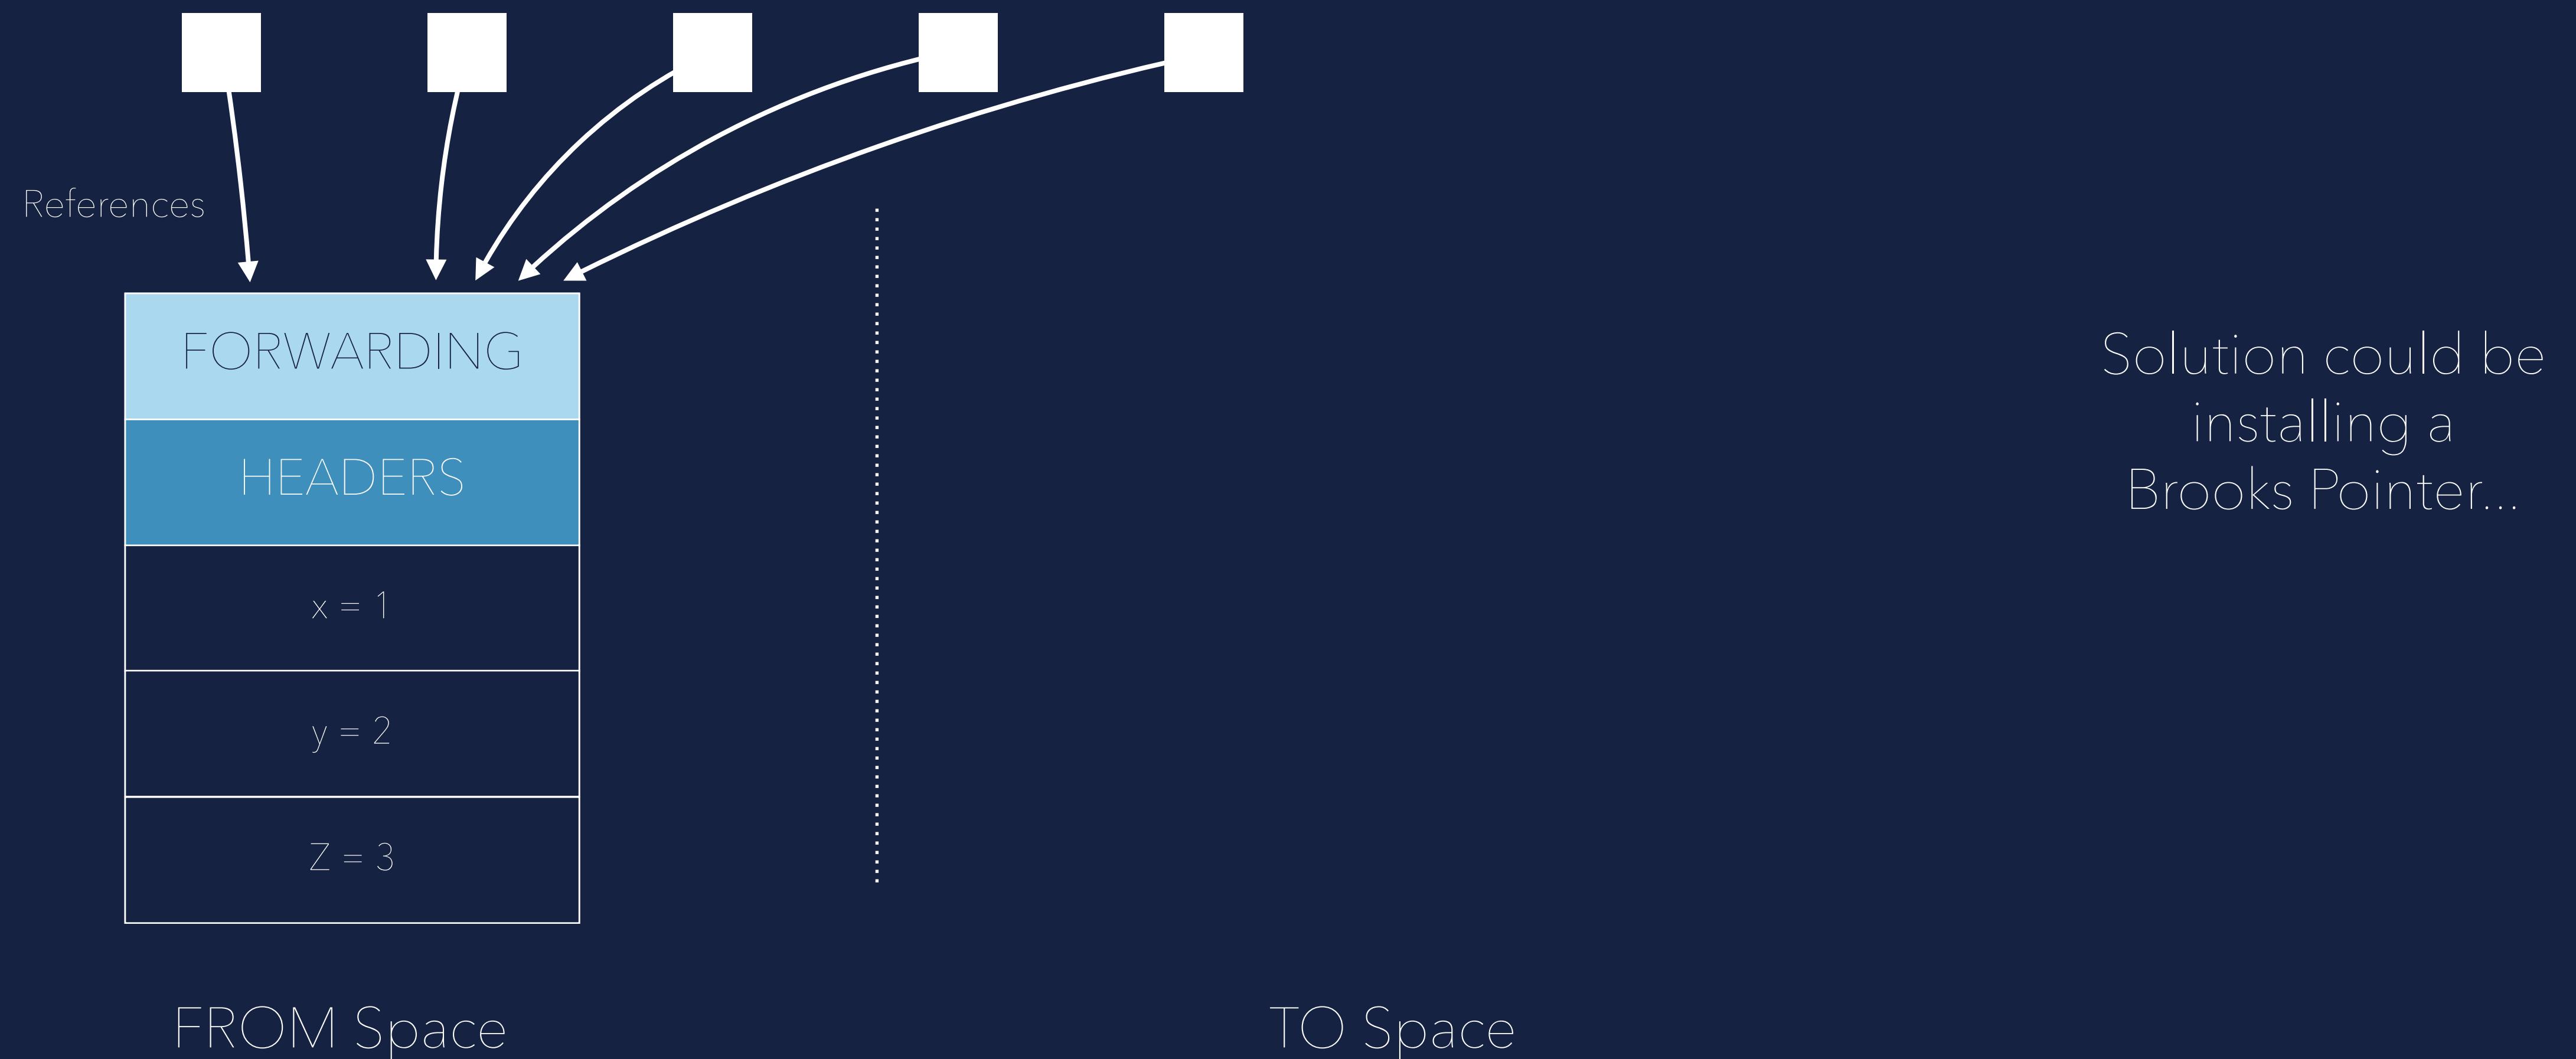
CONCURRENCY IS HARD...

Concurrent copying



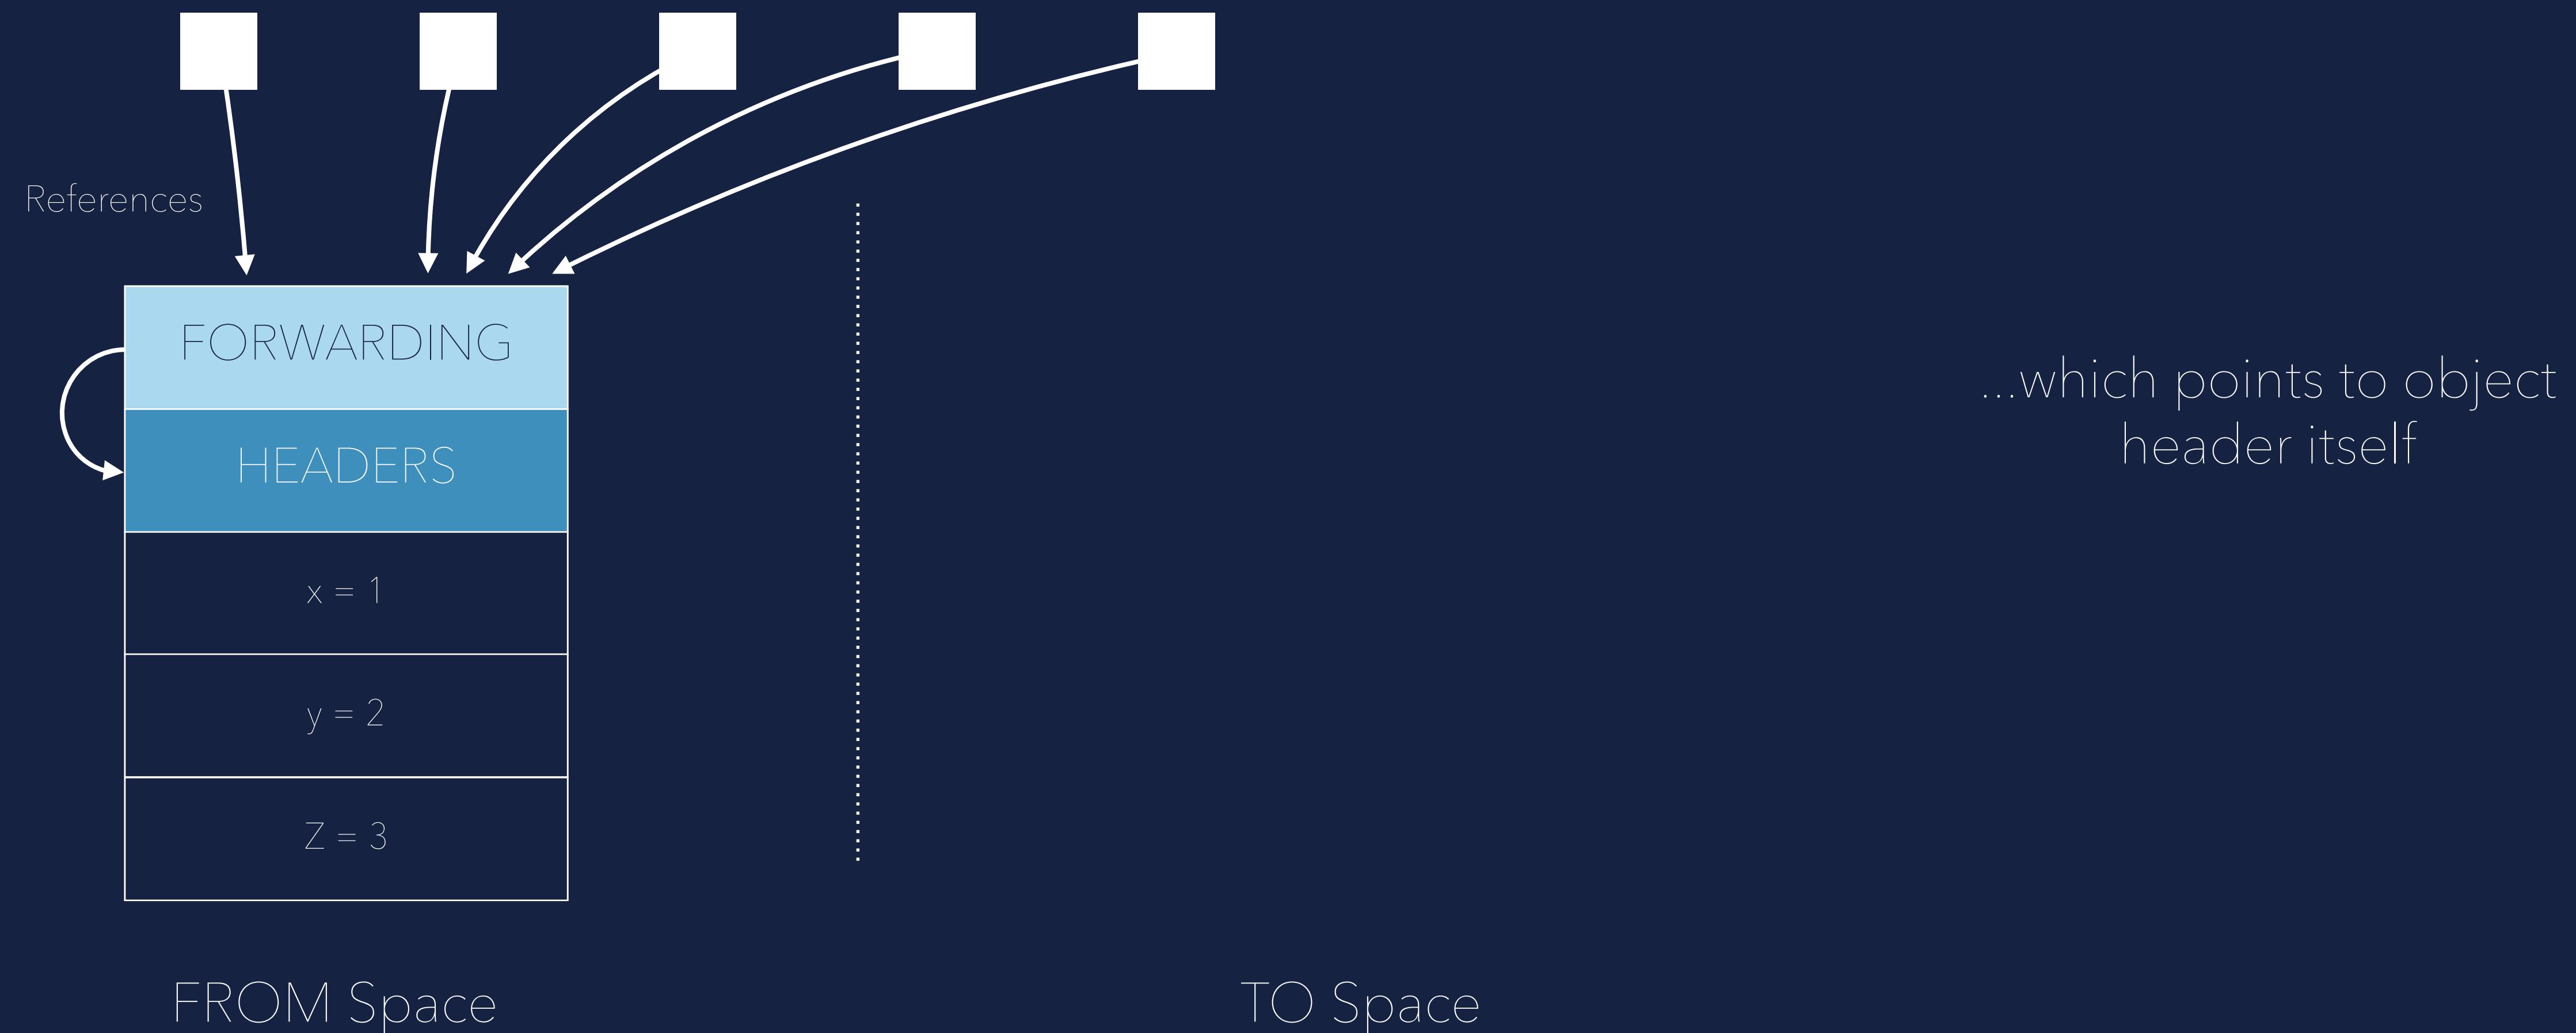
CONCURRENCY IS HARD...

Concurrent copying



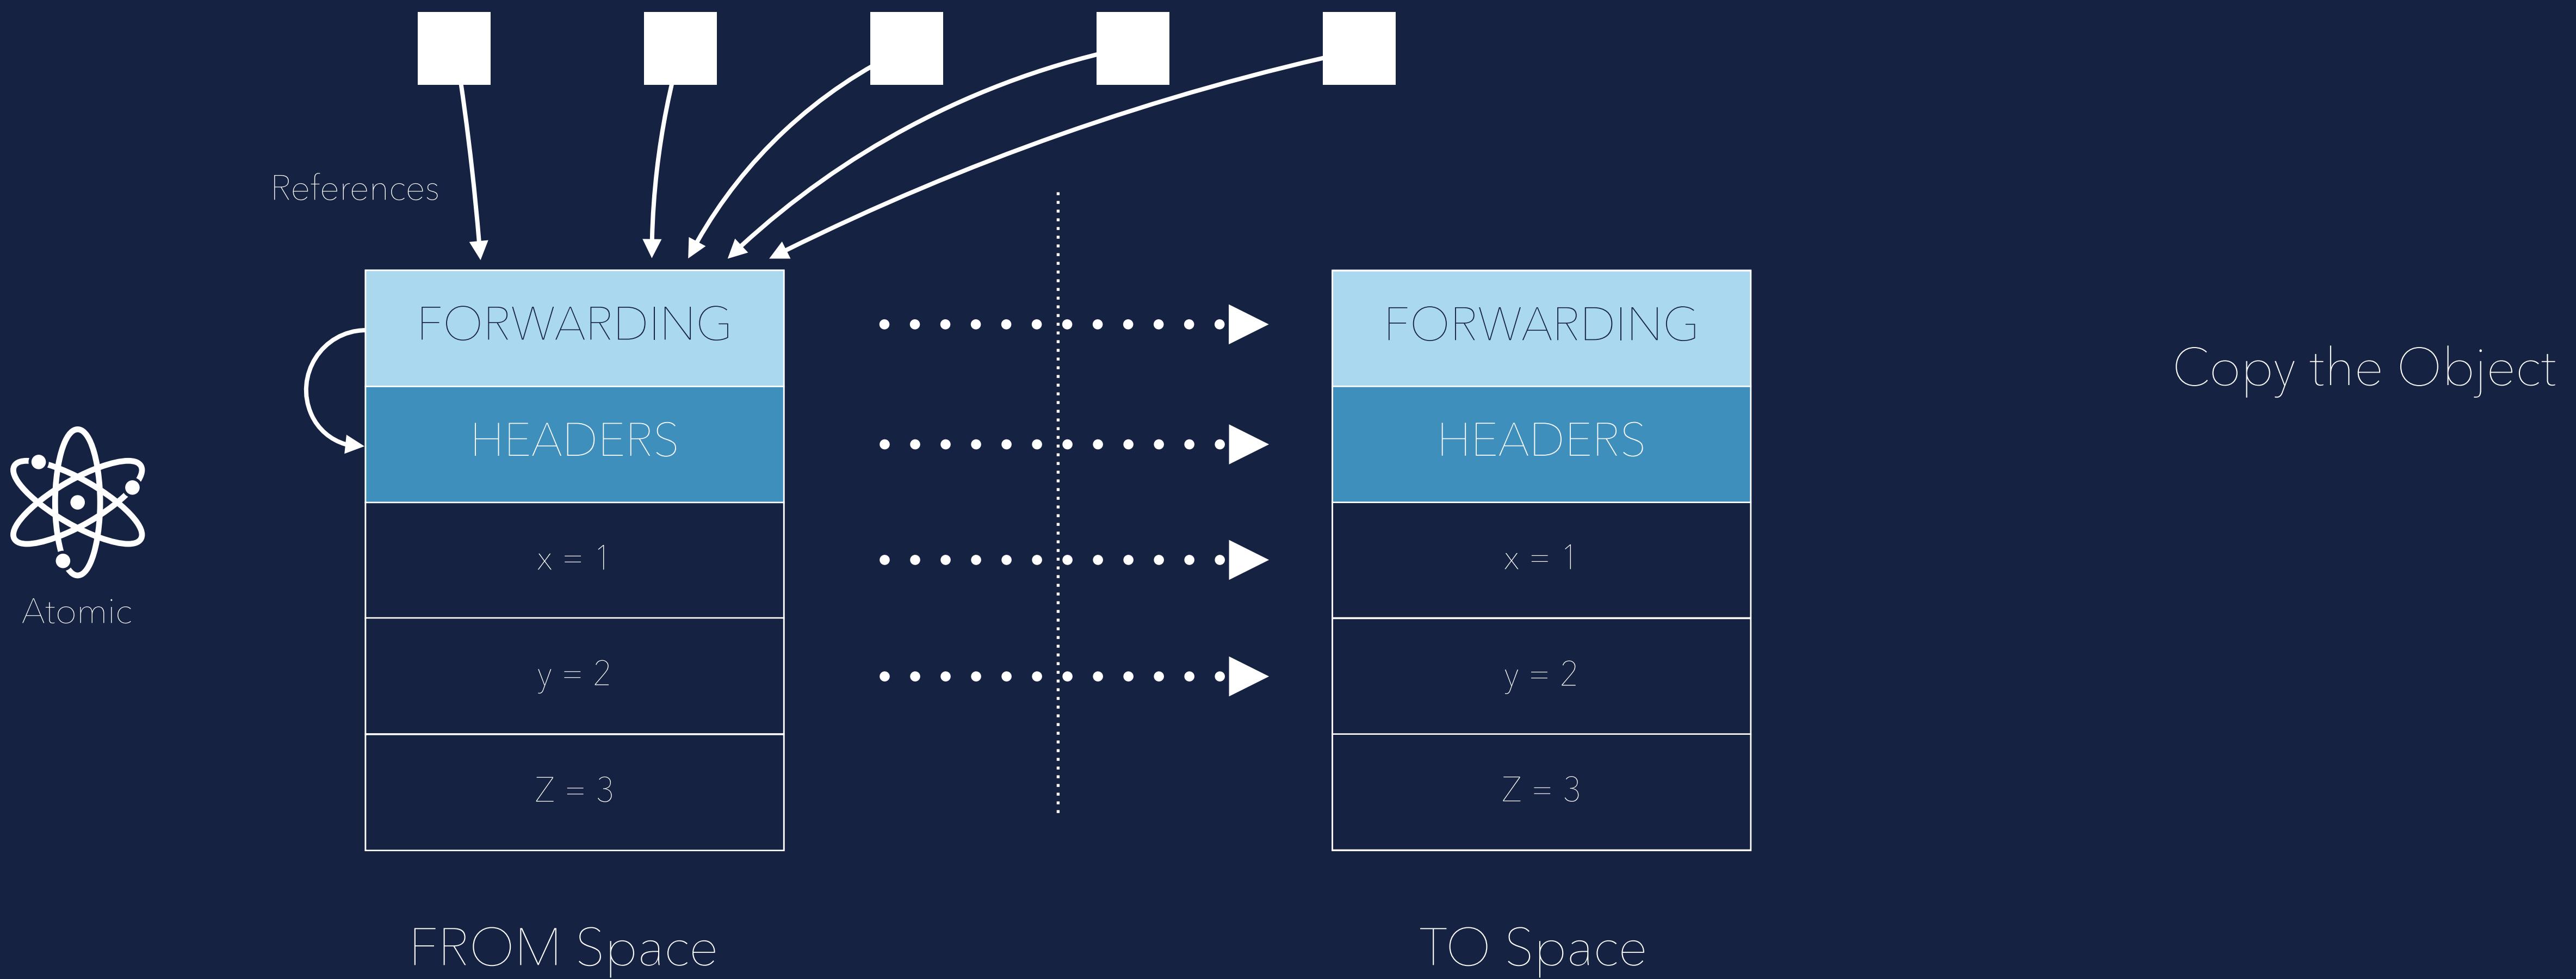
CONCURRENCY IS HARD...

Concurrent copying



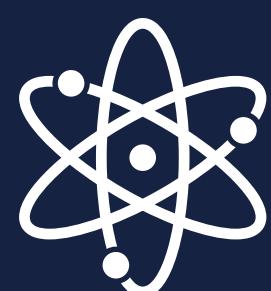
CONCURRENCY IS HARD...

Concurrent copying

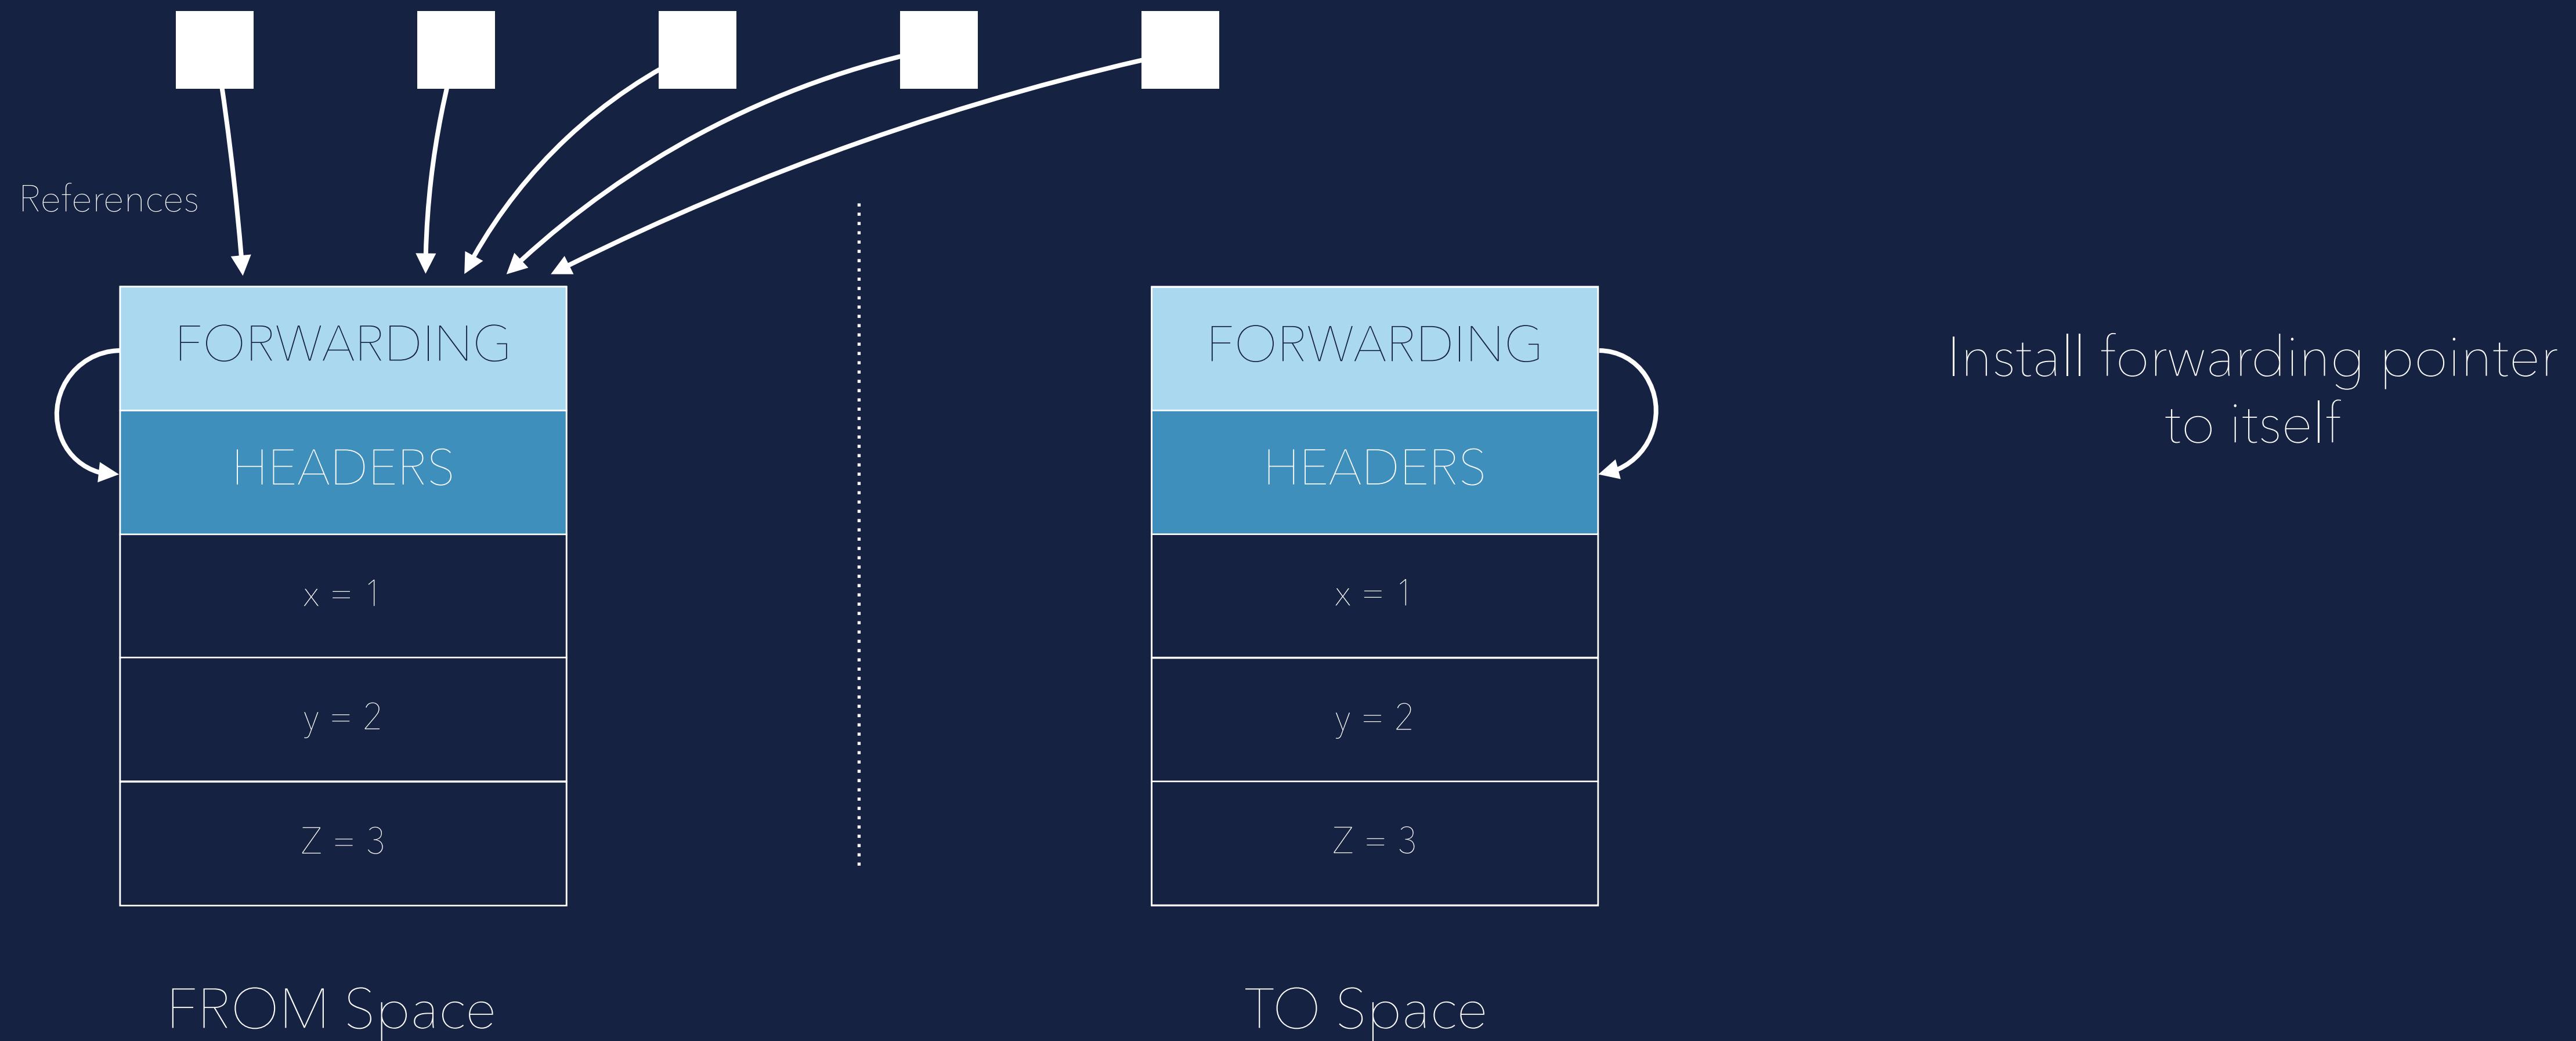


CONCURRENCY IS HARD...

Concurrent copying

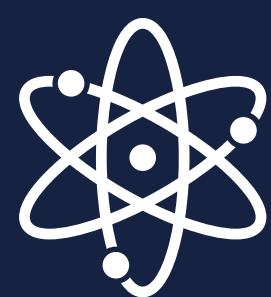


Atomic

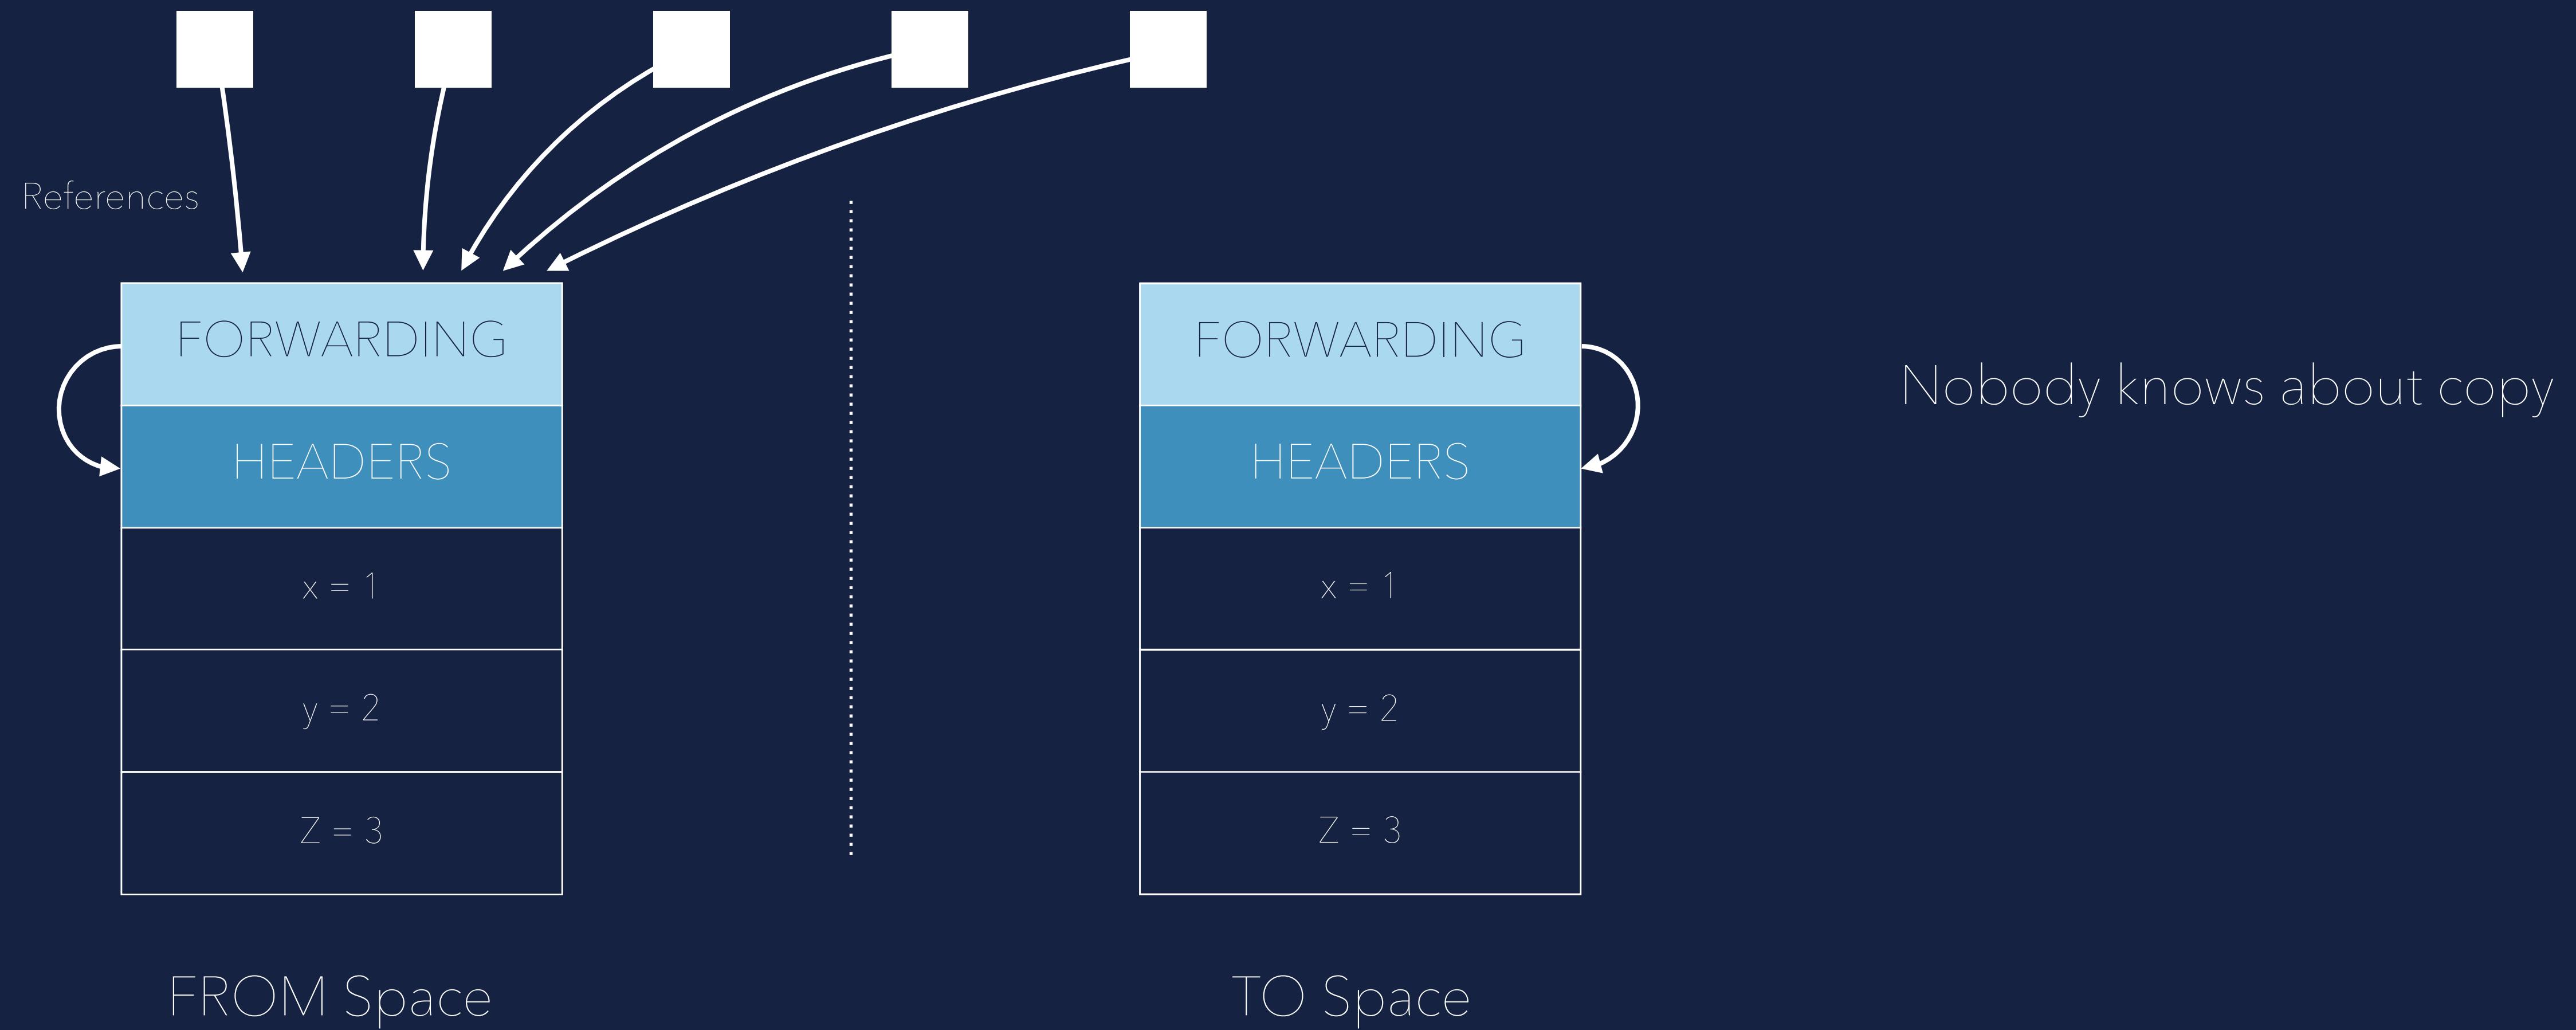


CONCURRENCY IS HARD...

Concurrent copying

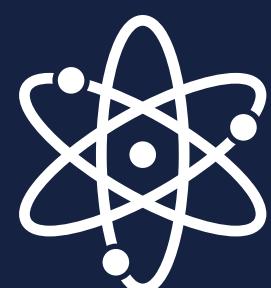


Atomic

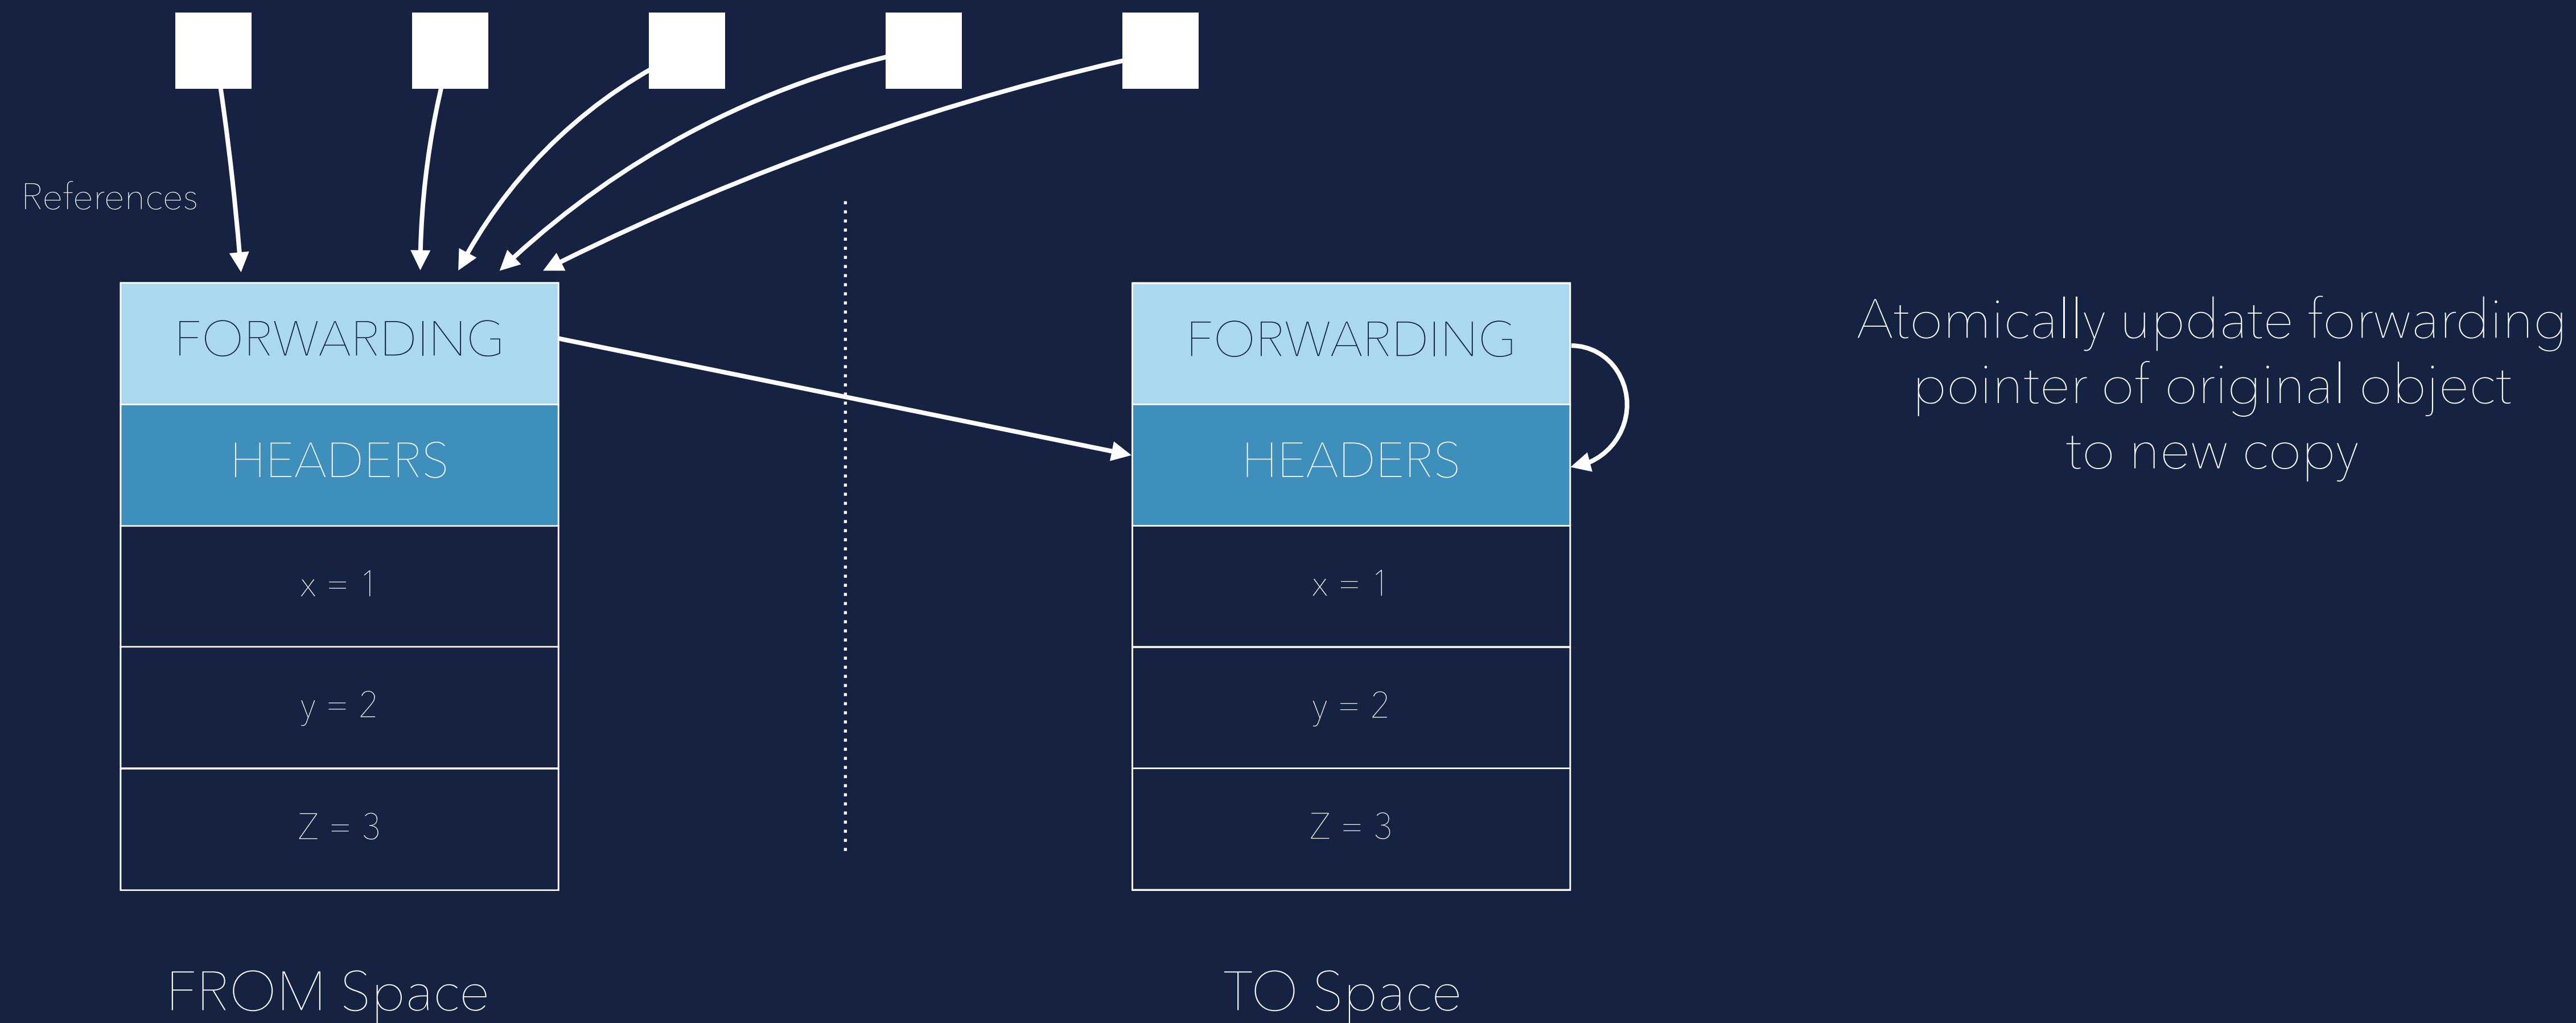


CONCURRENCY IS HARD...

Concurrent copying

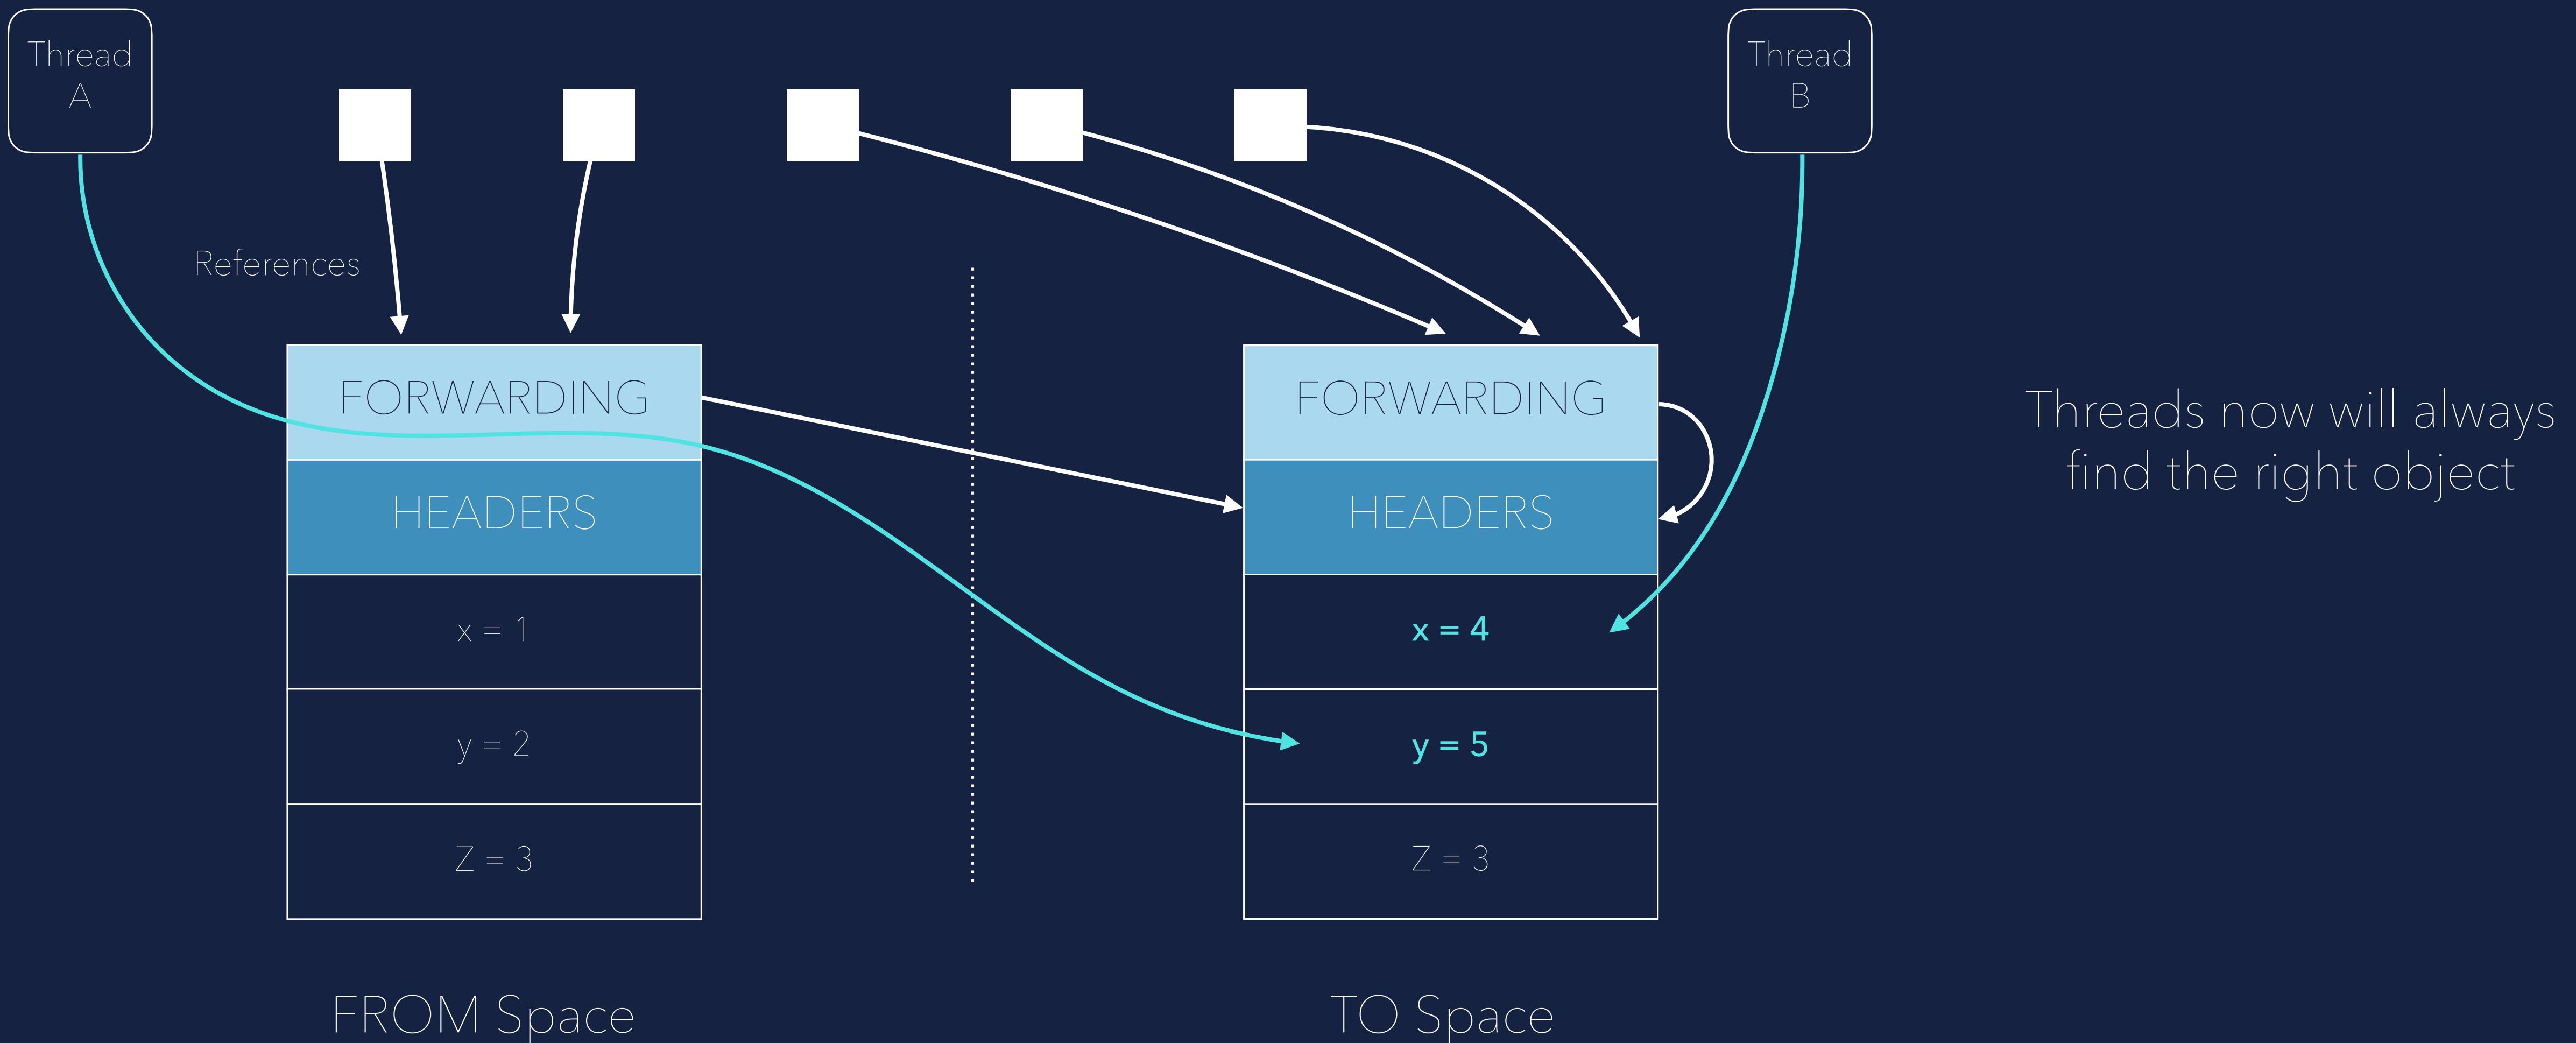


Atomic



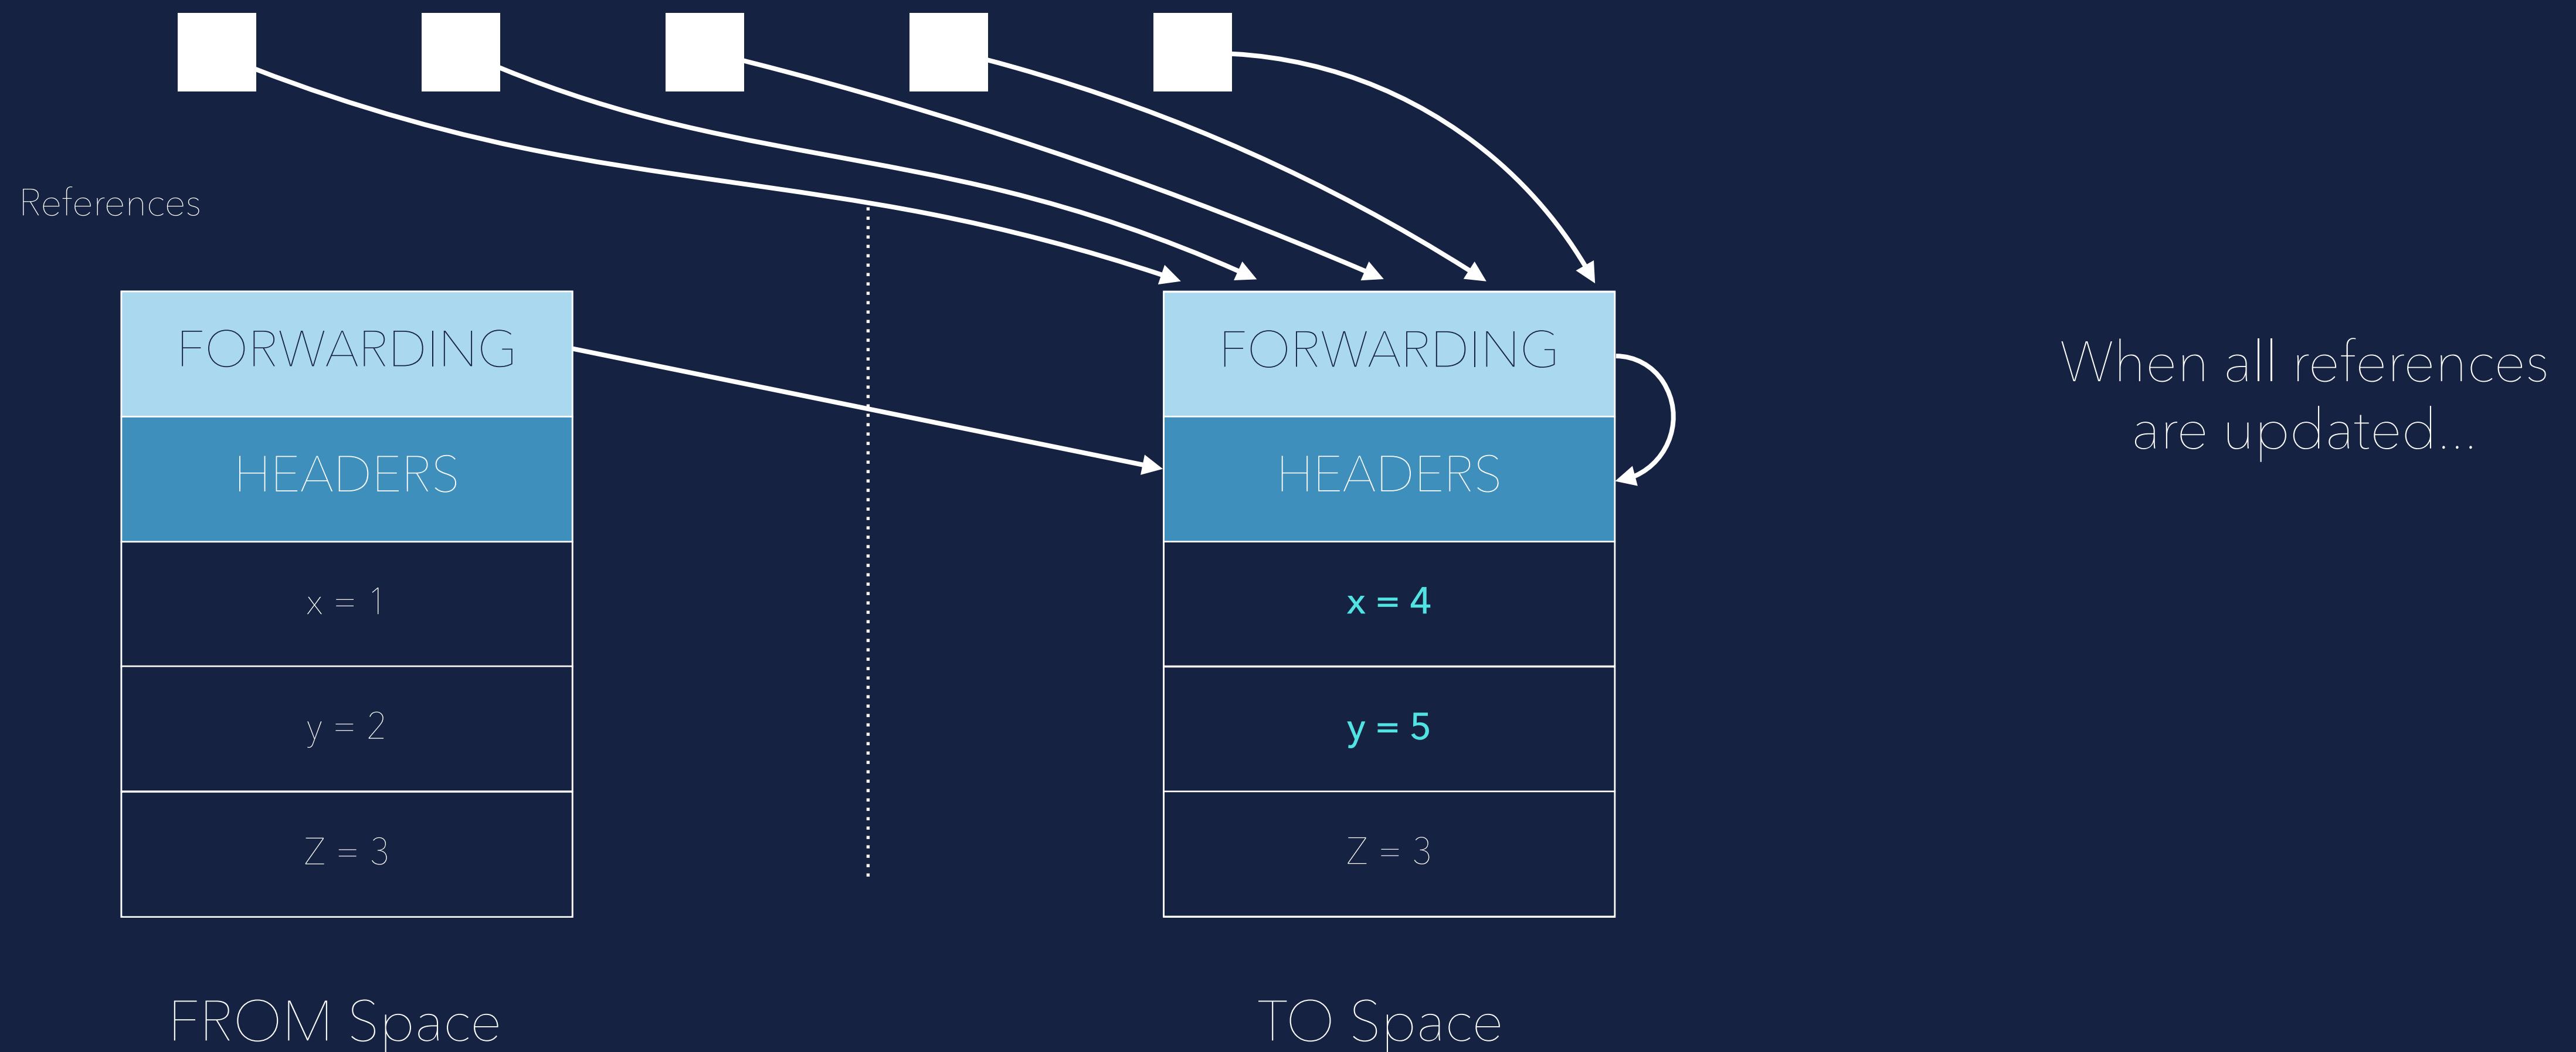
CONCURRENCY IS HARD...

Concurrent copying



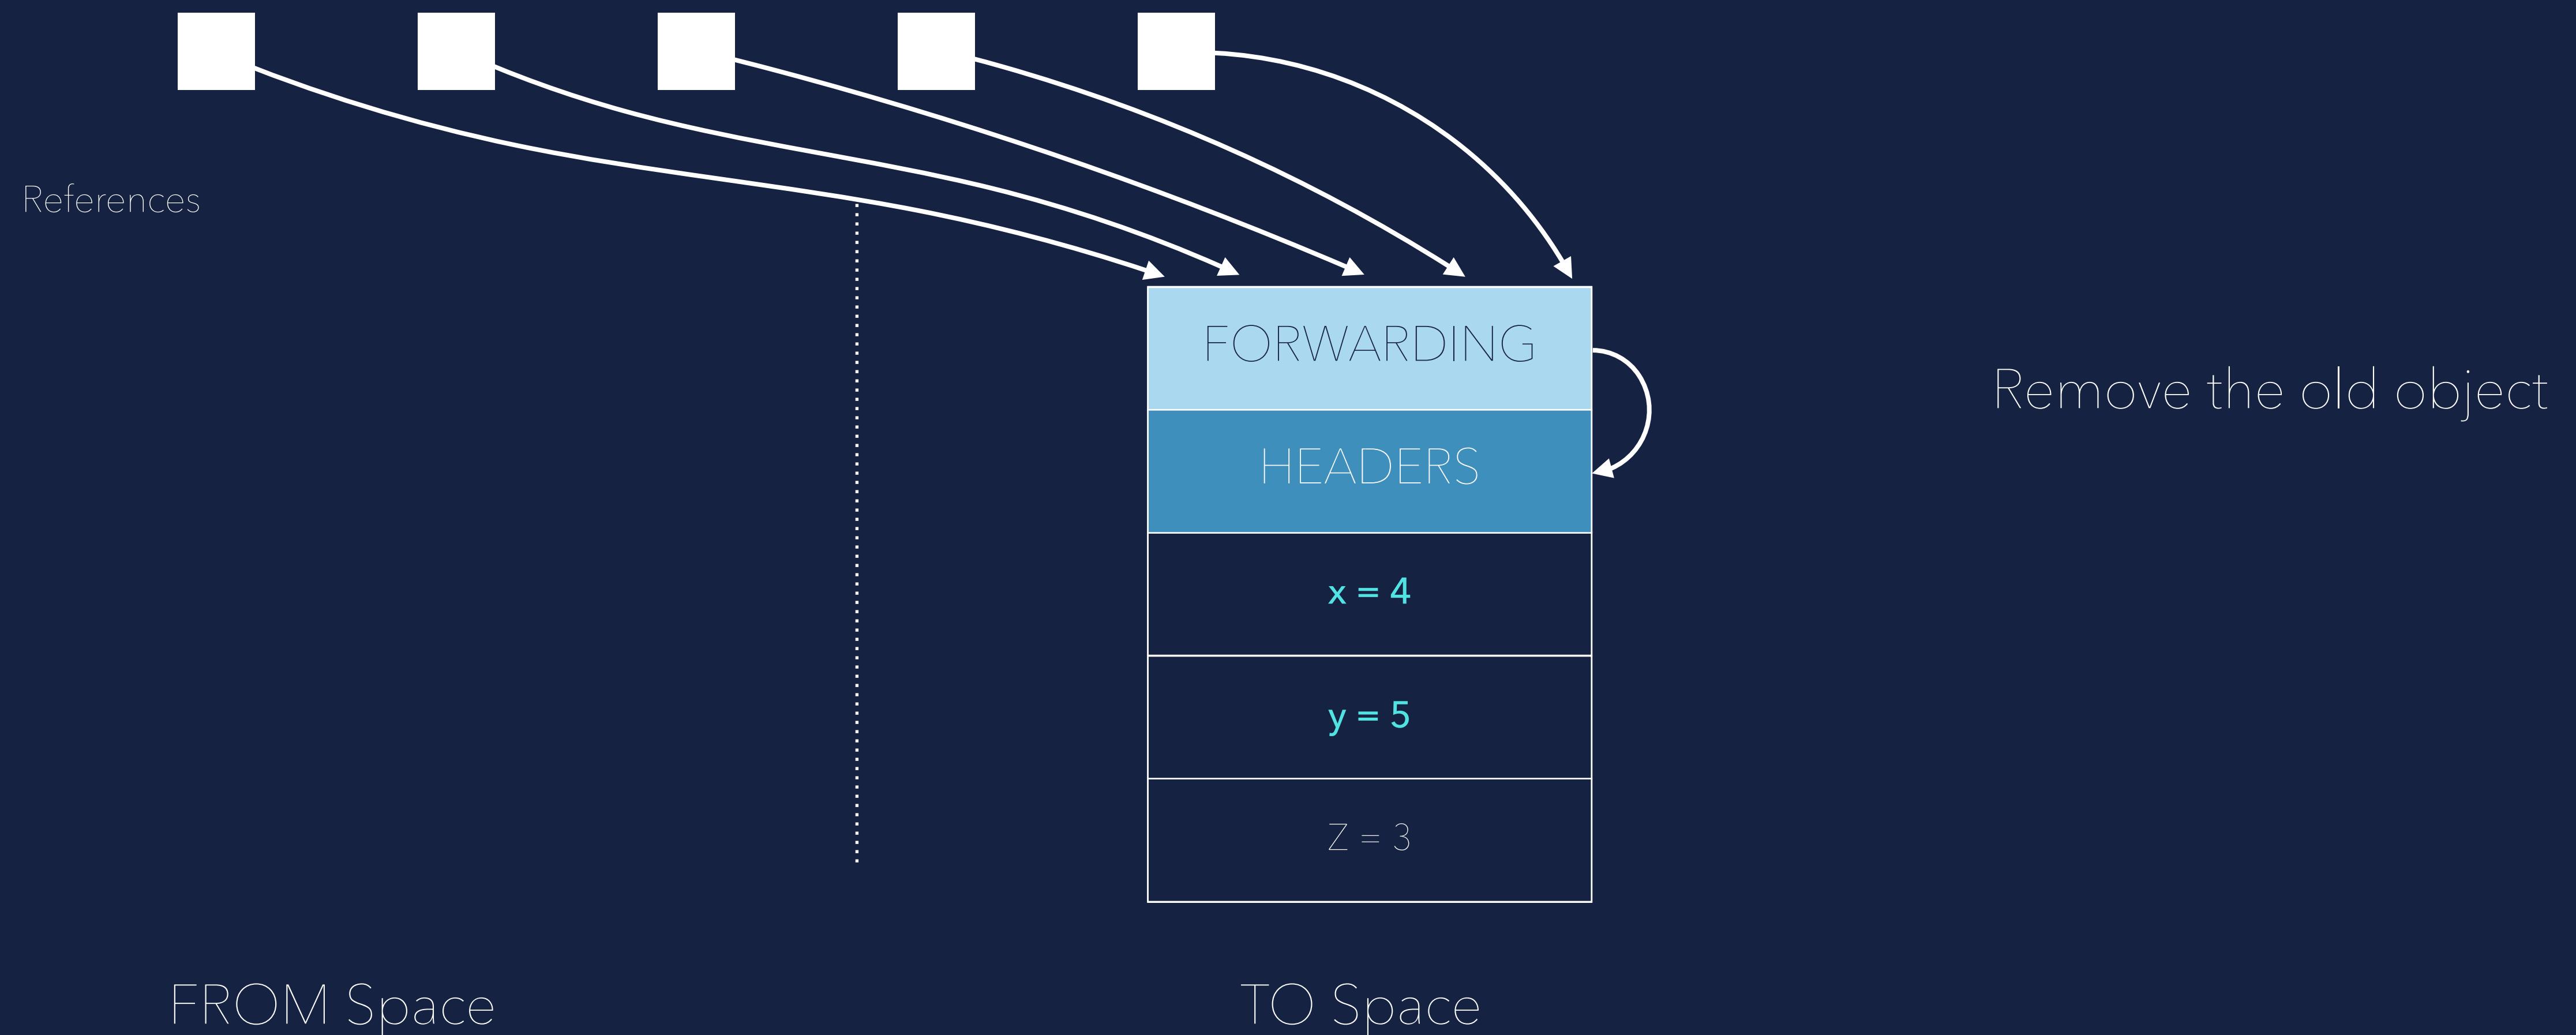
CONCURRENCY IS HARD...

Concurrent copying



CONCURRENCY IS HARD...

Concurrent copying



COLLECTORS IN THE JVM



SERIAL

azul



SERIAL

AVAILABILITY ALL JDK'S

PARALLEL NO

CONCURRENT NO

GENERATIONAL YES

HEAP SIZE SMALL - MEDIUM

PAUSE TIMES LONGER

THROUGHPUT LOW

LATENCY HIGHER

CPU OVERHEAD LOW (1-5%)

CHOOSE WHEN

- Single core systems with small heap (<4GB)
- No pause time requirements

BEST SUITED FOR

- Single threaded applications
- Development environments
- Microservices on small nodes



OS SUPPORT

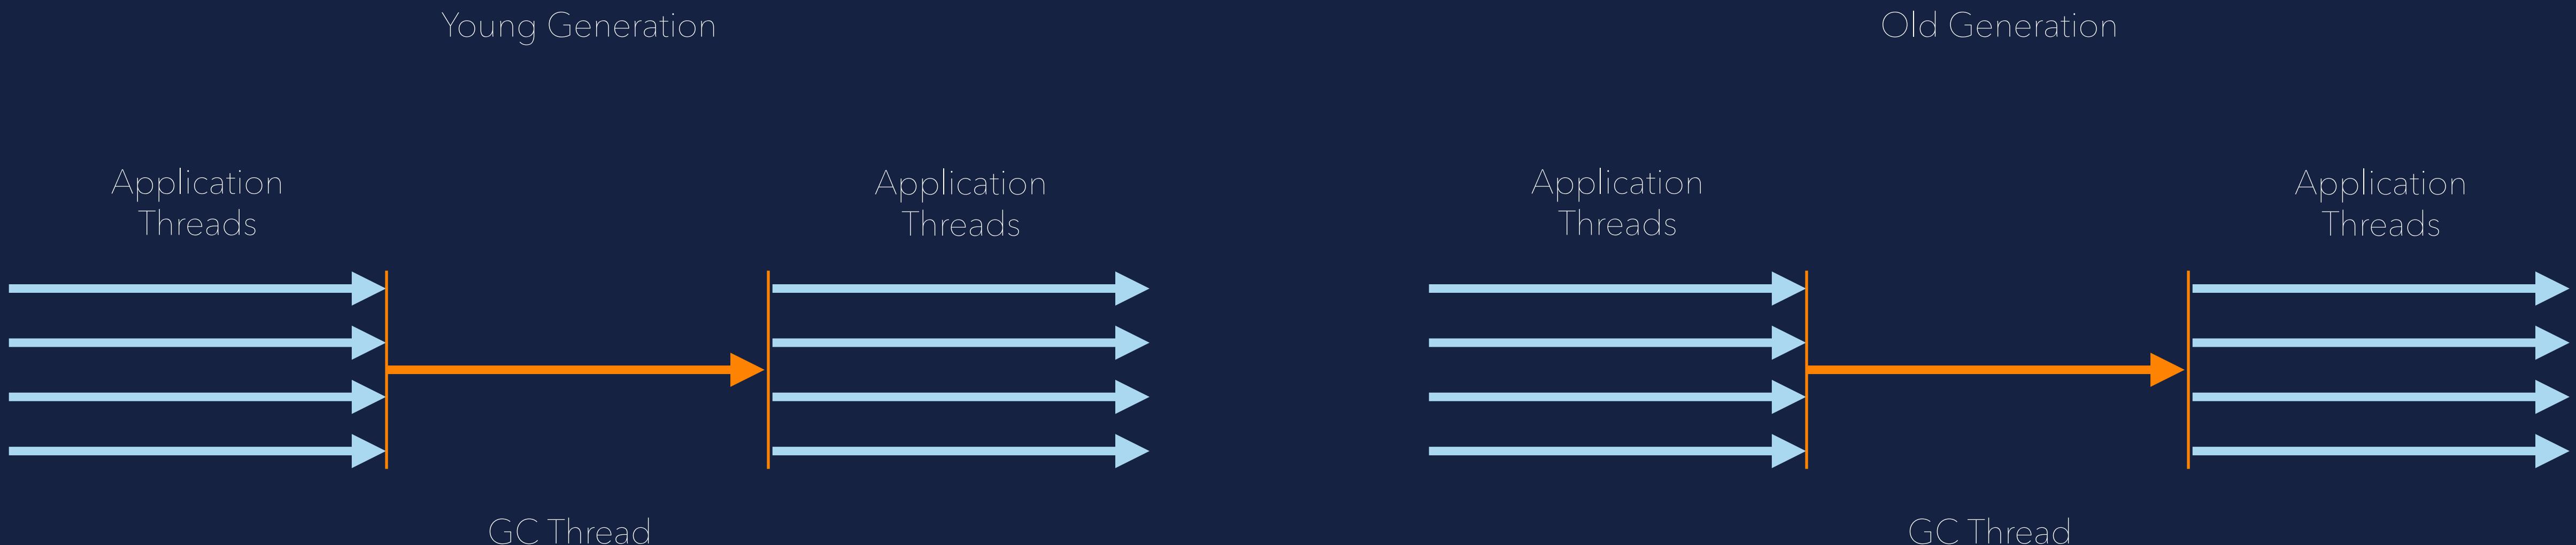
JVM SWITCH > java -XX:+UseSerialGC



SERIAL

NOTES

- Automatically selected if only a single processor is available
- Automatically selected if the avail. memory less than 1792 MB
- Mark and Compact





PARALLEL

azul



PARALLEL

AVAILABILITY

ALL JDK'S

PARALLEL

YES

CONCURRENT

NO

GENERATIONAL

YES

HEAP SIZE

MEDIUM - LARGE

PAUSE TIMES

MODERATE

THROUGHPUT

HIGH

LATENCY

LOWER

CPU OVERHEAD

MODERATE (5-10%)

CHOOSE WHEN

Multi-core systems with small heap (<4GB)

Peak performance is needed without pause time requirements

BEST SUITED FOR

Batch processing

Scientific computing

Data analysis

OS SUPPORT



JVM SWITCH

> java -XX:+UseParallelGC/-XX:+UseParallelOldGC



PARALLEL

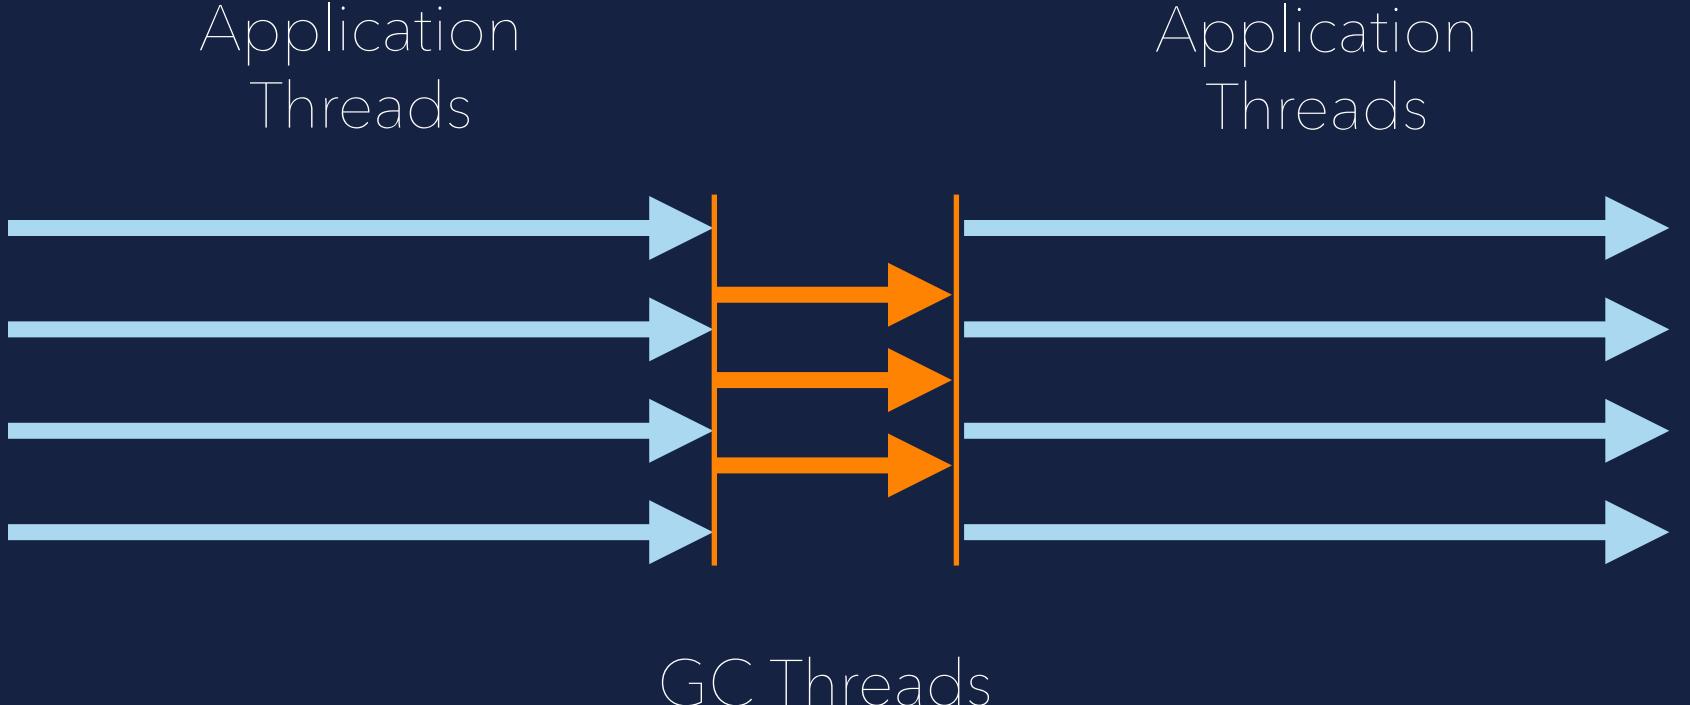
NOTES

Default garbage collector from JDK 5 to JDK 8

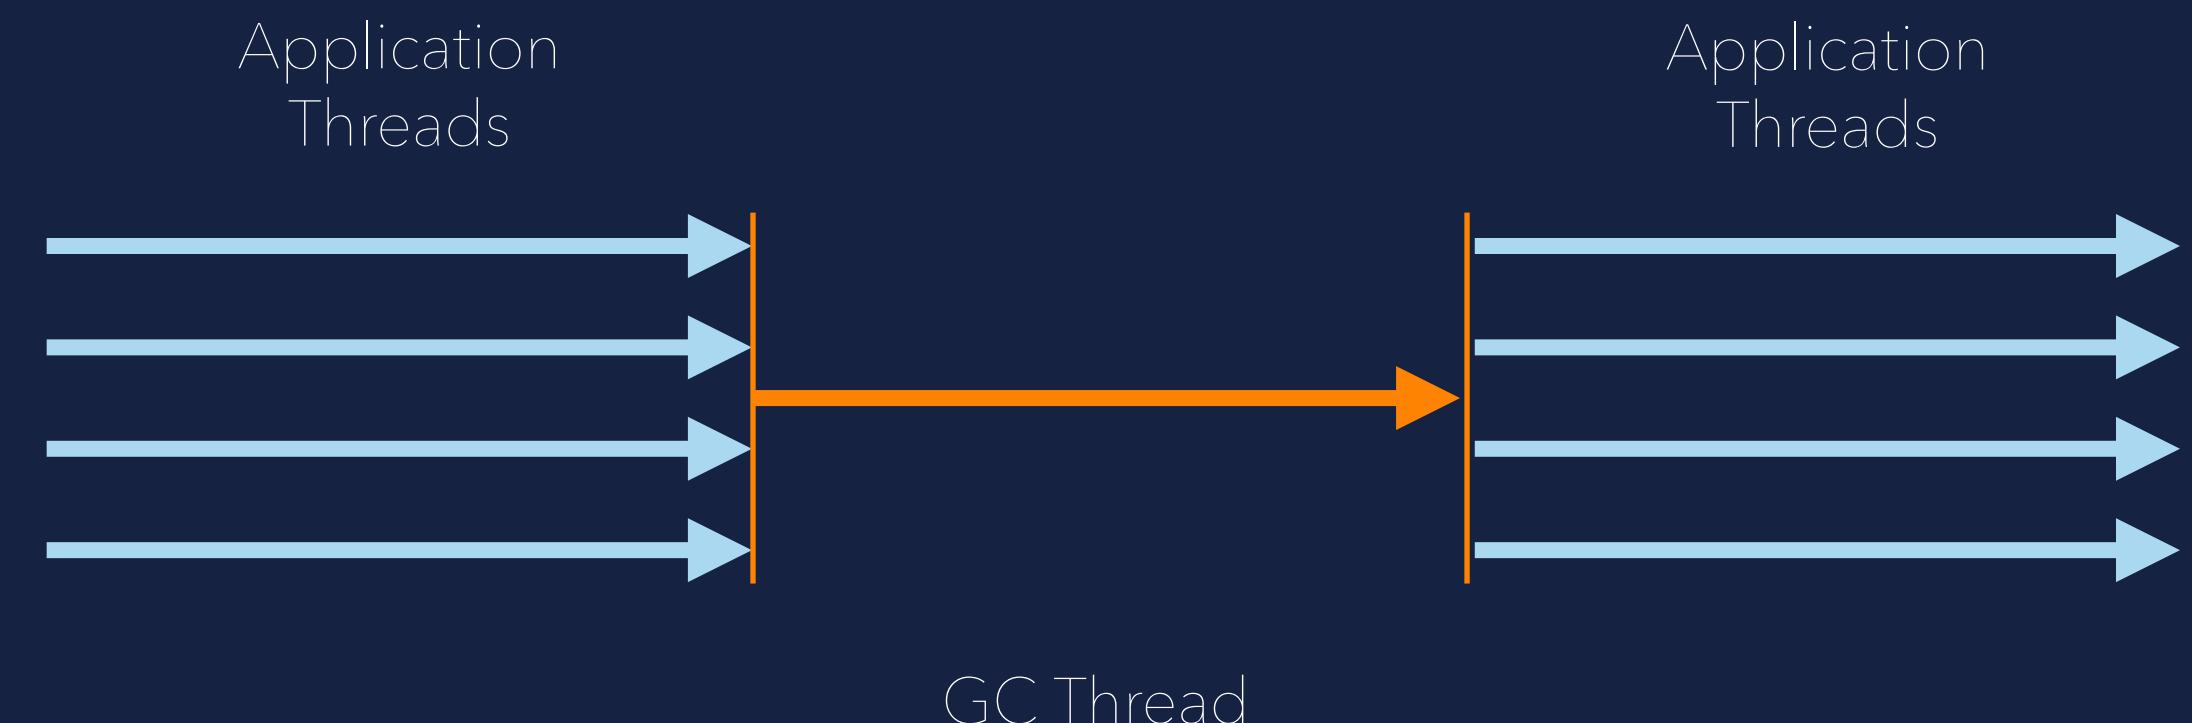
Mark and Compact

`-XX:+UseParallelGC`

Young Generation



Old Generation





PARALLEL

NOTES

Default garbage collector from JDK 5 to JDK 8

Mark and Compact

`-XX:+UseParallelOldGC`

Young Generation

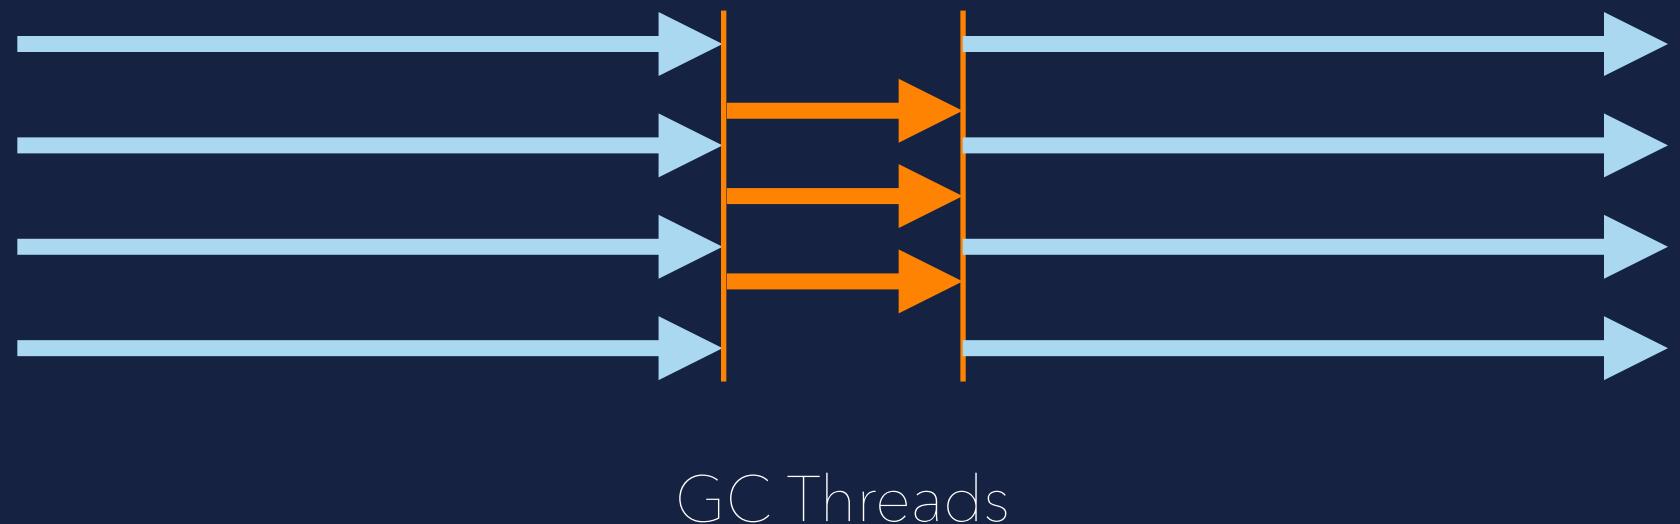
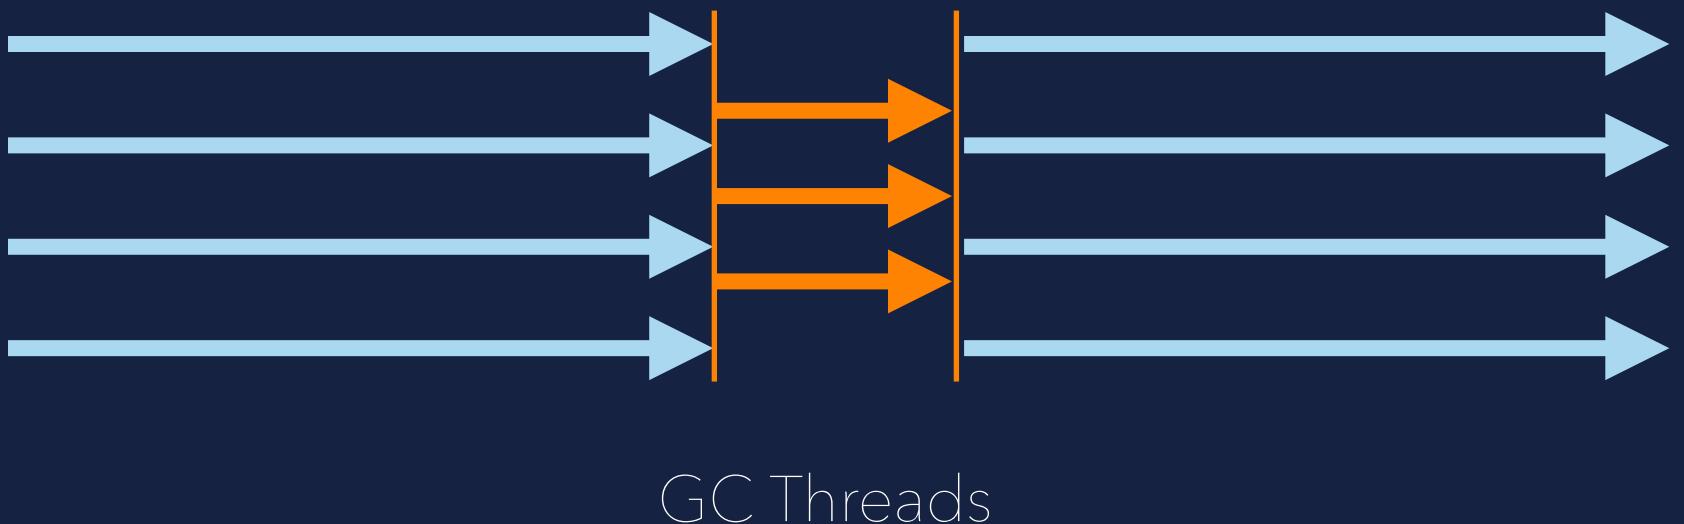
Old Generation

Application
Threads

Application
Threads

Application
Threads

Application
Threads





CMS

Concurrent Mark and Sweep

DEPRECATED

CMS



AVAILABILITY	JDK 1.4 - 13
PARALLEL	YES
CONCURRENT	PARTIALLY
GENERATIONAL	YES
HEAP SIZE	MEDIUM - LARGE
PAUSE TIMES	MODERATE
THROUGHPUT	MODERATE
LATENCY	MODERATE
CPU OVERHEAD	MODERATE (5-15%)

CHOOSE WHEN

- Recycling bin icon Response time is more important than throughput
- Recycling bin icon Pause time must be kept shorter than 1 sec

BEST SUITED FOR

- Recycling bin icon Web applications
- Recycling bin icon Medium sized enterprise systems

OS SUPPORT



JVM SWITCH

> java -XX:+UseConcMarkSweepGC

CMS

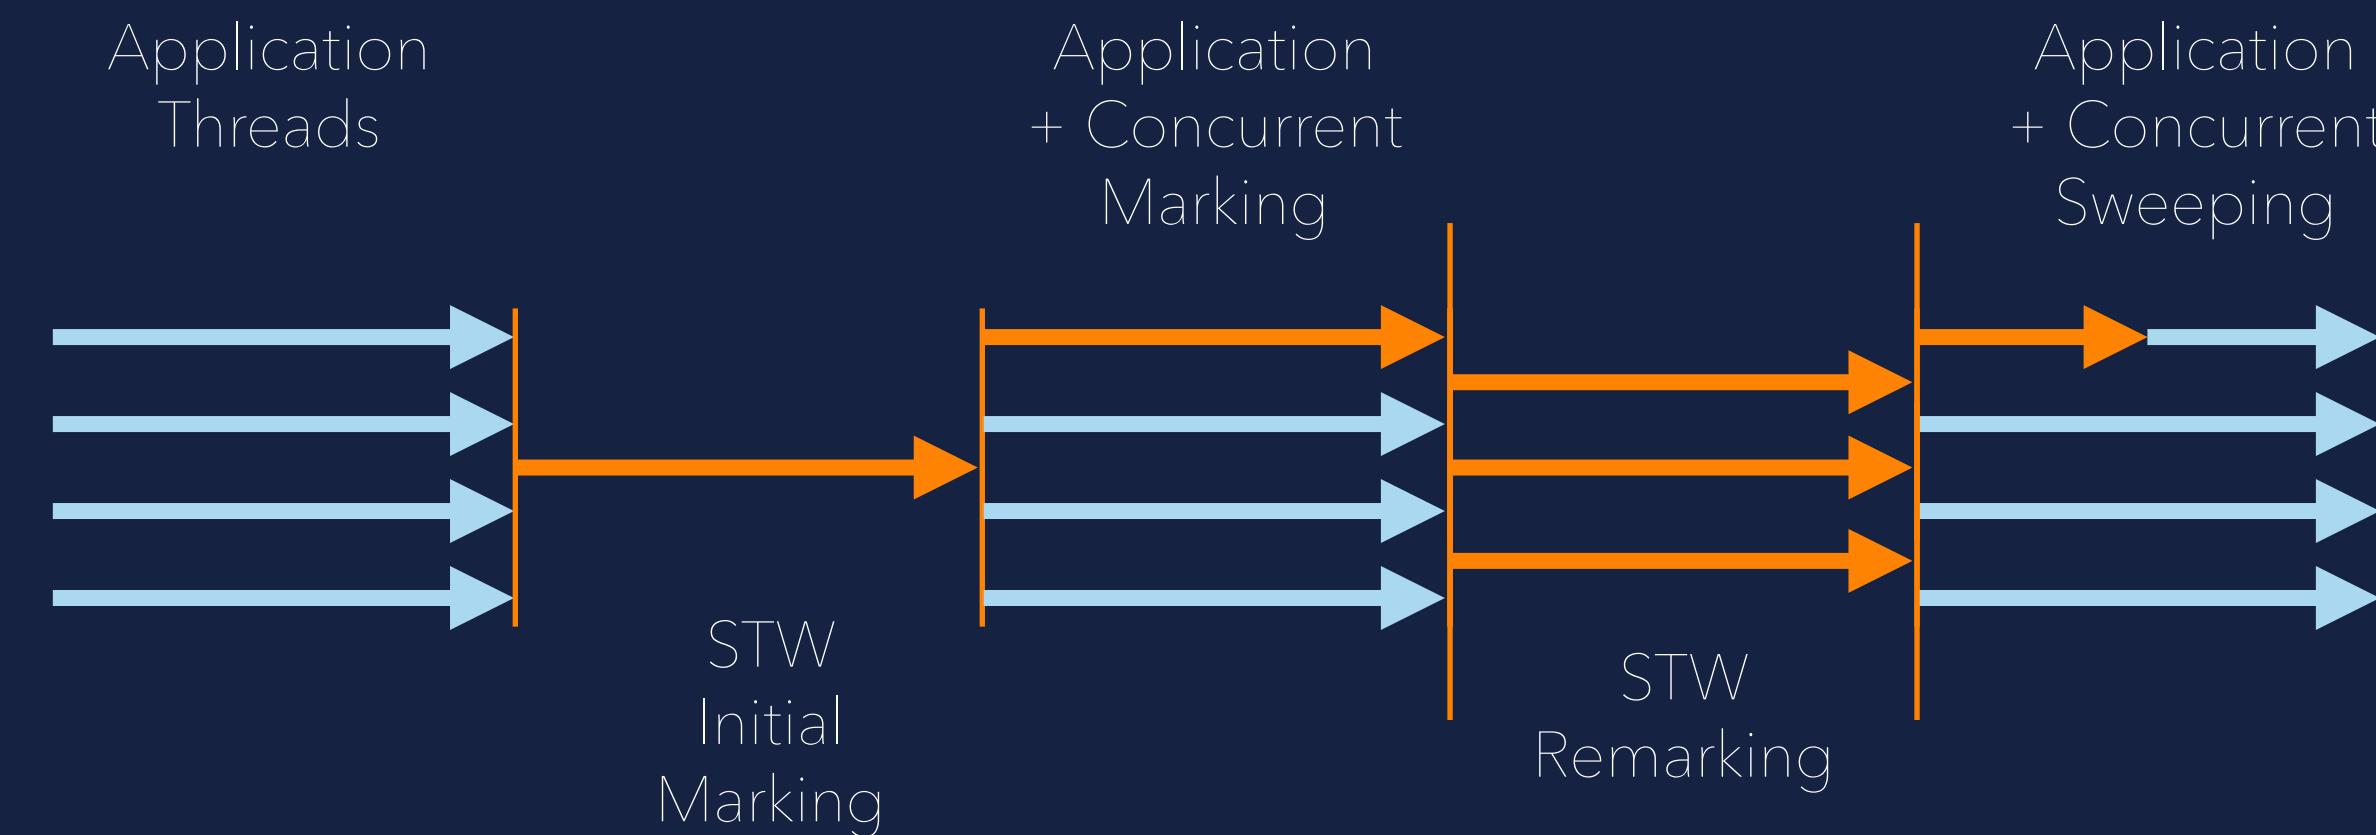


NOTES

☒ Deprecated as of JDK 9

☒ Removed from JDK 14

☒ Concurrent marking but no compaction -> Fragmentation





G1

Garbage First

azul



Heap-Layout

Region size 1 - 32 MB

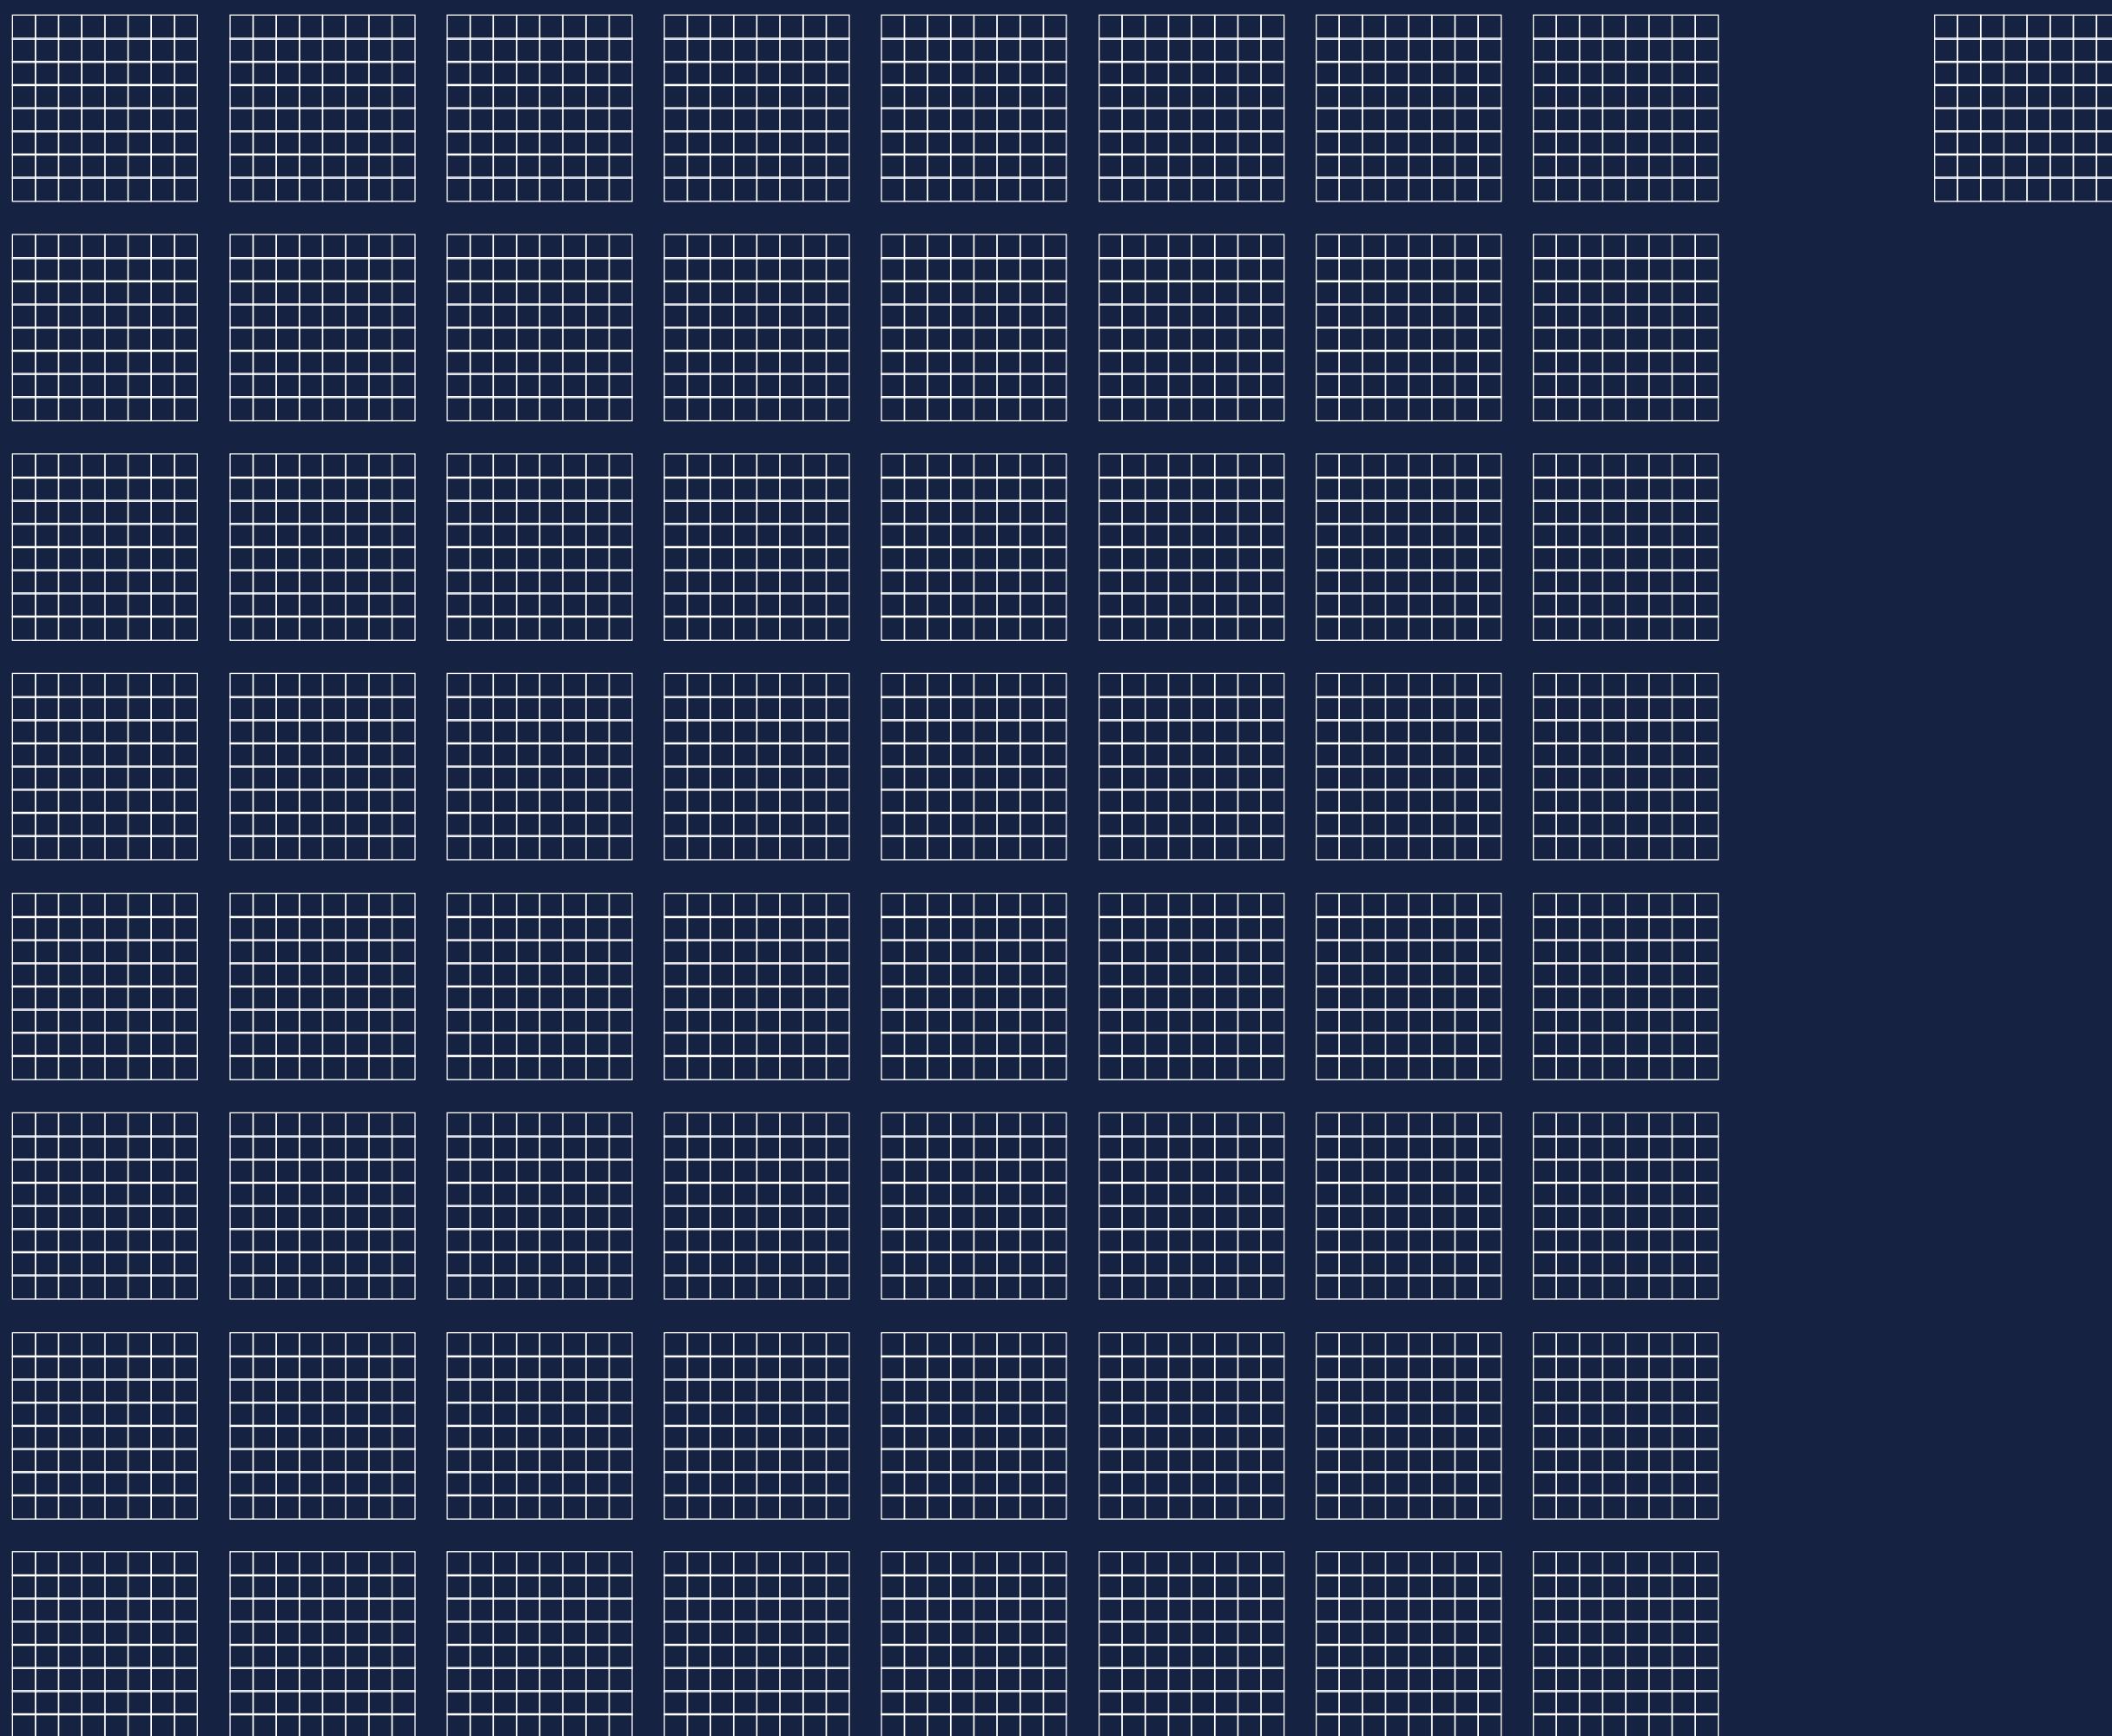
Max no. of region <= 2048

Heap	Region
< 4 GB	- 1 MB
< 8 GB	- 2 MB
< 16 GB	- 4 MB
< 32 GB	- 8 MB
< 64 GB	- 16 MB
> 64 GB	- 32 MB

Example 8GB Heap:

8 GB Heap = 8192 MB

8192 MB / 2048 = 4 MB region size



Unassigned region



Heap-Layout

Region size 1 - 32 MB

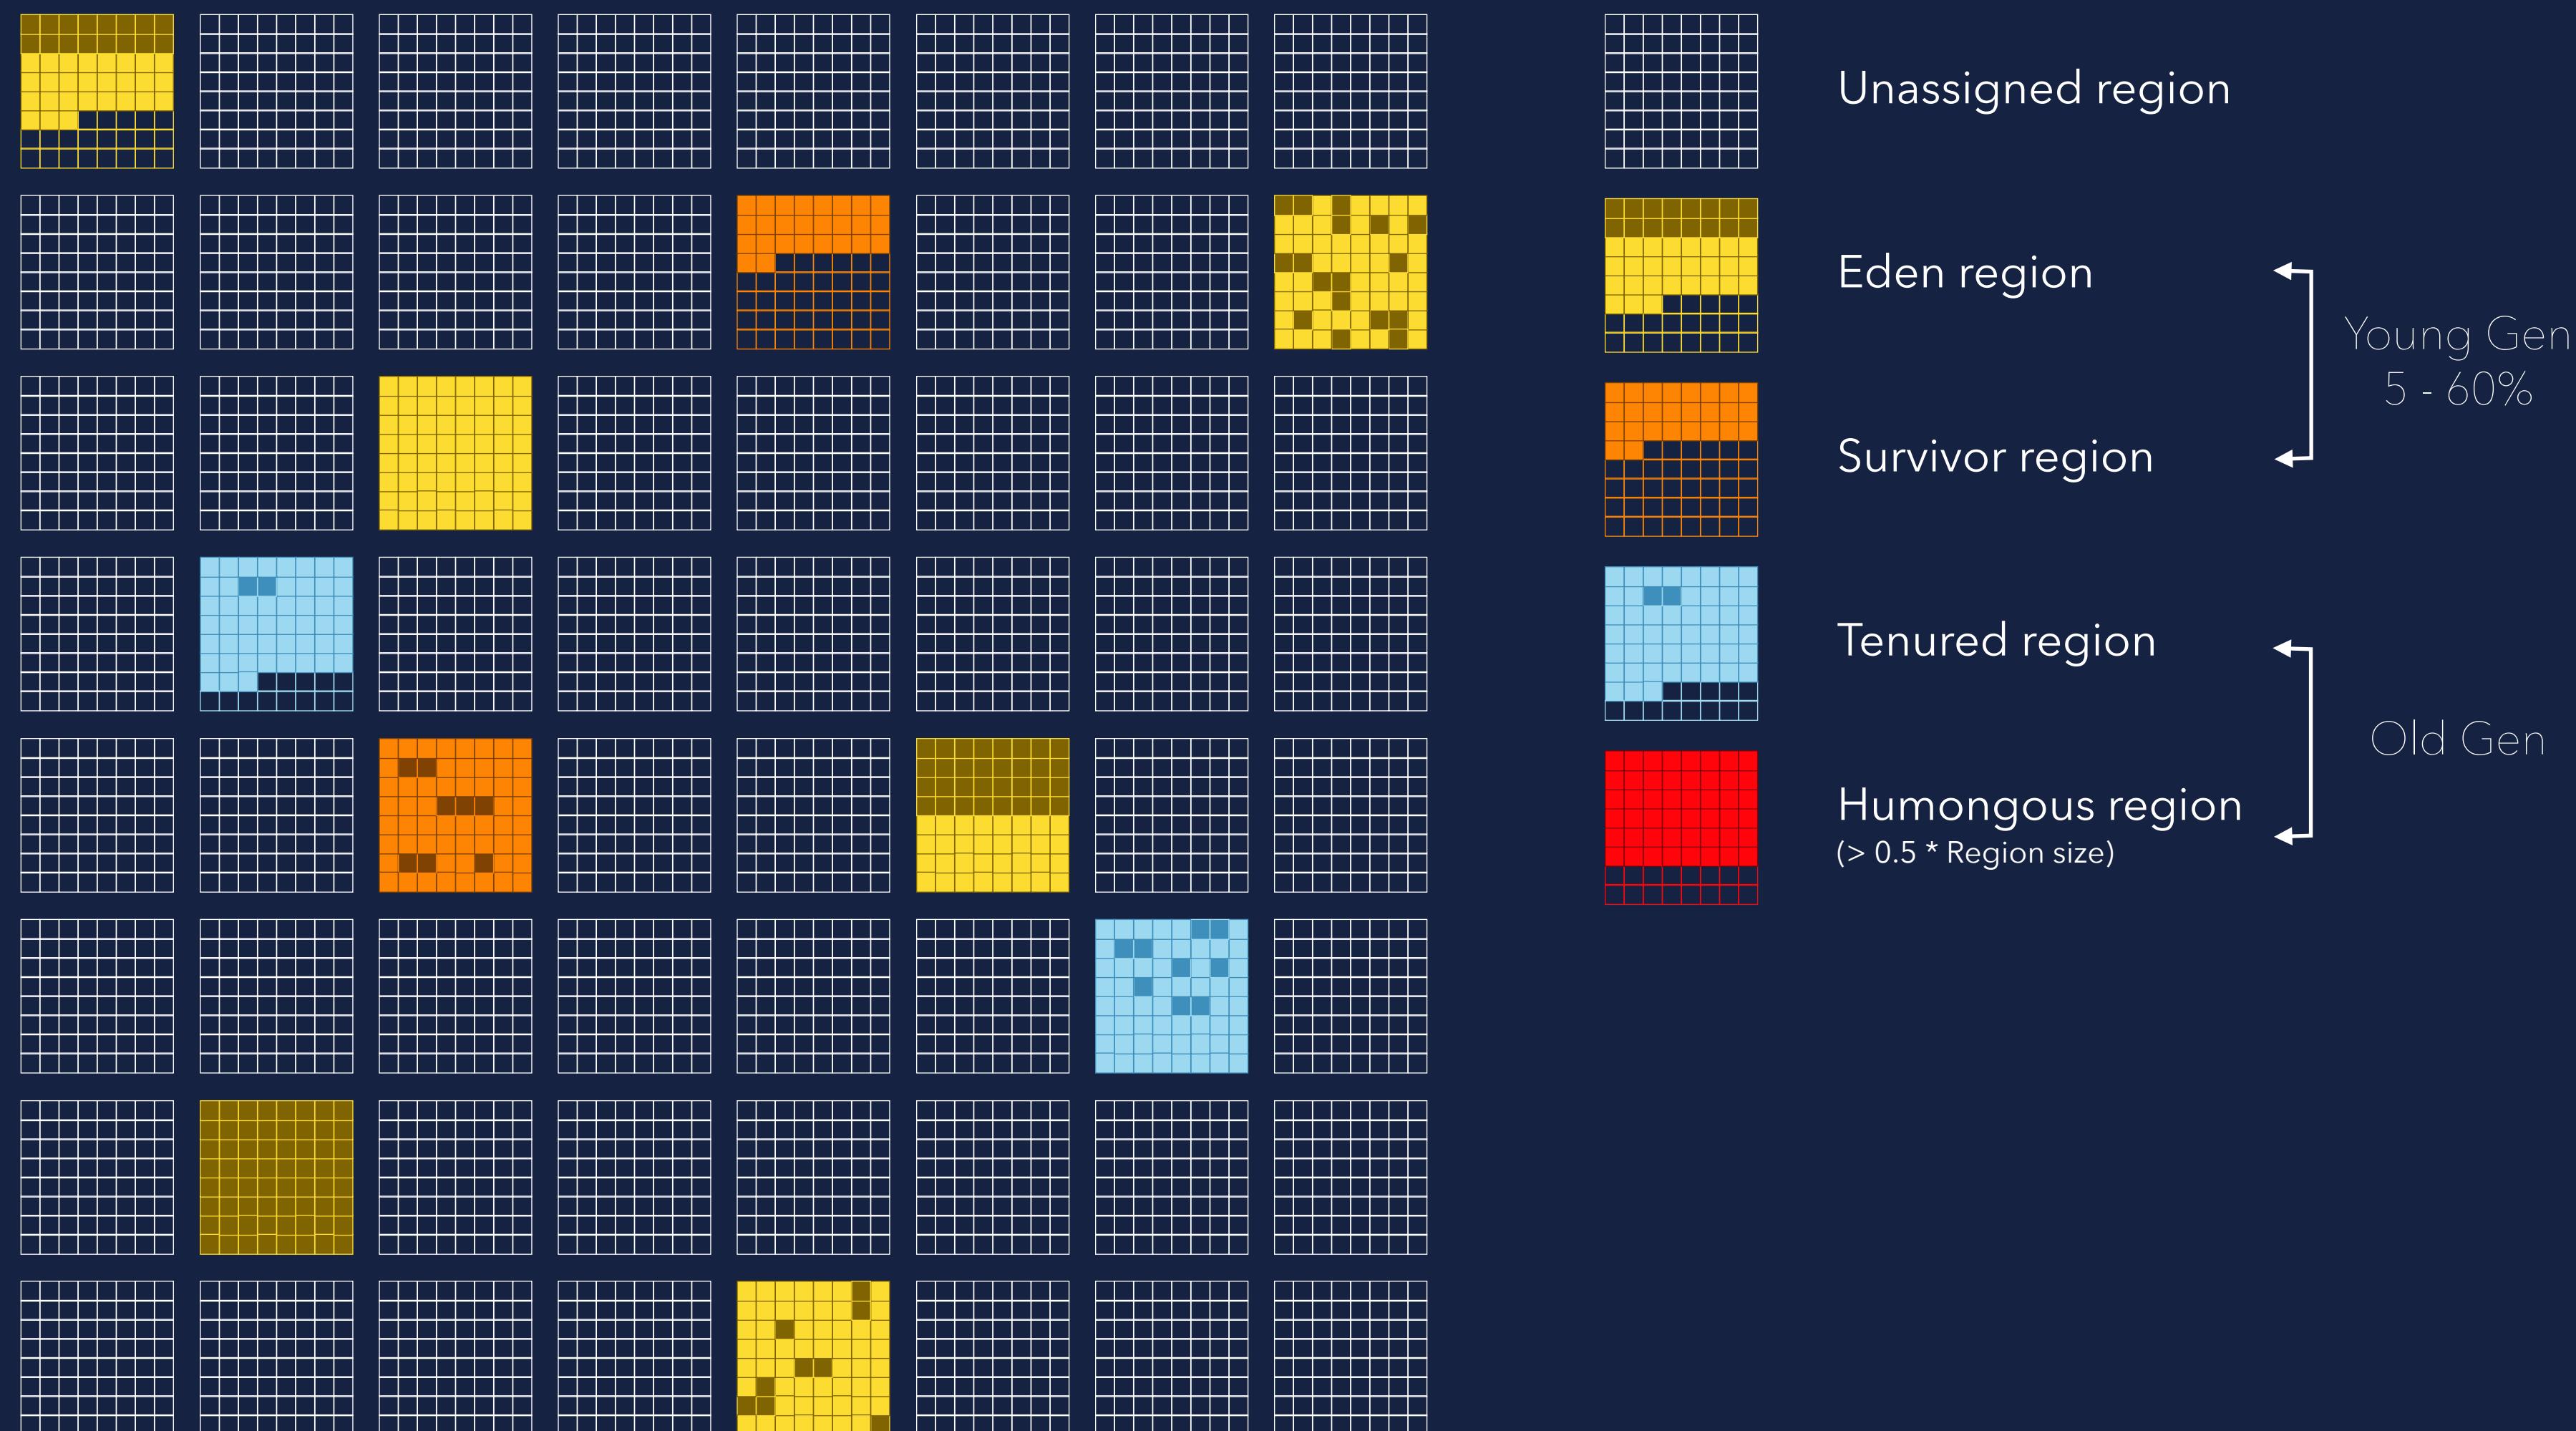
Max no. of region <= 2048

Heap	Region
< 4 GB	- 1 MB
< 8 GB	- 2 MB
< 16 GB	- 4 MB
< 32 GB	- 8 MB
< 64 GB	- 16 MB
> 64 GB	- 32 MB

Example 8GB Heap:

8 GB Heap = 8192 MB

8192 MB / 2048 = 4 MB region size





Heap-Layout

Region size 1 - 32 MB

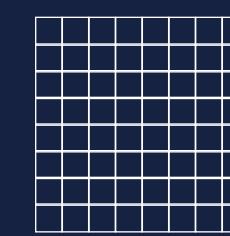
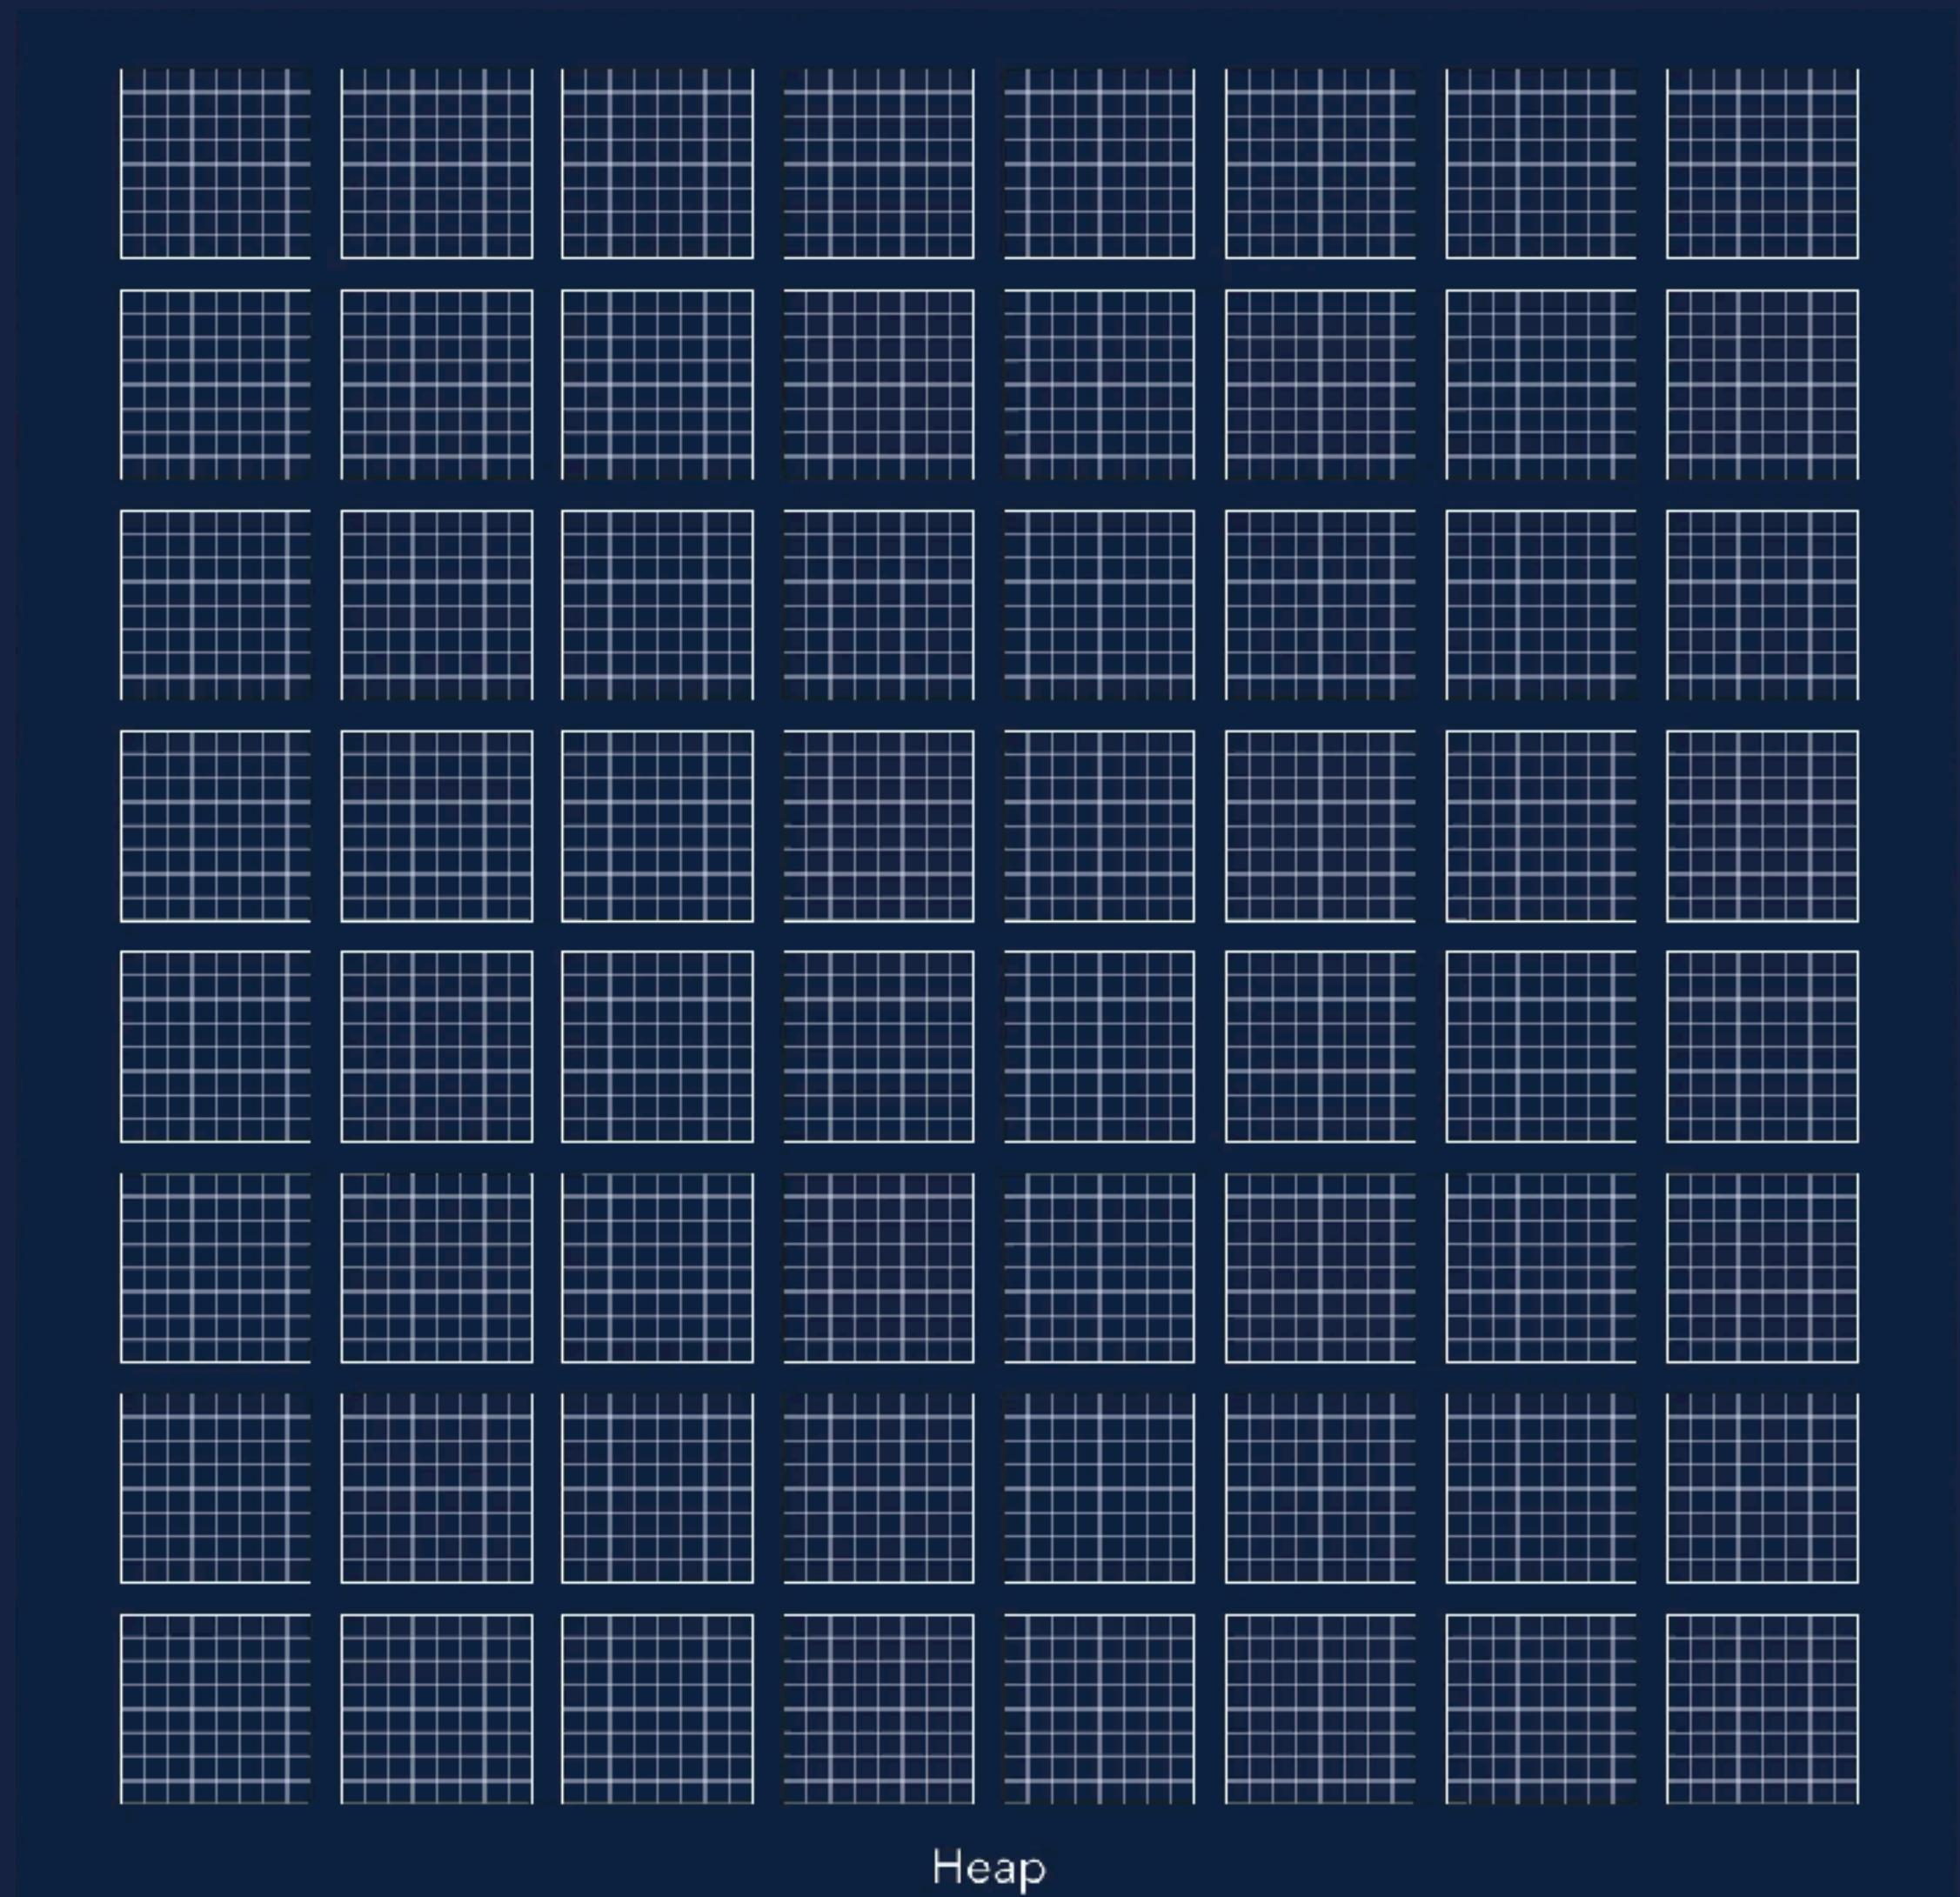
Max no. of region <= 2048

Heap	Region
< 4 GB	- 1 MB
< 8 GB	- 2 MB
< 16 GB	- 4 MB
< 32 GB	- 8 MB
< 64 GB	- 16 MB
> 64 GB	- 32 MB

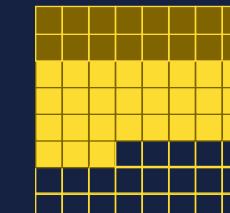
Example 8GB Heap:

8 GB Heap = 8192 MB

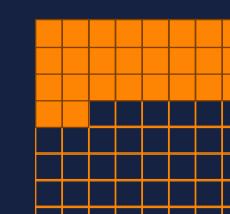
8192 MB / 2048 = 4 MB region size



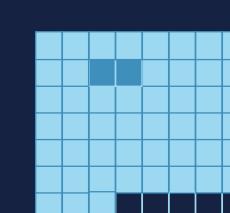
Unassigned region



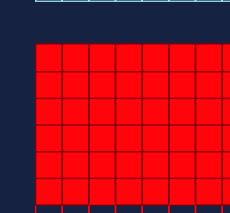
Eden region



Survivor region



Tenured region



Humongous region
(> 0.5 * Region size)

Young Gen
5 - 60%

Old Gen

Example:
6 Eden Regions
3 Survivor Regions

2 Regions with most garbage will
be collected/promoted



AVAILABILITY

JDK 7U4+

PARALLEL

YES

CONCURRENT

PARTIALLY

GENERATIONAL

YES

HEAP SIZE

MEDIUM - LARGE

PAUSE TIMES

SHORT - MEDIUM

THROUGHPUT

HIGH

LATENCY

LOWER

CPU OVERHEAD

MODERATE (5-15%)

CHOOSE WHEN

- Response time is more important than throughput
- Pause time should be around 200 ms
- Heap size is not larger than 16-32 GB

BEST SUITED FOR

- Mixed workloads
- Large sized enterprise systems
- Responsive in medium to large heaps

OS SUPPORT



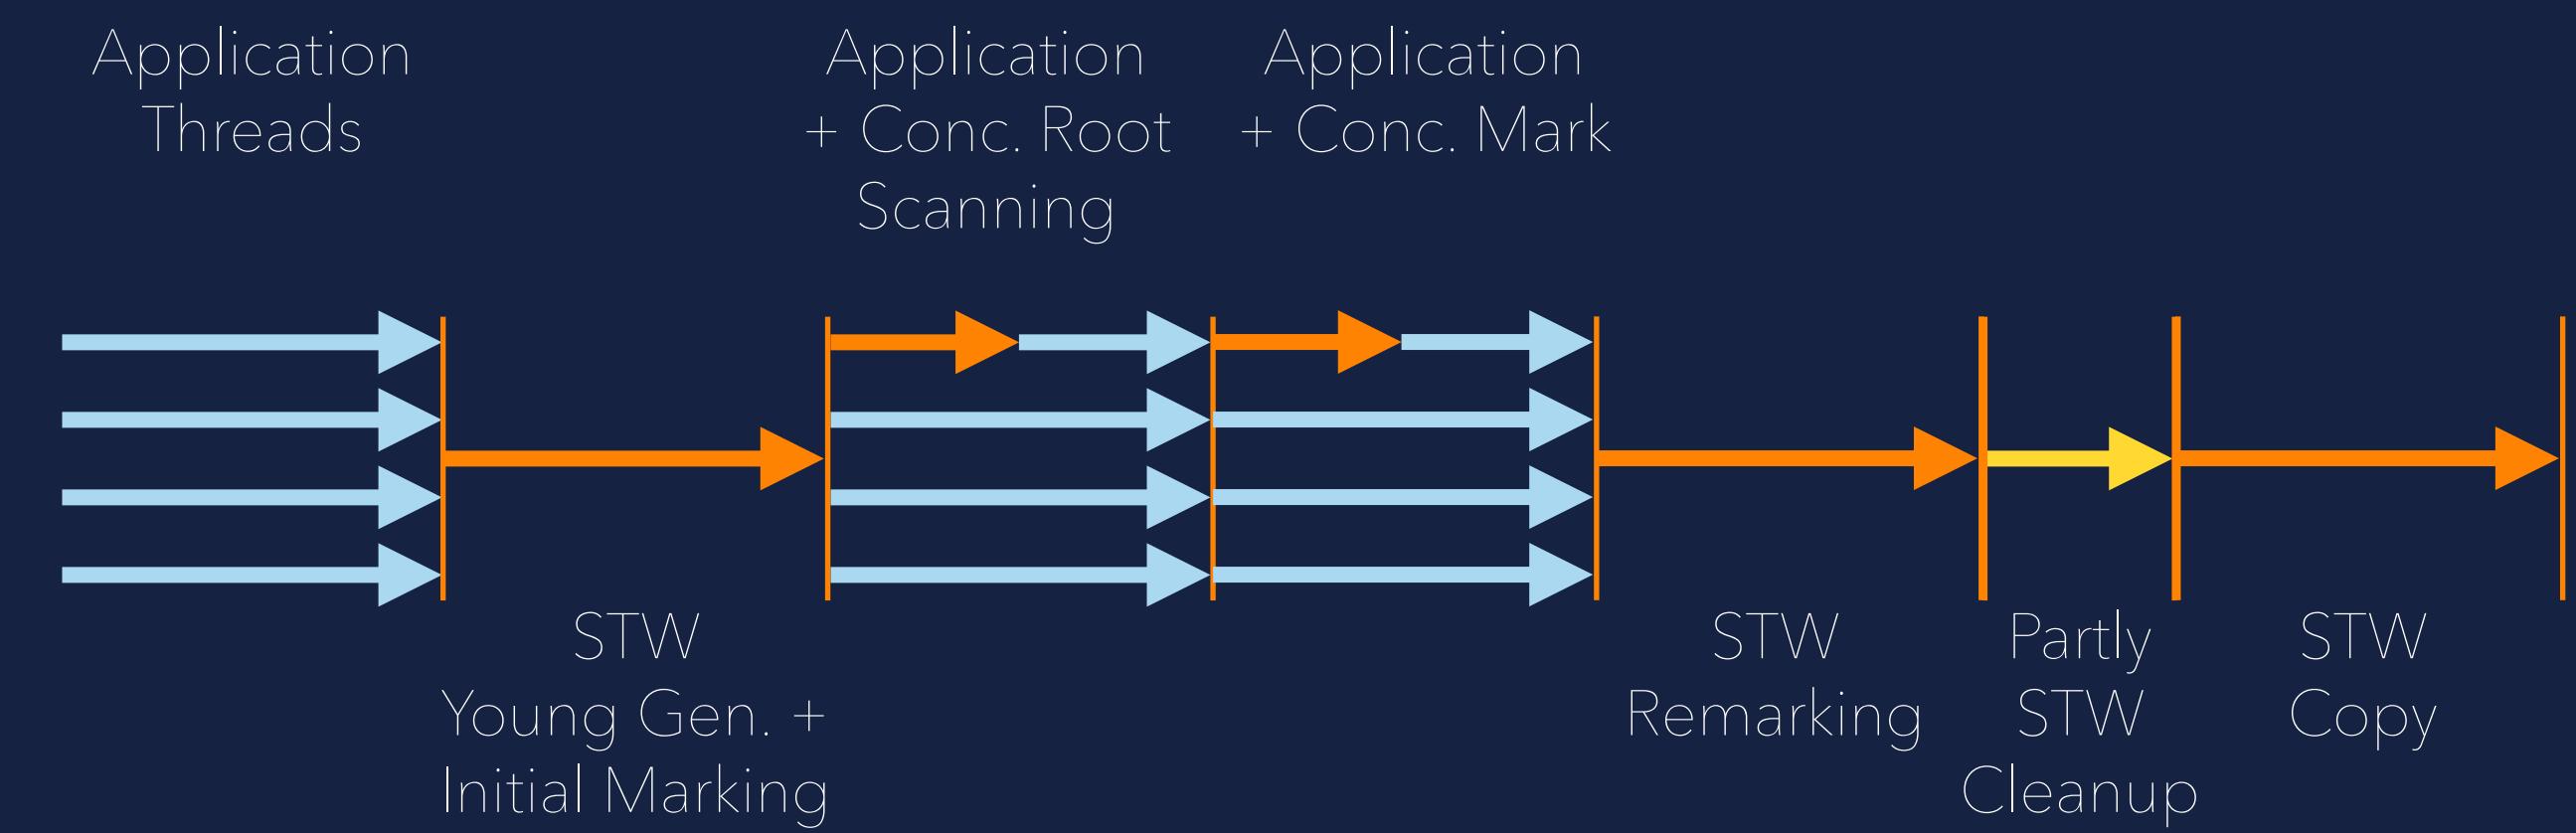
JVM SWITCH

> java -XX:+UseG1GC

NOTES

 Default collector from JDK 9 onwards

 Concurrent marking





EPSILON

azul



EPSILON

AVAILABILITY

JDK 11+

PARALLEL

-

CONCURRENT

-

GENERATIONAL

-

HEAP SIZE

-

PAUSE TIMES

-

THROUGHPUT

VERY HIGH

LATENCY

VERY LOW

CPU OVERHEAD

VERY LOW

CHOOSE WHEN

- Testing performance or memory pressure
- Highest performance is needed and nearly no garbage is created

BEST SUITED FOR

- Extremely short lived jobs
- Last drop latency improvements
- Last drop throughput improvements

OS SUPPORT



JVM SWITCH

```
> java -XX:+UnlockExperimentalVMOptions -XX:+UseEpsilonGC
```



SHENANDOAH

azul



SHENANDOAH

AVAILABILITY

JDK 11.0.9+ / JDK 24

PARALLEL

YES

CONCURRENT

FULLY

GENERATIONAL

NO / YES

HEAP SIZE

MEDIUM - LARGE

PAUSE TIMES

SHORT

THROUGHPUT

VERY HIGH

LATENCY

VERY LOW

CPU OVERHEAD

MODERATE (10-20%)

CHOOSE WHEN

- Response time is a high priority
- Using a very large heap (100GB+)
- Predictable response times needed

BEST SUITED FOR

- Latency sensitive applications
- Large scale systems
- Highly concurrent applications

OS SUPPORT



JVM SWITCH

> `java -XX:+UseShenandoahGC`



SHENANDOAH

NOTES

- Not available in Oracle JDK
- A bit reduced throughput due to concurrent GC
- Makes use of new barrier concept, load reference barrier
- First generational version in JDK 24



ZGC

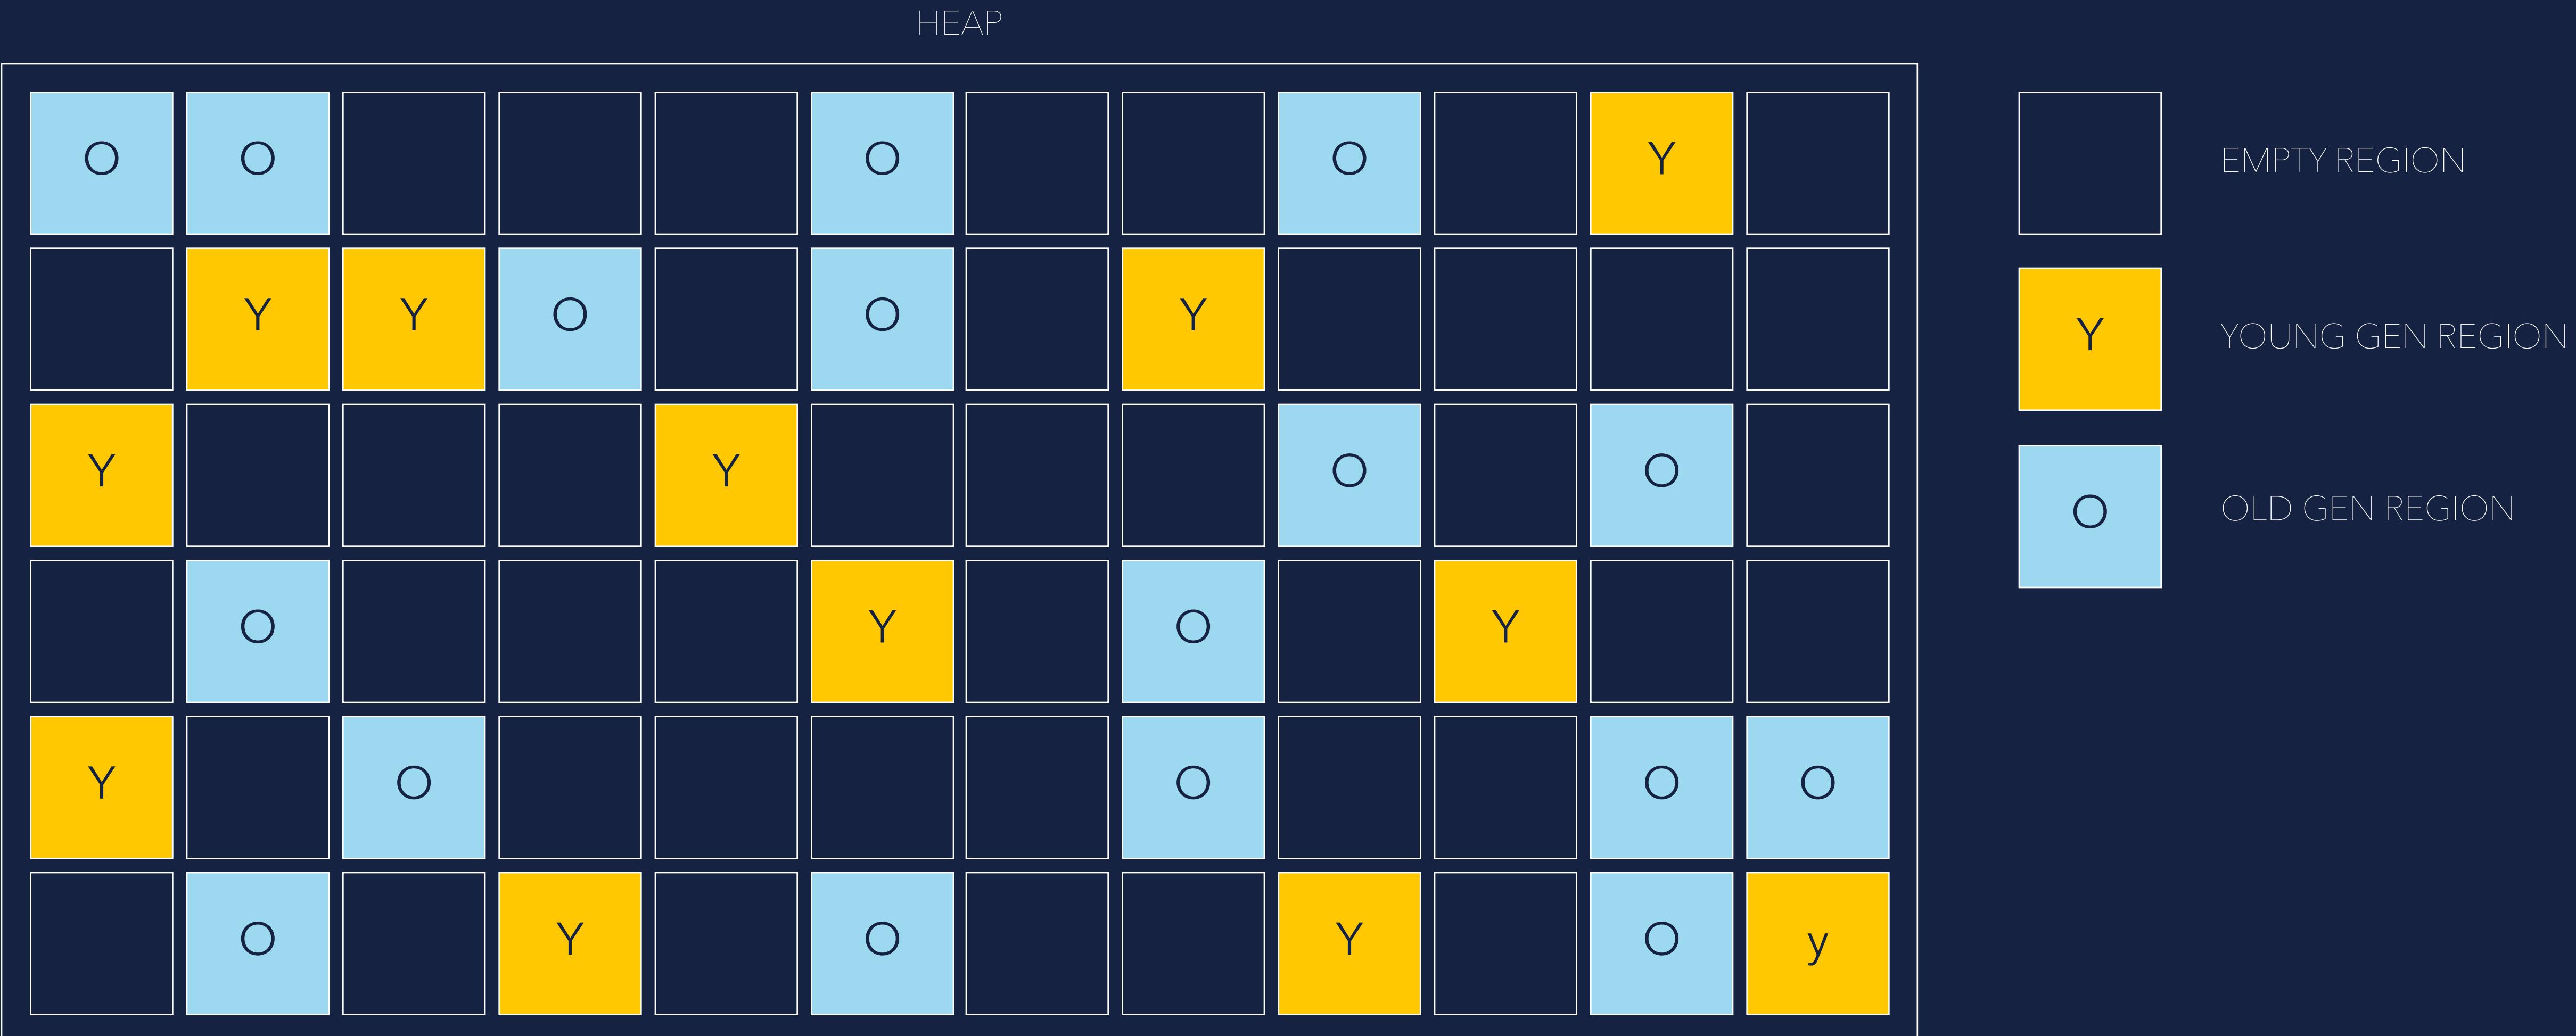
Z Garbage Collector

azul



ZGC

Heap-Layout



ZGC



AVAILABILITY

JDK 15-23 / JDK 21+

PARALLEL

YES

CONCURRENT

FULLY

GENERATIONAL

NO / YES

HEAP SIZE

LARGE

PAUSE TIMES

SHORT

THROUGHPUT

VERY HIGH

LATENCY

VERY LOW

CPU OVERHEAD

MODERATE (10-20%)

CHOOSE WHEN

- Response time is a high priority
- Using a very large heap (100GB+)
- Predictable response times needed

BEST SUITED FOR

- Low latency sensitive applications
- Large scale systems
- Highly concurrent applications

OS SUPPORT



JVM SWITCH

> `java -XX:+UseZGC`



NOTES



Non-generational version removed in JDK 24



C4

Concurrent Continues Compacting Collector

azul



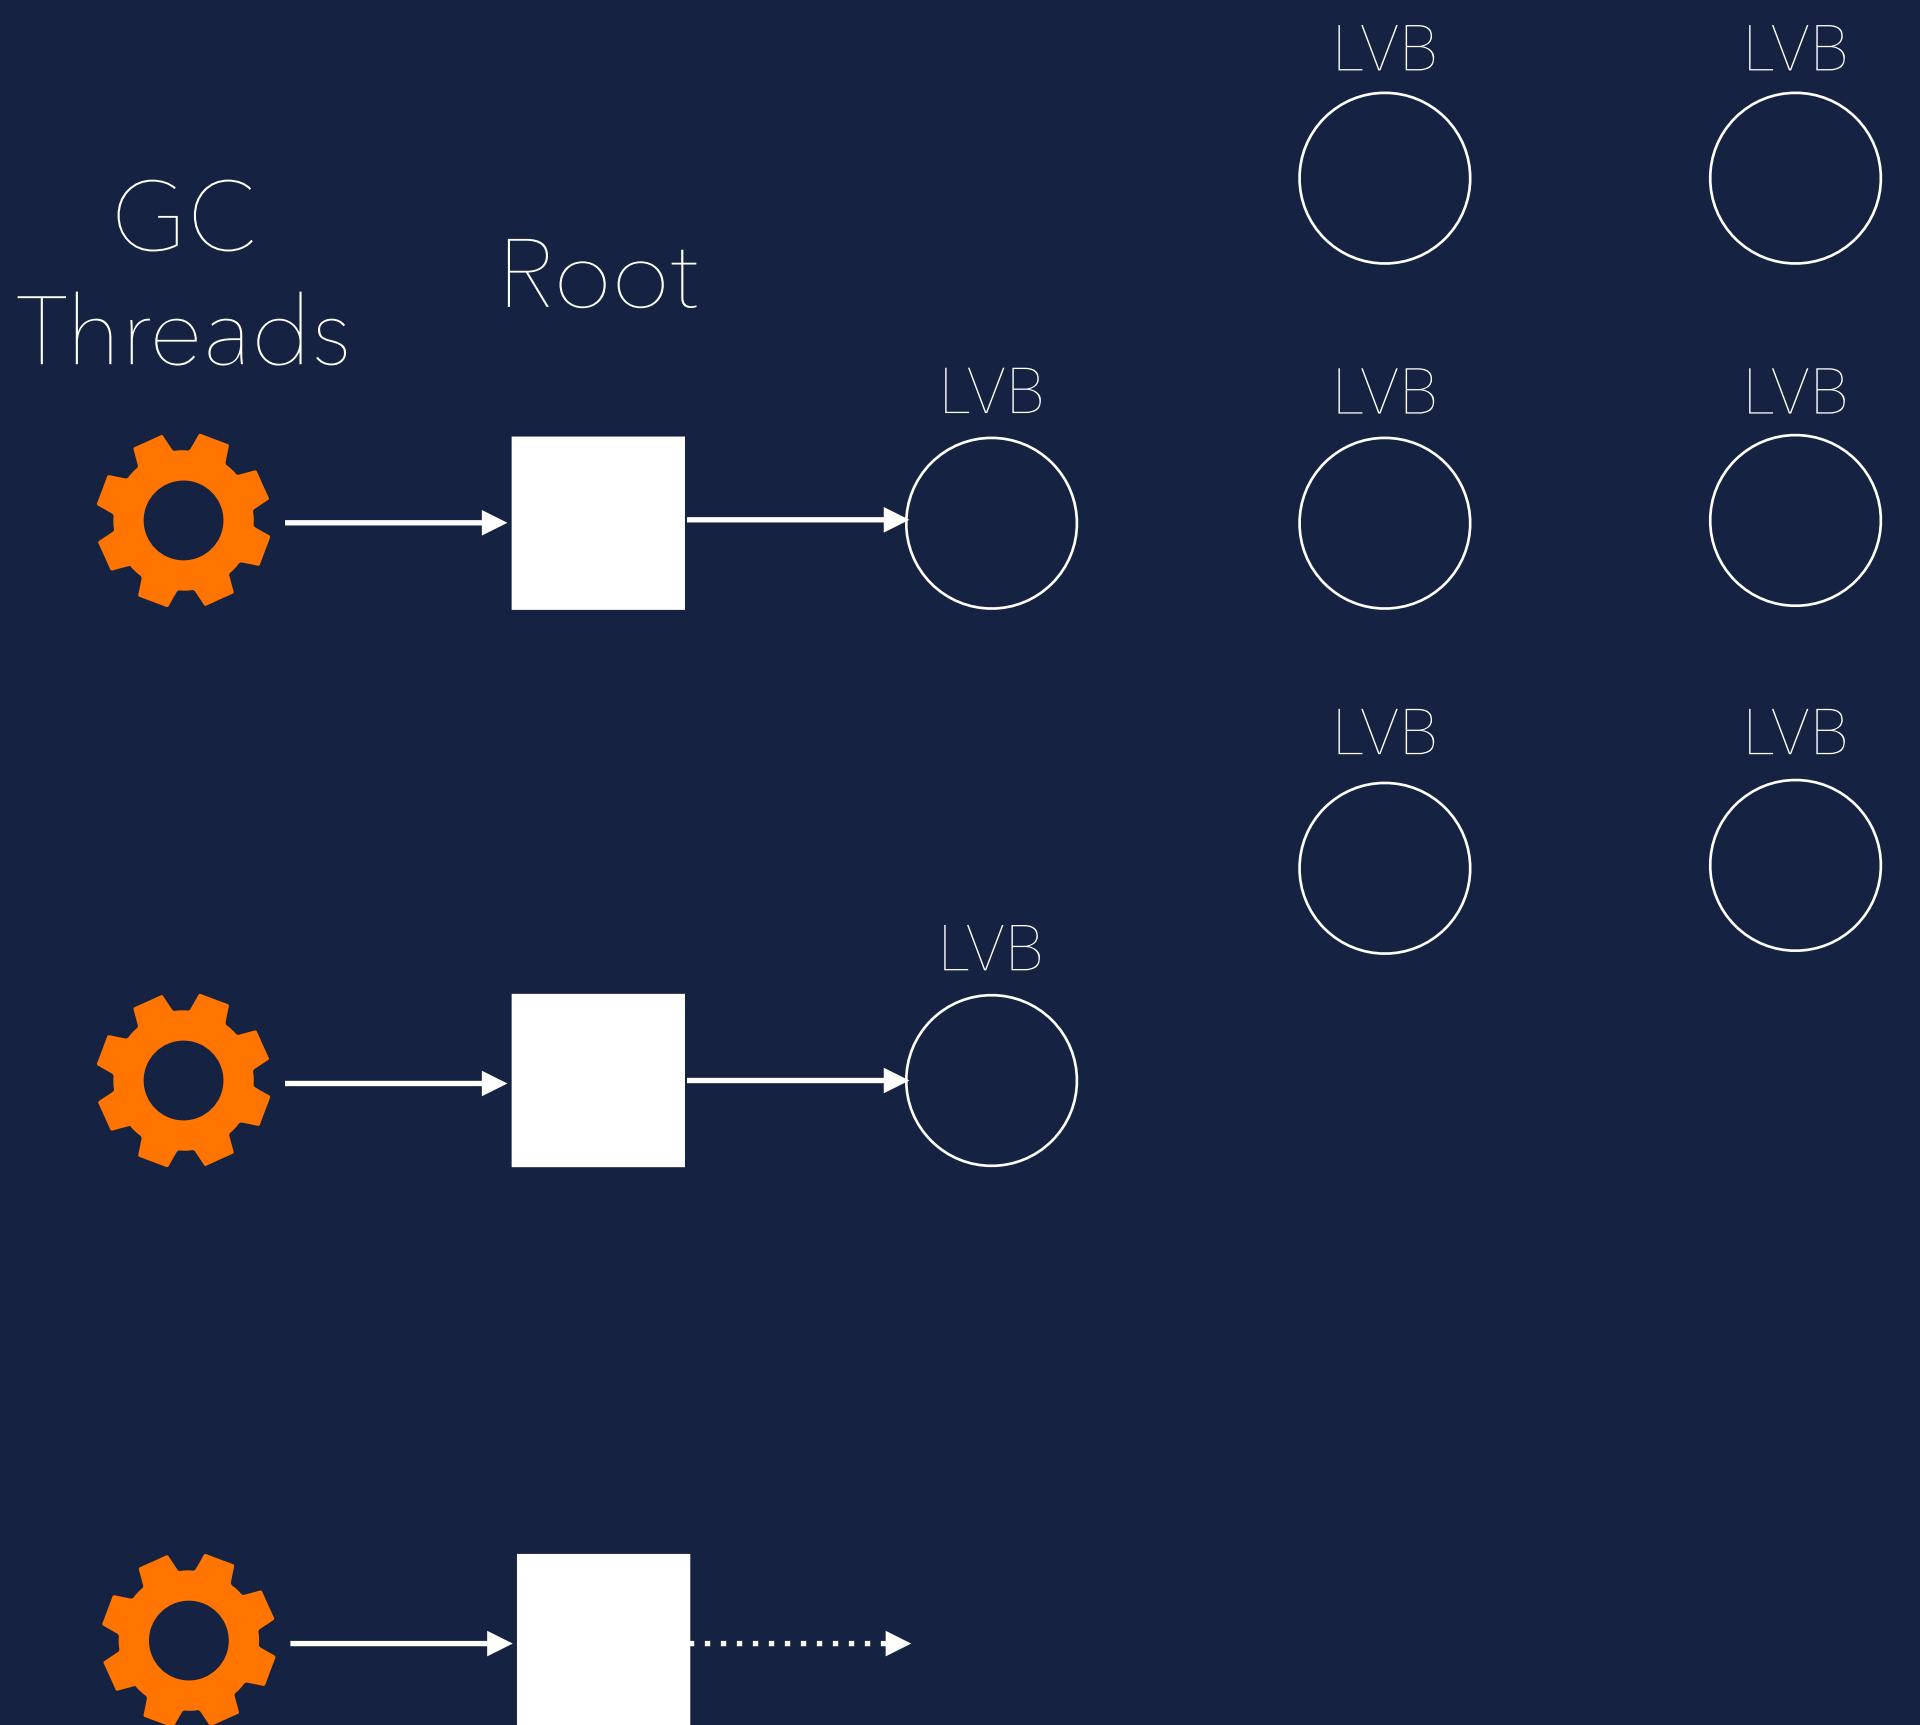
NOTES

- Part of Azul Zing JVM
- Makes use of Loaded Value Barrier (LVB) everywhere
(Test + Jump which only takes 1 cpu cycle -> very fast)
- LVB is read and write barrier
(guaranteed to be hit on every access)
- Best performance by using Transparent Huge Pages
(Normal page size 4kB, THP size 2MB)

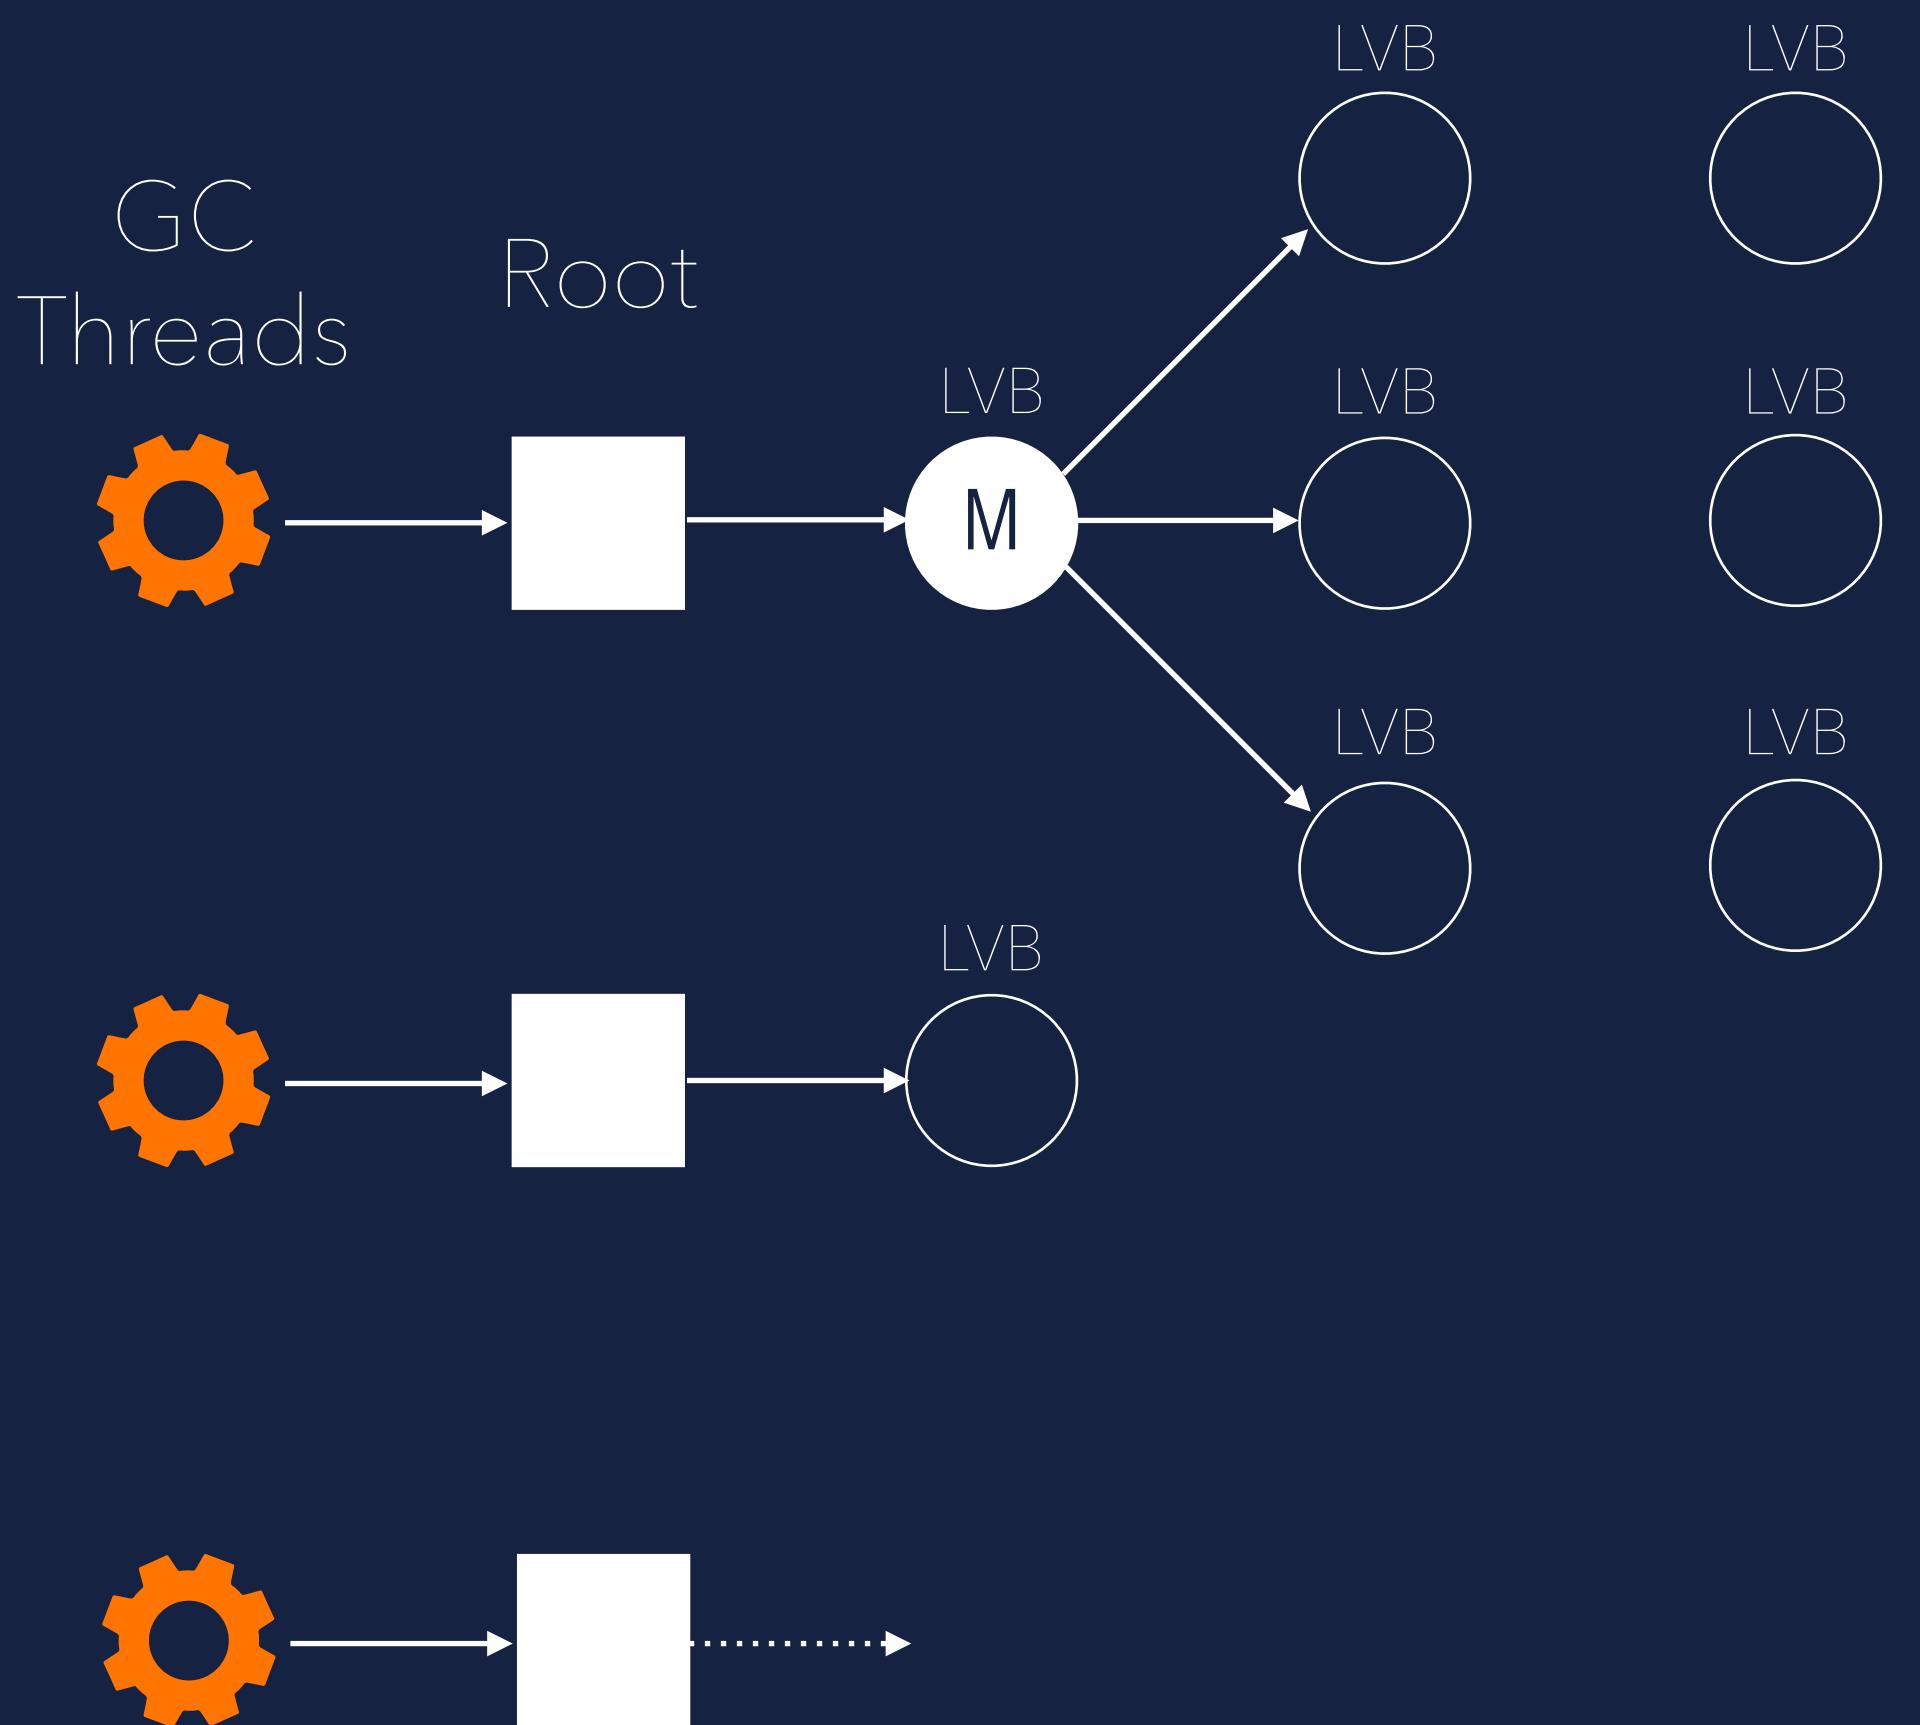


MARKING PHASE

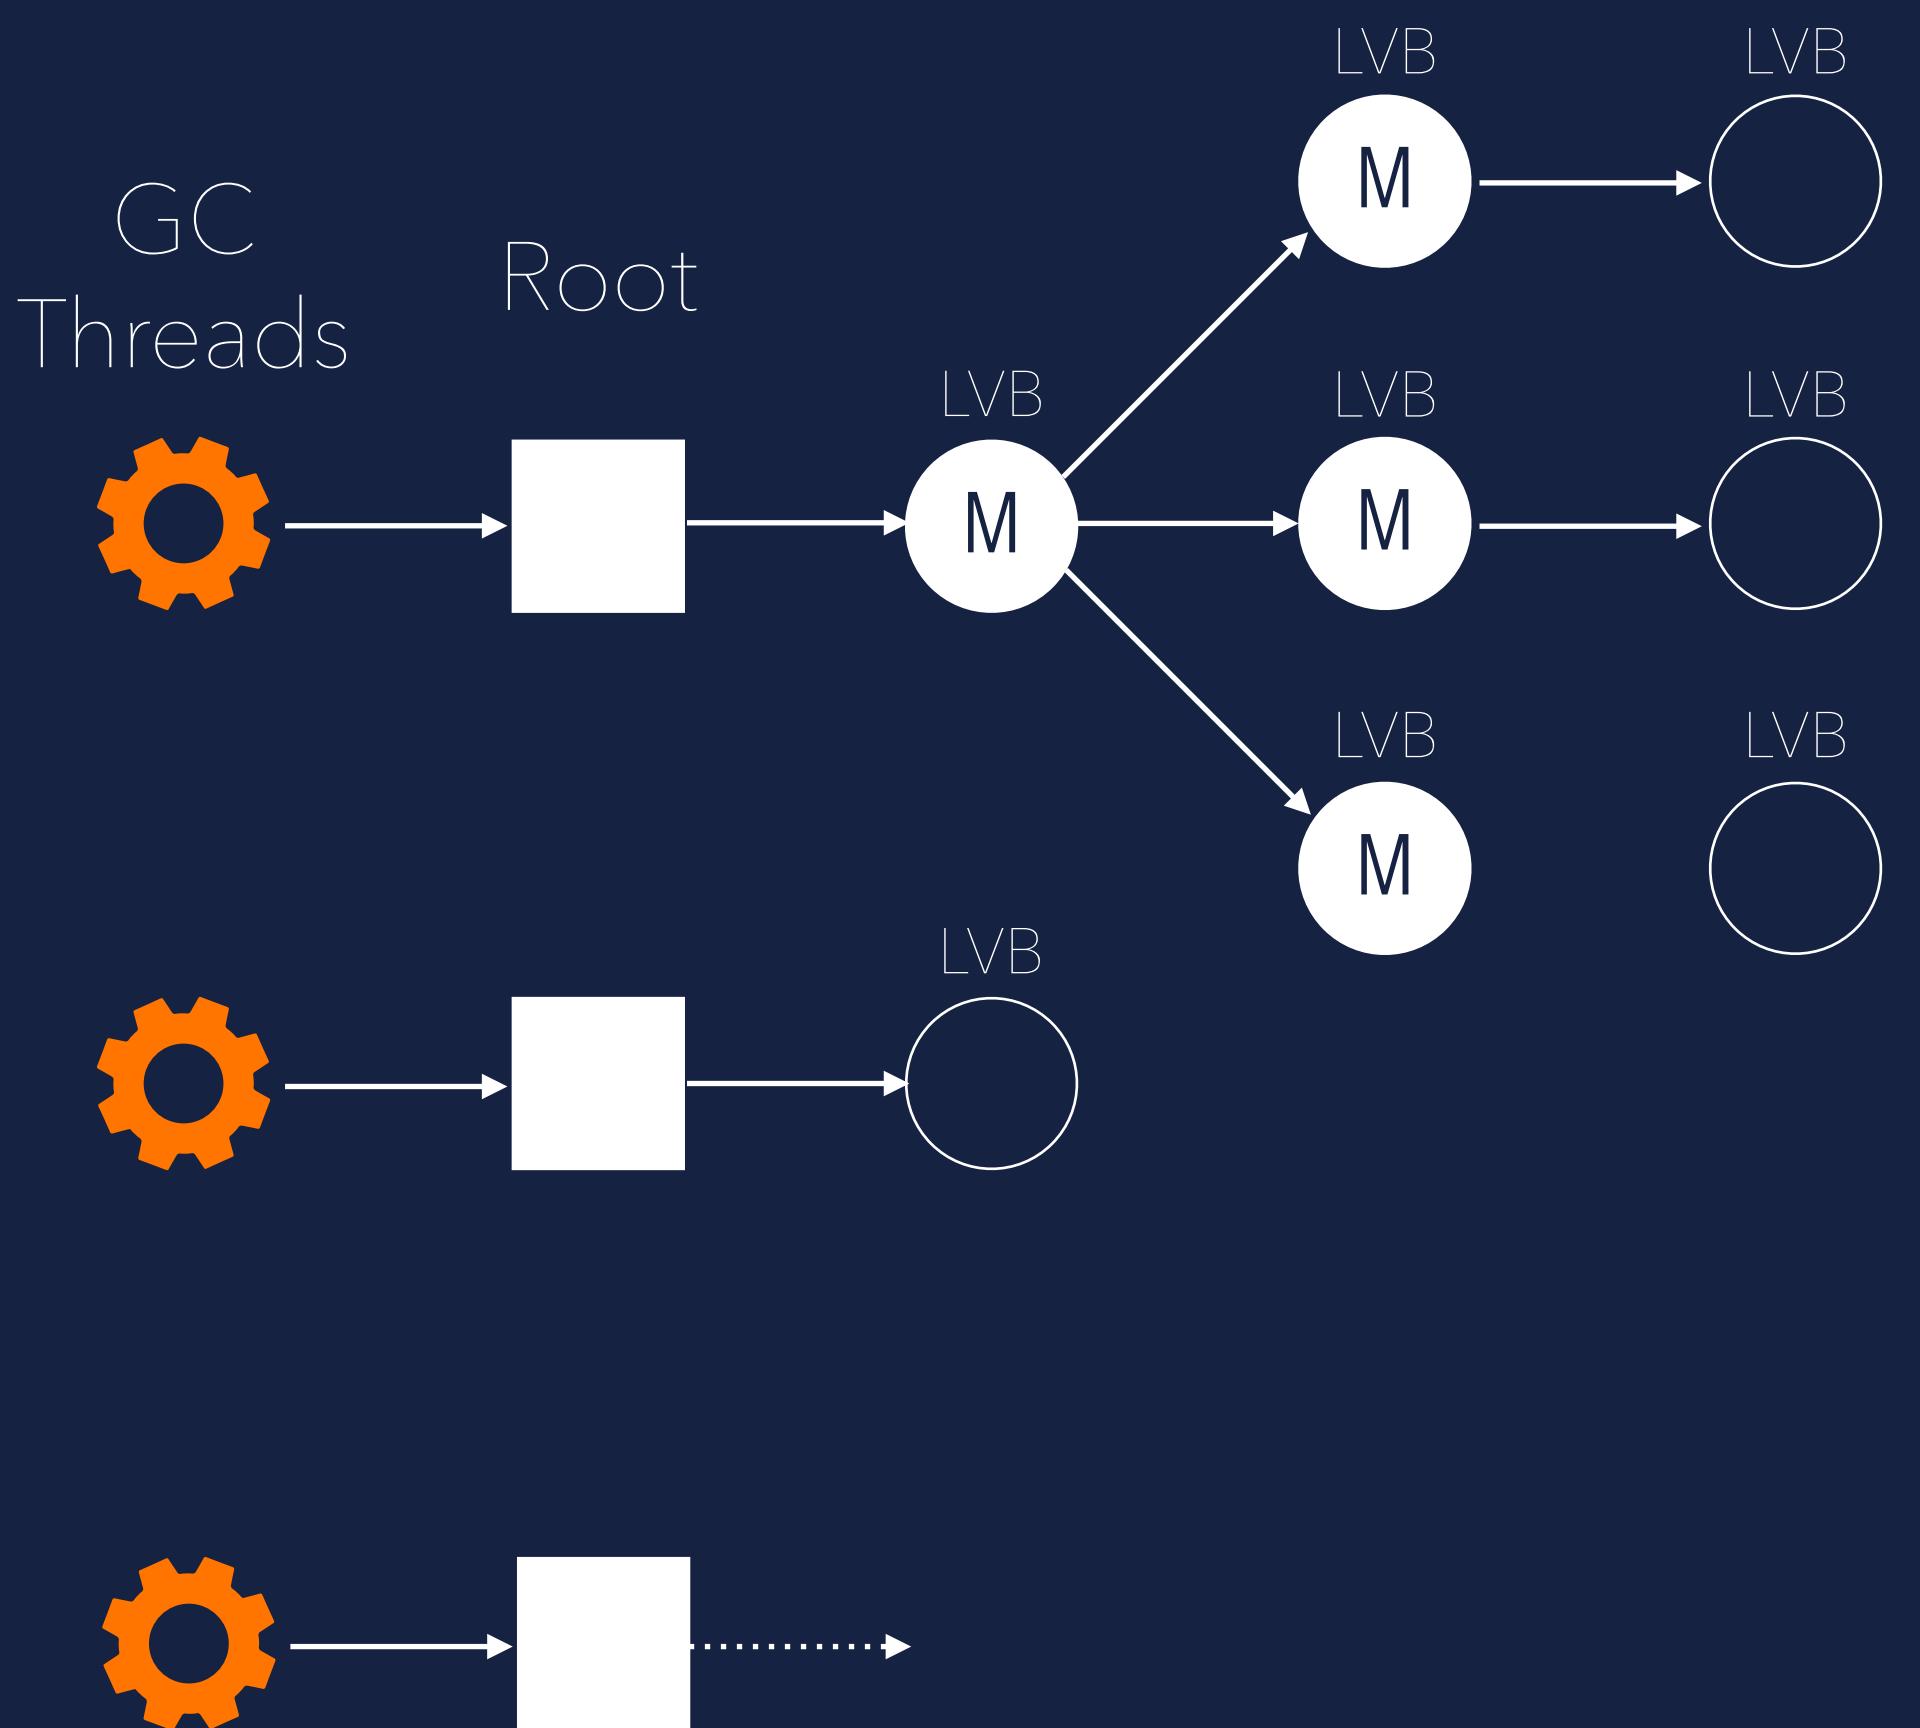
Marking Phase



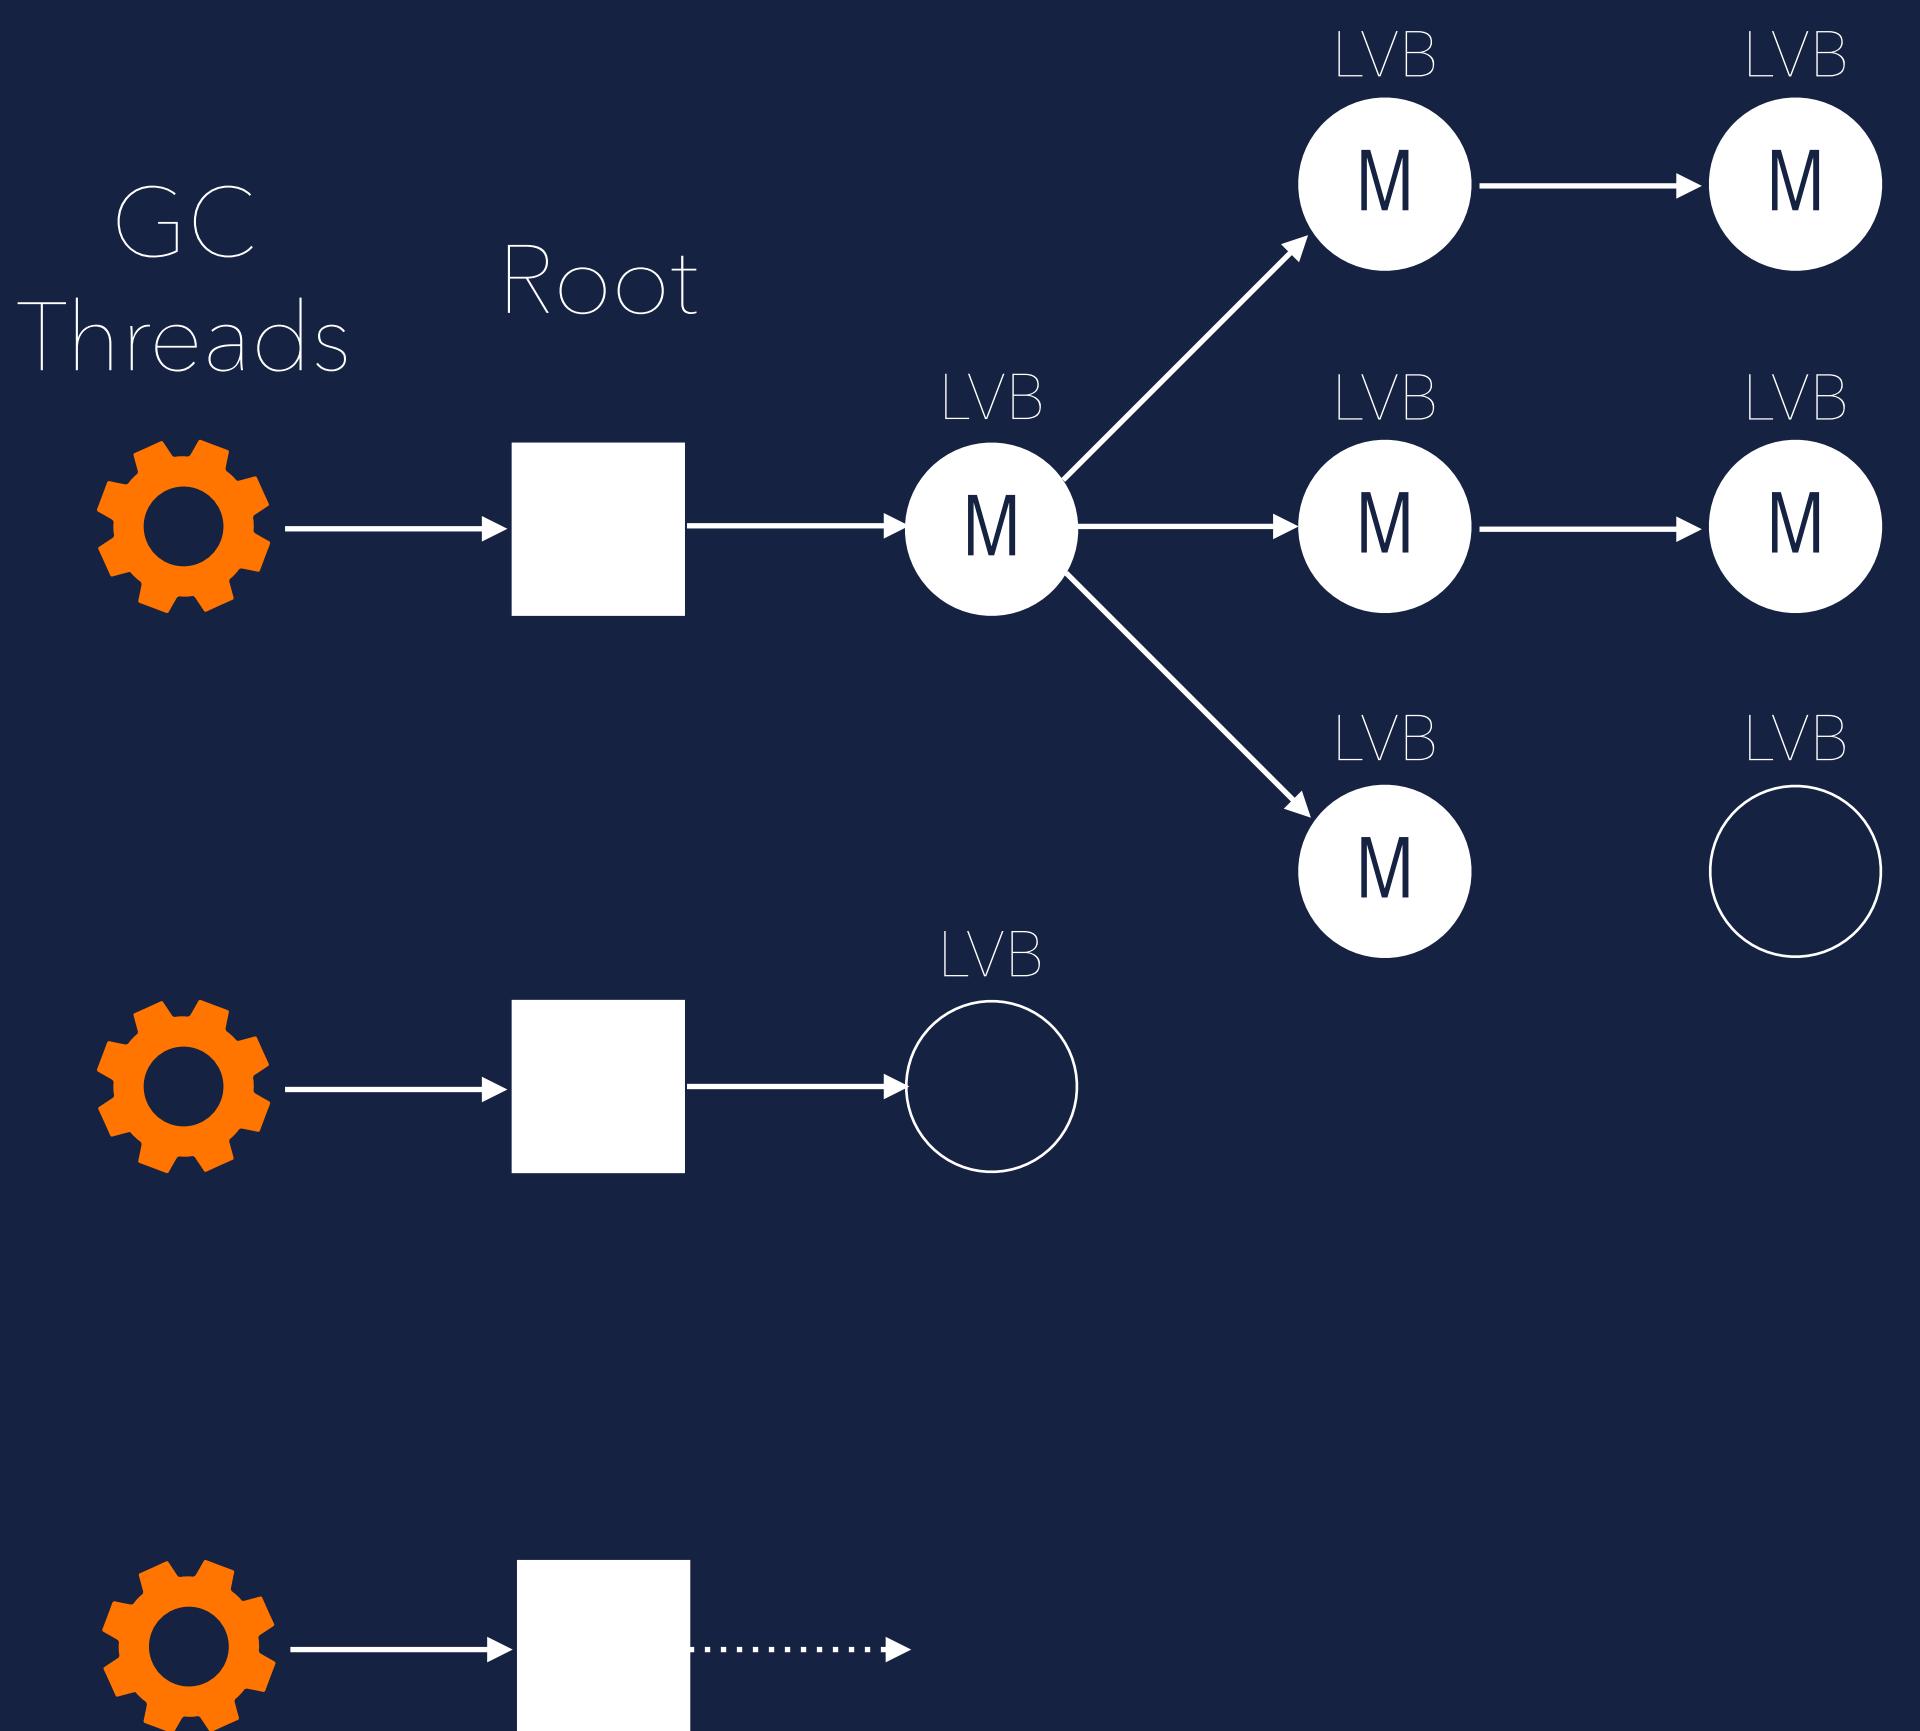
Marking Phase



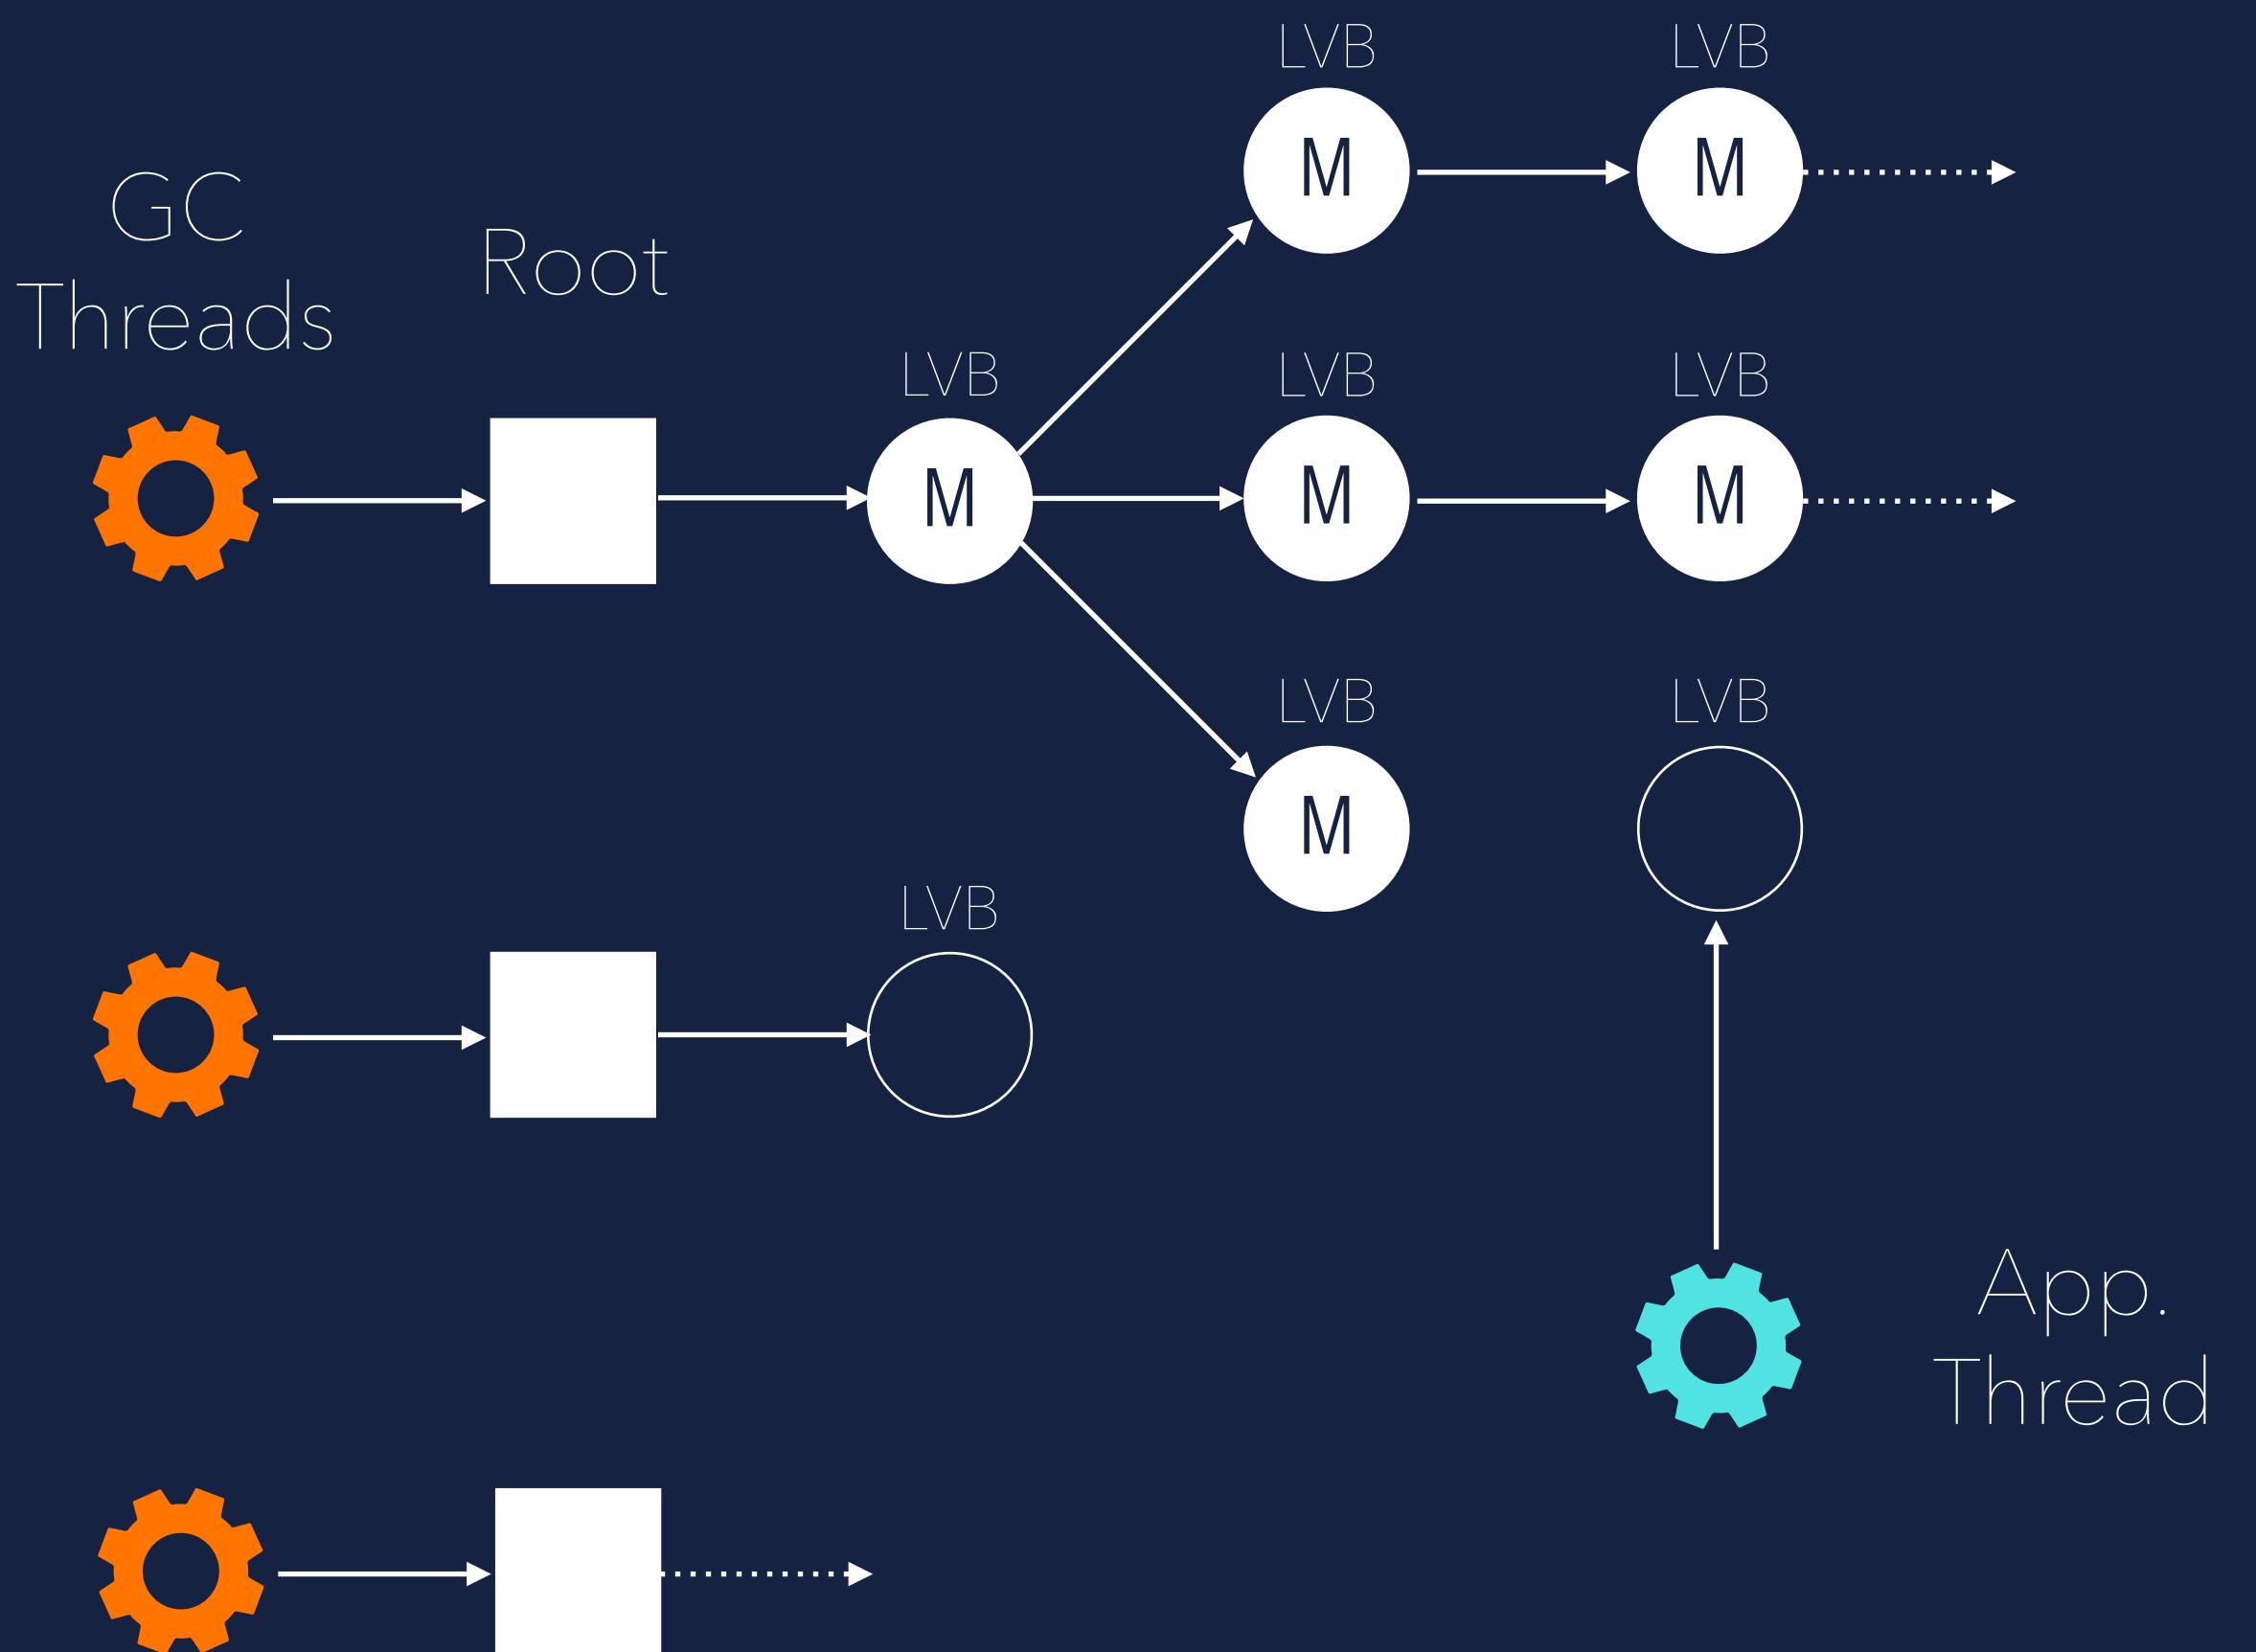
Marking Phase



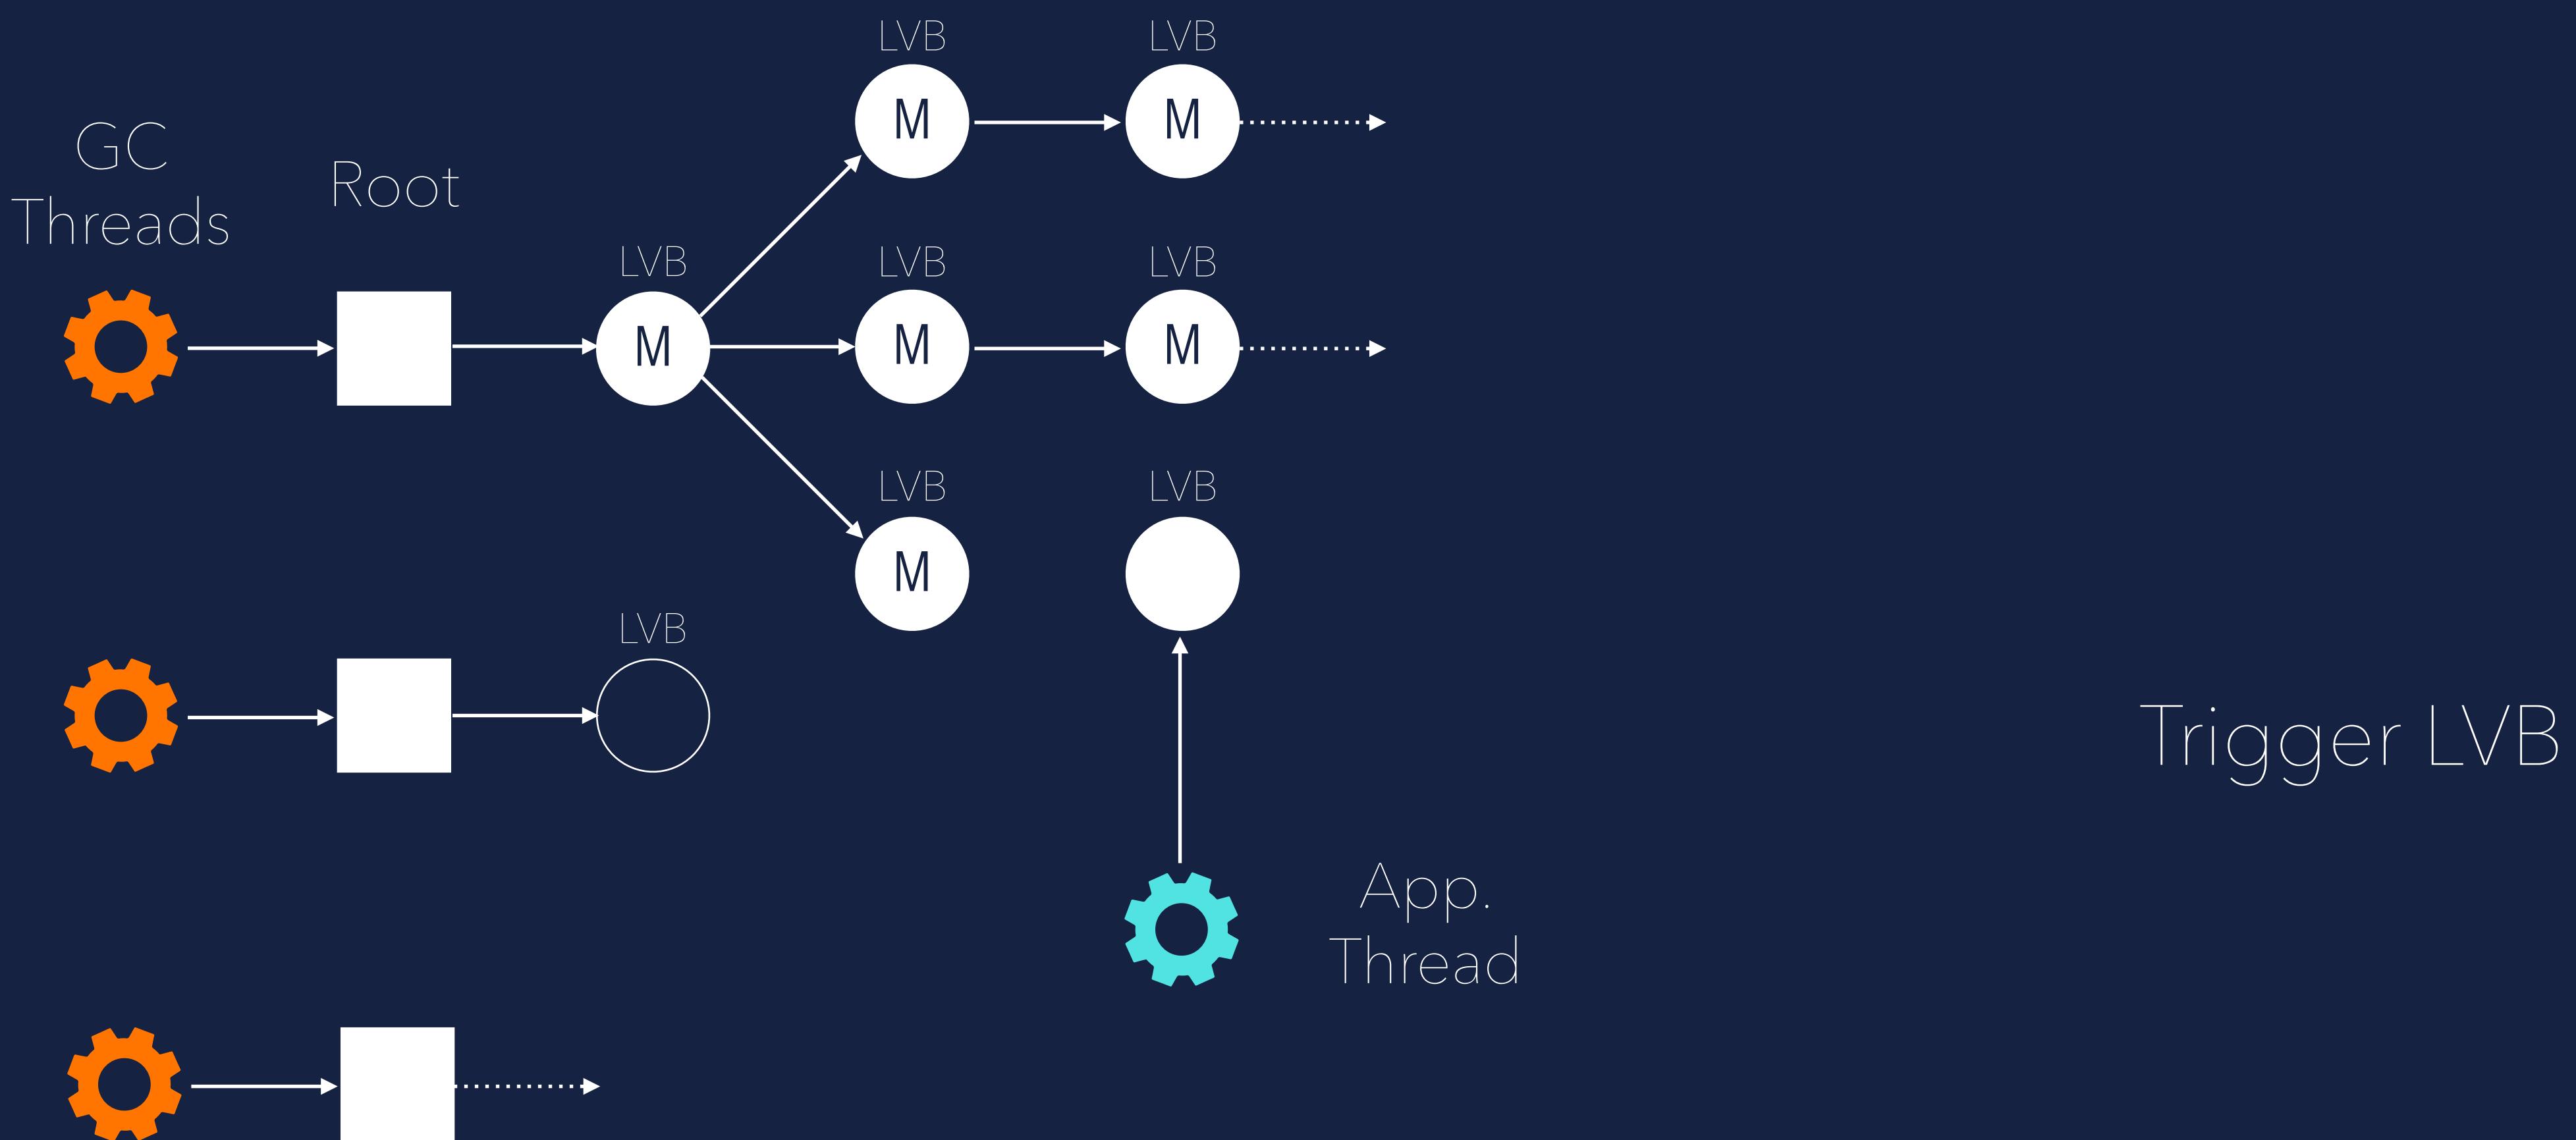
Marking Phase



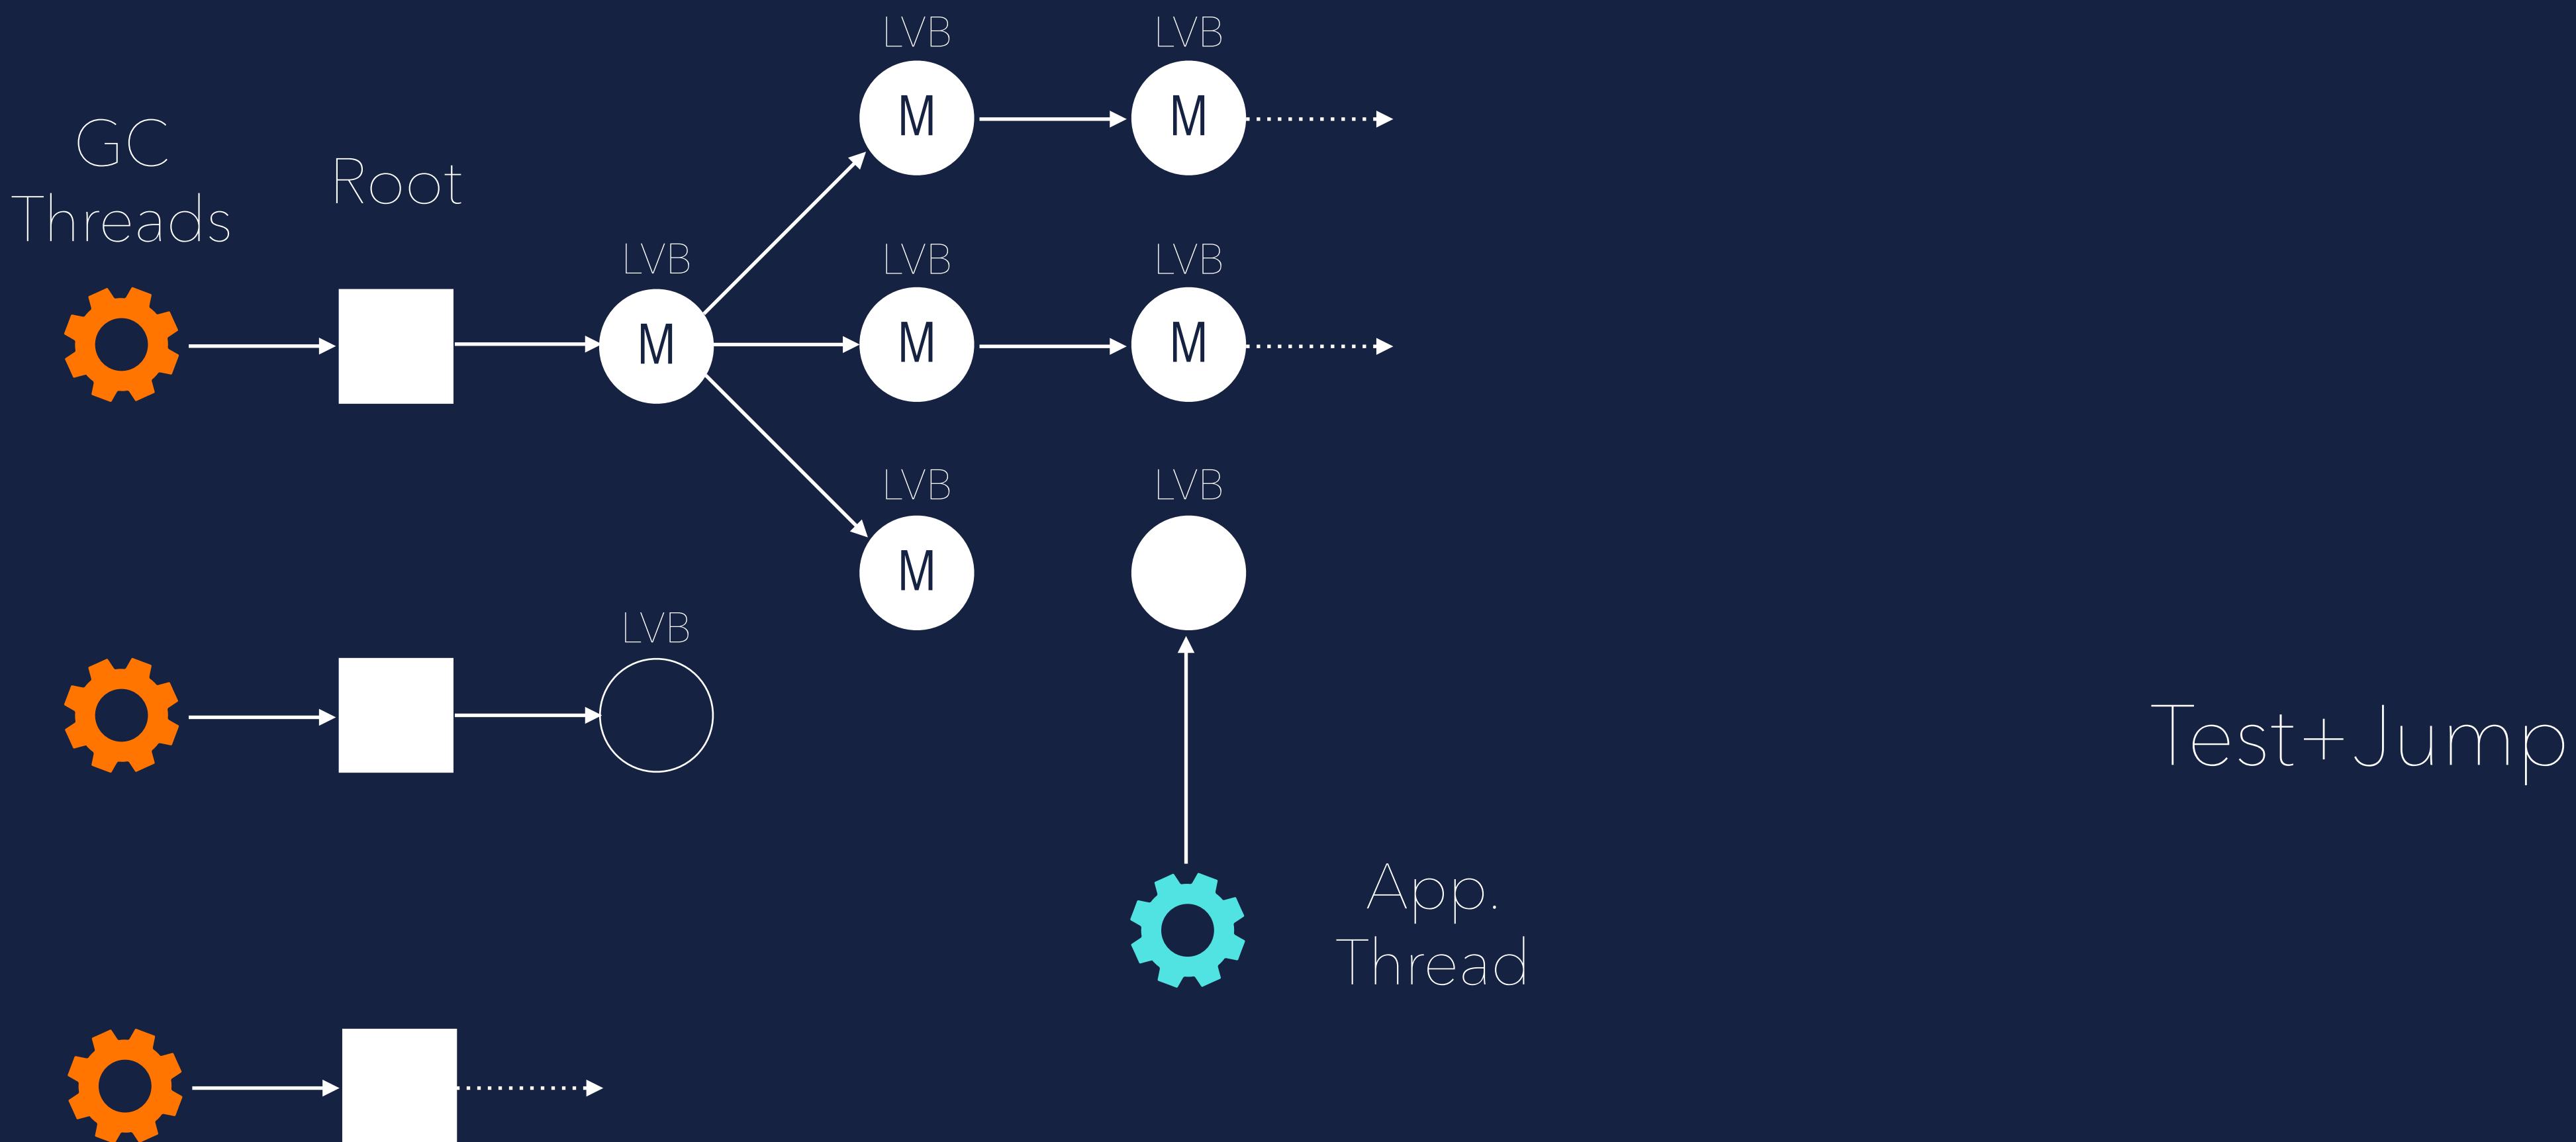
Marking Phase



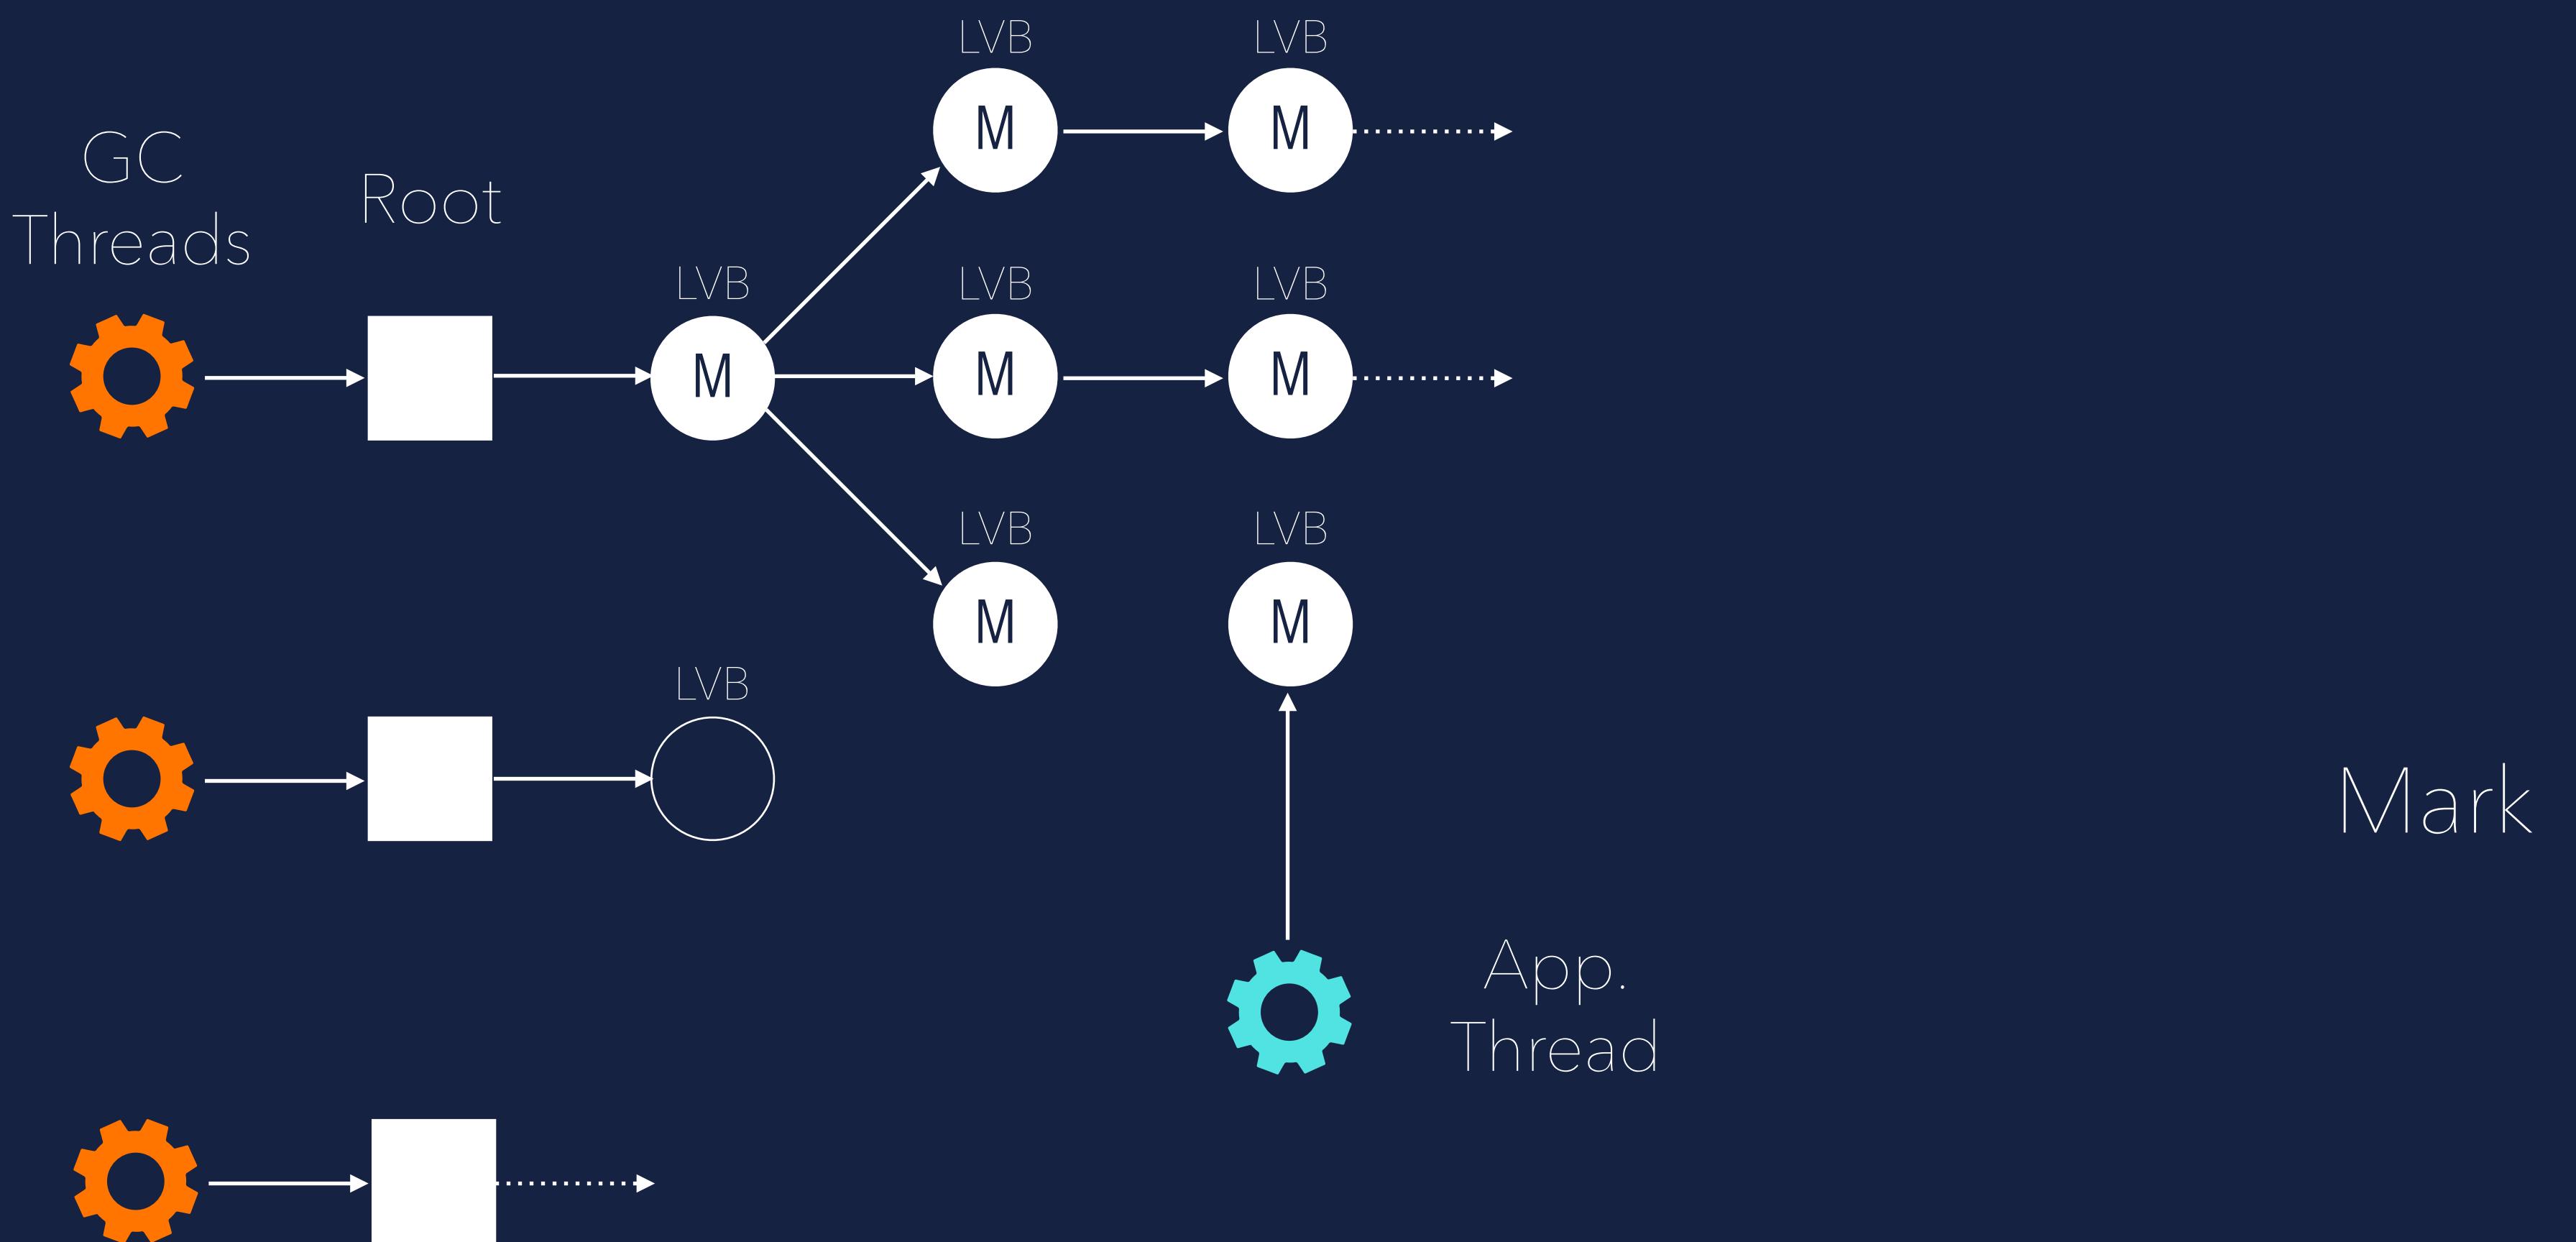
Marking Phase



Marking Phase

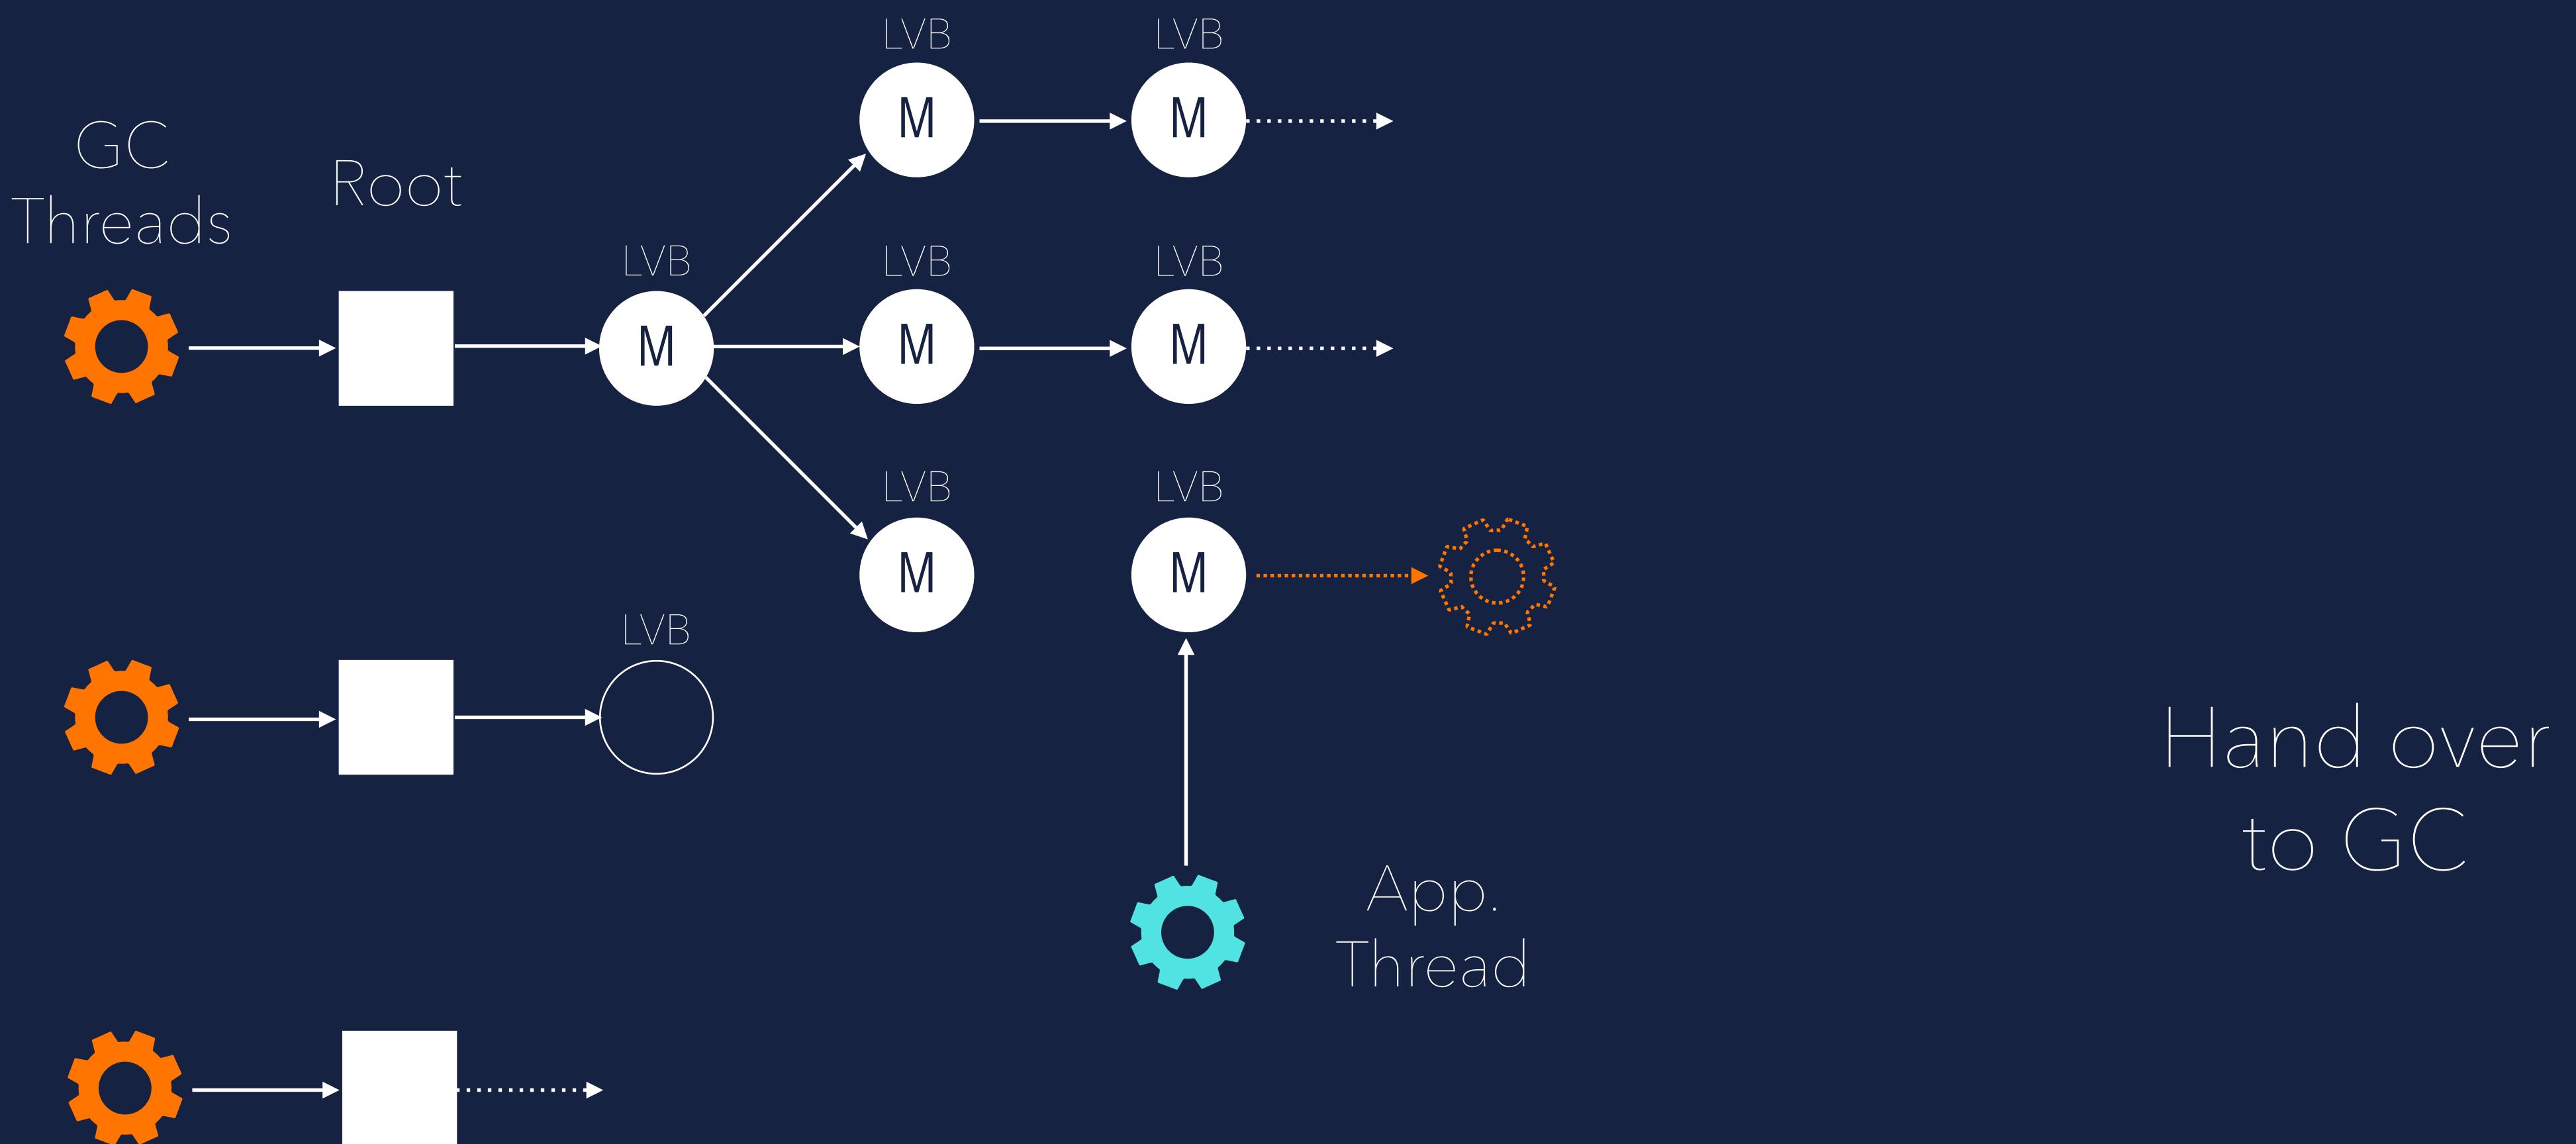


Marking Phase

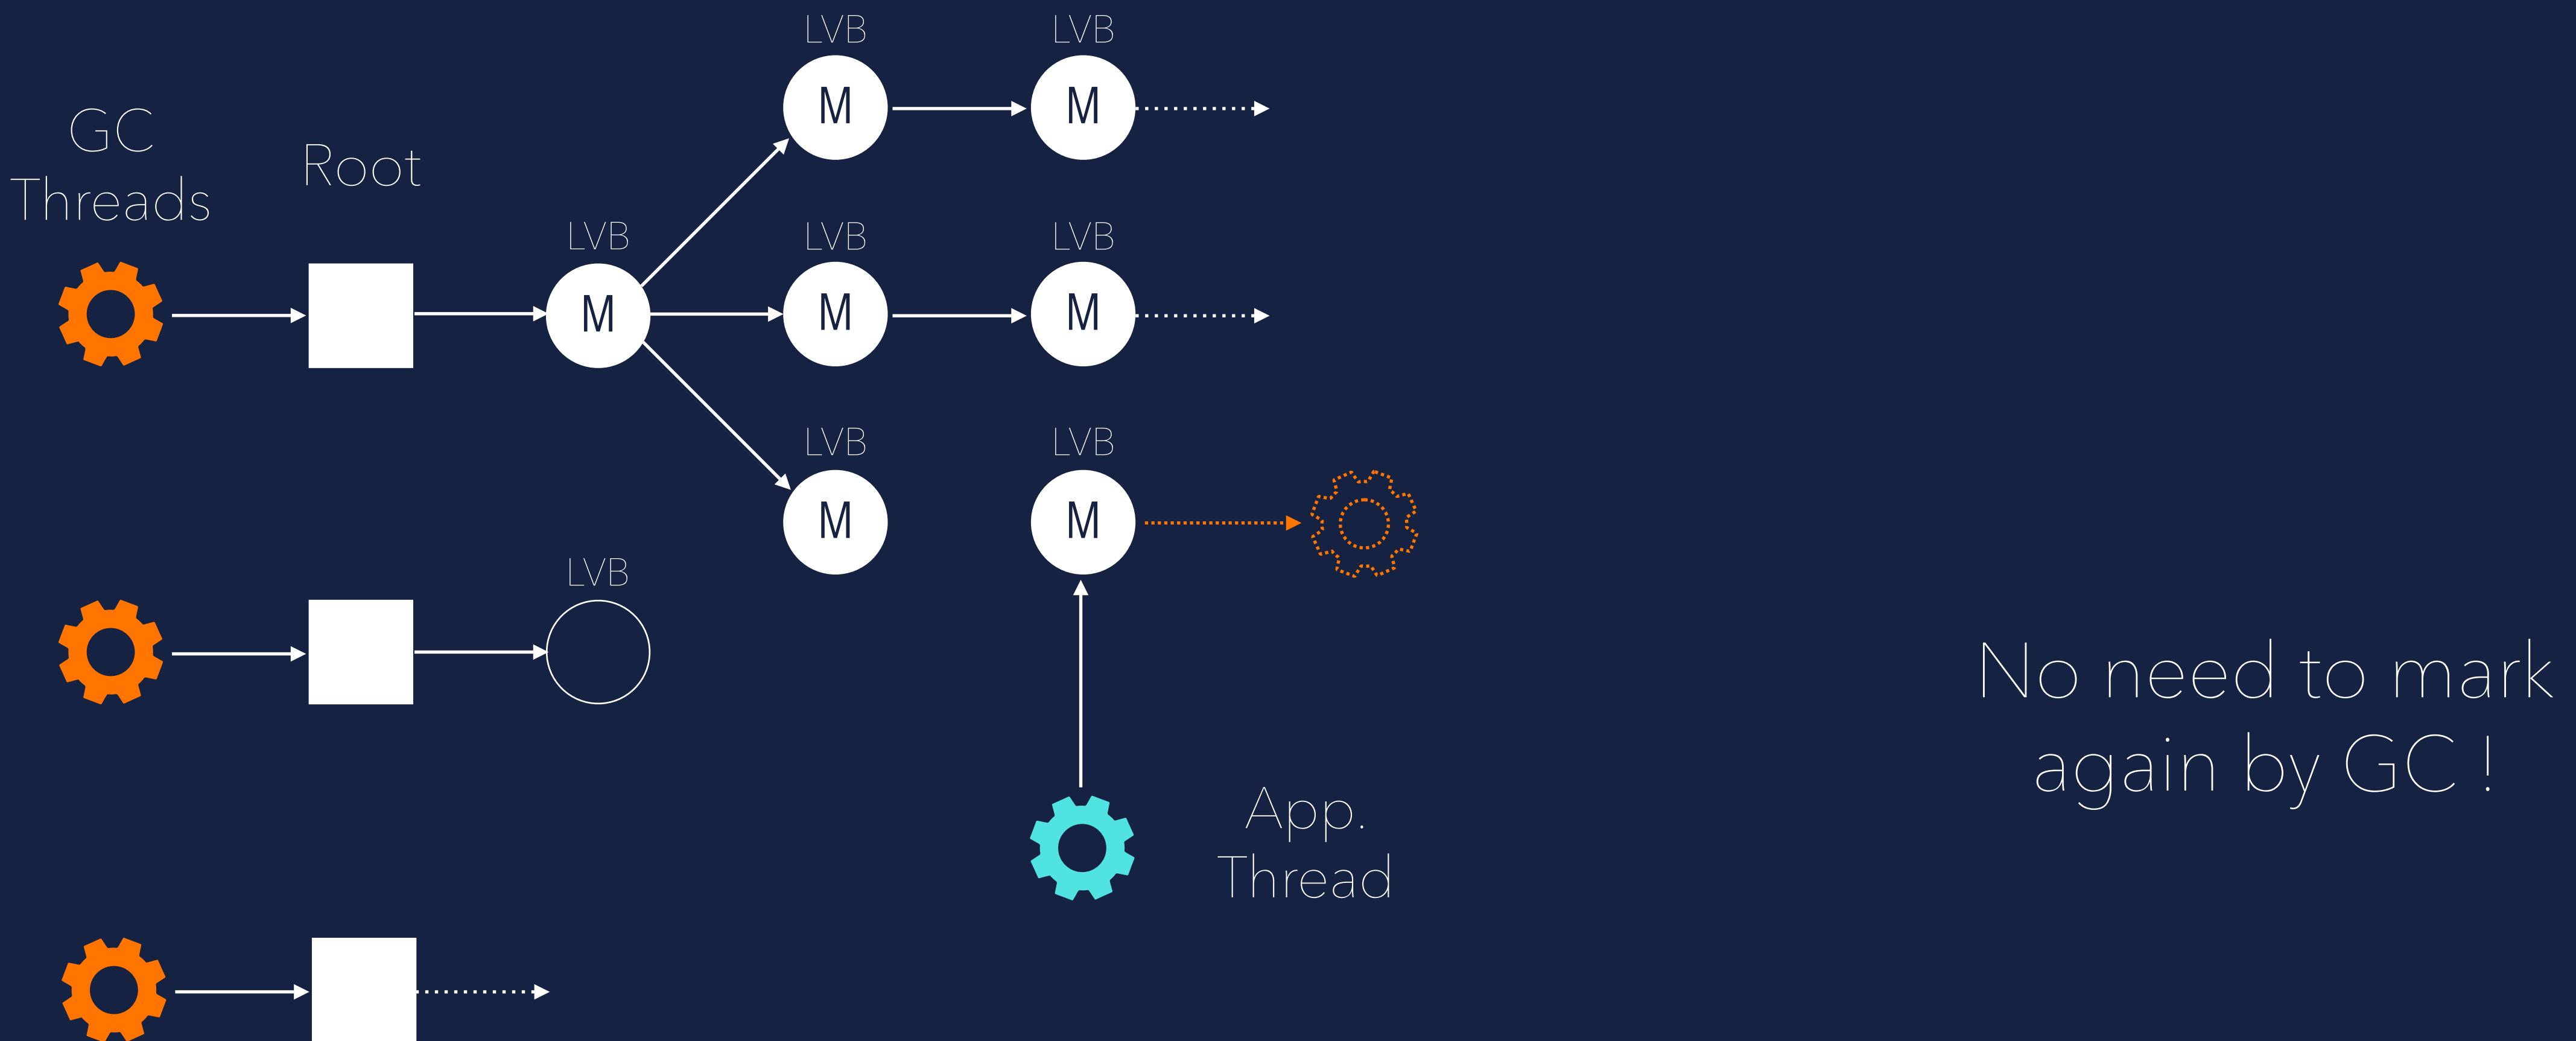


The logo consists of the letters 'C4' in a bold, white, sans-serif font. The letter 'C' is stylized with a thick vertical stroke on the right and a horizontal cutout on the left. The letter '4' has a vertical stem with a horizontal bar extending from its top.

Marking Phase



Marking Phase



C4

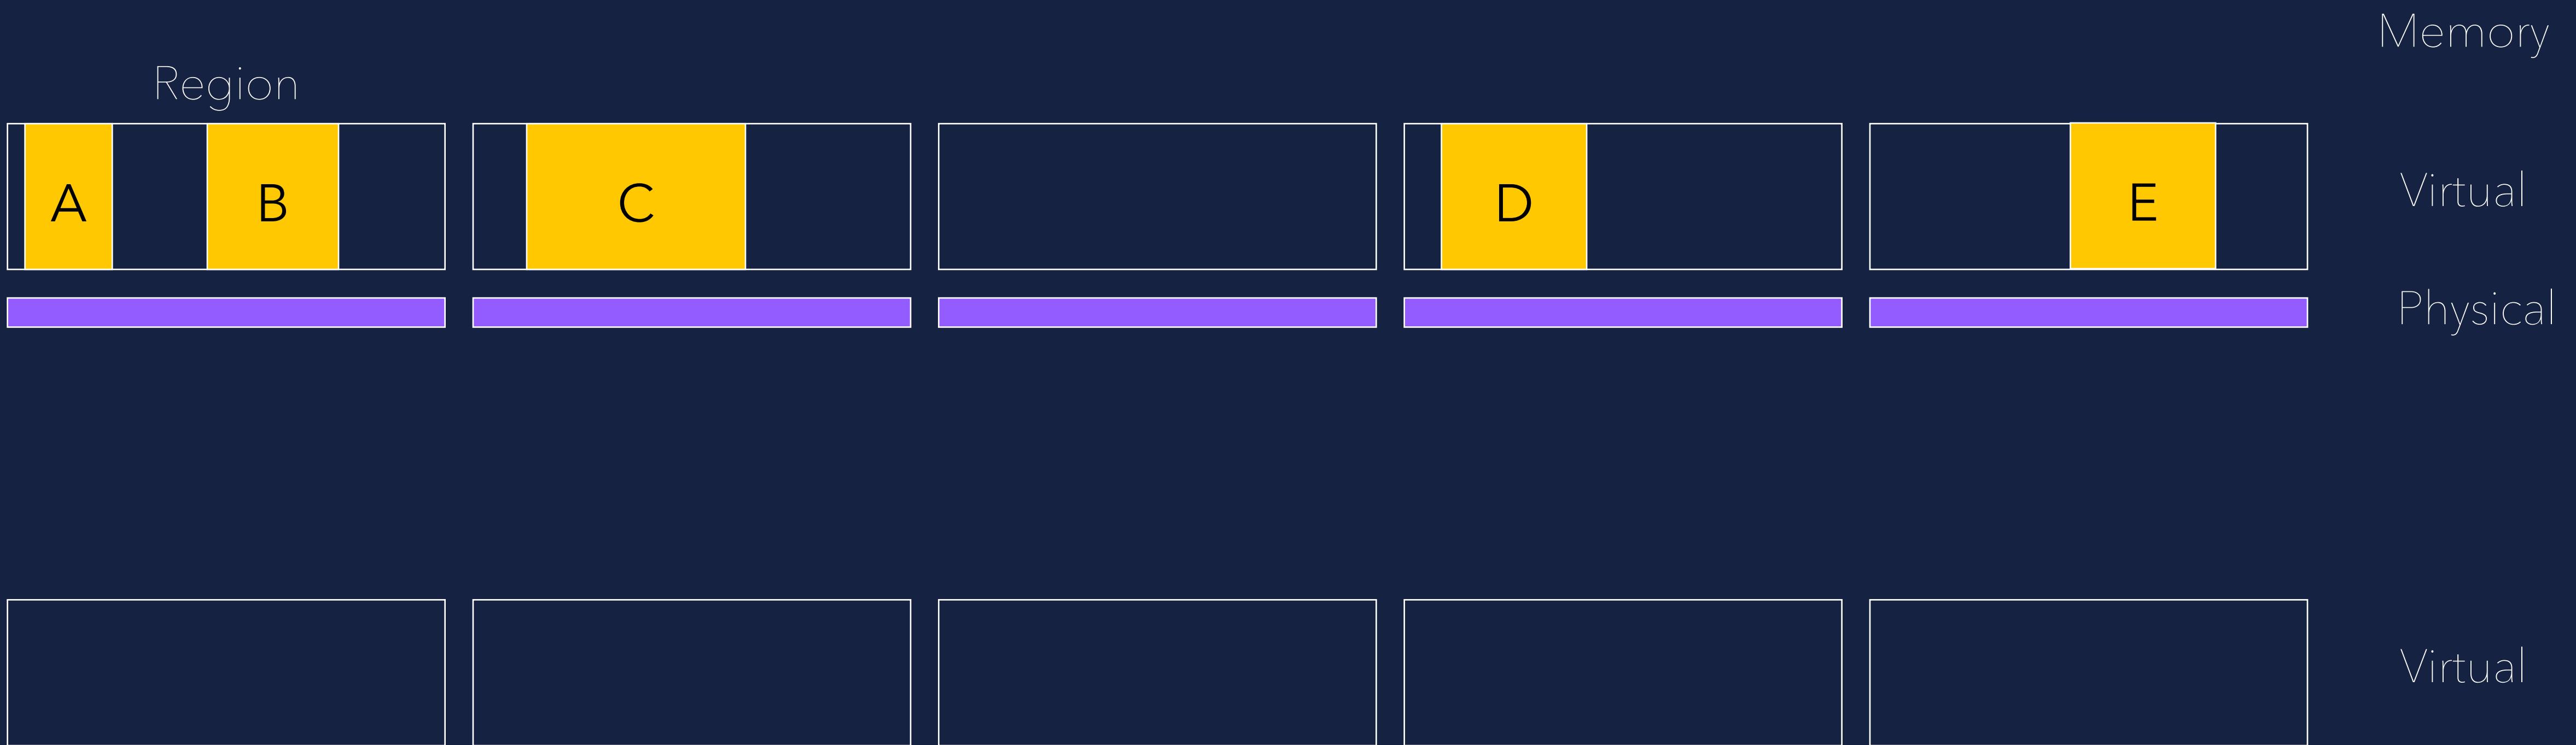


QUICK RELEASE

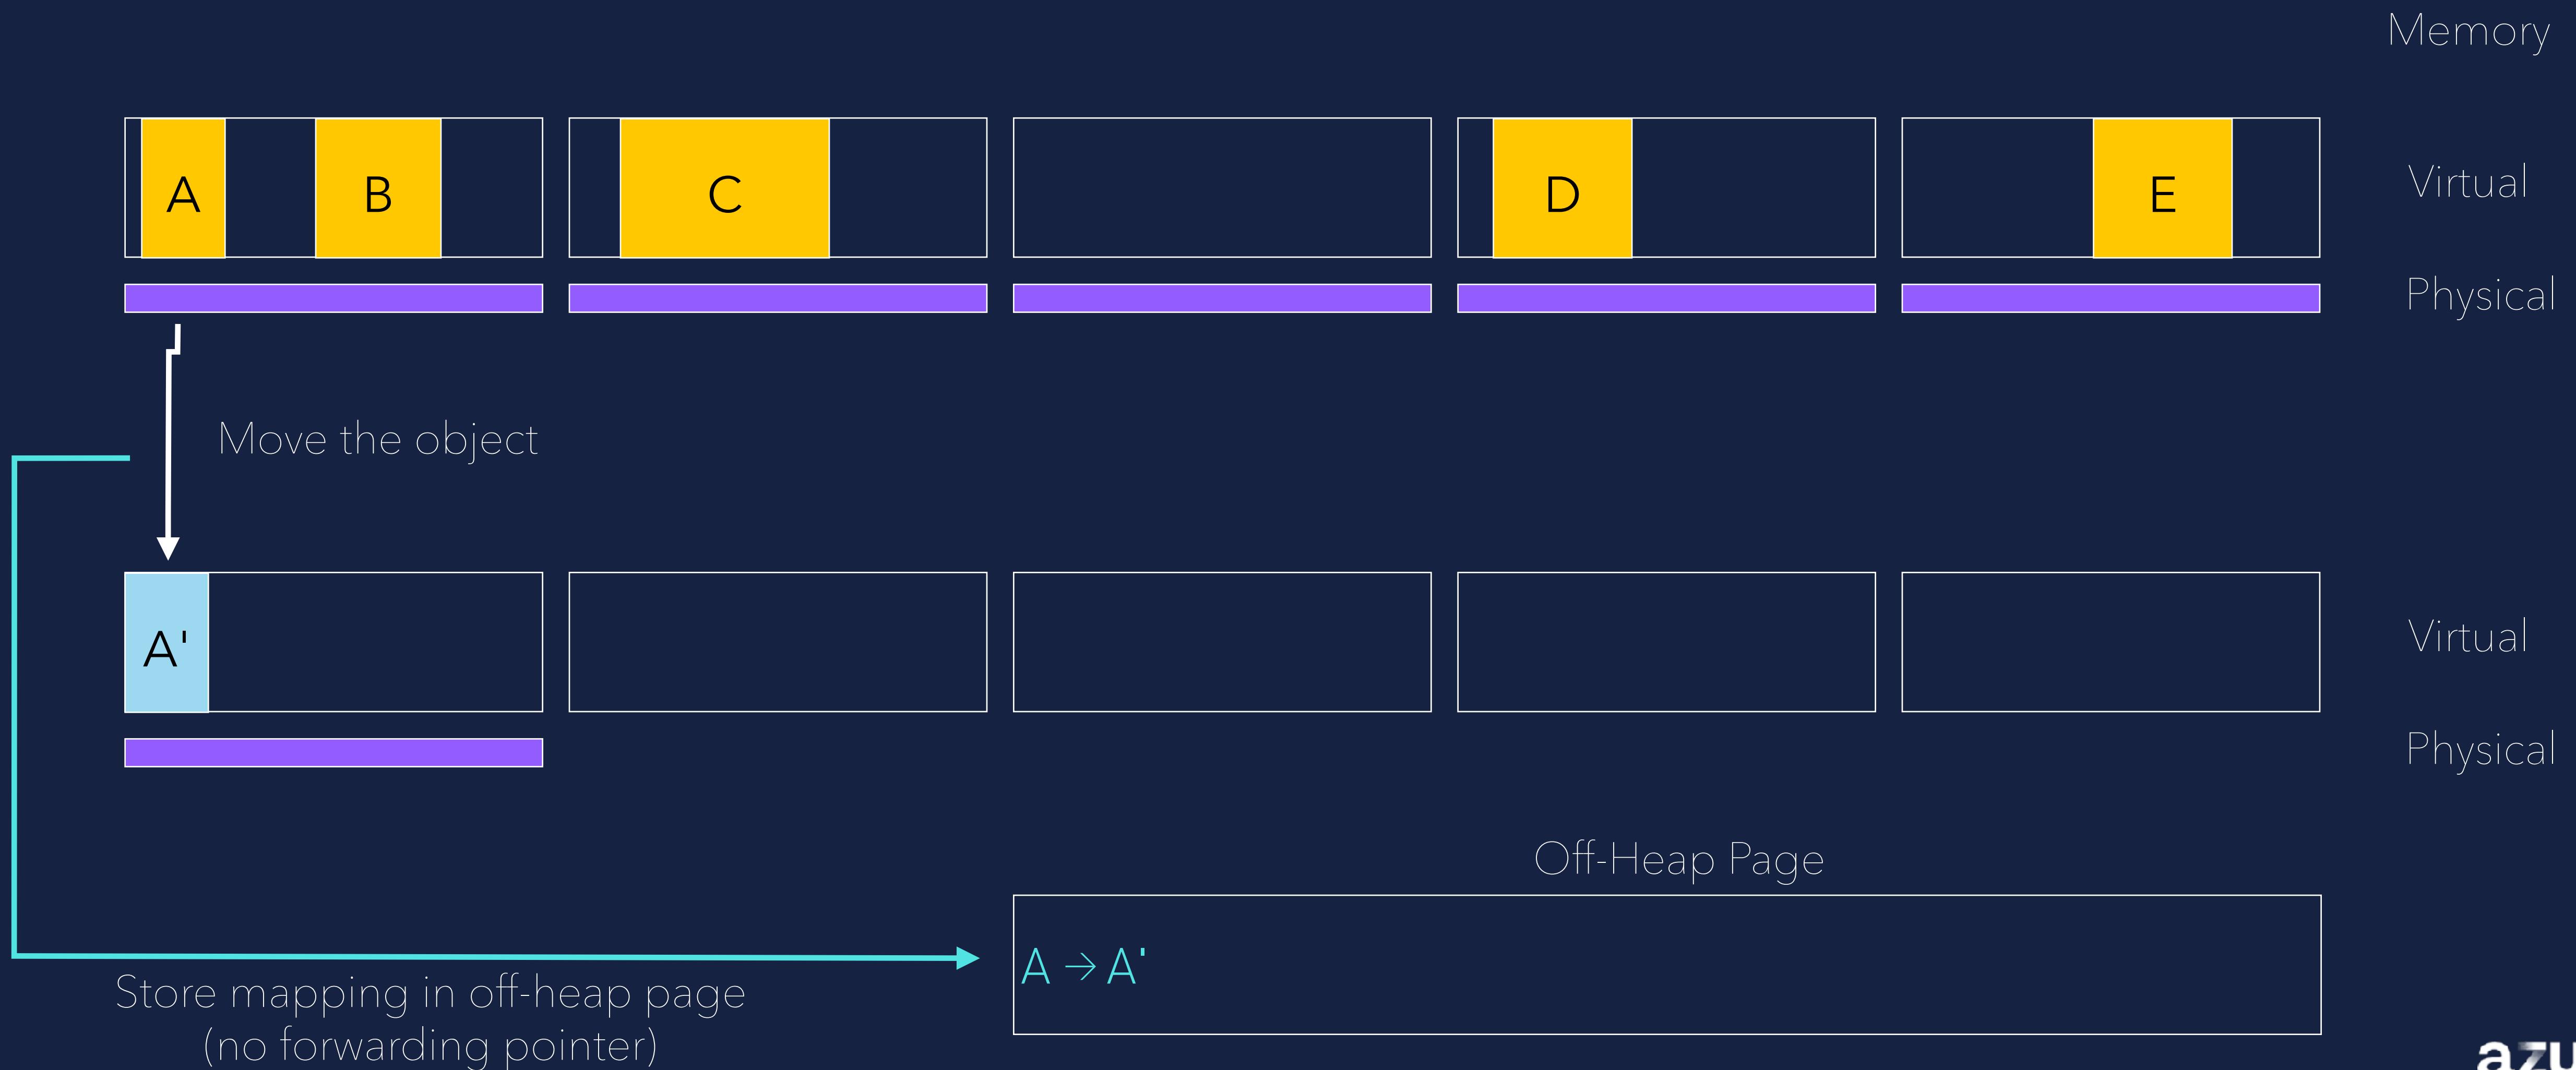
azul



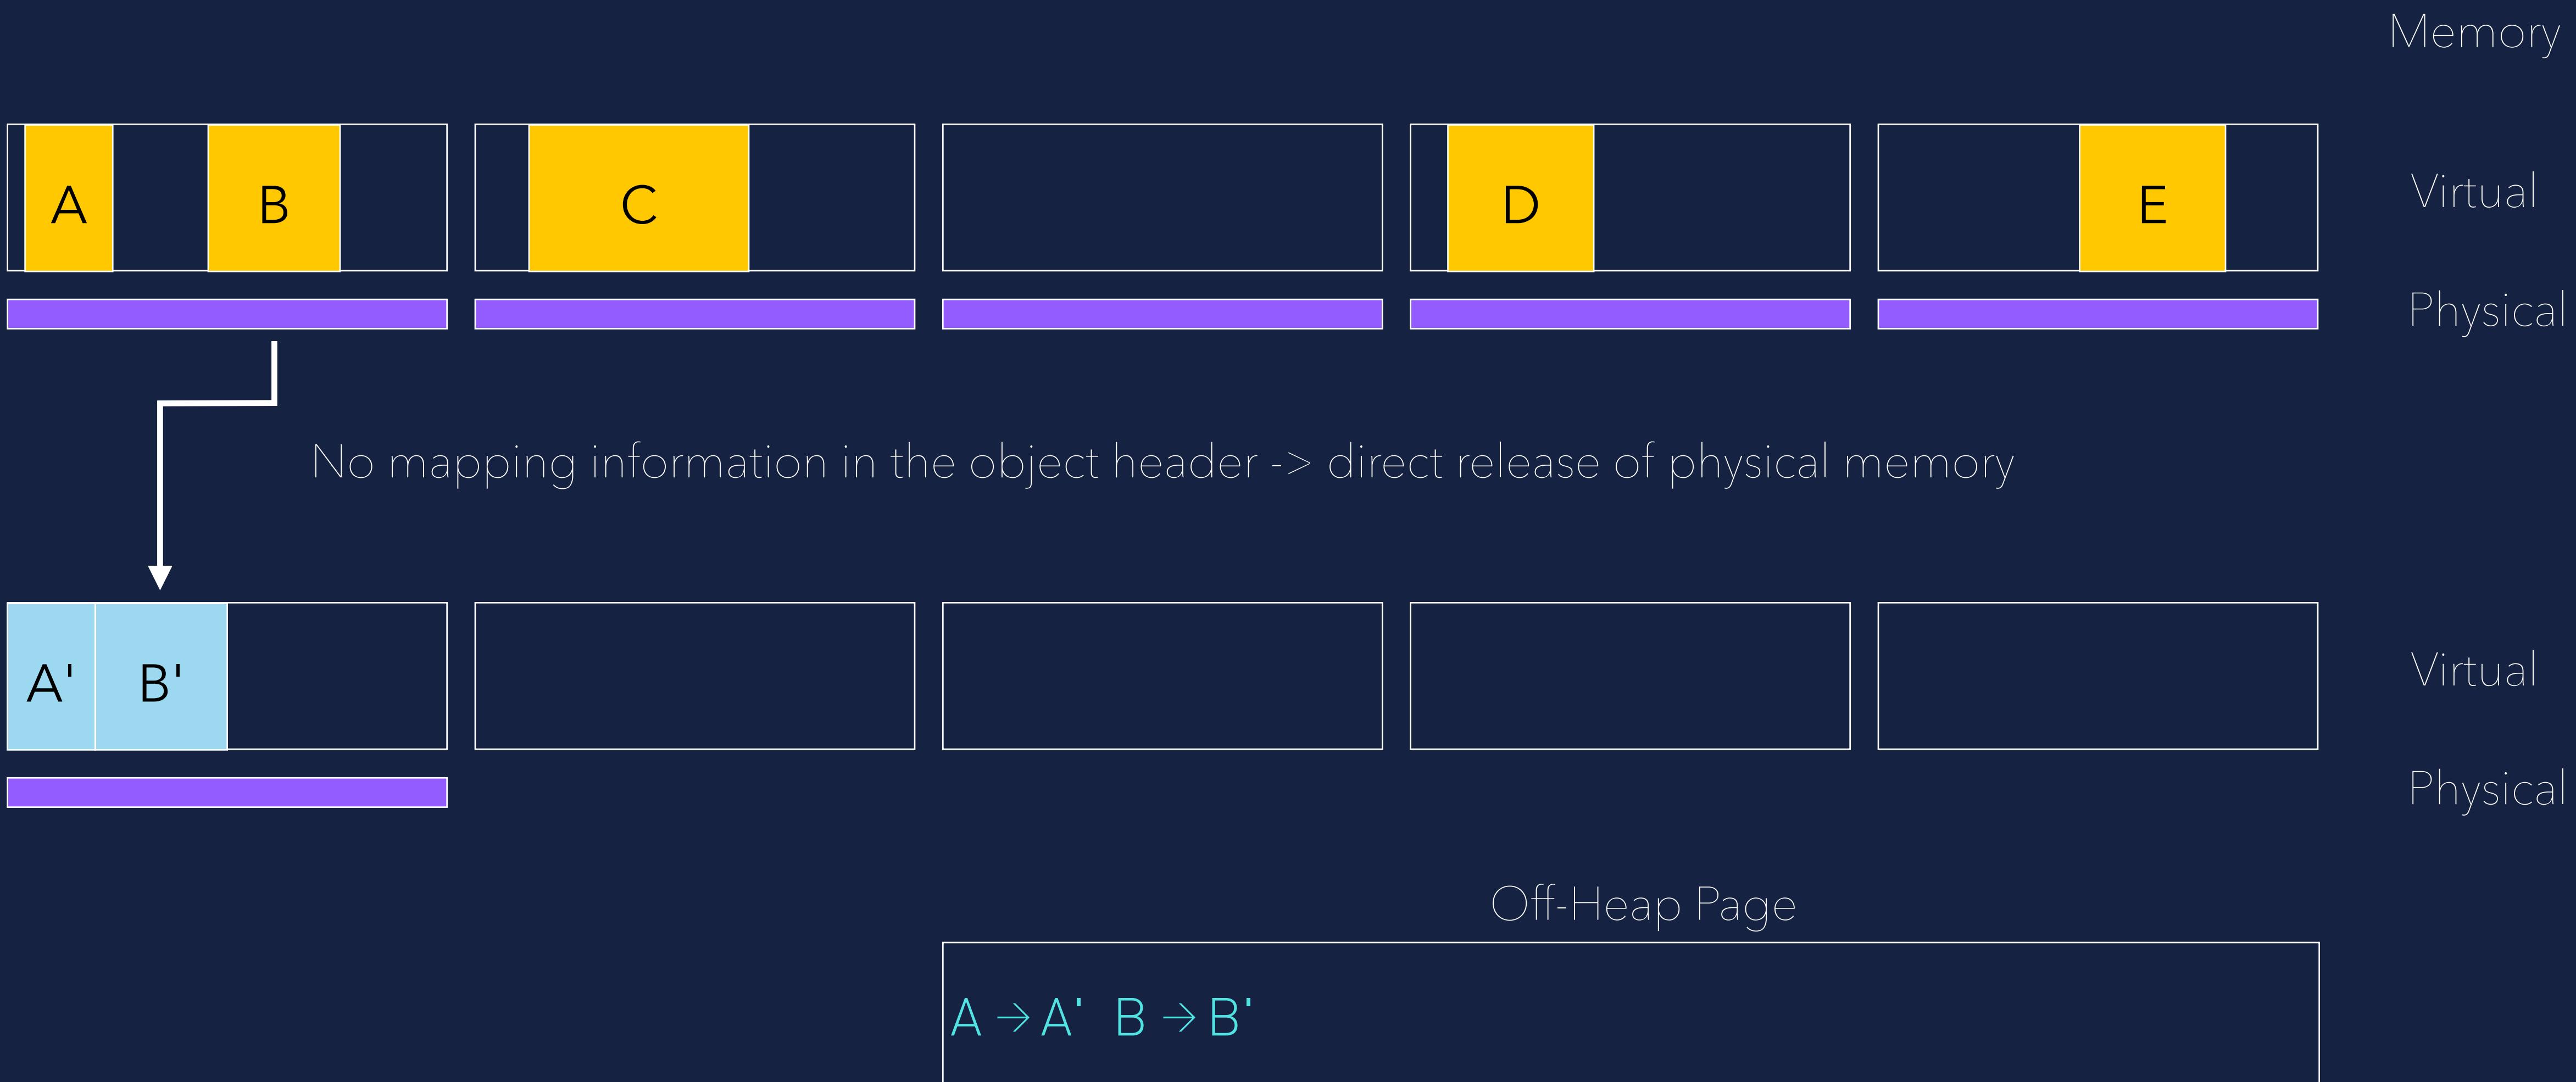
Relocation Phase



Relocation Phase (Compaction)



Relocation Phase (Compaction)



Relocation Phase (Compaction)



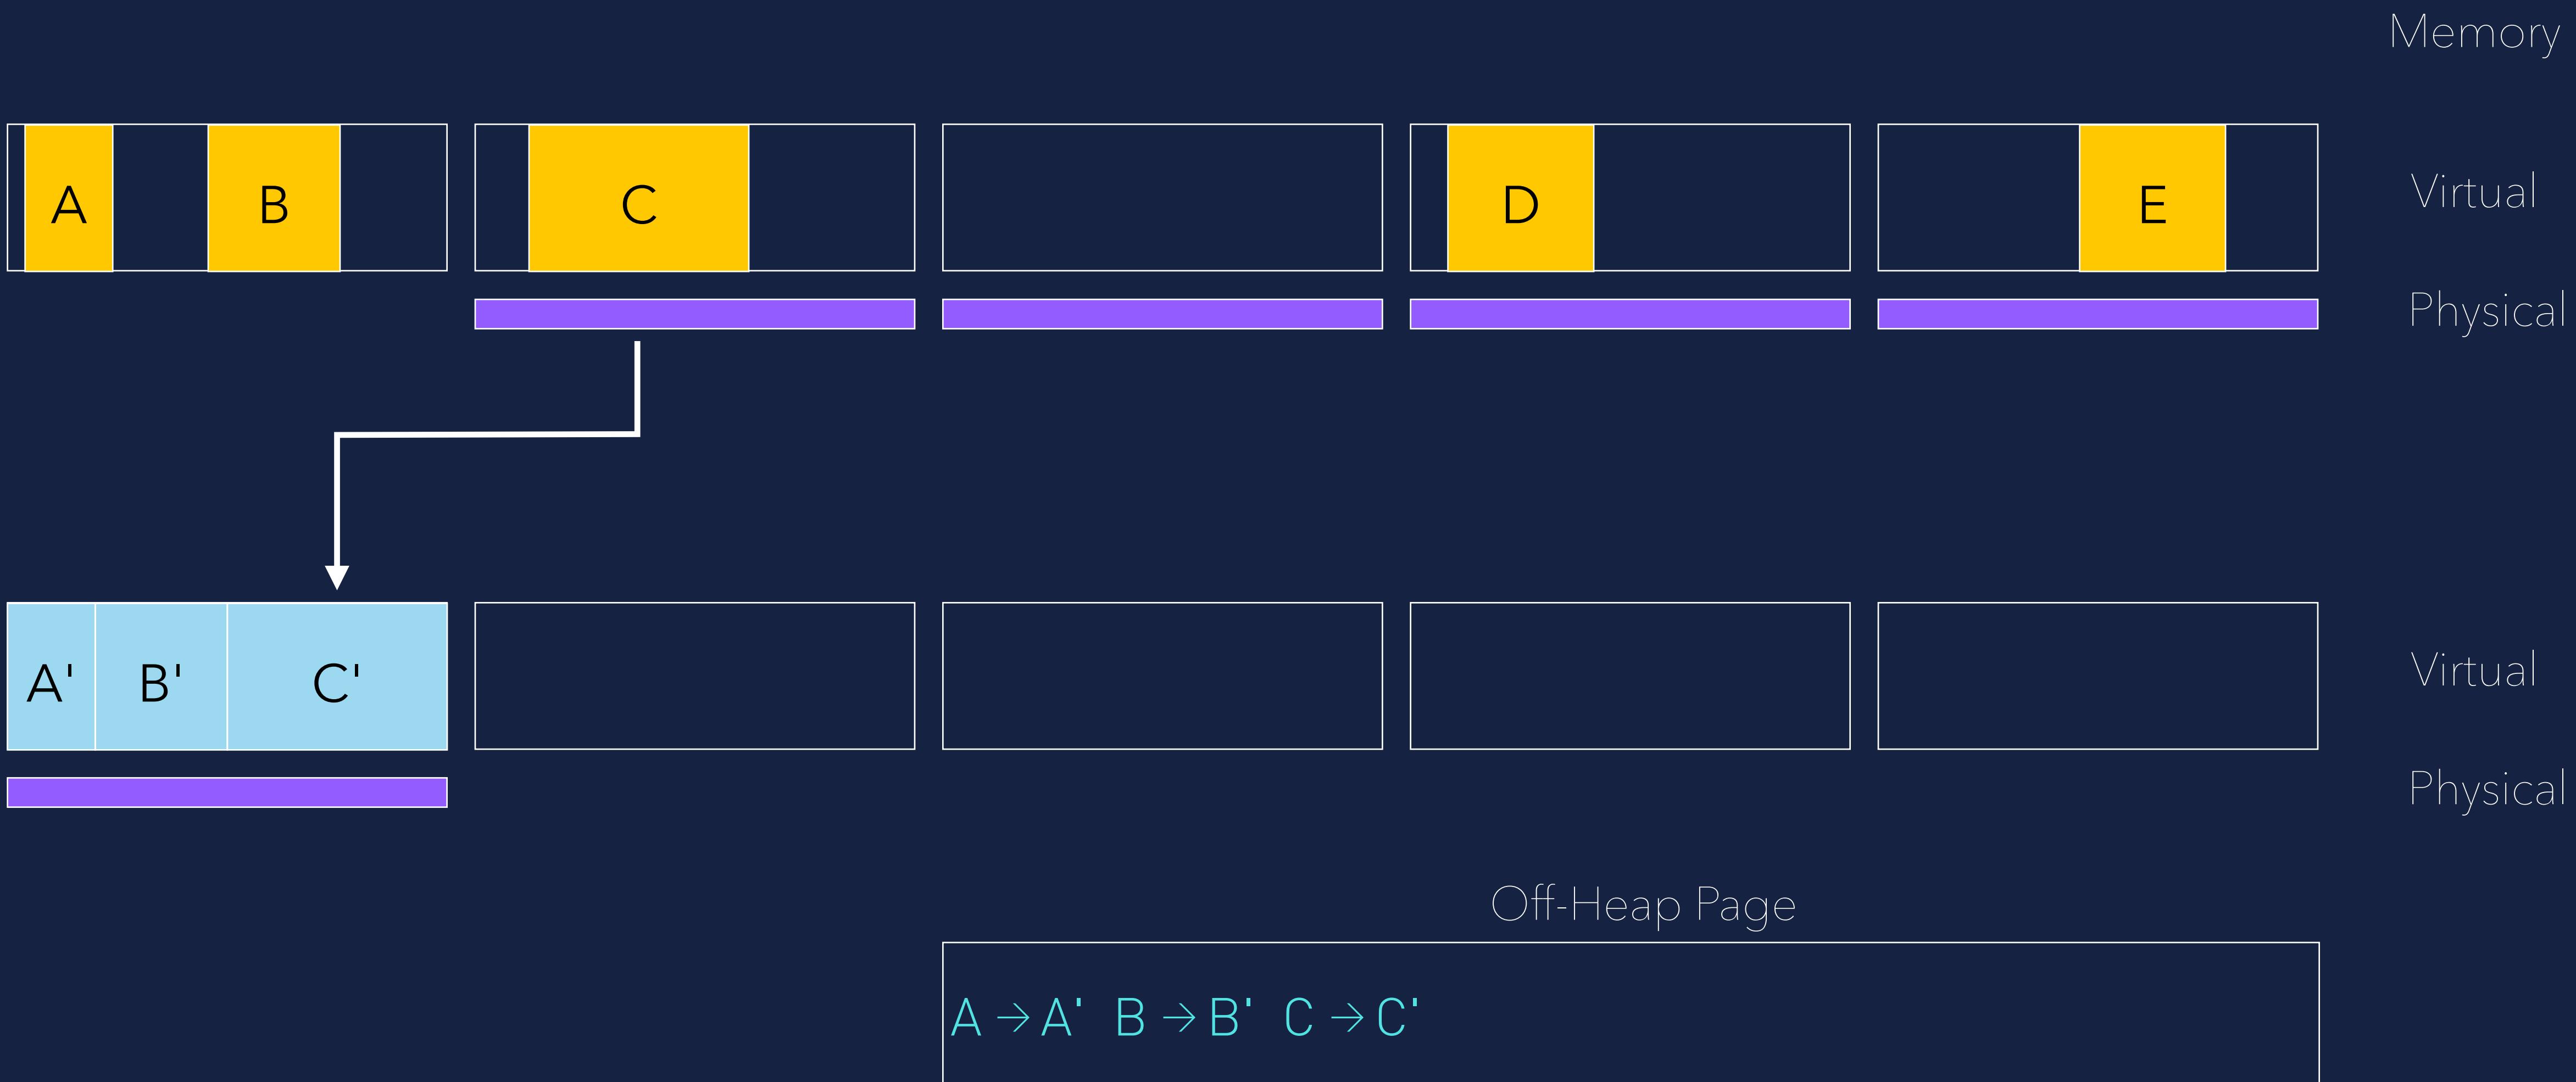
No mapping information in the object header -> direct release of physical memory



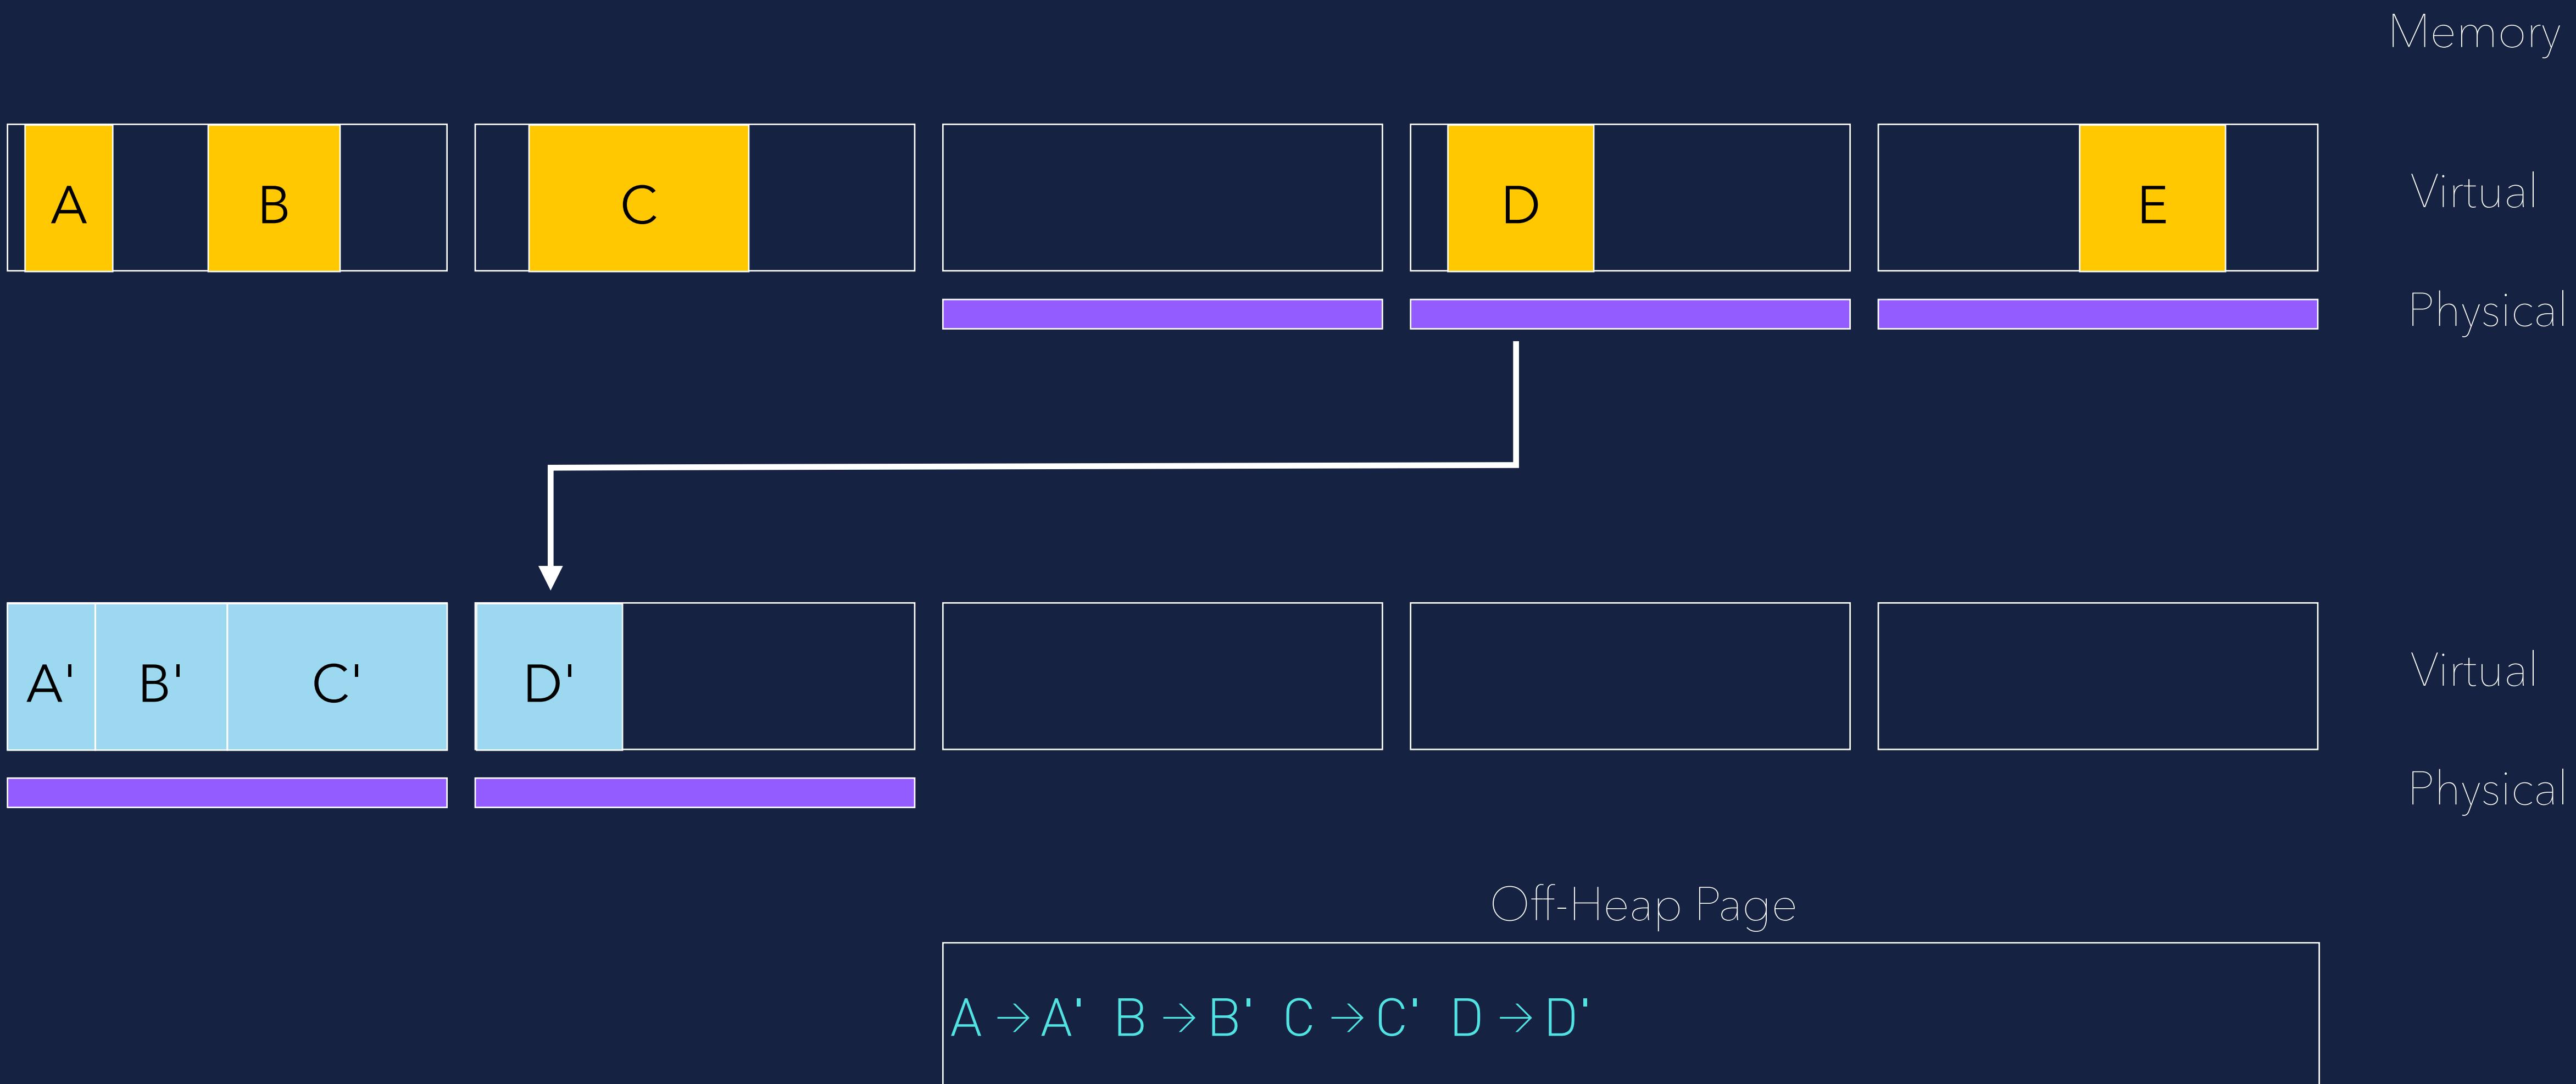
Off-Heap Page

$A \rightarrow A'$ $B \rightarrow B'$

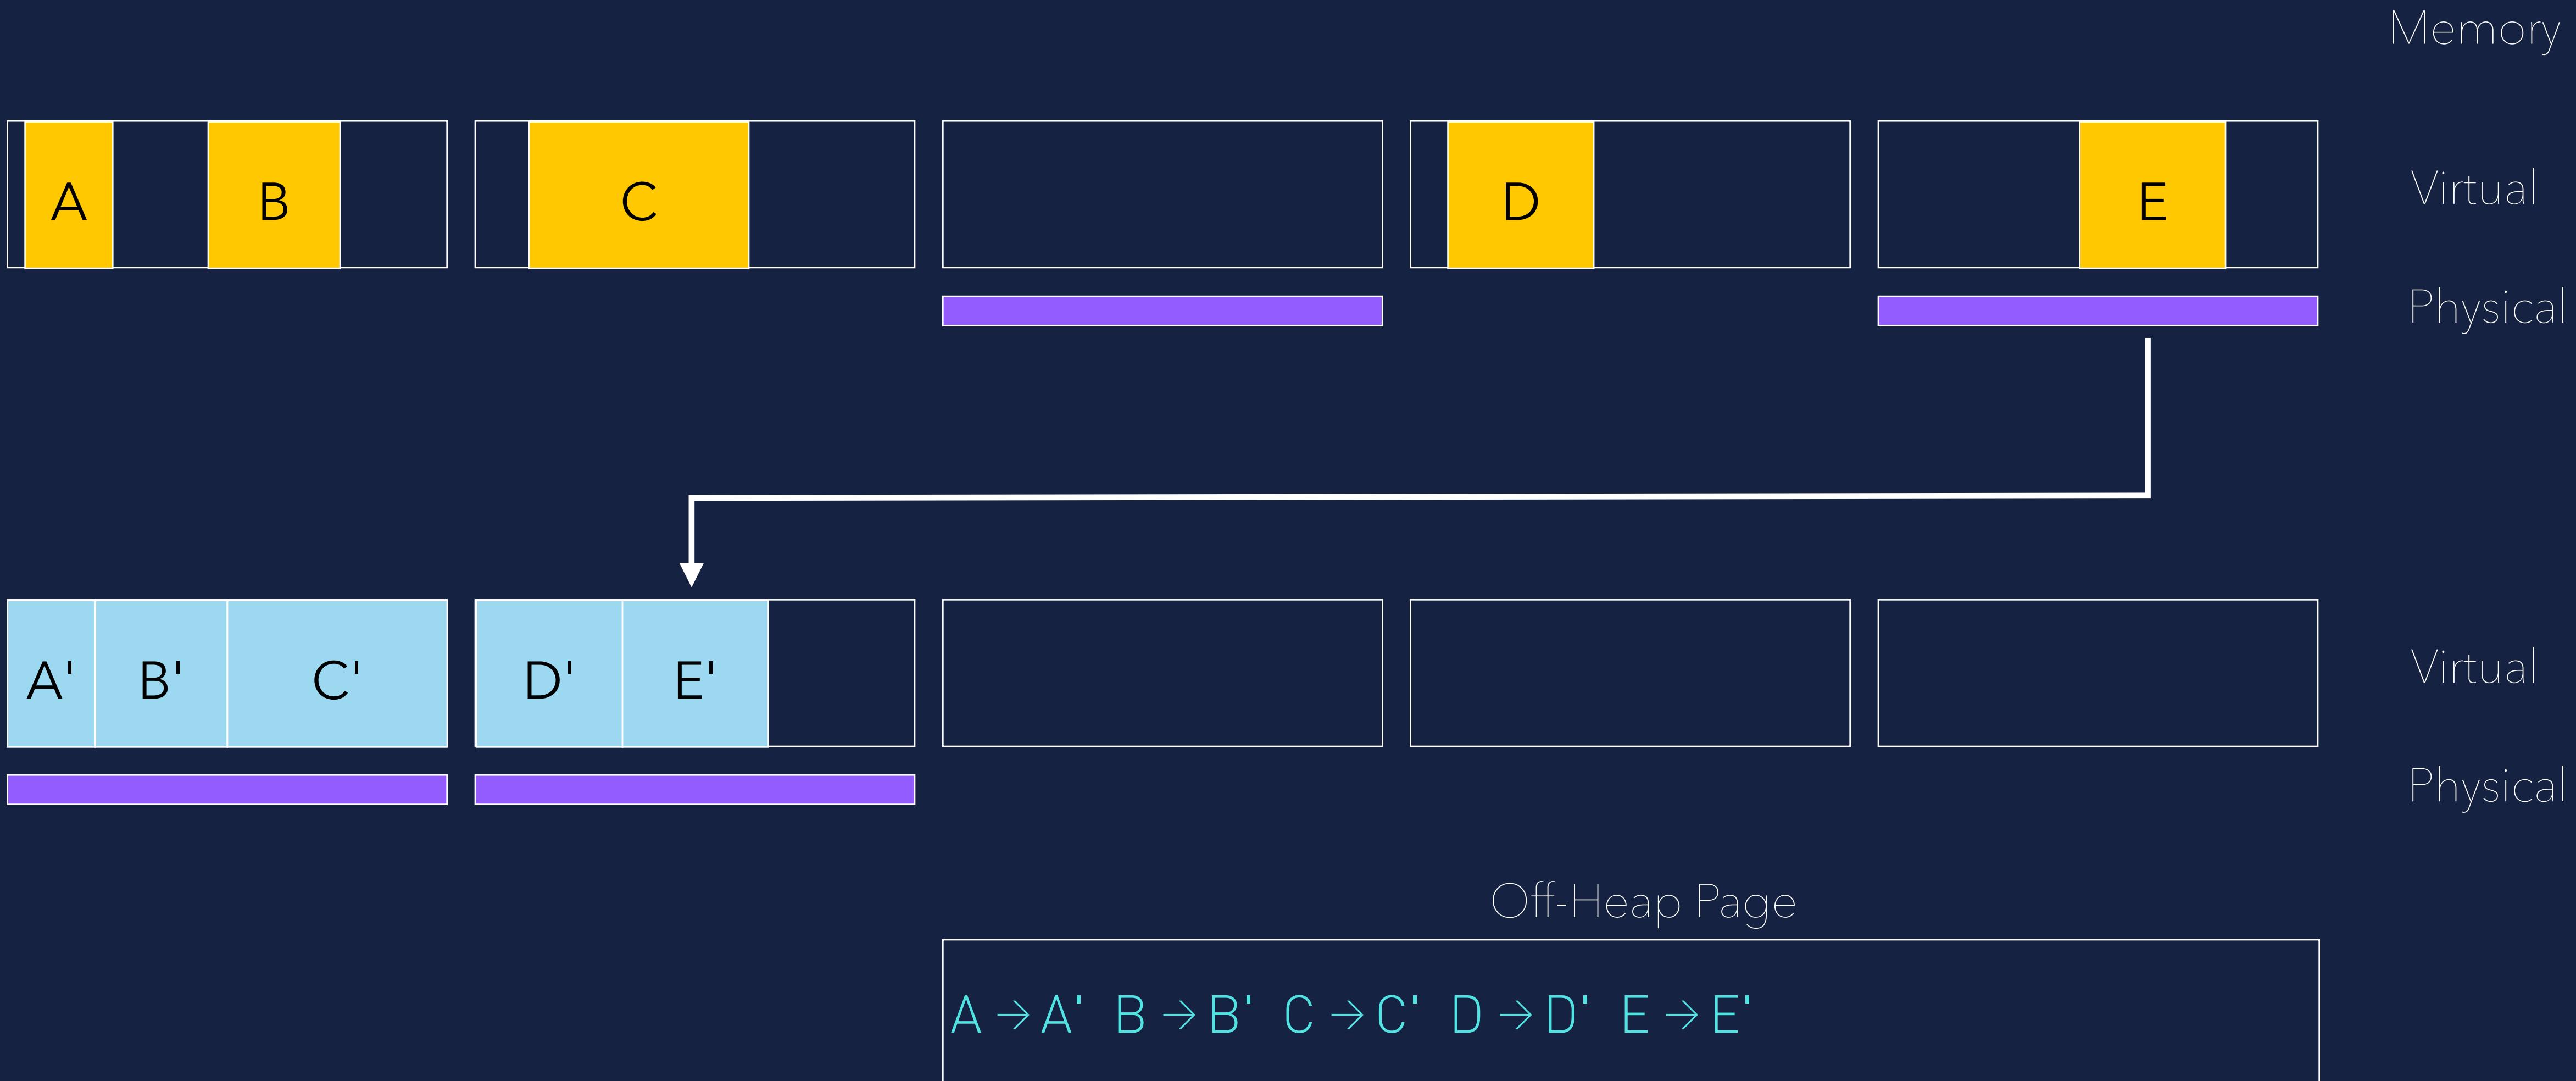
Relocation Phase (Compaction)



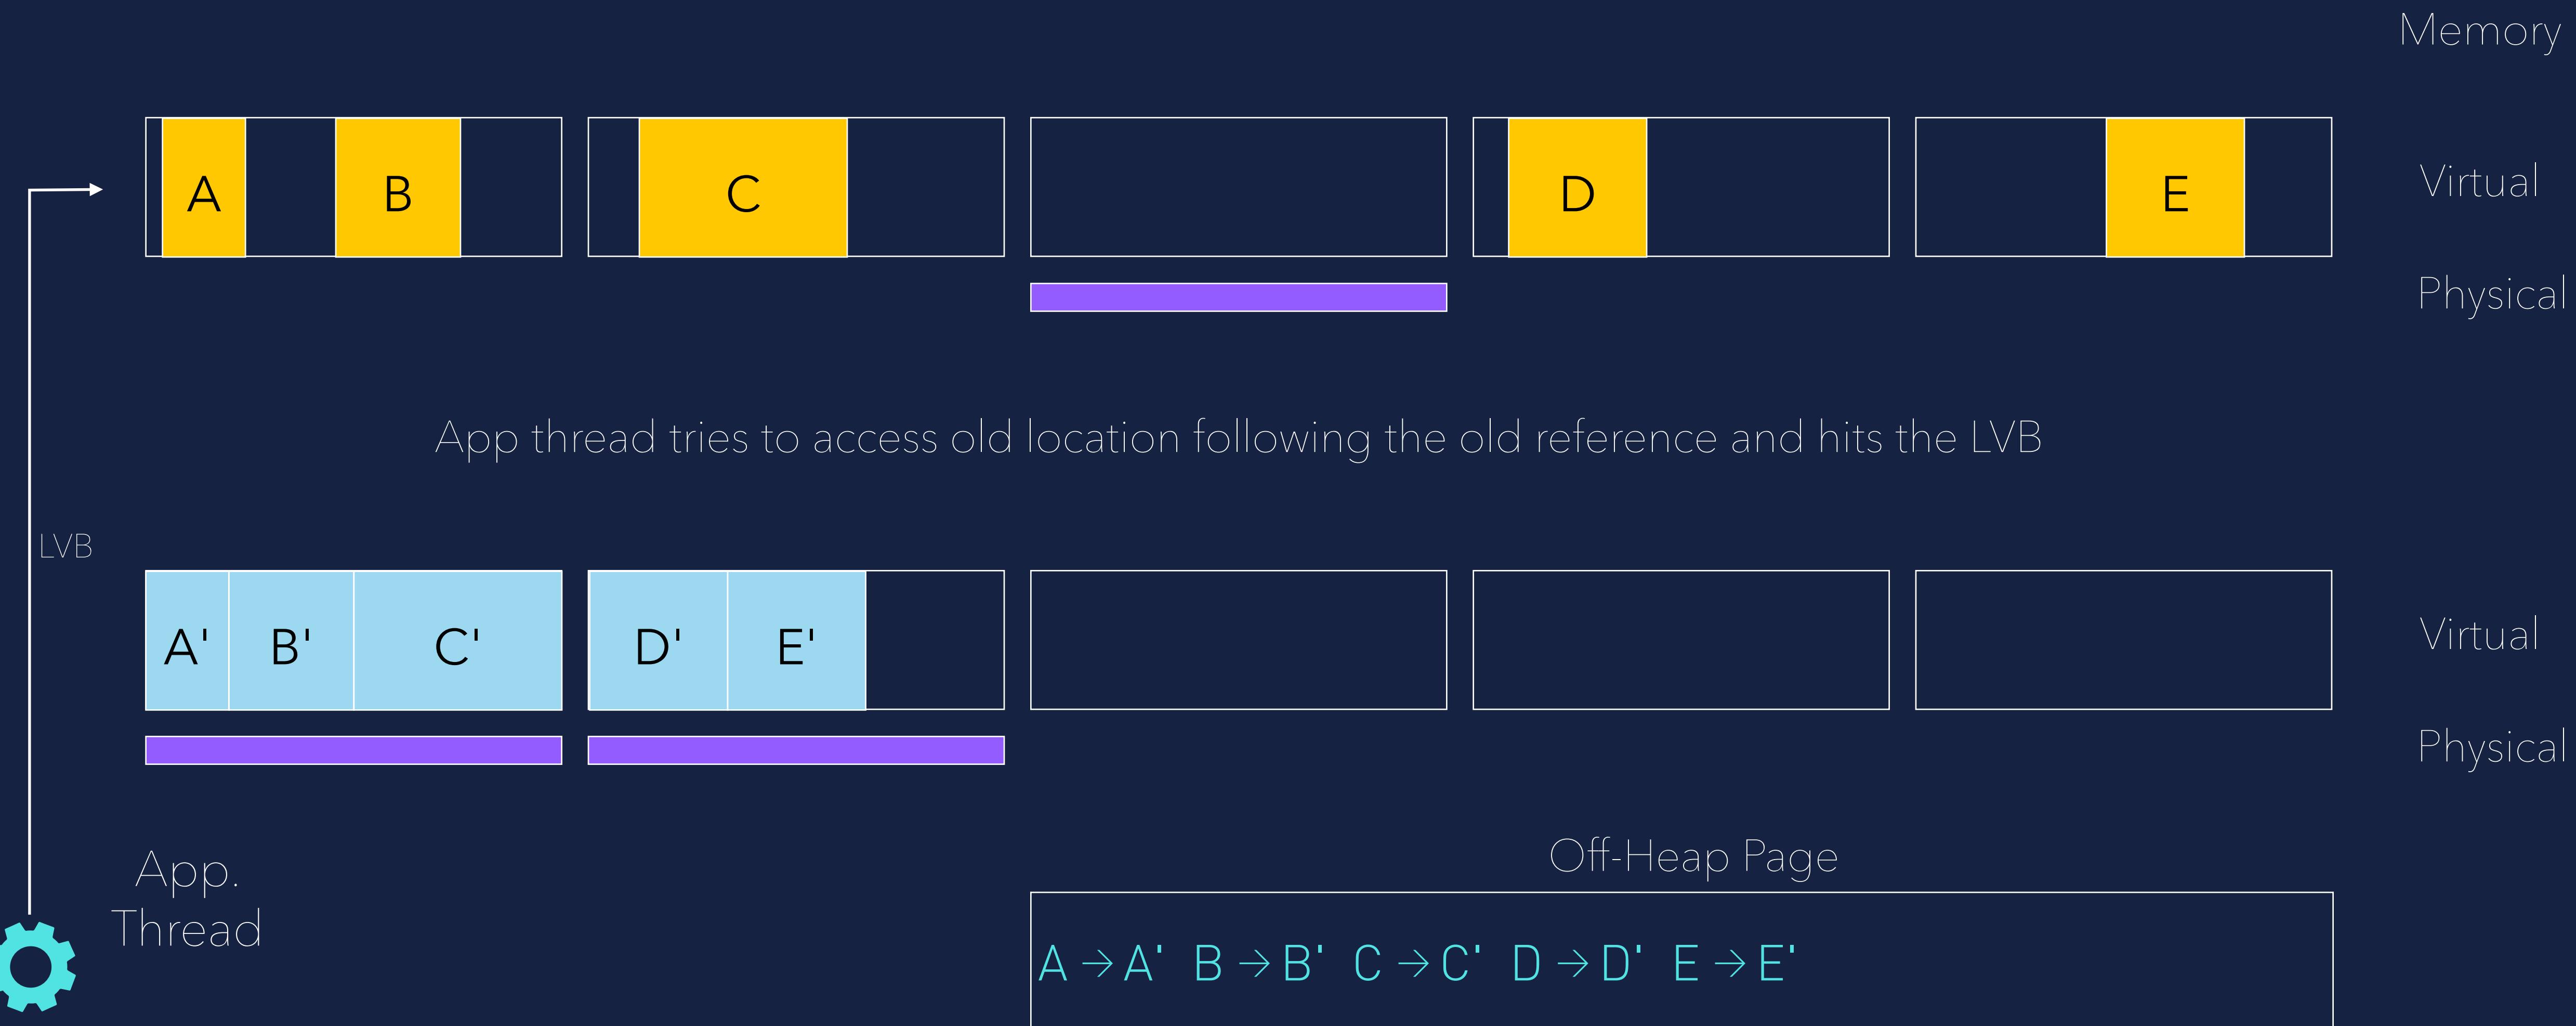
Relocation Phase (Compaction)



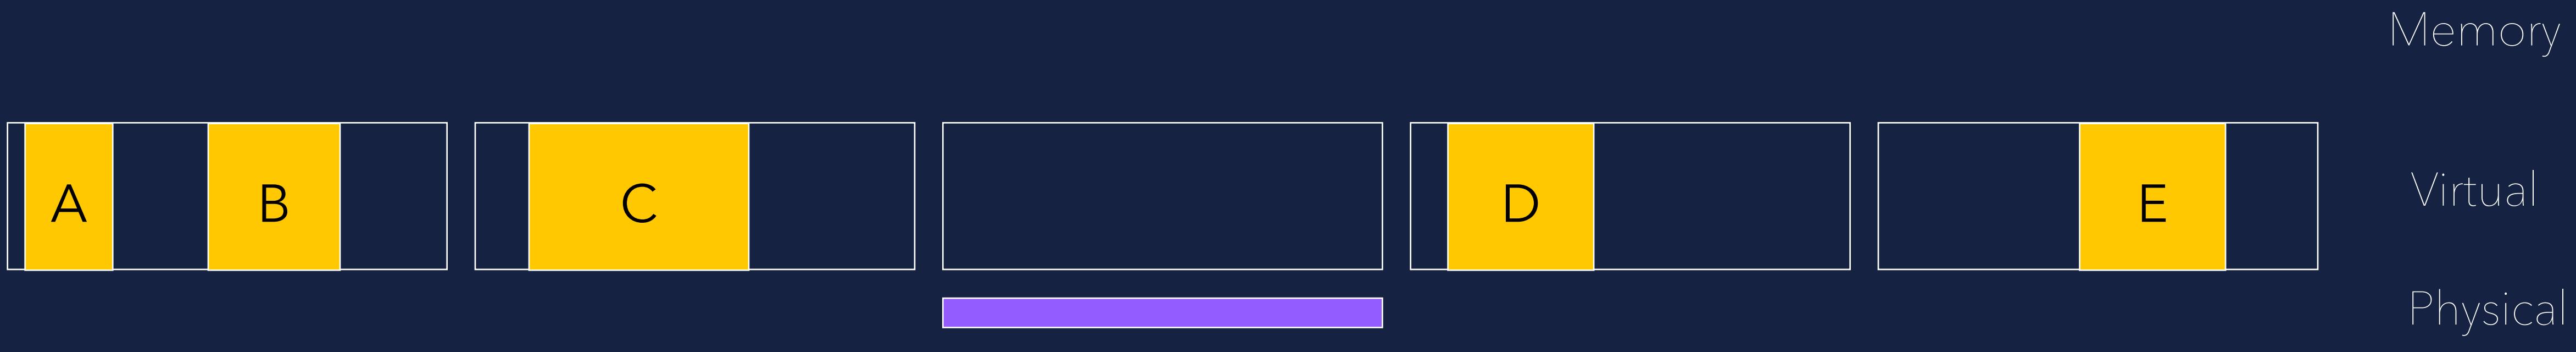
Relocation Phase (Compaction)



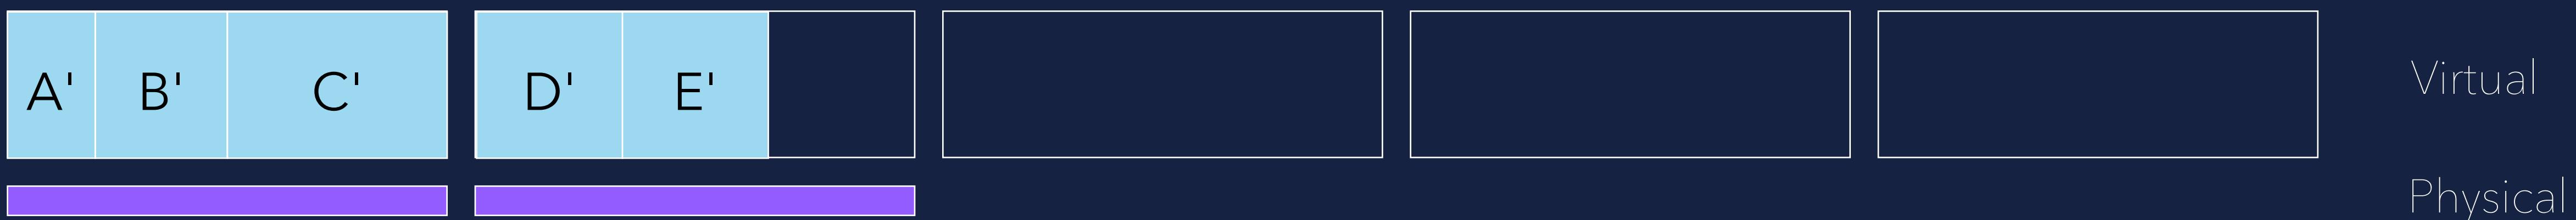
Relocation Phase (Quick Release)



Relocation Phase (Quick Release)



Gets new location from Off-Heap forwarding page



App.
Thread



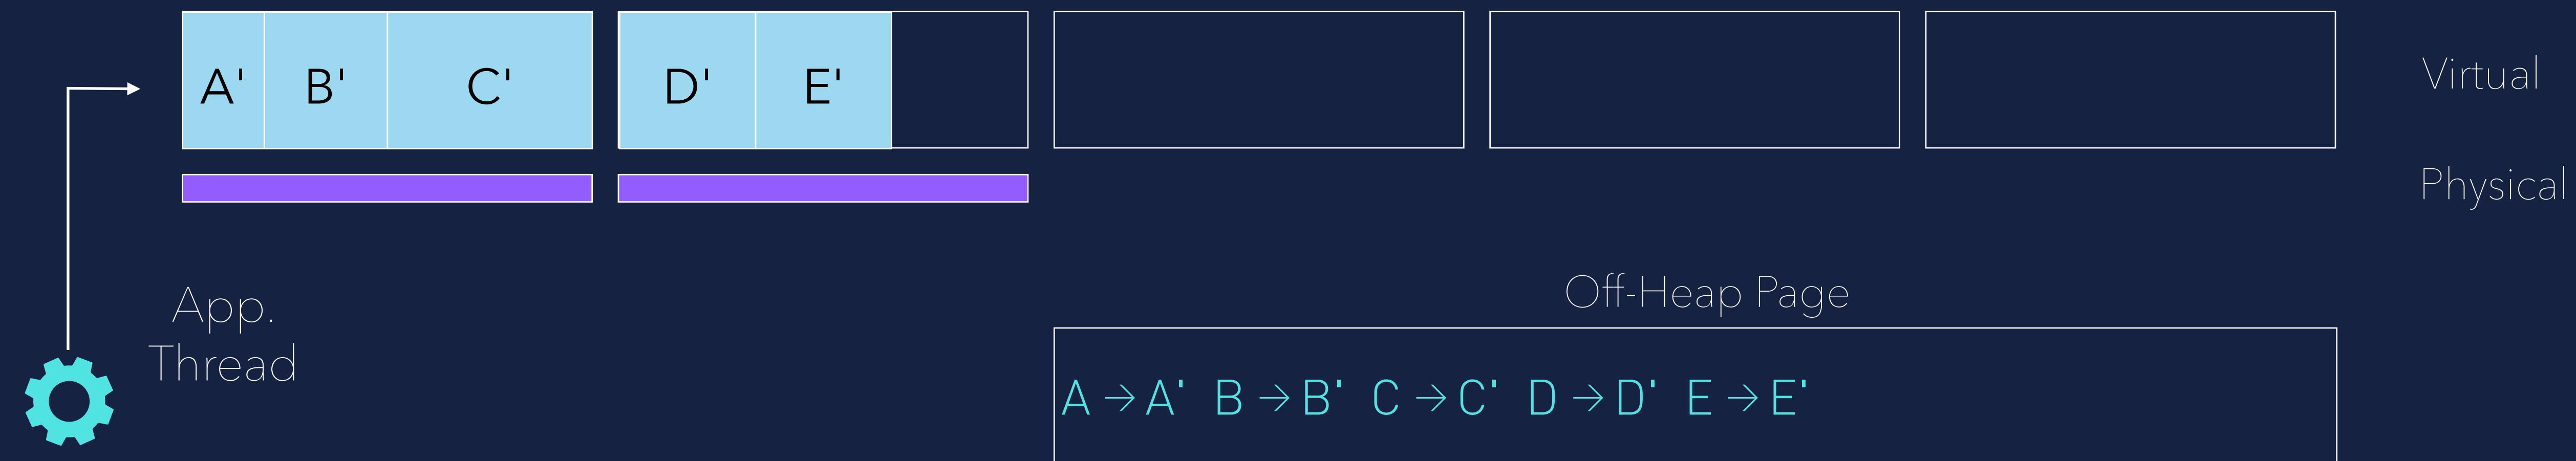
Off-Heap Page

A → A' B → B' C → C' D → D' E → E'

Relocation Phase (Quick Release)



Updates the reference and can access object at new location



AVAILABILITY	AZUL ZING JVM
PARALLEL	YES
CONCURRENT	FULLY
GENERATIONAL	YES
HEAP SIZE	LARGE
PAUSE TIMES	SHORT
THROUGHPUT	VERY HIGH
LATENCY	VERY LOW
CPU OVERHEAD	MODERATE (10-20%)

CHOOSE WHEN

-  Response time is a high priority
-  Using a very large heap (100GB+)
-  Predictable response times needed

BEST SUITED FOR

-  Low latency sensitive applications
-  Large scale systems
-  Highly concurrent applications

OS SUPPORT



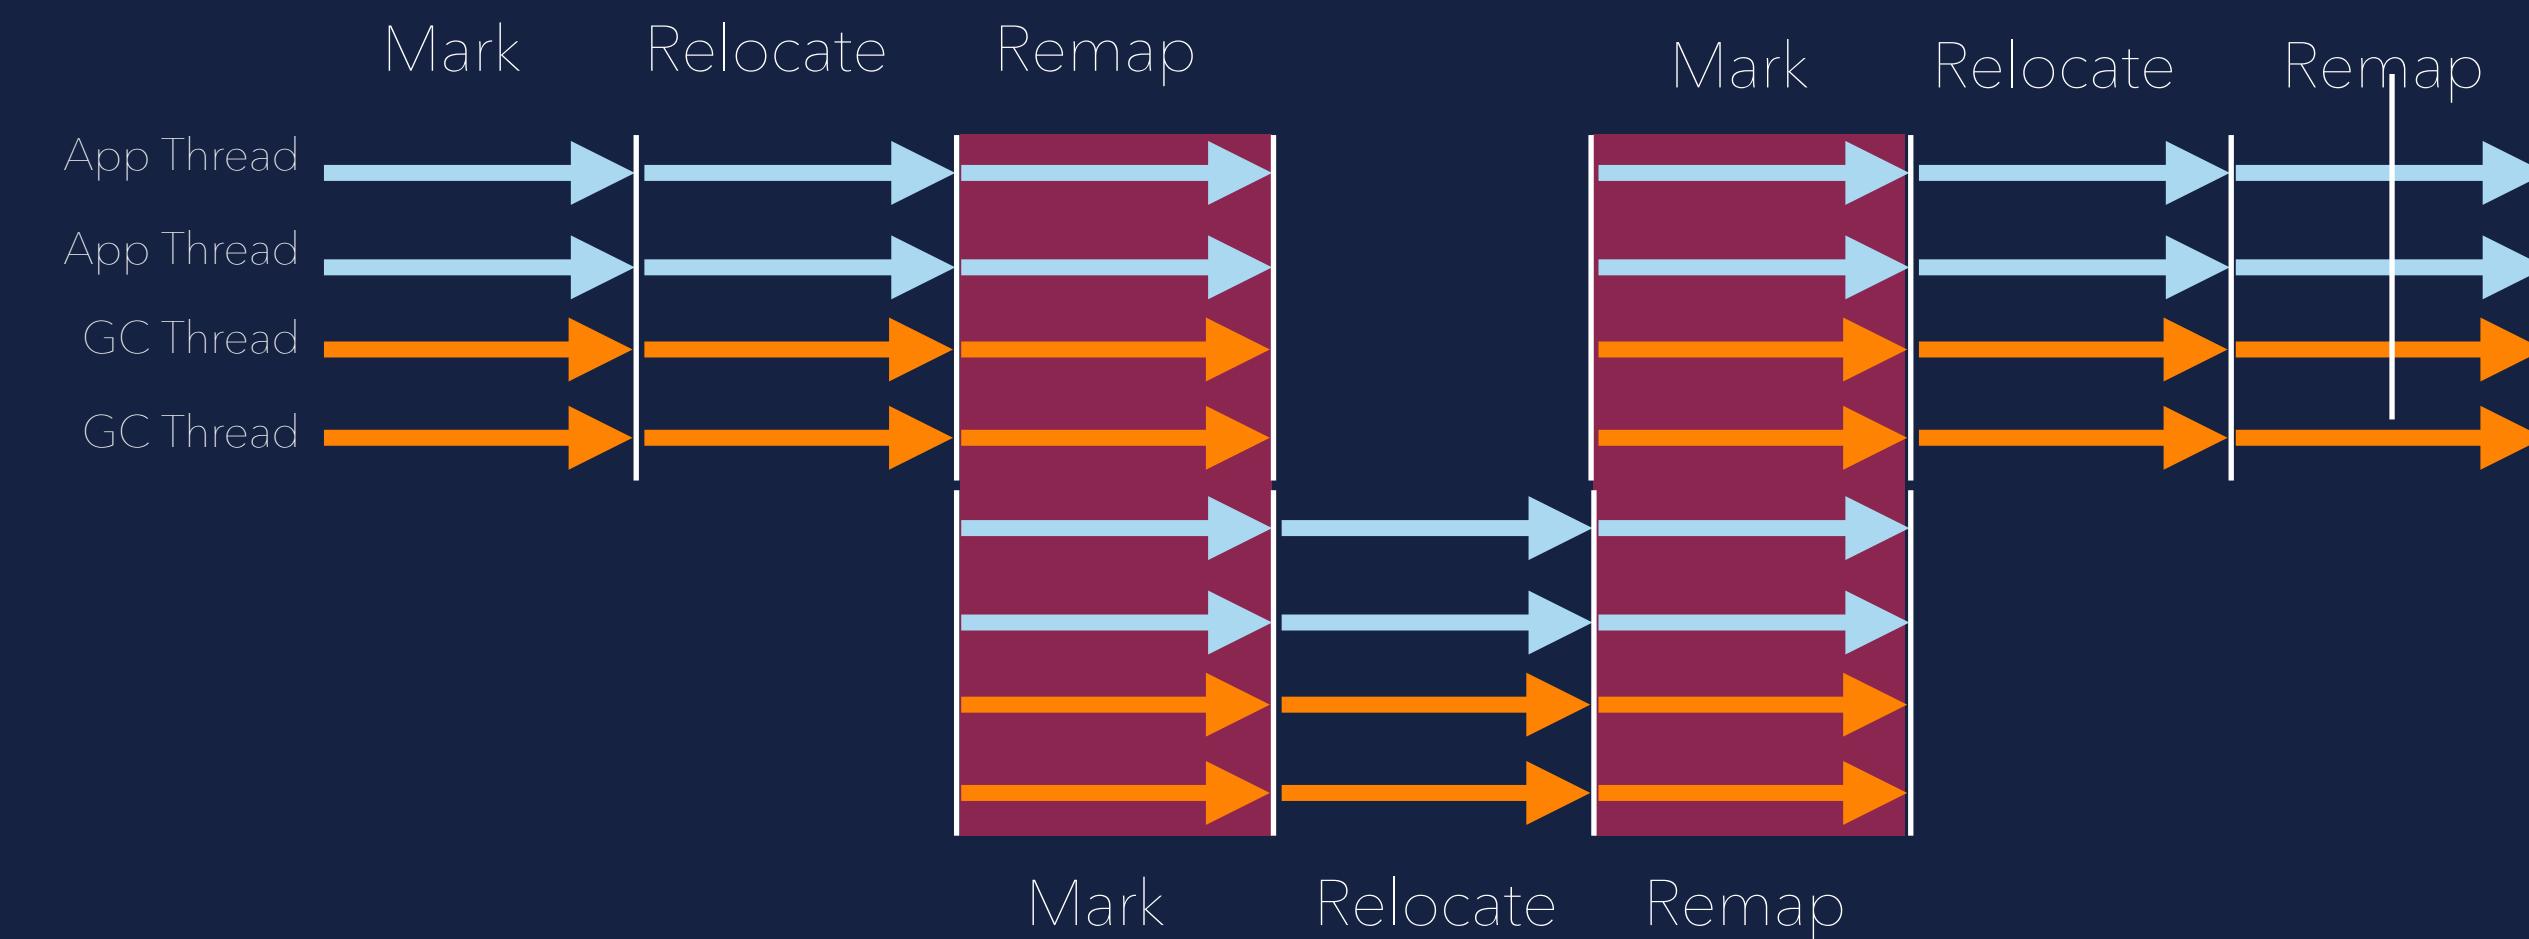
JVM SWITCH

> -

NOTES

Only available in Azul Zing JVM

No performance overhead because of faster Falcon compiler



WHICH ONE...?

WHICH ONE...?

Essential Criteria



Throughput

Percentage of total time spent in application vs. memory allocation and garbage collection

WHICH ONE...?

Essential Criteria



Throughput

Percentage of total time spent in application vs. memory allocation and garbage collection



Latency

Application responsiveness, affected by gc pauses

WHICH ONE...?

Essential Criteria



Throughput

Percentage of total time spent in application vs. memory allocation and garbage collection



Latency

Application responsiveness, affected by gc pauses



Resource usage

The working set of a process, measured in pages and cache lines

WHICH ONE...?

Essential Criteria



WHICH ONE...?

Choose dependent on your workload



* C4 has less overhead due to faster Falcon compiler

OVERVIEW

OVERVIEW

	Serial GC	Parallel GC	CMS GC	G1	Epsilon	Shenandoah	ZGC	C4
Availability	ALL JDK's	ALL JDK's	JDK 1.4-13	JDK 7u4+	JDK 11+	JDK 11.0.9+	JDK15 / 21+	Azul Zing 8+
Parallel	NO	YES	YES	YES	YES	YES	YES	YES
Concurrent	NO	NO	PARTIALLY	PARTIALLY	YES	FULLY	FULLY	FULLY
Generational	YES	YES	YES	YES	YES	NO	NO / YES	YES
Heap Size	SMALL - MEDIUM	MEDIUM - LARGE	MEDIUM - LARGE	MEDIUM - LARGE	YES	LARGE	VERY LARGE	VERY LARGE
Pause Times	LONGER	MODERATE	MODERATE	SHORT - MEDIUM	YES	VERY SHORT (<10ms)	VERY SHORT (<1ms)	VERY SHORT (<1ms)
Throughput	LOW	VERY HIGH	MODERATE	HIGH	VERY HIGH	VERY HIGH	VERY HIGH	VERY HIGH
Latency	HIGHER	LOWER	MODERATE	LOWER	YES	VERY LOW	VERY LOW	VERY LOW
Performance	LOWER	HIGHER	MODERATE	HIGHER	VERY HIGH	VERY HIGH	VERY HIGH	VERY HIGH
CPU Overhead	LOW	LOWER	MODERATE	MODERATE	VERY LOW	LOW - MODERATE	MODERATE - HIGH	MODERATE - HIGH
Tail latency	HIGH	HIGH	HIGH	HIGH	YES	Moderate	LOW	LOW

TOOLING...

TOOLING

 -Xlog:gc*:gc-%p-%t.log:tags,uptime,time,level:filecount=10,filesize=128m
Output of garbage collection details to log file

 jstat
Tool that provides info on performance and resource consumption of running applications

TOOLING

JITWatch

A tool for understanding the JVM JIT (<https://github.com/AdoptOpenJDK/jitwatch/wiki>)

jHiccup

A non intrusive tool to monitor platform "hicups" incl. JVM stalls (<https://github.com/giltene/jHiccup>)

VisualVM

All in one Java troubleshooting tool (<https://visualvm.github.io/>)

GCeasy

Universal GC Log Analyzer (<https://gceasy.io>)

JProfiler

All in one Java profiler (<https://www.ej-technologies.com/jprofiler>)

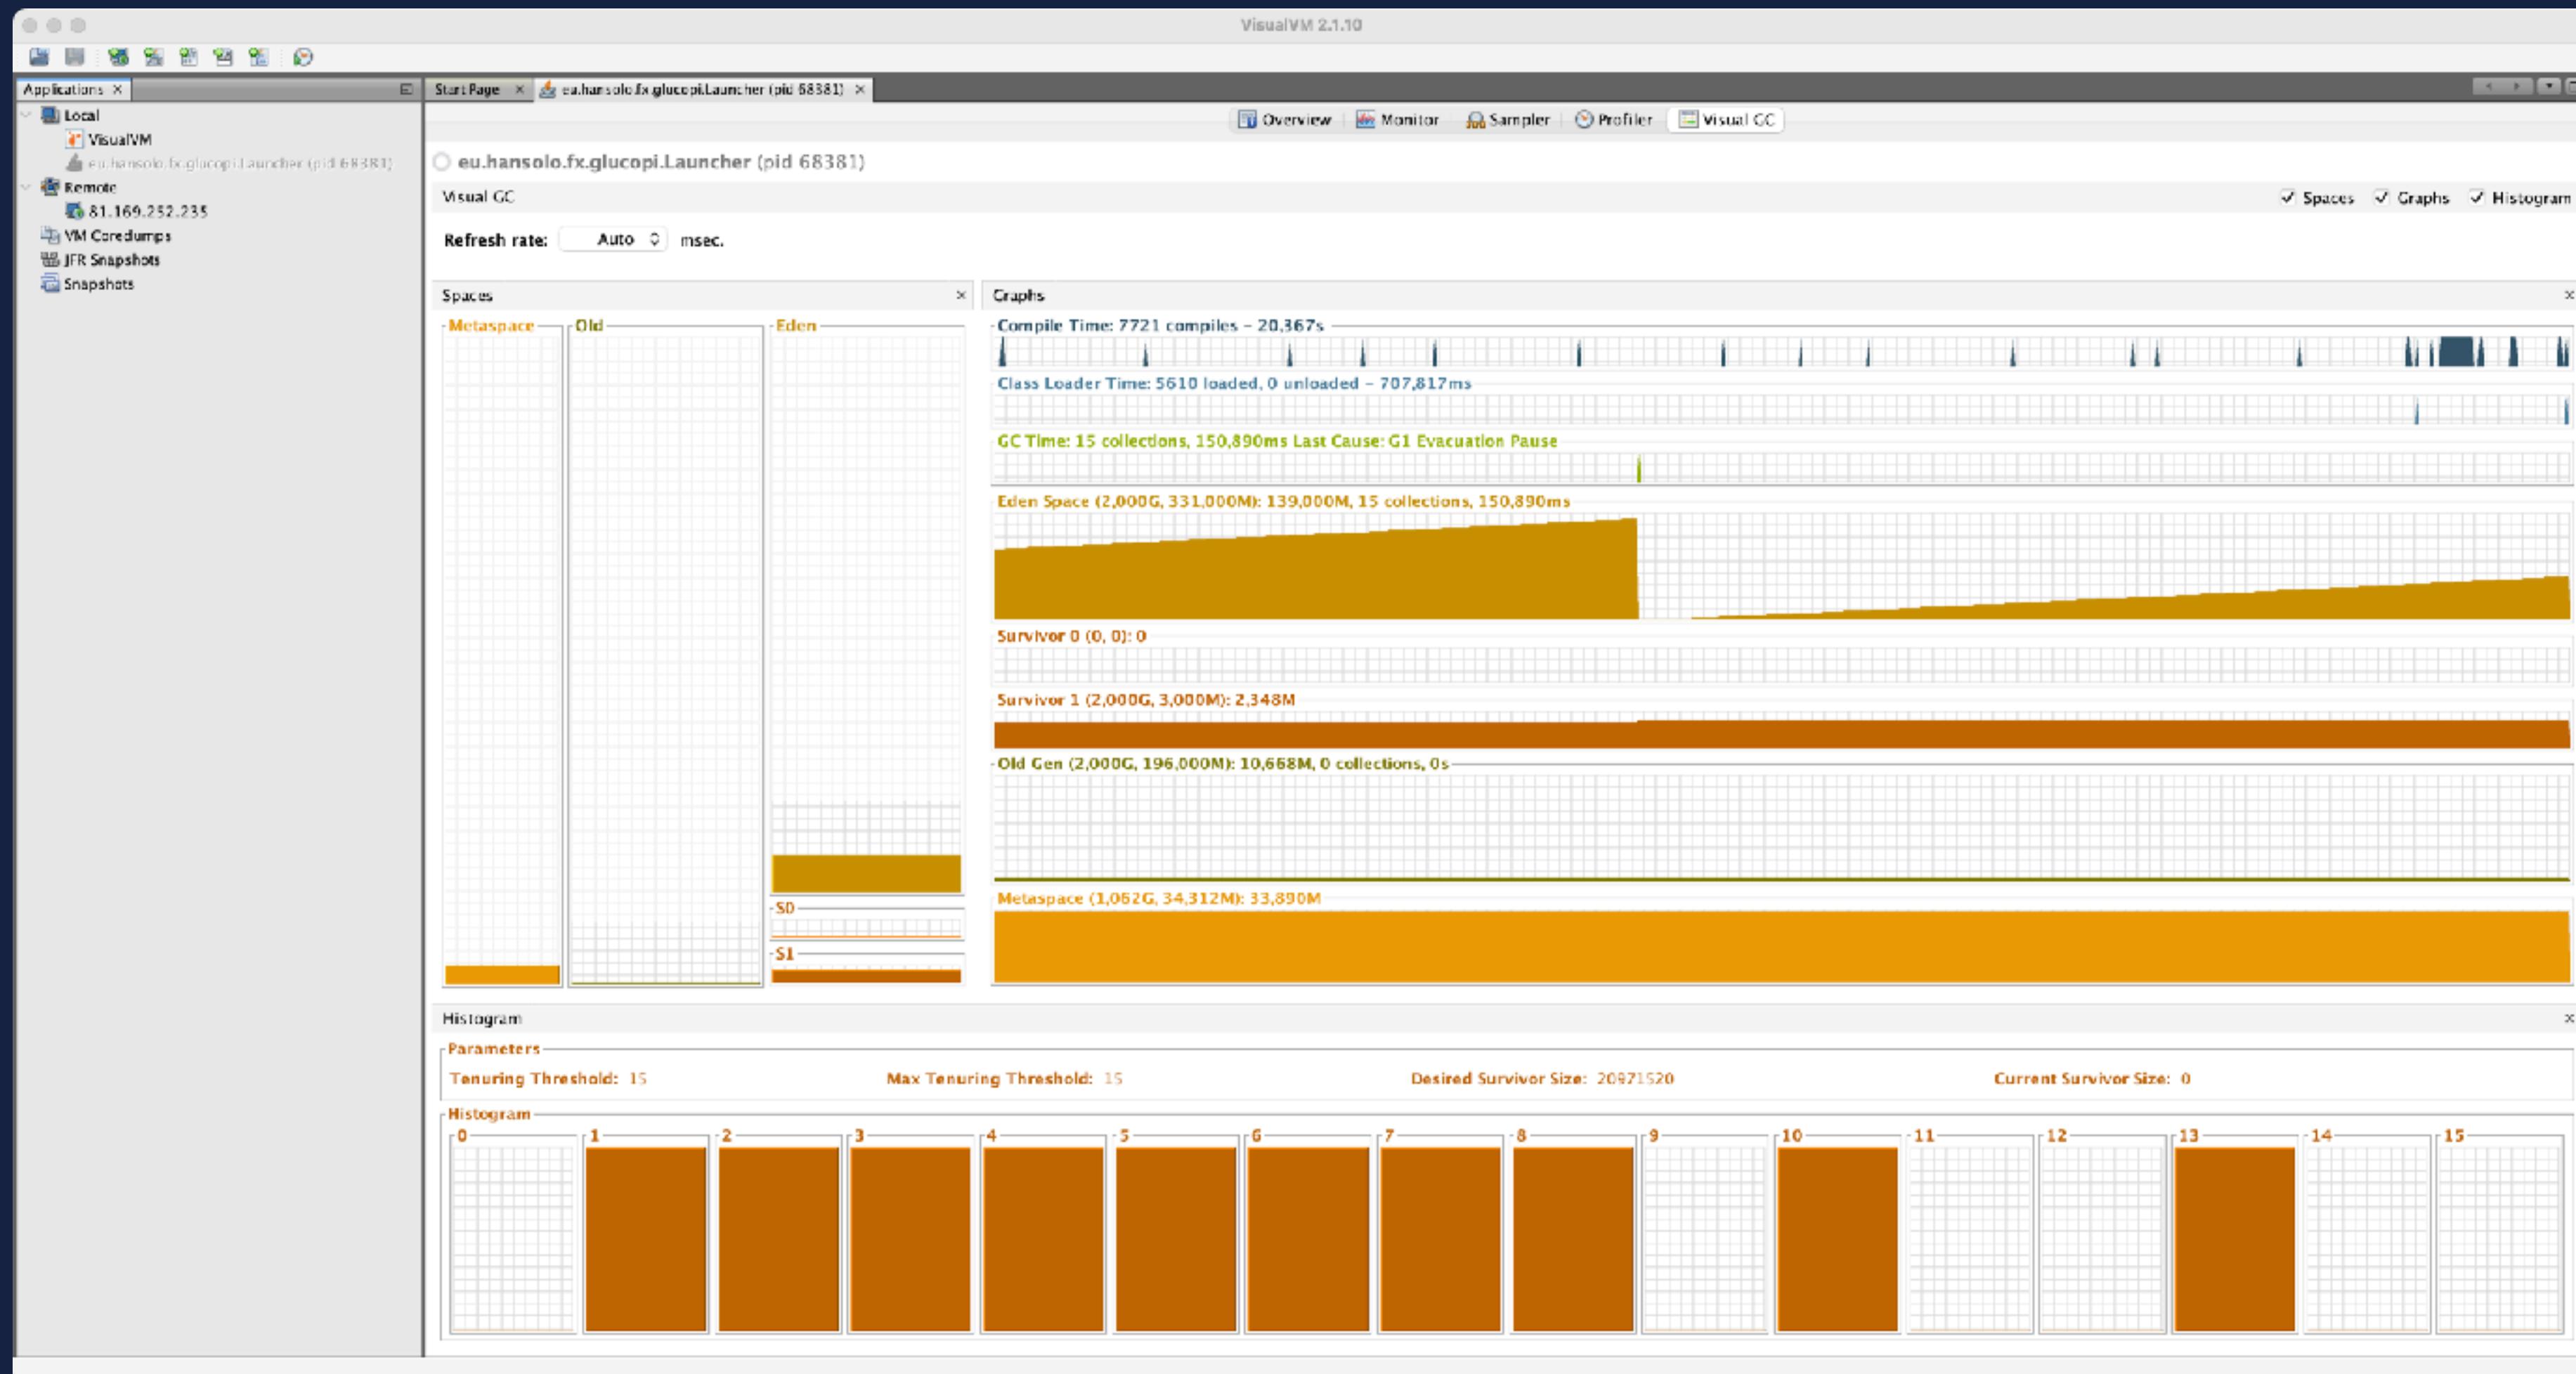
\$

YourKit Java Profiler

CPU and Java profiler (<https://www.yourkit.com/features/>)

TOOLING

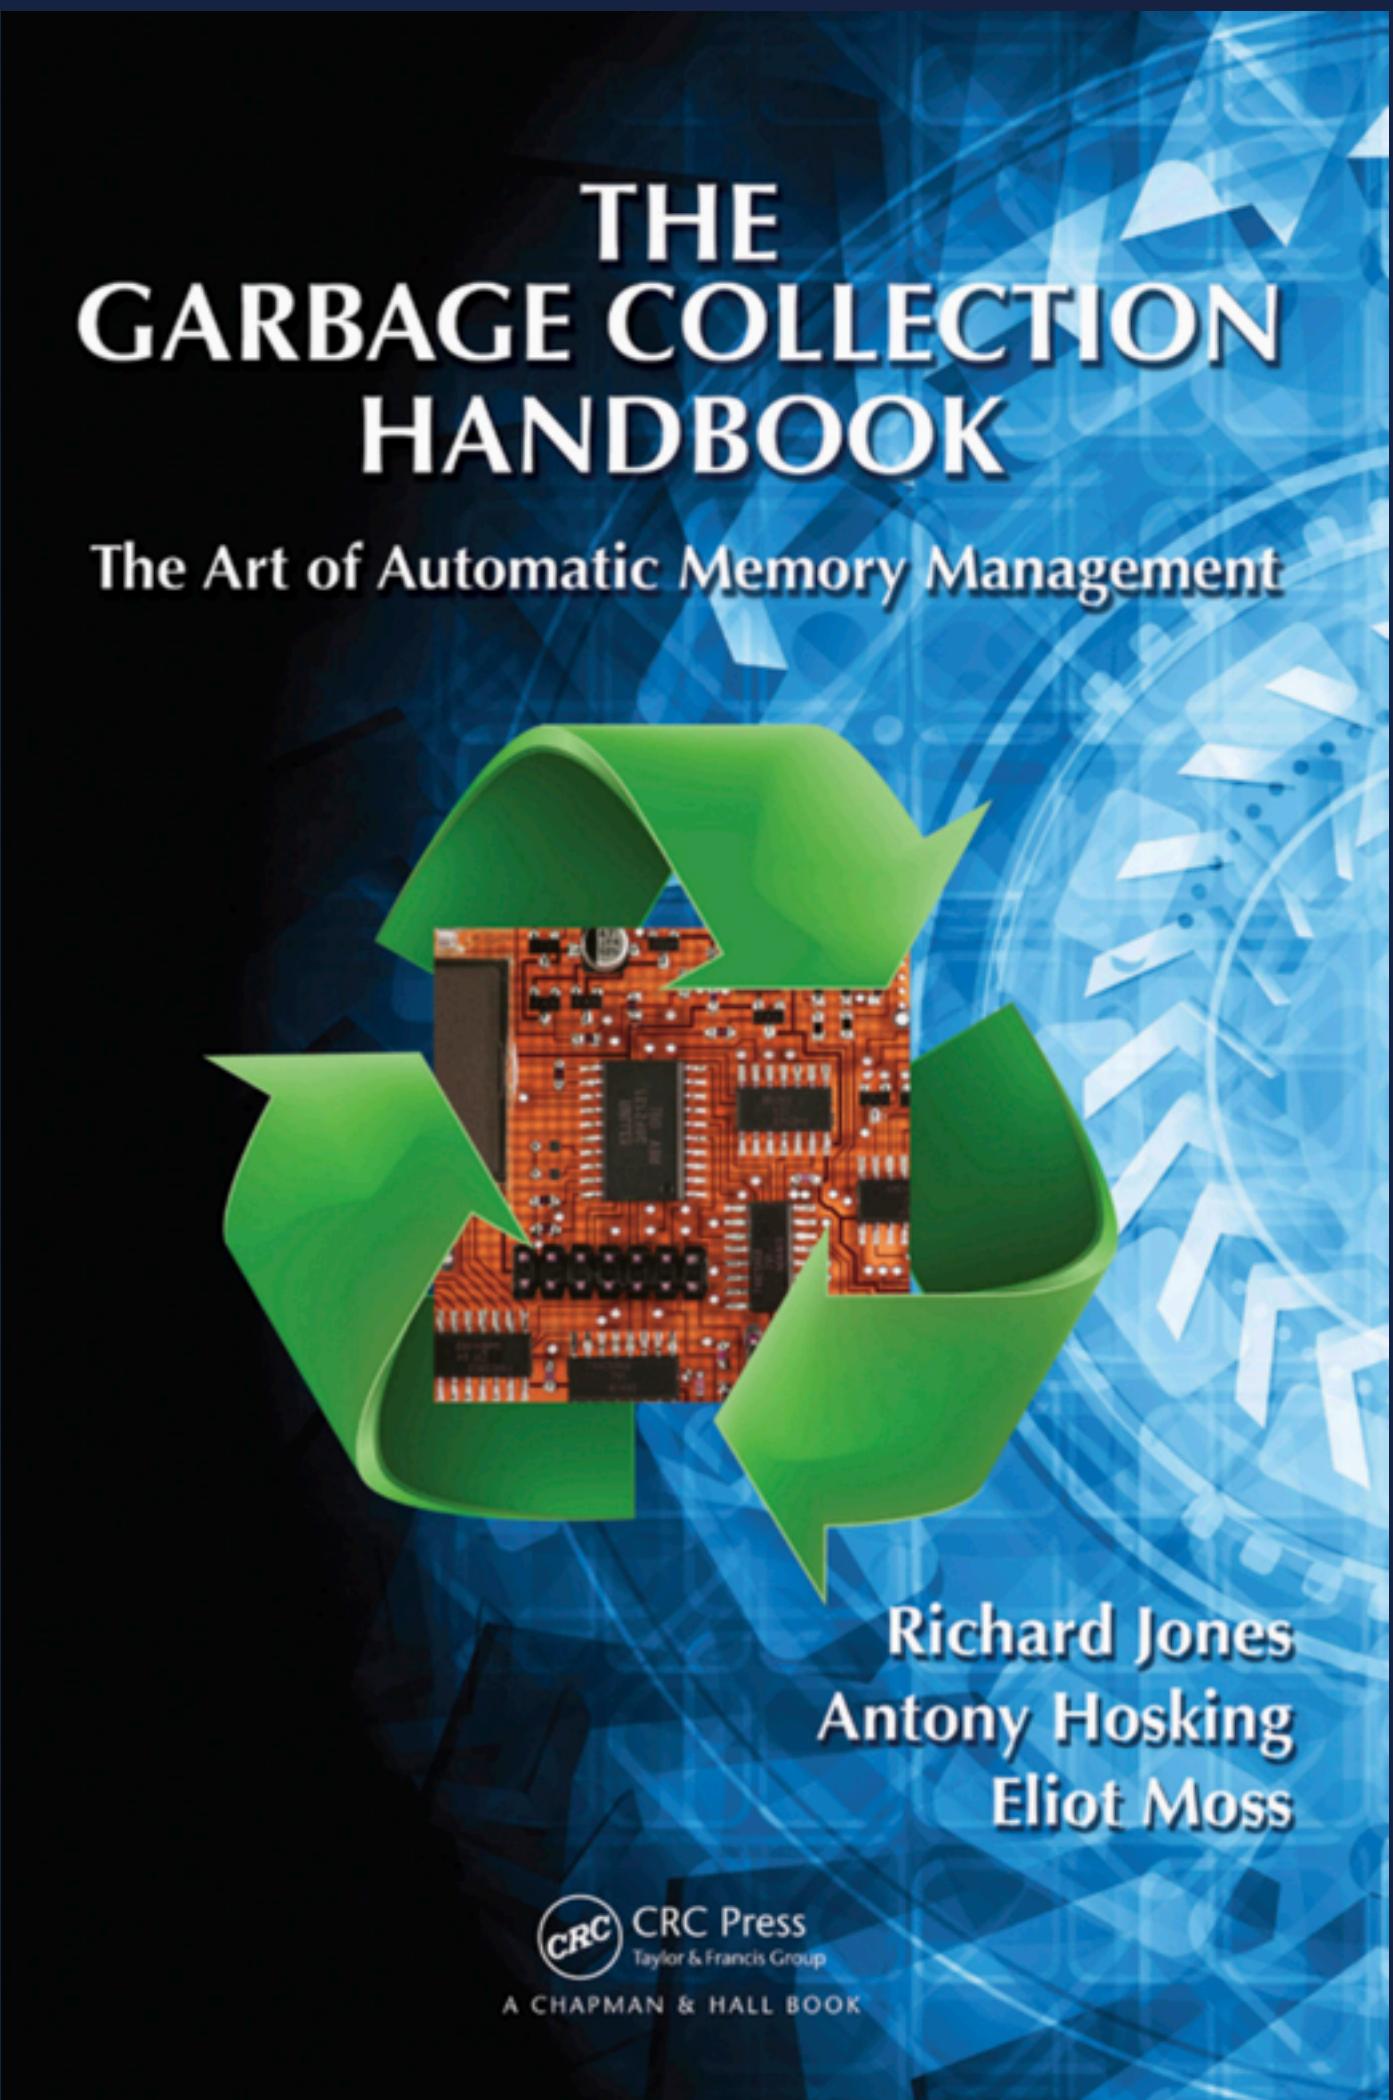
VisualVM + VisualGC plugin



WANNA KNOW
MORE?

WANNA KNOW MORE ?

R. Jones et al. "The Garbage Collection Handbook". Chapman & Hall/CRC, 2012



THANK YOU

